

A METAETHICAL EXPEDITION

THE OBJECTIVE MORALITY PROJECT

An Expedition to Build a Moral System
Everyone Has Reason to Obey

PHIL STILWELL

Copyright © 2026 Phil Stilwell

All rights reserved. No part of this book may be reproduced or distributed in any form without written permission from the copyright holder, except for brief quotations used in reviews, criticism, or scholarship.

The pressure cases in this book are philosophical constructions. Names, institutions, and places used in them are fictional unless otherwise indicated.

Contents

How to Use This Book	vi
Prologue — The Blank Constitution	1
 Part I — Specify the Thing You Hope to Build	 7
Chapter 1 — The Promise	8
Chapter 2 — What “Objective” Must Mean	17
Chapter 3 — What Are You Doing When You Say “Wrong”?	27
Chapter 4 — The Specification Sheet	37
Chapter 5 — One Argument, Eight Disagreements	47
 Part II — Give the System Content	 61
Chapter 6 — Draw the Circle	62
Chapter 7 — Decide What Receives a Verdict	70
Chapter 8 — Name What Matters	79
Chapter 9 — Turn Values into Guidance	87
Chapter 10 — Choosing Without a Common Measure	95
Chapter 11 — Like Cases, Exceptions, and Identity	103
 Part III — Find the Source of the “Ought”	 115
Chapter 12 — Three Different Things Called a Reason	116
Chapter 13 — Current Aims and Better-Informed Aims	124
Chapter 14 — Commitments, Promises, and Roles	133
Chapter 15 — Fair Terms and Constructed Reasons	141

Chapter 16 — Reasons That Apply Whether You Care or Not	150
Chapter 17 — From Approval to Obligation	159
Chapter 18 — God, Nature, and Rational Necessity	169
 Part IV — Ask How the System Could Be Known	 187
Chapter 19 — Does Evidence Have Authority?	188
Chapter 20 — Intuition and the Feeling of Obviousness	197
Chapter 21 — Where Values Come From	205
Chapter 22 — Disagreement and Convergence	214
Chapter 23 — Facts Enter; Values Select	224
Chapter 24 — Novel Cases, Uncertainty, and Correction	233
 Part V — Let the Great Candidates Make Their Best Case	 251
Chapter 25 — Count the Consequences	252
Chapter 26 — Respect the Rule and the Person	259
Chapter 27 — Character, Care, and Belonging	266
Chapter 28 — Receive the Moral Order	273
Chapter 29 — Build What No One Could Reasonably Reject	280
Chapter 30 — Let Standards Remain Local or Plural	287
Chapter 31 — The Hybrid Rescue	294
 Part VI — The Last Repairs and the General Result	 307
Chapter 32 — Criticize the World You Inherited	308
Chapter 33 — Build a Public Code Without the Missing Layer	315
Chapter 34 — Power, Dissent, and the Person Who Still Says No	324
Chapter 35 — Responsibility Without Cosmic Payback	332

Chapter 36 – Moral Emotions Without Moral Desert	340
Chapter 37 – No Candidate Returns Intact	348
 Part VII – After the Search	 361
Chapter 38 – Decide What to Do with Moral Language	362
Chapter 39 – Everything That Did Not Disappear	370
Chapter 40 – Begin with the Primitives	380
Chapter 41 – Stop Reconciling Moral Systems	393
Appendix A – The Candidate System Workbook	404
Appendix B – Pressure Cases	434
Appendix C – Common Repairs and Their Debts	448
Appendix D – The Civic Stakes Brief	457
Appendix E – Glossary and Position Map	480
Appendix F – Notes on the Limits of the Conclusion	493
Notes and Further Reading	501
Bibliography	509

How to Use This Book

This is an argument you perform.

You will begin by trying to build the thing many people mean when they speak of *objective morality*: a coherent, knowable, action-guiding system whose verdicts do not become true merely because someone approves of them and whose authority reaches even an informed rejecter. The book will not ask you to choose a favorite theory and cheer from the stands. It will put you at the drafting table.

Keep a notebook, a document, or a clean copy of Appendix A beside you. Your evolving construction is called the **Candidate System**. Whenever a premise, bridge, exception, or access route needs support, enter it in the **Debt Ledger**. Do not erase early answers. Date revisions and mark each as a clarification, discovery, repair, or surrender. The history of the attempt is part of the evidence.

The exercises are not tests of personality, empathy, or philosophical sophistication. A pressure case holds certain facts fixed so that one relation can be examined. Answer the exact question before improving the story. If you want to add a fact, write it as a controlled variation and note what changed. Cleverness is useful here only when it leaves fingerprints.

Four recurring devices organize the expedition:

- **The Candidate System** is the moral constitution you construct and revise.
- **The Debt Ledger** records what each claim still owes.
- **Pressure Cases** isolate seams by fixing tempting escape routes.
- **Inspection Reports** preserve what survived each movement and what was surrendered.

Two brief annotations help you recover the thread when the detail becomes dense. **Where this chapter is headed** appears after each opening scene and names the immediate task and its significance to the construction. **Purpose** appears in reusable worksheets and policy forms to state what the page produces and what completing it does not establish. These are route markers, not substitutes for the argument.

The recommended route is straight through the prologue and all seven parts. The final result depends on distinctions accumulated along the way. If you return later with a narrower question, three shorter routes are useful:

- For objective authority, read the Prologue; Chapters 1–5, 12–24, and 37–39; then Appendix F.
- For public policy after the search, read Chapters 5–11, 22–24, 29, and 32–41; then complete Appendix D.
- For responsibility, blame, and moral emotion, read Chapters 3, 7, 21, and 35–41; then consult Appendix E.

No background in metaethics is assumed. Unfamiliar terms are introduced when they begin doing work, and Appendix E provides a map when the terrain starts looking suspiciously like terrain.

One request governs the journey: give every candidate its strongest honest form. The aim is not to make a theory stumble. It is to discover what the theory must carry all the way home.

Prologue — The Blank Constitution

The first page is blank.

Around you, two hundred and forty-nine other delegates are pretending not to be intimidated by it.

You have been brought together to write the founding constitution of Port Meridian, a new coastal city assembled from communities displaced by drought, flood, and political collapse. Its residents have arrived with different religions, different languages, different family structures, different memories of government, and different ideas about what a decent life requires. Some escaped authoritarian states. Some escaped states too weak to protect them. Some regard individual freedom as the first political value. Others regard social fragmentation as the disaster from which freedom must be rescued. Some believe that moral law comes from God. Some believe that morality is a human construction. Some distrust the word *morality* altogether.

They nevertheless need rules.

Children must be educated. Water must be allocated. Contracts must be enforced or replaced by something else. Public officials must be selected. Dangerous conduct must receive a response. Religious communities must be protected, regulated, ignored, or some unstable combination of the three. The shoreline must be defended. Surveillance systems must be permitted or prohibited. Hospitals must decide whose needs take priority when equipment is scarce. People must know what they may do, what they may demand, and what will happen when they refuse.

The assembly has not asked you to draft a temporary bargain. It has asked for something more ambitious: a moral foundation.

The chair of the assembly explains the assignment.

“We do not want a code that is legitimate only because a majority happened to approve it. Majorities have approved slavery, conquest, caste, censorship, forced marriage, religious persecution, and the exclusion of whole classes of people. We do not want a code whose authority disappears whenever someone ceases to care. We do not want rules that become correct on one side of a border and incorrect on the other merely because different populations cheer for them. We want the right principles. We want to know what every person is actually obligated to respect.”

Several delegates applaud. A few look uneasy.

The chair continues.

“Our constitution must be coherent. It must tell us who counts, what matters, what may never be done, when exceptions are justified, how conflicts should be settled, and why any of this binds the person who rejects it. It must be public enough to guide us and objective enough to correct us. We must be able to say not only that an atrocity violates *our* code, but that the code permitting it was mistaken.”

Now everyone is looking at the same blank page.

What would you write first?

Perhaps you would begin with harm:

An action is right when it prevents the greatest suffering and makes life go best for everyone affected.

Perhaps you would begin with persons:

An action is right when it respects every person as an end and never uses anyone merely as a means.

Perhaps with freedom:

Every person has an equal claim to direct their own life, limited only by the equal freedom of others.

Perhaps with fairness:

The right principles are those that no one could reasonably reject under conditions of equal standing.

Perhaps with God:

Right action conforms to the will and nature of the perfectly good creator.

Perhaps with human nature:

Right action enables human beings to flourish according to the kind of beings they are.

Or perhaps you would resist beginning with any one of these. Suffering matters, but so do consent, loyalty, honesty, beauty, freedom, security, belonging, fairness, character, and trust. You might try to build a plural system in which several values constrain and correct one another.

That may be the sensible place to begin. But it is only the beginning.

If suffering controls the answer, whose suffering counts? How do you compare one person's terror with a thousand people's inconvenience? May a painless violation be ignored? May a person be sacrificed when enough others benefit?

If respect controls the answer, what counts as a person? What does respect require when two people make incompatible demands? If freedom controls the answer, which freedoms matter, and what happens when the exercise of one person's freedom creates another person's dependence? If fair agreement controls the answer, who participates, what makes the conditions fair, and why is the agreement authoritative for someone excluded from it?

If God controls the answer, how is the command known, how are rival interpretations resolved, and does goodness explain why the command has authority—or does the command define goodness? If flourishing controls the answer, who decides what counts as a flourishing life, and how does a fact about a species become a demand addressed to each member?

The blank page was easier.

Yet the difficulty does not show that the assignment is impossible. Every serious construction project begins with questions. A bridge engineer must specify loads, materials,

weather, maintenance, and acceptable risk. A constitution must specify offices, powers, limits, representation, amendment, and appeal. A scientific theory must specify what it explains, what evidence supports it, and what would count against it. Why should an objective moral system require less care?

This book asks you to attempt the construction.

You will not be asked merely to choose among famous theories. You will build the thing those theories are trying to be. You will decide who counts and why. You will distinguish counting for one's own sake from receiving equal weight or winning every conflict. You will decide what receives a moral verdict: actions, motives, consequences, rules, people, institutions, relationships, or all of them. You will identify the values that supply the system's content. You will define its central words. You will turn values into rules, specify exceptions, compare values that do not arrive in a common unit, and explain what makes one case relevantly similar to another.

Then you will face the hardest question: what gives any of those rules authority?

A person may love a rule without being obligated by it. A society may enforce a rule without making it correct. An agreement may create a responsibility among participants without binding every outsider. A punishment may give someone a powerful reason to comply without giving the punisher moral authority. A fact may explain why someone acted without making the action worthy of support. A desire may motivate without justifying. A fair procedure may produce stability while leaving a dissenter unconvinced. We will keep these possibilities separate long enough to see exactly what each one establishes.

We will also ask how objective moral truth could be known. Before relying on evidence, we will ask a prior question: what makes evidence count in favor of belief? Evidence can indicate truth, serve a truth-seeking project, improve practical success, follow rules of inquiry, or receive public procedural support. Whether any of those relations creates a stance-independent epistemic requirement is not settled by skepticism about morality. Confidence is not a method. Agreement is not necessarily discovery. Disagreement is not necessarily disproof. An intuition may be insight, trained perception, emotional salience, cultural inheritance, or some combination. A genealogy can explain a belief without refuting it. But a system that claims access to a stance-independent order must still explain how projection is distinguished from detection.

Finally, we will use what we have built. We will apply it to promises no one can enforce, suffering inflicted on enemies, sacrifice in emergencies, loyalty under pressure, social reform, artificial minds, dissenters, outsiders, punishment, and public policy. We will ask what becomes of guilt, shame, resentment, indignation, gratitude, and pride if suffering and esteem are not intrinsically deserved. We will allow the system to change. Every time a case exposes a weakness, you may repair the system rather than abandon it.

There is only one rule: repairs are not free.

If you add an exception, you must say what generates it. If you introduce a new value, you must say why it belongs. If you appeal to an ideal observer, you must say what makes the observer ideal. If you invoke human nature, you must say which features are normatively important. If you rely on practical wisdom, you must explain how wisdom

distinguishes perception from preference. If you appeal to self-evidence, you must say what justifies confidence when equally reflective people report incompatible certainties.

Every claim creates a debt of explanation. The Candidate System must keep a ledger.

That language may sound hostile, but it is not. Explanatory debts are simply the consequences of ambition. A system claiming little owes little. A household rule may be justified by the preferences of the people who share the house. A club may be governed by terms its members accepted. A law may identify the institution authorized to impose a penalty. A private commitment may organize one person's life. These practices can be clear and useful without being objective moral systems.

But a system that claims authority over every person, including the rejector, owes more. A system that claims to remain true if every mind in the universe opposes it owes more. A system that claims its verdicts are discoverable and action-guiding owes more. The size of the debt is set by the size of the claim.

You may discover that the debt can be paid. Perhaps one foundation will survive every pressure case. Perhaps several incomplete foundations will support one another. Perhaps objective morality will turn out to require less than its critics demand. Perhaps the construction will produce something more modest but still worthy of the name.

Do not decide in advance.

For now, take the blank page.

Where the prologue is headed. You will make a first attempt before the book supplies theory, then expose that attempt to one controlled pressure. The point is to create an honest starting record—not to pass a test or guess the author's conclusion.

The first entry

Purpose. Create an unrevised baseline in your own words; the entry records what first seems promising and is not graded for sophistication or agreement with the book.

Use a notebook, a document, or the margins of this book. Write the heading **Candidate System, Version 0.1**.

Complete this sentence without using the phrases “whatever is morally right,” “what a good person would do,” or “all things considered”:

An action is objectively right when ...

Do not worry about making the sentence sophisticated. Write what first seems promising. If you believe several conditions are needed, include them. If you believe objective rightness is grounded in God, reason, nature, welfare, freedom, agreement, dignity, or another source, say so directly.

Now add three short answers:

1. This standard is objective because ...
2. A person who rejects the standard is still bound because ...
3. We can know what the standard requires by ...

Leave space beneath each answer. You will return to them many times.

The first pressure

Suppose everyone in Port Meridian except one resident approves your standard. The city teaches it in schools, writes it into law, celebrates it in stories, and enforces it consistently.

The dissenter understands the standard and the reasons offered for it, but she still rejects it. She can violate one of its rules in complete secrecy. By stipulation, no one will discover the act, and it will bring her no punishment, shame, loss of trust, or later social consequence.

For the moment, set aside every other consideration. Do not ask whether the act causes harm, breaks a promise, or conflicts with one of her own aims. Any of those might give her a separate reason to refrain. We are testing one claim only:

Does the agreement of all the other residents, by itself, obligate the dissenter to comply?

Social agreement can accomplish a great deal. It can give a standard clarity, coordination, predictability, public legitimacy, stability, motivational force, and shared meaning. The question is whether any of those achievements also supplies the objective authority the assembly requested: authority that reaches someone who understands the standard, rejects it, and can evade every social response.

Record the distinction in the Debt Ledger:

Agreement provides _____.

To bind an informed dissenter, the system still needs _____.

Fill the first blank with whatever social agreement genuinely provides. Fill the second with whatever must connect that achievement to an obligation for the rejecter. If you think agreement itself supplies the connection, state why. That explanation becomes the first entry in the Debt Ledger.

The page is no longer blank.

Neither is the Debt Ledger.

PART I



Specify the Thing You Hope to Build

Before we can find an objective moral system, we must stop letting the word objective change costumes whenever the spotlight moves.

Chapter 1 — The Promise

Objective morality is attractive before it is philosophical.

Most people do not first encounter it as a theory about properties, truth conditions, or practical reasons. They encounter it as a promise. Some things, the promise says, are not merely disliked. Some people are not merely defeated. Some customs are not merely local. Some commands are not made legitimate by power. Some cruelties remain wrong when a ruler approves them, a population celebrates them, and a victim has no one strong enough to intervene.

That promise answers a deep political and personal fear: if morality is not objective, perhaps only preference and power remain.

Suppose a government removes children from a disfavored minority and distributes them to politically approved families. The policy is legal. It is popular. Its administrators sincerely believe it promotes national unity. The courts support it, the schools praise it, and most members of the majority experience the practice as normal.

To call the policy *wrong* seems to do more than report that we oppose it. We want to say the society is mistaken by a standard not created by its mistake. We want the separated families to have a claim that remains valid when no institution recognizes it. We want reformers to be correct before reform succeeds. We want the judgment to retain its status in an empty room.

Objective morality seems to promise all of that.

It also promises guidance. We do not only want a verdict about obvious atrocity. We want to know what to do when the claims are difficult: whether to tell a protective lie, how to distribute a scarce treatment, whether loyalty can justify an exception, and how much present sacrifice is required for future strangers. We also need to know when punishment becomes cruelty, which beings have standing, how much weight that standing receives, and how unlike values can be compared when every available policy imposes serious costs on someone.

And it promises authority. A moral system is expected to do something that a prediction, preference, or threat does not. “People will punish you” may influence behavior. “I hate what you did” may express a stance. “Our rules forbid it” may classify an act. “You promised” may invoke a relationship. A moral demand seems to claim more: that a consideration genuinely counts in favor of acting, whether or not the agent is moved by it.

These are different promises. They are often bundled into one word.

Before we build anything, we need to know which promises the finished system must keep.

Where this chapter is headed. We will unpack the promise hidden inside *objective morality* and identify what a successful system would have to explain. Until that promise is explicit, agreement, law, emotional force, and objective authority can look interchangeable.

The sentence that sounds complete

Consider this exchange:

“Why should I return the money?”

“Because keeping it would be wrong.”

The second sentence sounds like an answer. Grammatically, it is an answer. But what information has it supplied?

Perhaps the money belongs to someone else. Perhaps returning it would preserve trust. Perhaps it was received under a clear agreement. Perhaps keeping it would expose another person to hardship. Perhaps the law requires its return. Perhaps the speaker condemns taking what was not freely given. Perhaps the agent values honesty and wants to remain the kind of person who can be trusted.

Each of those considerations may matter. The word *wrong* does not identify which one is doing the work.

Now ask again:

“Why should I return the money?”

“Because it would be wrong not to.”

“What makes it wrong?”

“Because you ought to return what is not yours.”

“What makes that something I ought to do?”

“Because keeping what is not yours is wrong.”

The conversation has made a circle without moving. The terms *wrong* and *ought* pass authority back and forth. Neither explains where it came from.

This does not show that the money may be kept. It shows that a verdict and a justification are different things.

The difference matters because familiar moral sentences can borrow their apparent completeness from the practical considerations surrounding them. If keeping the money will cause an identifiable hardship, we can describe the hardship. If the person gave it under a condition, we can describe the agreement. If property law applies, we can identify the jurisdiction and the likely remedy. If trust will be damaged, we can explain the dependence that makes trust important to the relationship. If the agent cares about honesty, we can trace the action to a commitment.

Once these things are said, the moral status may appear to summarize them. But summary is not grounding. The word *wrong* may gather the considerations into one urgent label. It does not, merely by being spoken, add an objective property that explains why they count.

The Candidate System must therefore do more than produce verdicts. It must expose the route to them.

A verdict is not a system

Write three plausible judgments:

- Cruelty is wrong.
- People deserve equal respect.
- Promises should be kept.

They seem like the beginnings of a moral system. In fact, each is a bundle of unanswered design decisions.

“Cruelty is wrong” does not yet tell us what cruelty is. Does cruelty require an intention to cause suffering? Can indifference be cruel? Can an institution be cruel when no participant intends the total effect? Is humiliating someone cruel if it causes no lasting injury? Is frightening a child during necessary medical treatment cruel? Does cruelty make an act forbidden, blameworthy, punishable, regrettable, or all four?

“People deserve equal respect” leaves different questions. Who qualifies as a person? Does a fetus? Does a person in irreversible unconsciousness? Does a conscious artificial agent? Is the ground felt experience, agency, interests, relationship, membership, a constructed inclusion rule, or an approval-independent status? Does equal respect require identical treatment, equal rights, equal consideration, equal resources, or an explanation for relevant differences? Can respect for one person’s agency conflict with protection of another person’s security? What settles the conflict?

“Promises should be kept” also needs an architecture. Does a coerced promise count? A deceived promise? A forgotten promise? A promise made to someone now dead? A promise whose fulfillment would cause a disaster? Does the duty arise from the speaker’s act of commitment, the listener’s reliance, a beneficial social practice, an objective property of promising, or a rule no one could reasonably reject? When may a person release themselves?

A verdict becomes part of a system only when the system can answer questions like these without changing its standard to protect the preferred answer.

This gives us the first construction principle:

The strength of a moral conclusion cannot exceed the support supplied by its stated grounds.

If the ground is “our community condemns cruelty,” the conclusion may accurately report a community standard. If the ground is “cruelty predictably causes terror and pain,” the conclusion may identify consequences that matter to people who care about terror and pain. If the ground is “we agreed to prohibit cruelty,” the conclusion may invoke the terms of an agreement. None of these conclusions is trivial. None is yet the claim that cruelty possesses an objective moral status independent of every stance.

The distinction is easy to lose because the same practical decision may be supported in several ways. A society can prohibit assault because citizens fear injury, because mutual

restraint benefits almost everyone, because victims demand protection, because a religious tradition condemns violence, because officials want public order, and because lawmakers believe assault is objectively wrong. The prohibition may remain in place even if those grounds conflict.

But if we are building an objective moral system, we must identify which ground makes the moral claim true and authoritative. A coalition can make a law. A shared fear can sustain it. A threat can enforce it. What makes it objectively correct?

The five hopes

The promise of objective morality usually contains at least five hopes.

Independence

An objective moral truth should not become false merely because someone disapproves of it. Stronger versions say it should survive universal disapproval. If every society embraced torture for entertainment, torture would still be wrong.

This is the hope that protects criticism from a popularity contest.

Scope

An objective system should say whom its truths reach. Some moral claims appear to address everyone: no person may torture an innocent child for amusement. Others are role-specific: a physician must protect patient confidentiality. A system needs to explain both the universal reach and the local condition.

This is the hope that prevents authority from expanding and shrinking opportunistically.

Authority

An objective moral consideration should count in favor of action even for the person who rejects it. The torturer's amusement does not erase the child's claim. The corrupt official's indifference does not make bribery permissible. A moral reason seems capable of applying without first borrowing a desire from the agent.

This is the hope that moral demands are more than advice addressed to people who already care.

Knowledge

A moral order that no one can distinguish from projection cannot perform the work expected of it. The system needs some account of access: perception, intuition, reason, revelation, natural facts, ideal reflection, convergence, or another method.

This is the hope that correction is possible.

Guidance

The system should help with more than cases described to trigger immediate agreement. It should define terms, distinguish verdicts, handle exceptions, prioritize conflicts, and respond to new facts. When values collide, it must say whether it uses a common scale, priority to the worse off, a protected boundary, an emergency threshold, lexical rank, plural judgment, fair procedure, or an openly chosen settlement. If it says every value matters but never explains how action is settled, it preserves sensitivity by surrendering direction.

This is the hope that moral truth can organize action.

These hopes are not automatically compatible. A system may be objective but inaccessible. It may be universal but too indeterminate to guide. It may be precise because it follows local rules, but local dependence may cost objectivity. It may be flexible enough to handle context, but the flexibility may conceal inconsistent standards. It may bind everyone by definition while leaving the source of that authority unexplained.

The aim of this book is not to demand perfection from every moral practice. It is to see whether a system can keep the promises for which objectivity is wanted.

The unheard song

Port Meridian includes a district populated by the descendants of two communities that fought a bitter war. The current majority forbids the funeral songs of the defeated community. Officials say the songs glorify the old conflict. Most residents experience them as threatening. Public performance brings condemnation, loss of employment, and a substantial fine.

Lio's mother belongs to the defeated community. When she dies, he enters a sound-proof memorial room and quietly sings one of the prohibited songs. No microphone is present. No one hears him. No record is made. The act changes no later event.

Suppose, for this exercise, that the majority's disapproval is the only ground offered for the prohibition. We are not yet told whether the song contains a threat, whether singing comforts Lio, whether honoring his mother matters to him, whether the rule violates a constitutional promise, or whether the community accepted a fair procedure. The case isolates one question:

What does the majority's disapproval itself establish when no one discovers the act?

It clearly establishes that the act violates a standing social rule. A rule can exist even when a particular violation is secret.

It may establish something about classification: by the community's standards, the act is forbidden.

It does not create an enforcement-based reason in this instance. No punishment or exclusion will occur.

It may have shaped Lio's psychology. He could feel shame or fear even in private. But the case does not specify that response, and an internalized reaction would still be a psychological effect of the social code.

Does the disapproval make the song objectively wrong?

If your answer is yes, write the bridge explicitly:

Whenever a community strongly condemns an action, that condemnation gives the action an objective moral status even when the agent rejects the standard and no one is affected by the violation.

You may not want that bridge. It would make opposing community codes objectively correct in their respective locations solely because each population approves its own. It would make reformers objectively wrong while the old rule remained popular. It would extend a community's authority over outsiders who never accepted its terms.

Perhaps only some forms of community approval create objective authority. Then we need the selector. Approval produced through fair procedures? Approval of rules that reduce harm? Approval consistent with equal standing? Approval a rational person would endorse?

Each option may improve the rule. Each introduces a ground other than actual approval. Fairness, harm, equality, or rational endorsement now carries the argument.

That is not a failure. It is an entry in the debt ledger.

Exercise: Separate the claims

Purpose. The word *wrong* can compress several different claims into one verdict. This exercise asks you to pull those claims apart. You are not yet deciding whether Lio should have sung, whether the prohibition is defensible, or what the final moral verdict should be. You are auditing what each kind of statement establishes.

Guardrails

- Use only the facts stipulated in **The unheard song**. Do not add a threat in the lyrics, an effect on a listener, a promise by the city, or a motive for Lio.
- Give one short, concrete answer to each prompt. If the case does not tell you, write **not specified**. An acknowledged gap is better than an invented fact.
- Do not replace *wrong* with a disguised equivalent such as *immoral*, *unethical*, *bad*, *unjustified*, *objectively forbidden*, or *something he ought not do*. You may write *forbidden by Port Meridian's code* because that names a social classification rather than assuming an objective verdict.
- Keep the levels separate. What officials **can punish** is not yet what they **have authority to punish**. What Lio **fears or values** is not yet a reason that applies to every person. What most residents **approve** is not yet what an informed rejecter is obligated to accept.
- In item 7, state a bridge principle or an additional ground. Do not write merely that the rule "has authority" or "is objectively correct." Those phrases restate the claim instead of supporting it.

Complete the audit

1. **Classification—what rule exists?** Under Port Meridian's current code, Lio violated _____.

2. **Prediction—what might happen psychologically?** If the act were discovered, Lio might experience _____. If the case gives you no basis for a prediction, write **not specified**.
3. **Institutional power—what can officials do?** The community's institutions have the power to _____.
4. **Personal commitment—what might matter to Lio?** Lio may care about _____, but the case establishes _____ about whether he actually does.
5. **Effects and interests—who or what is implicated?** The prohibition affects _____. The secret performance itself is stipulated to change _____.
6. **Social support—what does approval accomplish?** The majority's approval gives the prohibition _____.
7. **Missing bridge—what has not yet been established?** The prohibition would bind Lio even in perfect secrecy only if _____. State the ground or principle that would make the rule authoritative for an informed person who rejects the community's standard.

Check your work

- Could a reader accept your answers to items 1–6 and still ask whether the prohibition binds Lio? If not, you may have smuggled the conclusion into an earlier answer.
- Does item 7 state a premise that could be examined and challenged? If it says only that the rule binds Lio because it is binding, the circle remains.
- Did you turn *may* into *does* anywhere the case is silent? Replace the invented fact with **not specified**.

A successful page contains several concrete claims and one exposed bridge. It does not require a verdict on the song. Its purpose is to prevent classification, prediction, power, commitment, effect, social support, and objective authority from borrowing one another's credentials.

The reformer's time

The unheard song tested what approval contributes when one person rejects a rule. The next case tests a different hope: that an objective standard can vindicate a reformer before the surrounding society changes its mind.

Port Meridian inherits a technical academy that once prohibited girls from studying engineering. Asha secretly taught them. At the time, the exclusion was lawful, popular, and religiously defended. She lost her employment and was denounced by nearly everyone in her district. Twenty years later, the academy admits its error and names its largest laboratory after her.

What can we say about Asha at the time she acted?

We can say she violated the rule then in force. We can say later residents celebrate the same act. We can identify a social change.

We can praise her from our present commitments. We can say her action served education, agency, equality, economic independence, or reduced domination. We can identify

the interests of the excluded students. If the academy had publicly promised equal citizenship, we can say its rule conflicted with its own commitment.

All of these are genuine forms of criticism. None requires us to pretend that the past majority secretly agreed.

But the promise of objective morality says more. It says Asha was correct even then, whether or not anyone recognized it.

That claim is one reason many people will not accept simple majority relativism. Yet the reformer's case does not itself establish the objective standard. It creates a demand for one. We still have to identify what made Asha correct: an objective value of equality, a right to education, the wrongness of domination, the welfare of the students, a rational requirement of consistency, a divine command, an ideal agreement, or something else.

The case tells us what we want the system to do. It does not build the mechanism.

This is a pattern we will encounter repeatedly. A disturbing consequence can reveal a cost of a theory. It does not automatically prove the rival foundation. If majority relativism cannot support the criticism we want, we may reject majority relativism. We must still pay for whatever replaces it.

When seriousness masquerades as objectivity

Some people worry that separating these claims weakens condemnation. If we stop saying an atrocity is objectively wrong and instead describe terror, injury, domination, grief, coercion, broken trust, and the values that move us to resist, have we said less?

We have said more about the world and less about an alleged status.

Whether that is a loss depends on what the status contributes. Perhaps objective wrongness explains why the terror matters and why anyone has reason to stop it. If so, the system must provide that explanation. Perhaps the status merely intensifies the speaker's commitment. In that case, direct language may be at least as serious:

This practice terrifies children, separates families, entrenches control by one group, and destroys forms of trust on which their lives depend. We will organize to stop it, protect those targeted, and change the institutions that enable it.

Nothing in that statement is casual. It names effects, values, commitments, and intended action. It does not ask a moral property to carry the emotional weight.

We are not entitled to assume yet that this translation is sufficient. The objectivist may reasonably answer: "You have explained why *you* oppose the practice. You have not explained why its supporters have a reason to stop."

Correct. That is the central problem of authority, and we will not evade it.

For now, the point is narrower. Emotional intensity, political resolve, and objective truth are different. A claim does not become objective because the stakes are high. Nor do the stakes become low if objectivity is not established.

Construction exercise: The promise you require

Purpose. Decide which aspirations make objective morality worth seeking for you; this sets the standard of success but does not show that the aspirations can be jointly satisfied.

Return to Candidate System 0.1. Create a page titled **Promises**.

For each property below, mark it **essential**, **desirable**, or **unnecessary**.

1. Its basic truths remain true if every actual person rejects them.
2. Its verdicts apply to people who lack any desire to comply.
3. It supplies determinate answers to most serious conflicts.
4. Its truths can be known with reasonable reliability.
5. Its central rules can be applied by ordinary people, not only ideal reasoners.
6. It distinguishes actions, motives, outcomes, characters, and institutions.
7. It explains exceptions without inventing them case by case.
8. It applies across cultures, eras, and forms of life.
9. It can criticize its own interpreters and institutions.
10. It supplies reasons for obedience, not merely motives produced by reward or punishment.

Now choose the three properties for which you most want objective morality. Write why.

Do not choose what seems easiest to defend. Choose the promise that matters to you. The construction should be judged by the job you actually want it to perform.

Inspection note

At the end of this chapter, the Candidate System has not been refuted. It has barely been specified.

But several distinctions are now on the table:

- a verdict is not a justification;
- a social rule is not the same as an objective status;
- enforcement supplies motives without necessarily supplying authority;
- reform creates a need for criticism beyond present approval without proving a particular foundation;
- seriousness does not establish objectivity;
- the strength of a conclusion cannot exceed the support supplied by its grounds.

Add the first line to the debt ledger:

***Debt 1 — Added status:** What does an objective moral verdict add to the concrete facts, values, commitments, rules, and responses already identified?*

In the next chapter, we decide exactly what the word *objective* is allowed to promise.

Chapter 2 — What “Objective” Must Mean

The word *objective* is useful partly because it is generous.

It can praise evidence, fairness, distance from self-interest, independence from fashion, agreement among observers, freedom from bias, stability under reflection, or reality itself. A journalist tries to be objective by representing evidence fairly. A measurement is objective when different competent observers can reproduce it. A procedure is objective when it does not let the judge alter the standard for friends. A fact is objective when its truth does not depend on what anyone wants it to be.

These are related achievements. They are not the same achievement.

If we allow them to blur, an objective moral system can appear to succeed by moving between them. When stance-independent truth is challenged, the system retreats to impartial procedure. When the procedure’s selectors are challenged, it points to broad agreement. When agreement is challenged, it points to what people would endorse under better information. When idealization is challenged, it calls the result objective because it is not merely one person’s preference.

Each move may identify something valuable. The question is whether it preserves the project written on our specification sheet.

Where this chapter is headed. We will separate six forms of independence that the word *objective* often bundles together. Later candidates should receive credit only for the kind they actually establish, not for every favorable association carried by the word.

Six kinds of independence

Begin with six claims that might be expressed by saying a moral judgment is objective.

1. It is not impulsive

The judgment survives a passing mood. I remain opposed to humiliating prisoners after my anger subsides, after I sleep, and after I consider the matter calmly.

This establishes stability within a person. It does not establish independence from every person’s stance.

2. It is not selfish

The judgment does not simply favor the speaker. I support a rule even when it costs me, or I apply the same evidentiary standard to my ally and my opponent.

This establishes a form of impartiality or integrity. It does not establish that the rule is true independent of anyone’s evaluative commitments.

3. It is not idiosyncratic

Many people can recognize and apply the judgment. Agreement may cross families, parties, religions, and cultures.

This establishes intersubjective reach. It does not establish why the agreement is truth-tracking rather than produced by shared biology, common interests, education, power, or coincidence.

4. It is not altered by ordinary disagreement

The judgment is made under a standard whose application does not change when an opponent objects. A court can apply a public rule despite a defendant’s rejection. A game official can make a correct call even when a player protests.

This establishes objectivity relative to a rule or practice. It does not establish the objective authority of the rule or practice itself.

5. It is what suitably improved agents would endorse

The judgment is not tied to what people happen to want now. It is tied to what they would want if informed, coherent, vivid in their understanding of others, free from manipulation, and placed under fair conditions.

This establishes independence from actual attitudes. It may not establish independence from attitudes as such. It also creates questions about which improvements are permitted and why.

6. It is true independently of every stance

The judgment’s truth is not constituted by any person’s or group’s actual or hypothetical approval, preference, commitment, convention, command, or procedure. It would remain true if no one believed it and if every agent opposed it.

This is stance independence. It is the kind of objectivity most directly suited to the promise that a whole society can be morally mistaken.

The Candidate System may use several of these. But it may not collect the authority associated with the sixth while defending only the first five.

A fact, a measurement, and a verdict

Suppose a reservoir contains 11.8 million cubic meters of water.

The fact does not depend on what residents want. The measurement does depend on human practices: units, instruments, calibration, sampling, and conventions for reporting uncertainty. Yet those practices are designed to track a quantity that constrains them. A faulty instrument does not make the reservoir fuller. A unanimous vote does not add water.

This example makes objective truth seem straightforward. There is a world that pushes back.

Now consider a policy that distributes sixty percent of the water to households, twenty percent to agriculture, fifteen percent to industry, and five percent to ecological maintenance.

The physical quantity constrains the policy. So do facts about crop failure, health, leakage, employment, migration, and the wetland. For anyone engaged in a truth-responsive inquiry, better evidence can expose false predictions. A later chapter will ask whether this truth-indicating relation also supplies a stance-independent epistemic reason; moral skepticism alone does not answer that question. The distribution still depends on standing and priority. Whose need counts directly? How much current hardship may be imposed to prevent future loss? Does the wetland matter only through its effects on people and sentient animals, or does it have standing of its own? Should equal need receive equal water, or should certain uses take priority because they sustain the system on which everyone depends?

If the assembly reaches an impartial agreement, it has accomplished something substantial. It has not necessarily discovered a distribution written into the water.

The difference is not that physical facts are real and human stakes are fake. Thirst is real. Fear is real. A destroyed livelihood is real to the person who loses it. Ecological collapse has observable effects. The difference concerns what makes the allocation *correct*. The reservoir's quantity constrains every stance. The claim that one distribution has objective moral status requires an additional account.

This is why moral objectivity cannot be established merely by pointing to objective nonmoral facts. A moral system must explain how facts enter and what selects among their possible practical implications.

The unanimous chamber

Imagine a society whose members all approve the same code. Their agreement is not superficial. They have discussed the code for decades. They understand every consequence. They can explain the rules. No one is coerced.

The code places a hereditary group in permanent service to the rest. Members of the servant group approve because, before birth, all citizens underwent an intervention that makes hierarchy feel safe, meaningful, and beautiful. The intervention does not impair factual reasoning. It gives everyone the same terminal values.

Call this the Unanimous Chamber.

The society's code is stable, informed, coherent, and universally endorsed. Is it objectively correct?

Most readers will say no. Engineered agreement does not make hierarchy morally true.

That response reveals something important about the desired objectivity. We want the system to be independent not only of current disagreement but also of current agreement. Consensus can be mistaken. A perfectly shared value can still fail to constitute moral truth.

Now remove the intervention. Imagine the same population arrived at the same values through ordinary upbringing, shared history, mutual dependence, and reflection. Is the code now objective?

The suspicious causal history has disappeared. The relation between agreement and truth has not changed. If approval constitutes correctness, the code is correct because it is approved. If truth is stance-independent, approval alone does not settle it.

You may object that a population reasoning under genuinely fair conditions would never endorse hereditary service. That may be right. But notice the repair. Actual unanimity is no longer sufficient. The system now appeals to *suitable* unanimity.

We need to know what makes the conditions suitable.

The idealized chamber

Let us improve the agents.

They receive every relevant fact. They can imagine every role vividly. They reason consistently. They are free from coercion. They deliberate as equals. They do not know which social position they will occupy. They understand the emotional lives of everyone affected. They are given unlimited time.

This procedure has eliminated many causes of bad judgment. It protects against ignorance, inconsistency, bargaining advantage, narrow self-interest, and failure of imagination.

Suppose all participants reject hereditary service.

The result has several attractive properties:

- it is not a product of misinformation;
- it is not tailored to known personal advantage;
- it represents every participant symmetrically;
- it is stable under disciplined reflection;
- it supplies a public basis for a rule.

Calling the result objective may be reasonable in an institutional or procedural sense. But has it become stance-independent?

The procedure still begins with valuing agents. Information alone does not supply terminal ends. Consistency alone does not choose among coherent aims. Vivid understanding may activate compassion in agents who are capable of compassion, but it does not logically force every possible agent to care. Equal representation matters only within a standpoint that treats representation as important. Ignorance of one's future position restrains advantage because participants want safeguards across the possible positions they may occupy.

The procedure can transform, discipline, and coordinate values. It may create reasons for people who accepted its terms. It may identify principles that actual people can defend to one another. Those are major achievements.

But if the original project requires a truth whose correctness is not constituted by any actual or hypothetical stance, idealized endorsement is a different foundation.

This is not a verbal trick. It marks the difference between two important projects:

1. finding standards that no one could reject under specified conditions; and
2. finding stance-independent moral facts.

The first may be possible when the second is not. We should not treat that as failure unless the second is what we promised to build.

Universal is not objective

A rule can apply universally without being objectively true.

Port Meridian may adopt a traffic rule requiring everyone to use the right side of the road. The rule reaches residents, visitors, officials, dissenters, and people who would prefer the left. Its universal application is essential to coordination. A visitor cannot safely exempt themselves because they never approved it.

The rule is authoritative within a legal jurisdiction. It may give nearly everyone strong practical reasons to comply. It can be enforced impartially. It can be the correct answer to “Which side must I use here?”

None of this shows that right-side driving is objectively morally correct independent of human practices. Another city can drive on the left without making a moral error.

The example separates scope from grounding. Universality answers *who is governed*. Objectivity answers *what makes the standard correct*. A social practice can have broad scope because coordination requires it. A legal institution can claim jurisdiction over nonconsenting visitors. Whether the jurisdiction is legitimate raises further questions, but no stance-independent traffic fact is needed to explain why mixed conventions are dangerous.

Moral rules are usually thought to be different. “Do not torture children for amusement” does not appear conventional in the way traffic direction is. We do not think a coordinated society may choose the opposite rule as an equally acceptable solution.

Why not?

Because children experience terror and pain? Because amusement cannot justify those effects? Because persons have rights? Because a rational agent could not will the rule? Because no one would accept the risk under fair conditions? Because compassion responds appropriately to suffering? Because God forbids cruelty?

Any of these may become part of the Candidate System. The point for now is that universal scope does not choose among them or establish their objectivity.

Impartial is not viewless

The assembly appoints a water-allocation judge. She owns no farm, factory, or water-front property. Her family receives the same household allotment as everyone else. She discloses every conflict of interest. She applies published standards. She gives reasons for each decision and permits appeal.

She is impartial in several important senses. Yet she does not make decisions from nowhere. The published standards prioritize drinking water, then sanitation, then food production, then economic use, then recreation. That order expresses a substantive ranking. Her impartiality consists partly in applying it without favoritism.

If the ranking is challenged, pointing to her neutrality does not answer the challenge. A neutral official can administer a disputed standard faithfully.

The same is true of an ideal observer. Making an observer fully informed and unbiased removes familiar distortions. It does not by itself give the observer a terminal value. We must specify what the observer cares about, which emotional capacities remain, whose perspective they take, and what counts as bias rather than appropriate partiality.

Suppose the ideal observer is compassionate. Compassion will influence the verdict. Why is compassion included while a taste for hierarchy is corrected away? Perhaps compassion tracks suffering, and suffering has objective moral importance. Then the importance of suffering grounds the decision, not the ideality of the observer. Perhaps compassion is constitutive of the standpoint we choose to call moral. Then the system is defining a valued standpoint. Perhaps every suitably informed human would retain compassion. Then the system appeals to facts about human psychology, not yet to stance-independent truth.

Again, none of these answers is worthless. The ledger asks only that the decisive work be named.

Response-dependent facts

Not every fact involving a response is unreal or arbitrary.

Whether a joke is funny depends in some way on the responses of beings capable of amusement. Whether food is nauseating depends on creatures with digestive and sensory systems. Whether a room is terrifying depends on a subject’s experience and situation. These can be perfectly real facts. A person can be mistaken about what will frighten someone. A community can ignore real distress. Experts can investigate patterns of response.

Some philosophers understand value in a similarly response-dependent way. Perhaps what is valuable is what suitably situated agents would value. Perhaps moral truth concerns the responses warranted under ideal conditions.

This approach can preserve truth and error. An actual person may be mistaken about what they would endorse if informed, coherent, and vividly aware. A society may build practices around distorted or incomplete responses. Objective inquiry can investigate the relevant conditions.

But response dependence differs from the strongest stance independence. The truth is connected constitutively to a kind of response, even if not to anyone’s present response.

The Candidate System must decide whether that is objective enough.

There is no law requiring one use of the word. The danger lies in asking response dependence to perform two incompatible roles: to ground value in human responsiveness

when metaphysical access is challenged, and to claim independence from every response when universal authority is desired.

Objectivity as a bundle

We can now separate several dimensions that ordinary moral language compresses:

Dimension	Question	Possible success without stance-independent moral truth
Stability	Does the judgment survive mood and time?	Yes: durable commitment or character
Impartiality	Is the standard applied without favoritism?	Yes: public procedures and integrity
Intersubjectivity	Can many people share and apply it?	Yes: common values, language, and institutions
Standard-relative correctness	Is there a correct answer under a stated rule?	Yes: law, games, roles, and local codes
Idealized endorsement	Would improved agents converge?	Possibly: constructivist or response-dependent standards
Universality	Does the rule govern every person in scope?	Yes: universal legal or institutional rules
Stance independence	Is the truth constituted independently of every stance?	This is the remaining realist claim

The table is not a ladder from weak to strong in every respect. An impartial democratic procedure may be politically more valuable than an inaccessible stance-independent fact. A stable commitment may organize a life better than a disputed metaphysical claim. Standard-relative truth can be precise where abstract moral principles are vague.

But the table does show how a project can change while keeping the same word.

Pressure exercise: The convergence machine

Purpose. Separate what idealized convergence can establish—information, impartiality, legitimacy, or a procedure-relative result—from the further claim of stance-independent truth.

Port Meridian develops a deliberation system. It can provide every resident with verified facts, correct logical inconsistency, simulate the experience of every affected party, remove threats and bargaining advantages, and permit a century of subjective reflection in one hour. It does not implant a specific conclusion.

After using the system, every resident endorses Policy A.

Answer these questions separately:

1. Is Policy A supported by actual consensus?
2. Is the consensus unusually informed?
3. Is the procedure impartial?
4. Does the result have democratic legitimacy if all residents agreed in advance to use the system?
5. Does a resident who refused the system have a practical reason to comply?
6. Is Policy A true relative to the decision procedure?
7. Would Policy A remain correct if a different population went through the same procedure and converged on Policy B because its terminal values differed?
8. Is Policy A stance-independently morally true?

It is possible to answer yes to the first six and remain uncertain about the last two. That is not inconsistency. The questions concern different achievements.

Now imagine that the machine cannot produce convergence. Everyone understands everyone else and all factual disagreement disappears, but two terminal priorities remain. One group ranks individual control over one’s body above aggregate survival. Another ranks aggregate survival above bodily control in emergencies.

What should the machine do next?

- Add a rule for priority?
- Search for a deeper shared value?
- Choose the policy with the least irreversible loss?
- Use a vote?
- Use a lottery?
- Protect a right?
- Maximize an outcome?
- Refuse to decide?

Each answer introduces a value or procedure not supplied by information and consistency alone.

Add it to the Candidate System. Then add its debt.

Construction exercise: Set the threshold

Purpose. Fix the strongest kind of objectivity your project requires and the weaker achievements you would still count as success; the exercise defines the target rather than proving it exists.

Return to the Promises page from Chapter 1. Under the heading **Objectivity**, write definitions for these terms:

- objective truth;
- impartial judgment;
- universal rule;
- public standard;
- idealized agreement;

- stance-independent moral fact.

Keep each definition to two sentences. None may use another term in the list without explaining the relation.

Now complete the following:

My Candidate System ceases to count as objective morality if its central verdicts depend constitutively on _____.

Possible answers include actual individual preference, current majority approval, any actual human stance, any possible stance, any response including idealized response, a divine stance, or something else.

Then complete:

I would still consider the project successful if its verdicts were not stance-independent but were _____.

You may answer “I would not.” You may also discover that what you most want is impartiality, defensibility, stability, humane content, or public legitimacy rather than stance independence. Record that discovery. Do not rename it objective truth merely to preserve the original label.

A fair standard for the search

We now have enough precision to avoid two unfair conclusions.

The first would be:

Humanly constructed standards are not stance-independent, so they are arbitrary and useless.

That does not follow. Construction may be informed, reciprocal, transparent, stable, protective, and open to correction. A constructed bridge is not an imaginary bridge. A constructed legal procedure can constrain officials and protect people. Human origins do not erase human purposes.

The second unfair conclusion would be:

Impartial, informed, and widely shared standards are objective in some sense, so objective moral truth has been established.

That does not follow either. A procedure can earn one kind of objectivity without establishing another.

The inquiry will therefore use a simple discipline:

Name the kind of independence achieved, and claim no more authority than that achievement supplies.

This discipline protects both sides. It prevents the skeptic from dismissing every human standard as whim. It prevents the realist from using the successes of public reason as evidence for an additional stance-independent status.

Inspection note

Add a second line to the debt ledger:

***Debt 2 — Independence:** Which parts of the Candidate System remain true under universal rejection, and what makes their truth independent of the attitudes used to identify or apply them?*

Beside it, add a warning:

Consensus is not yet objectivity. Impartiality is not a view from nowhere. Universality is not a ground. Idealization needs selectors.

The Candidate System now has a clearer target. In the next chapter, we ask what kind of claim its central moral sentences are making in the first place.

Chapter 3 — What Are You Doing When You Say “Wrong”?

Wrong is a small word with impressive luggage.

Four letters can report a supposed fact, condemn a person, warn a friend, invoke a rule, demand a law, end a conversation, begin a protest, express disgust, coordinate a group, or spoil an otherwise pleasant family dinner. Few words work so hard while explaining so little.

That is not an insult to moral language. Ordinary language is efficient because context does much of the work. If someone shouts “Fire!” in a crowded theater, we do not request a dissertation on combustion. We look for smoke and an exit. If a friend says, “That was cruel,” tone, history, and circumstance may tell us exactly what needs to be addressed.

The trouble begins when a word that works well in conversation is recruited as the load-bearing term of an objective system. Engineers are allowed to say “thingamajig” over lunch. They may not put *thingamajig* in the bridge specifications.

Your Candidate System currently uses moral vocabulary. Its first sentence probably contains *right*, *wrong*, *ought*, *good*, *duty*, or a close relative wearing formal clothes. Before the assembly places those words in Port Meridian’s constitution, it needs to know what they claim.

The question is not yet which acts are wrong. It is what a speaker is doing by calling an act wrong.

Where this chapter is headed. We will identify the different jobs a moral sentence can perform before asking whether it succeeds. The distinction matters because two speakers can share a verdict while making different claims—or appear to disagree while doing different things with the same word.

The missing medicine

A refrigerated medicine arrives at Port Meridian’s central clinic. One dose remains. A clerk named Sora secretly diverts it to her father, who is seriously ill but not in immediate danger. The dose had been allocated by lottery to a stranger whose condition is equally serious. Sora knows the allocation. She alters the record and says nothing.

When the diversion is discovered, one sentence appears in meetings, messages, and news reports:

What Sora did was wrong.

Six residents agree on every ordinary fact. They know what Sora did, why she did it, how the lottery worked, and what happened to both patients. They agree that the clinic

should prevent a repeat. They even agree that Sora should return to work only after trust is rebuilt.

When asked what the widely used sentence commits a speaker to, they offer importantly different accounts. Some use it literally. One thinks its literal use contains an error. Another retains it only as a conscious fiction.

The reporter

The first speaker means:

Sora’s act had a moral property—wrongness—that it would have possessed even if everyone approved of the diversion.

This is a factual interpretation. The sentence is intended to describe reality. On this view, “Sora diverted the medicine” and “Sora acted wrongly” report two facts about the event, though facts of different kinds. The first is institutional and causal. The second is moral.

The speaker may believe wrongness is natural, nonnatural, rationally necessary, divinely grounded, or constituted in some other objective way. Those are different accounts of the property. They share the claim that the sentence is true independently of actual approval.

This family of views is usually called **moral realism**. The name covers substantial disagreement, so it is better to understand the move before memorizing the label: moral sentences report facts, and at least some of those facts are objective.

The reporter owes the debts already entered in our ledger. What is this property? What makes the sentence true? How is the property known? Why does it have authority? The reporter has made a clear claim. Clarity is not yet payment, but it tells us where to send the invoice.

The mistaken reporter

The second speaker thinks ordinary moral language *tries* to work as the reporter says. When people sincerely say that Sora acted wrongly, they attempt to describe an objective moral status.

The difference is that this speaker thinks no such statuses exist.

“Sora acted wrongly” purports to report an objective fact, but there is no objective wrongness for it to report.

This is **moral error theory**. It is more radical than uncertainty about one case. It says ordinary moral claims contain a systematic mistake. People speak as if actions possess objective moral properties or as if categorical moral requirements apply, but the world supplies no such properties or requirements.

The error theorist need not approve of Sora. They may care intensely about reliable lotteries, equal access, truthfulness, and protecting patients. They may demand dismissal, repair, or institutional redesign. Their claim about the sentence’s meaning is narrower:

if “wrong” literally reports objective wrongness, the sentence is false because its required property is absent.

This can sound like a person saying, “The dragon is dangerous, but dragons do not exist.” That would be a peculiar warning. Yet the error theorist’s point is not that the diversion did not happen or that nothing about it matters. The point is that ordinary descriptions, reactions, rules, and values do all the practical work while the additional objective status fails to refer.

The mistaken-reporter view therefore creates a later question: should we stop using the sentence, revise its meaning, or keep it for some practical purpose? We will not settle that here. First we are learning what the sentence would commit us to.

The opponent

The third speaker is not trying to detect a moral property. The sentence voices opposition and helps form a practical stance:

I oppose diversions like Sora’s. Let us reject them, protect the lottery, and expect one another not to do this.

This is not merely a report of private dislike, like “I dislike fennel.” Moral opposition can be reasoned, consistent, informed, and addressed to others. The speaker may cite the stranger’s loss, the deception, the need for trust, and the danger of letting staff privately rank their families above patients. The sentence gathers those concerns into a practical commitment.

Views in this family are often called **expressivist** or **noncognitivist**, although contemporary versions are more sophisticated than the old caricature “Wrong means boo.” A sports fan shouting “Boo!” is not normally proposing a public standard, committing to a pattern of judgment, or inviting structured reasoning. Moral language often does all three.

One useful analogy is a vote. A vote expresses a choice, but it is not a random noise. It can be informed by facts, governed by consistency, connected to plans, and coordinated with other votes. Casting the vote does not describe a mysterious property of the candidate. It helps decide what the voter and the group will do.

The opponent’s challenge is to explain how moral reasoning works when the conclusion expresses a stance rather than reports a fact. How can attitudes be inconsistent? What makes evidence relevant? How can one condemnation be better informed or less arbitrary than another? Expressivist theories have detailed answers. For our construction, the essential point is that a practical commitment and an objective report are different jobs.

The rule user

The fourth speaker means:

Sora’s act is forbidden by the valid standards of this clinic or community.

The statement may be straightforwardly true. The clinic has an allocation rule. Staff accepted it as part of their role. The record shows a violation. No objective moral property is needed to establish the institutional classification.

This is a **standard-relative** interpretation. The truth of the sentence depends on a specified framework: the law, a profession, a culture, an agreement, a game, a relationship, or a personal code.

Framework-relative truth is not the same as “whatever anyone feels.” A chess move can be illegal even if both players like it. A court can misapply its own law. A clinic can violate its written allocation policy. Once a standard is identified, there may be objective facts about what it permits.

The harder questions concern the standard itself. Why this clinic’s rule? Does it bind a staff member who never freely accepted it? Can the rule be criticized? What if the clinic openly favors wealthy patients? A relative classification answers “What follows under this standard?” It does not automatically answer “Why should this be the standard?”

The rule user may be completely clear about that limit. Trouble begins only when a truth relative to a framework is presented as a truth whose authority is independent of every framework.

The builder

The fifth speaker means something like:

Sora’s act is prohibited by the principles that suitably informed and fairly situated people would construct for governing access to medicine.

Unlike the rule user, the builder does not simply accept whichever standard happens to exist. Actual institutions can be corrupt. The relevant standard is generated through a procedure: perhaps equal representation, full information, freedom from coercion, impartial choice, reciprocal justification, or what no one could reasonably reject.

This family is called **constructivism**. Some versions are realist in some sense; others are not. The classification dispute matters less here than the underlying proposal. The builder does not find a complete rule floating beyond human deliberation. The rule is the outcome of the right kind of construction.

That gives the verdict critical power. Port Meridian’s actual clinic may favor officials, while a defensible procedure would reject that privilege. The builder can criticize existing practice without appealing to whatever the current majority approves.

But the construction has selectors. Who participates? What information do they receive? Which inequalities are removed? What counts as reasonable rejection? Are any outcomes excluded before deliberation begins? The procedure may organize and discipline values brilliantly. It does not choose its own design.

The builder therefore owes an account of why the selected procedure has authority and whether its result is objective, response-dependent, or simply the most defensible human standard available.

The conscious storyteller

The sixth speaker does not believe in objective wrongness but continues to speak *as if* it exists because the practice is useful:

Let us treat Sora's diversion as wrong within a deliberately retained moral story, because this way of speaking helps protect trust and coordinate action.

This is a form of **moral fictionalism**. We already use useful fictions. A corporation is not literally a biological person, but treating it as a legal person helps organize ownership and liability. A nation is not a single mind, but saying “the nation decided” can summarize an institutional process. A line of longitude is not painted on the Earth, yet navigators are advised not to throw away the maps.

The fictionalist proposes that moral language might remain useful after belief in objective moral facts is withdrawn. Speakers could knowingly participate in the practice without treating its surface grammar as a literal inventory of the universe.

This approach owes a disclosure policy. If some people believe the story literally while others use it consciously, are they engaged in the same practice? When is pretense manipulative? Could direct language about harm, trust, rules, and commitment perform the useful work with less confusion? Fictionalism preserves familiar language, but familiarity is not free if it conceals what speakers actually believe.

One verdict, six interpretations

The six residents may support the same response. Most may say, “Sora acted wrongly”; the error theorist may instead quote the sentence while refusing literal assent. They disagree about what would make it true, whether it is literally true, and what job its use performs.

This gives us an important form of disagreement:

People can agree about what to do while disagreeing about the meaning and status of their moral language.

That fact complicates the usual picture. We often imagine moral disagreement as two armies facing one another across a battlefield of conclusions. Sometimes the disagreement is inside one camp. Everyone wants the medicine rule protected, but one person thinks the prohibition tracks an objective fact, another treats it as a fair construction, another invokes an actual institutional rule, and another expresses a shared commitment.

Their practical alliance can last for years. It may even write a constitution. The deeper disagreement becomes visible when someone asks whether the rule binds a rejector, whether another society is objectively mistaken, or whether moral language should continue after realism is abandoned.

The reverse can also occur.

One meaning, opposite verdicts

Two moral realists can agree completely about what *wrong* means and disagree about whether Sora’s act has that property. One may think family loyalty justifies the diversion. The other may think the stranger’s equal claim controls. Their dispute is first-order: which action has the agreed moral status?

Two standard-relative speakers can use the same framework and disagree about its application. Did Sora’s father qualify for an emergency exception? Did accepting employment constitute agreement to the lottery? Which rule has priority?

Two expressivists can agree that moral claims express practical stances while expressing opposite stances. One opposes all private diversion. Another endorses family partiality.

Two constructivists can agree on the nature of moral construction and disagree about the correct procedure or its output.

So another important fact is:

People can agree about the meaning of moral language while disagreeing about the verdict.

Meaning and conclusion are separate axes. Confusing them is like assuming that two mathematicians who disagree about an answer must be using the word *number* differently. Sometimes they are. Usually they are disputing the calculation.

The smuggling problem

Using moral language in more than one way is not automatically a defect. Context changes what words do. “You must be exhausted” can express a confident prediction, sympathetic recognition, or a tactful invitation to go home. Competent speakers navigate such flexibility constantly.

The Candidate System faces a stricter standard because its central terms carry authority. It cannot change meanings in order to evade pressure.

Imagine this conversation:

“Why is Sora’s act wrong?”

“Because our fair procedure prohibits it.”

“Would it remain wrong if everyone rejected that procedure?”

“Of course. Wrongness is objective.”

“What makes the objective fact true?”

“The fact that no one could reasonably accept a system of private diversion.”

“What makes that rejection unreasonable?”

“Because private diversion is wrong.”

The speaker has moved from a constructed standard to stance-independent truth and back again. The movement is quick enough to look like progress. It is actually a philosophical airport walkway: you feel as though you are speeding forward while the same terminal remains suspiciously nearby.

The problem is not that the speaker must choose one theory forever. A plural account may connect objective facts, procedures, attitudes, and conventions. The problem is that each connection must be stated. A procedure cannot borrow the authority of objective truth and then serve as the only evidence that the truth exists.

We will call this **status smuggling**:

A moral term is defined modestly when its existence is questioned and interpreted strongly when its authority is used.

Examples include:

- “Wrong just means that our society condemns it”—until the speaker criticizes another society as mistaken.
- “Rights are principles we constructed”—until a dissenter asks why the construction binds them and the rights become self-evident facts.
- “Good means conducive to flourishing”—until someone asks why flourishing matters and *good* becomes an unexplained source of importance.
- “I am only expressing my opposition”—until the opponent refuses and the expression becomes a reason that supposedly applies independently of every attitude.

The cure is not to ban moral words. It is to disclose their current job.

Case study: The captain's promise

Port Meridian operates a rescue ferry. Before a dangerous voyage, Captain Imani promises the engineer, Vale, that if one emergency suit remains, Vale will receive it. Vale accepts the assignment partly because of that assurance.

A collision occurs. One suit remains. Five passengers without promises will die unless one of them receives it. Whoever receives the suit survives. Everyone else in this small group dies. There is no time for a lottery, no way to share the suit, and no future precedent because the ferry will be retired.

Imani says:

It would be wrong to break my promise.

Before deciding who receives the suit, translate the sentence.

1. **Objective report:** Promise-breaking in these circumstances has a stance-independent moral property.
2. **Systematic error:** The sentence attempts an objective report, but no such property exists.
3. **Practical stance:** Imani opposes breaking this relied-upon promise and commits herself to keeping it.
4. **Standard-relative judgment:** The maritime code or promise practice classifies breaking the assurance as prohibited.
5. **Constructed verdict:** A suitably fair procedure would protect relied-upon rescue promises, perhaps with or without an emergency exception.

6. **Useful fiction:** Imani knowingly treats the promise as a binding moral fact because that practice supports reliability and resolve.

Now notice what the translation does *not* decide. It does not decide whether Vale or one of the five should receive the suit. Semantics does not settle priority. The objective reporter still needs a substantive theory. The practical stance may be defeated by another stance. The maritime code may contain an emergency clause. A fair procedure may prioritize lives. The useful fiction may become harmful in this case.

Clarifying meaning is not a magic method for resolving every dispute. It is a way to stop resolving one dispute while pretending it was another.

Exercise: The six translations

Purpose. Render one moral sentence under six semantic interpretations so their different commitments become visible; do not use the exercise to decide which interpretation is true.

Choose one sentence:

- Deception is wrong.
- Animals deserve moral consideration.
- Wealth inequality is unjust.
- Parents ought to protect their children.
- Punishment should be proportional.
- People have a right to control their bodies.

Write six translations in plain language:

1. **Fact:** What objective feature is the sentence reporting?
2. **Error:** What would make the apparent report systematically false?
3. **Stance:** What opposition, endorsement, or plan is the speaker expressing?
4. **Standard:** Relative to which actual rule or practice is the sentence true?
5. **Construction:** Which idealized procedure generates the standard?
6. **Fiction:** What useful pretense is being maintained, and for what purpose?

Do not choose among them yet. The aim is to feel the difference in their commitments.

Exercise: Agreement in disguise

Purpose. Determine whether a remembered dispute concerned action, meaning, or ground; locating a shift does not establish bad faith or settle the practical controversy.

Return to a moral argument you have witnessed. It may concern abortion, policing, speech, war, education, diet, sexuality, religion, wealth, or family obligation.

Ask:

- Did the speakers want different actions, or did they mostly want the same action while holding different views about what moral claims mean or what makes them true?

- Did anyone treat a local rule as if it were an objective fact?
- Did anyone describe an objective-sounding judgment and later defend it as mere opposition?
- Would the practical disagreement remain if every speaker translated their verdict?

Do not accuse anyone of bad faith. People inherit moral vocabulary long before they inherit metaethics textbooks, which is fortunate because very few infants would finish one. A shift in meaning may reflect uncertainty, not deception. Record the shift before judging it.

Construction exercise: Give the system a voice

Purpose. Commit the Candidate System to a stable account of what its moral terms do; choosing a meaning exposes the corresponding truth and authority debts rather than paying them.

Open the Candidate System and add a section titled **Semantic Rule**.

Complete one of these, or write a more precise alternative:

When this system says an act is “wrong,” it reports _____.

When this system says an act is “wrong,” it expresses _____ and proposes _____.

When this system says an act is “wrong,” it means that the act violates _____.

When this system says an act is “wrong,” it means that _____ would be selected under _____.

This system uses “wrong” as a conscious fiction whose purpose is _____ and whose disclosure rule is _____.

If you think the system uses more than one interpretation, state when each applies. For example:

In personal deliberation, “wrong” expresses a settled commitment. In legal discussion, it abbreviates a public standard. It never reports a stance-independent property.

Or:

“Wrong” always reports an objective moral status. Expressions of anger and institutional rules may accompany the report but do not constitute its truth.

The second answer incurs a larger debt. That is not a reason to avoid it. The size of the debt should match the ambition of the claim.

The third debt

Add this entry to the ledger:

Debt 3 — Meaning: *What job do the Candidate System’s moral terms perform, and do they keep that job when truth, authority, disagreement, and practical usefulness are tested separately?*

Under it, write:

A shared verdict may hide a disagreement about meaning. A shared meaning may permit opposite verdicts. Translation clarifies the dispute; it does not decide it.

We now know what kind of thing we are trying to build, what kind of objectivity it promises, and what its central words claim to do.

That is enough to write a specification sheet.

It is also enough to discover how many ways a system can fail inspection.

Chapter 4 — The Specification Sheet

Port Meridian's assembly has now received three pages from you.

The first contains a candidate principle. The second lists the kinds of objectivity the principle promises. The third explains what words such as *right* and *wrong* are supposed to mean.

The assembly's chief engineer reads them twice.

"This is philosophically ambitious," she says.

You wait for the compliment to finish.

It has finished.

"What does the completed system have to do?"

You point to the pages.

"Those tell me what you hope it is. I need to know how we determine whether it works."

This is an irritatingly good question.

If we were buying a kettle, we would know what success looks like. It must contain water, heat it, stop before setting the kitchen on fire, and pour without reenacting a minor flood. A moral system is expected to do more complicated work, yet people often test it with a single question: "Does it produce my preferred answer in this case?"

That is not an inspection. It is an audition for a conclusion already cast.

A serious construction needs a specification sheet written before the favorite theory arrives. Otherwise, every theory passes by redefining the examination around its strengths. Consequentialism points to outcomes. Deontology points to constraints. Virtue ethics points to judgment. Constructivism points to defensible procedures. Divine-command theory points to authority. Each may be answering a genuine need. The question is whether any one of them—or any combination—does the whole job for which objective morality was requested.

We therefore need success conditions.

Where this chapter is headed. We will turn the book’s opening aspirations into explicit conditions of success. The specification sheet prevents us from adding a requirement only when a rival fails it or quietly dropping one when a favored theory cannot supply it.

Nine demands on the finished system

The following nine demands do not come from nowhere. They are drawn from the promises identified in Chapter 1 and the meanings separated in Chapters 2 and 3. Not every reader will require all nine. That choice is part of the exercise.

The crucial rule is this:

Decide what counts as success before you know which requirements your preferred foundation can satisfy.

This is the philosophical equivalent of checking the target before firing the arrow. Painting a target around the arrow afterward is excellent for morale and poor for archery.

1. Truth

At least some of the system’s central verdicts must be capable of being true or false.

If “torturing a child for amusement is wrong” merely expresses opposition, it may be sincere, stable, and practically important. But a system promising objective truth must explain what makes the sentence correct rather than incorrect.

Truth requires more than grammatical shape. “The present king of Port Meridian is left-handed” looks capable of being true or false, but Port Meridian has no king. A sentence can have the form of a factual report and still fail because the thing it presupposes is absent. The error theorist thinks ordinary moral claims face a problem of this kind.

Inspection question: What in reality makes a central verdict true?

2. Stance independence

The truth of at least some central verdicts must not be constituted merely by actual approval, preference, commitment, convention, command, or procedure.

This condition preserves the promise that every actual person can be mistaken. It does not require moral facts to be ghostly, supernatural, or located in a celestial filing cabinet. A naturalist may claim that moral facts are ordinary natural facts. A rationalist may ground them in requirements of reason. The issue is dependence: would the truth remain if every stance changed?

Inspection question: Which change in attitudes could alter the verdict, and why?

3. Scope

The system must say whom its verdicts reach.

Some rules appear universal. Others depend on roles, relationships, agreements, capacities, or institutions. A physician may have reasons a stranger lacks. A parent may

properly give special attention to their child. A citizen may be governed by a traffic rule that does not apply on another continent.

Scope prevents a system from expanding when it wants authority and contracting when it encounters a counterexample. “Everyone must obey this rule” requires an account of *everyone*. “Members are bound” requires an account of membership, entry, exit, disclosure, and unfair terms.

Inspection question: Who is addressed, under what conditions, and who is not?

4. Authority

The system must explain how its considerations count in favor of action for a person who does not already endorse them.

This is stronger than causal influence. A threat can motivate. Propaganda can motivate. A drug can motivate. None becomes morally authoritative by being effective.

It is also stronger than classification. A rule may correctly classify an act as prohibited without explaining why an outsider has reason to accept the rule.

Authority is where many systems perform an elegant costume change. A human concern enters the dressing room and returns as a categorical command. The audience applauds. The ledger asks who tailored the suit.

Inspection question: What, exactly, gives the rejector a reason?

5. Accessibility

Agents must have a defensible route to the system’s central truths.

A completely objective moral order that no one can identify may exist, but it cannot reliably guide a constitution. If access depends on intuition, revelation, rational reflection, perception of natural facts, idealized agreement, or expert interpretation, the system must explain how success is distinguished from projection and error.

Accessibility does not require certainty. Science works through fallible methods. Courts use evidence without omniscience. Navigation existed before perfect maps. The demand is for a method whose errors can at least sometimes be diagnosed.

Inspection question: How would a sincere inquirer improve the odds of getting the answer right?

6. Coherence

The system’s claims must fit together without contradiction, equivocation, circular support, opportunistic exceptions, or hidden master principles.

Coherence is not truth. A novel can be internally coherent and entirely fictional. A legal code can be coherent and brutal. A paranoid conspiracy can explain every objection by adding another conspirator.

Yet incoherence matters. If the system says loyalty is decisive when it protects our friend and irrelevant when it protects an opponent, it needs a stated difference. If rights are absolute in one chapter and outweighed by small benefits in the next, the transition

needs an ordering rule. If “wrong” reports a fact when authority is desired and expresses a stance when facts are questioned, the meaning has drifted.

Inspection question: Can all of the system’s central claims be accepted at once under stable definitions?

7. Determinacy

The system must settle enough cases to perform the work promised of it.

No sane standard requires an instant answer to every possible situation, including dilemmas invented by graduate students at three in the morning. Real systems can acknowledge uncertainty, ties, tragedy, and the need for judgment.

But a system that says “all relevant values must be balanced appropriately” has not yet supplied guidance. *Appropriately* is where the entire problem has gone to hide. The same is true of “do what is reasonable,” “respect dignity,” or “choose the all-things-considered best action” when no account of the comparison follows.

Inspection question: Which conflicts can the system settle, and what does it say when it cannot?

8. Usability

Finite, biased, emotional, time-limited agents must be able to use the system.

A principle may be correct yet demand information no one can gather, calculations no one can perform, emotional neutrality no one can sustain, or foresight no institution possesses. A rule for perfectly informed beings can help clarify an ideal while failing as an operating manual.

Usability also concerns communication. Can affected people understand why a rule applies? Can officials be trained? Can errors be appealed? Can ordinary agents identify the difference between an exception and a convenient excuse?

Human limits are not an embarrassment accidentally attached to moral life. Human beings are the creatures asking for the guidance.

Inspection question: What could a conscientious but imperfect person actually do with the system?

9. Correctability

The system must provide a way to identify and repair errors in belief, interpretation, application, or institutional design.

An objective truth need not change when people learn. But human access to it does. A religious tradition may distinguish fixed revelation from revisable interpretation. A rights theory may preserve a principle while revising its application to artificial minds. A public code may amend both principle and procedure.

The difficult balance is disciplined revisability. A system that never changes can insulate error. A system that changes whenever insiders are inconvenienced has no stable

standard. Correctability asks what may be challenged, by whom, with what evidence, through which procedure, and with what repair after acknowledged failure.

Inspection question: How can the system learn without becoming whatever current power wants?

Is humaneness a tenth condition?

Many readers will want to add another demand:

The system must not be cruel.

The impulse is understandable. Why construct an objective system if it licenses terror, domination, or casual suffering?

But adding *humaneness* to the specification requires care. What does it mean? Reduced suffering? Respect for agency? Compassionate motivation? Protection of the vulnerable? A package of values we already endorse?

If we define humaneness through those values, then requiring a humane moral system builds substantive content into the test. That may be entirely appropriate for a system we intend to live under. It cannot by itself prove the values objective.

We need two inspections:

1. **Does the system meet its claimed structural standards and its claims about moral truth and authority?**
2. **Does the system serve the beings and values we care about?**

A system can pass one and fail the other. Keeping the inspections separate prevents a useful human code from being dismissed because it lacks metaphysical objectivity. It also prevents an allegedly objective system from being praised merely because it is coherent.

You may add humaneness as an essential practical condition. Just label it correctly: it states what you want the system to protect. The later value chapters will ask what supplies that content and how conflicts among humane concerns are resolved.

Three cities

The assembly sends your specification sheet to three experimental settlements. Each claims to possess the moral system Port Meridian seeks.

Clockwork

Clockwork's code is perfectly coherent.

Every person belongs to one of four hereditary classes. The rules specify duties, privileges, education, marriage, work, punishment, and appeal for each class. Definitions never shift. Similar cases receive similar treatment. Exceptions are public. Children learn the system easily. Officials apply it reliably. The code has operated for two centuries without internal contradiction.

It is also horrifying.

The lowest class exists for dangerous labor and involuntary medical experimentation. Clockwork’s philosophers explain the arrangement consistently: social stability is the master value; each class serves the whole; birth fixes function; suffering among the lowest class is regrettable but permitted when system stability requires it.

Does Clockwork possess a coherent moral system?

Structurally, perhaps. Does it possess a *true* objective moral system? Coherence alone cannot answer. Does it possess a system Port Meridian should adopt? Not if Port Meridian values agency, equal standing, and protection from severe suffering.

Clockwork teaches:

Coherence can organize cruelty. Internal order does not certify the selectors placed at the foundation.

Harborlight

Harborlight’s residents are compassionate, generous, and quick to repair harm. Their code proclaims five principles:

1. Never deceive.
2. Always prevent the greatest available harm.
3. Never use a person without consent.
4. Give family members priority.
5. Treat every person impartially.

The principles sound admirable. They also collide almost immediately.

A violent pursuer asks where a victim is hiding. A patient refuses a treatment needed to prevent an epidemic. A rescue worker can save either a sibling or five strangers. Harborlight has no priority rules. In every conflict, residents announce that “all five principles remain absolute” and then do what seems best.

Their judgments may often be humane. Their formal system does not explain them. The actual tie-breaker lives in unrecorded intuition, emotion, custom, or power.

Harborlight teaches:

Good values do not become a usable system merely by standing together for a group photograph.

High Oracle

The third settlement claims that one objective moral principle exists and correctly settles every possible case. The principle is carved on a crystal inside a mountain.

No person can reach it. No signal escapes. No one knows its content. The settlement’s founders discovered conclusive evidence that the crystal contains the truth, but the evidence and every record of it were destroyed. Current residents sincerely believe the story for incompatible reasons.

Suppose the claim is true. The principle exists. It is objective. It is complete. It is perfectly coherent.

What can the settlement do with it?

Nothing.

Officials still decide by custom, intuition, bargaining, and power. People can say “the crystal has an answer,” but the sentence does not distinguish one policy from another.

High Oracle teaches:

A truth can exist without becoming guidance. Accessibility is not existence, but a system built for agents needs both.

Tensions inside the specification

The nine demands do not always cooperate.

Determinacy and sensitivity

Precise rules increase predictability. They can also ignore context. Flexible judgment responds to detail. It can also conceal favoritism.

“Never lie” is determinate and brittle. “Lie when the situation warrants it” is sensitive and unfinished. A successful system needs either principled exceptions or a disciplined method of judgment.

Universality and role

Universal rules prevent arbitrary exclusion. Role-sensitive rules recognize promises, dependence, expertise, and relationship.

“Treat everyone identically” ignores that a doctor has responsibilities to patients that strangers lack. “Fulfill your role” becomes dangerous when roles are imposed or corrupt. The system needs an account of when difference is relevant.

Independence and accessibility

The more completely a truth is independent of human responses, the harder it may be to explain human access. The more tightly it is connected to idealized endorsement or cultivated perception, the more accessible it becomes—but its stance independence may weaken.

This is not a decisive objection. Physical facts are independent of belief and accessible through instruments. The moral system needs its analogue of instruments and an explanation of why they track the proposed domain.

Stability and correction

Stable principles protect against fashion and pressure. Revisability protects against error and exclusion.

“Our foundational rule can never be questioned” blocks correction. “Our foundational rule means whatever our best current judgment says” risks making the foundation ornamental. The system needs to distinguish truth, interpretation, application, and institution.

Authority and usability

A categorical demand may claim authority over everyone. Yet ordinary agents may need rules connected to motives, relationships, incentives, and institutions. A system could be universally authoritative in theory and practically inert. Another could be motivationally powerful because it expresses shared values while lacking authority over rejectors.

The Candidate System must not substitute practical effectiveness for objective force or dismiss effectiveness as irrelevant to guidance.

The anti-rigging declaration

Before testing any famous theory, write an **anti-rigging declaration**:

I will not count a requirement as essential only when a rival theory fails it. I will not quietly drop a requirement only when my preferred theory fails it. When I revise the specification, I will mark the revision and state what prompted it.

This does not freeze your starting assumptions. Learning should change them. The declaration requires only an audit trail.

Three labels will help:

- **Clarification:** The requirement was ambiguous, and the revision states what you meant all along.
- **Discovery:** A case showed that the project needs a condition you had overlooked.
- **Surrender:** The candidate cannot meet a desired condition, and you choose a more modest project.

Surrender is not necessarily defeat. Discovering that what you need is a defensible public code rather than a stance-independent moral truth may be progress. The label prevents us from calling the new project a completed version of the old one without explanation.

Construction exercise: Your bill of specifications

Purpose. Freeze the Candidate System's essential, negotiable, and unnecessary conditions before theory testing begins; your rankings define success for this project, not success for every possible moral practice.

Create a table with the nine conditions.

For each one, mark:

- **Essential** — without it, the project is no longer the objective moral system you intended to build;
- **Negotiable** — some loss is acceptable if the remaining system performs the central work;
- **Unnecessary** — the system need not provide this feature.

Then complete the associated sentence.

Truth

A moral verdict counts as true when _____.

Stance independence

The verdict would remain true despite changes in _____, but could depend on _____.

Scope

The system addresses _____ because _____.

Authority

A rejecting agent has a reason to comply because _____.

Accessibility

A sincere inquirer improves their judgment by _____.

Coherence

When two principles conflict, the system _____.

Determinacy

The system may leave a case unsettled when _____, but not merely when _____.

Usability

An ordinary person can apply the system with access to _____.

Correctability

A claim or application may be challenged by _____ using _____, and acknowledged error requires _____.

If you added humaneness or another substantive condition, define it without using *good, right, wrong, just, or moral* as an unexplained shortcut.

Pressure exercise: Choose the failure

Purpose. Reveal which failure matters most when no option satisfies every specification; preference among the options identifies priorities but does not establish which system is objectively true.

Imagine you must choose one of these systems for Port Meridian:

- a stance-independent truth that no one can know;
- a highly reliable and humane procedure whose standards depend on human values;
- a universal and determinate code that has no authority beyond enforcement;
- a flexible plural system that usually produces judgments you endorse but has no explicit tie-breaker;
- a coherent realist theory with clear access but several verdicts you find cruel.

Rank them for two different purposes:

1. Which comes closest to the objective moral system specified at the beginning?
2. Which would you prefer as Port Meridian's public practice?

If your rankings differ, do not force them together. The difference shows why meta-physical success and practical desirability need separate columns.

The fourth debt

Add this entry to the ledger:

***Debt 4 — Joint performance:** Can the Candidate System satisfy its essential conditions at the same time, or does each repair preserve one promise by surrendering another?*

Then copy your essential conditions onto a clean page titled **Bill of Specifications**.

The system will be judged by this bill, not by familiarity, prestige, emotional comfort, or its ability to produce one dramatic verdict.

We are almost ready to begin construction.

First, however, we need to learn how to listen when the assembly argues. A single moral dispute can contain several different conflicts, and solving the wrong one with great confidence is a Port Meridian tradition the city would prefer not to formalize.

Chapter 5 — One Argument, Eight Disagreements

The first full meeting of Port Meridian’s constitutional assembly begins at nine in the morning.

At 9:07, two delegates accuse one another of ignoring suffering.

At 9:12, someone says the argument is really about freedom.

At 9:18, someone else says it is really about corporate influence.

At 9:23, a physician objects that nobody has defined the central term.

At 9:31, a theologian announces that definitions are a distraction from objective truth.

At 9:34, the chair calls a fifteen-minute recess and considers a new career in reservoir maintenance.

The subject is assisted dying. The proposed law would allow a mentally competent adult with an incurable condition and less than six months to live to request medication that ends life. Two independent physicians must confirm the diagnosis and capacity. A waiting period and repeated request are required. No clinician may be compelled to participate.

Delegate Arin says:

Forcing a competent person through unbearable dying is morally wrong. The law should respect autonomy.

Delegate Bela replies:

Helping someone die is morally wrong. The law should protect vulnerable life.

The sentences look like mirror images. It is tempting to imagine one disagreement with two answers.

That appearance is misleading.

Arin and Bela may disagree about the meaning of a term, a factual prediction, whose interests count, or which admitted value takes priority. They may instead disagree about what attitude to express, which public rule to adopt, who should decide, whether an objective moral fact exists, or several of these at once.

Until we know which layer is active, an argument can continue indefinitely while each side answers a question the other side is not asking.

Moral disagreement is less like a boxing match than a building with noise coming through the walls. Before striking back, it helps to find the room.

Where this chapter is headed. We will separate eight kinds of disagreement that often occupy one argument at once. The map matters because evidence, clarification, negotiation, procedure, and grounding answer different problems; success at one layer cannot be billed as success at all eight.

The first discipline: do not invent the dispute

Suppose all you know is this:

- Arin says assisted dying is permissible.
- Bela says assisted dying is impermissible.

What do they disagree about?

The only safe answer is: their verdicts.

You do not yet know why. Perhaps Arin predicts careful safeguards and Bela predicts abuse. Perhaps they make the same predictions but rank autonomy differently. Perhaps Arin expresses compassion and Bela expresses reverence. Perhaps both share every value but interpret *competence* differently. Perhaps one makes an objective claim and the other proposes a public compromise.

Familiar controversies encourage us to fill missing information with familiar scripts. We have heard versions of these arguments, so we supply motives, facts, and priorities. That is efficient when predicting television panels. It is poor philosophical method.

The rule is simple:

Classify only the disagreement supported by what the speakers have actually said.

This restraint matters beyond thought experiments. People are often told what they “really” care about by opponents who have not asked them. A safety concern becomes disguised hatred. A liberty concern becomes selfishness. A religious commitment becomes fear. An empirical disagreement becomes rejection of science. Sometimes those diagnoses are correct. They require evidence.

We will now open the eight rooms one at a time.

Room 1: Words and frames

Arin calls the practice **medical aid in dying**. Bela calls it **physician-assisted suicide**.

They may be referring to the same act. The phrases place it in different conceptual neighborhoods.

Medical aid invokes care, relief, professional support, and the process of dying. *Suicide* invokes self-destruction, crisis, prevention, and the danger of normalizing death as a solution. Neither phrase is a neutral window. Each selects features and associations.

A definitional disagreement can be narrower:

- Does “terminal” mean death expected within six months, or an incurable condition that will eventually cause death?

- Does medication administered by the patient differ morally or legally from medication administered by a clinician?
- Does refusing life support count as causing death, allowing death, or both?
- Does “coercion” include pressure created by poverty and inadequate care, even without a threatening person?

Clarifying words can dissolve some disputes. It can also reveal that the wording was carrying a value conflict. If Arin and Bela agree to use the neutral description “a clinician provides lethal medication that a competent terminal patient voluntarily takes,” they may continue to disagree completely.

This gives us our first caution:

A verbal dispute may conceal a substantive dispute, but not every substantive dispute is verbal.

Changing a label can improve the conversation. It cannot guarantee agreement about the act.

Checkpoint

Ask:

- Are the speakers describing the same practice?
- Which definitions differ?
- What does each frame make vivid or obscure?
- Would the disagreement remain under a mutually accepted description?

Room 2: Predictions about the world

Arin says jurisdictions with strict safeguards can provide assisted dying without widespread abuse. Bela says eligibility will expand, subtle coercion will be missed, and disabled or poor patients will feel pressure to die.

This is an empirical disagreement. Evidence matters: records from existing programs, rates of expansion, error detection, patient reports, health-care access, family pressure, clinician behavior, and the effects of palliative care.

Calling the disagreement empirical does not mean a spreadsheet will settle it by Friday. Evidence may be incomplete. Jurisdictions differ. Rare coercion is hard to measure. Future institutions may not resemble current ones. The parties may reasonably interpret uncertain data differently.

Evidence-responsive does not mean **guaranteed to converge**.

It does mean that a claim can be improved, weakened, or defeated by facts. If Bela predicts that ninety percent of patients will be coerced and careful evidence shows a fraction of one percent, that matters. If Arin predicts perfect safeguards and hidden pressure repeatedly appears, that matters too.

Even complete agreement about effects would not settle the policy. Suppose both delegates agree that a thousand people will receive desired relief, three patients will die under

undetected pressure, and public trust will change by a known amount. They may still rank those effects differently.

Facts enter the argument. They do not choose their own priority.

Checkpoint

Ask:

- Which claims predict observable effects?
- What evidence would change each prediction?
- Are probabilities or causal relations disputed?
- If the factual dispute disappeared, what disagreement would remain?

Room 3: Who has standing?

Arin centers the competent patient's experience and agency. Bela centers patients who are vulnerable to pressure. A disability advocate asks whether disabled people are being treated as a group whose lives are more readily judged burdensome. A physician asks whether clinicians have standing to refuse participation. A family member describes dependence and grief. A taxpayer discusses the public health system.

These parties are all affected in some way. The standing question is more precise:

Whose interests, agency, experience, commitments, or claims enter the decision for their own sake?

The patient's standing seems obvious to most participants. What about a family that desperately wants more time? Their grief is real, but does it give them authority over the patient's body? What about clinicians? May a profession be required to provide a lawful service some members reject? What about future patients whose trust may be altered? What about religious communities that experience the practice as a violation of a shared order? What about citizens who are merely offended that the practice exists?

Including a party does not determine the answer. Standing is admission to the deliberative room, not a guarantee of the largest chair.

This distinction prevents two common mistakes:

1. "Their interests count, therefore their preferred policy must win."
2. "Their preferred policy should not win, therefore their interests do not count."

A family member's grief may deserve consideration without becoming a veto. A clinician's integrity may justify nonparticipation without justifying obstruction. A patient's choice may receive great weight without making safeguards irrelevant.

Checkpoint

Ask:

- Which beings, persons, groups, relationships, or institutions are included?
- Do they count directly or only through effects on someone else?

- Is anyone missing because the speakers assume rather than defend the boundary?
- Has inclusion been confused with equal weight or final victory?

Room 4: Priority among shared values

Arin and Bela may agree that autonomy, relief of suffering, protection from coercion, equal standing, medical trust, and care all matter.

They disagree about what leads when the values cannot all be fully protected.

Arin gives priority to a competent patient's control over the final stage of life. Bela gives priority to protecting life under conditions in which error is irreversible. Neither denies the other's value. The dispute concerns ordering, thresholds, and acceptable risk.

This kind of disagreement is often misdescribed as one side "not caring" about something. Arin must not care about vulnerable people. Bela must not care about suffering. Perhaps. But people can share concerns and assign them different weights.

Priority questions include:

- Does bodily agency function as a protected boundary or one value among others?
- How much risk of coercion is enough to restrict voluntary choice?
- Should the worse-off or most vulnerable receive extra priority?
- Is death ranked first, so that no amount of prolonged suffering can outweigh it, or can suffering defeat the normal priority of survival?
- Does a fair procedure settle action when no common measure exists?

Evidence can tell us how much coercion occurs. It cannot, without an evaluative bridge, tell us how much risk is acceptable.

Checkpoint

Ask:

- Which values do the speakers share?
- Which value controls in a direct conflict?
- Are priorities absolute, weighted, threshold-based, fixed in rank, contextual, or procedural?
- What would make the priority change?

Room 5: Attitudes and commitments

Arin says, "I cannot accept a law that forces someone to endure this." Bela says, "I refuse to make intentional killing part of medicine."

These statements express practical identities and commitments. Arin may feel compassion, horror at control, or loyalty to a person who suffered. Bela may feel reverence for life, fear of abandonment, or commitment to a healing profession.

An attitude is not merely a belief about an attitude. Compare:

Bela strongly opposes assisted dying.

and:

I strongly oppose assisted dying.

The first reports a psychological fact about Bela. The second takes a stand.

Attitudes can respond to evidence. If Arin's compassion rests on a mistaken picture of pain, new facts may change it. If Bela's fear rests on a false prediction, evidence may reduce it. But terminal commitments can survive factual agreement. One person may continue to value self-direction more; another may continue to recoil from intentional killing.

Calling these "mere feelings" understates their role. Commitments organize lives, institutions, loyalties, and political action. Calling them objective truths overstates what has been established.

Checkpoint

Ask:

- What does each speaker oppose, endorse, protect, or refuse?
- Which emotions make certain features salient?
- Are they reporting an attitude, expressing it, or claiming the attitude tracks an objective fact?
- Which factual changes could revise the stance, and which conflict would remain?

Room 6: A rule for living together

Arin and Bela may never agree about the best death. They still need a public rule.

Possible settlements include:

- permit assisted dying under strict conditions;
- prohibit it;
- allow patient self-administration but not clinician administration;
- permit individual clinicians to abstain while requiring referral;
- allow regional variation;
- create a narrow pilot program with review;
- leave the decision to private medical associations;
- tolerate civil disobedience under specified circumstances.

A coordination rule answers:

What practice will govern us while substantive disagreement remains?

This is not identical to moral agreement. People drive on the same side of the road without believing that side objectively superior. A ceasefire coordinates enemies who continue to judge each other harshly. A legal compromise can permit conduct some citizens condemn while protecting their refusal to participate.

Practical settlement may depend on stability, mutual advantage, tolerable loss, enforceability, or a shared desire to avoid domination. Its success is measured partly by whether people can live under it, not by whether it proves one moral theory true.

The distinction cuts both ways. Passing a law does not make a moral proposition true. Failing to pass a law does not prove the prohibited act immoral. Law determines institutional permission and response. It may express values and affect behavior. It does not function as a metaphysical voting machine.

Checkpoint

Ask:

- What rule could coordinate action despite continuing disagreement?
- Which losses would each side be asked to tolerate?
- Does the settlement depend on shared values, bargaining, power, or all three?
- Has practical acceptance been mistaken for moral agreement?

Room 7: Who decides?

Suppose Arin and Bela agree that a public decision is needed. They may still disagree about procedure.

Should the issue be decided by:

- a popular referendum;
- the elected legislature;
- constitutional courts;
- medical regulators;
- individual patients and clinicians;
- a citizens' assembly;
- local jurisdictions;
- an expert commission followed by legislative review?

This is a procedural disagreement. It concerns authority, competence, representation, accountability, appeal, and the protection of minorities.

A referendum may represent citizens broadly while reducing a complex case to a slogan. Experts may understand medicine while lacking democratic authority. Courts may protect rights while removing decisions from public revision. Individual choice may protect agency while missing structural pressure.

Procedure does not merely count preferences. Its design determines who speaks, what evidence enters, how burdens are distributed, and what may be challenged.

Yet a fair procedure can produce a substantively bad result. It can be misinformed. Participants can share prejudice. A majority can burden a minority. Procedural legitimacy and moral truth are distinct even when both matter.

Checkpoint

Ask:

- Which institution has jurisdiction?
- Who is represented, and who is affected without representation?
- What expertise, transparency, reasons, appeal, and revision are required?
- Are any outcomes protected from majority decision, and what grounds those limits?

Room 8: The objective claim

After every other layer is clarified, Arin says:

Even if every person opposed assisted dying, a competent patient's autonomy would make it objectively permissible.

Bela replies:

Even if every person approved, intentionally helping an innocent person die would be objectively wrong.

Now the disagreement is explicitly **metaethical**: it concerns what could make a moral claim true, not merely which policy to choose. Both speakers claim a truth independent of actual approval. They may agree about definitions, effects, standing, values, policy, and procedure while disagreeing about objective status.

Or one may make the objective claim while the other does not:

I am not claiming that assisted dying has a stance-independent status. I support permission because I value agency and relief, and I am willing to defend a rule with safeguards.

That speaker may be politically firm and metaethically modest. The opponent cannot establish an objective disagreement merely by repeating an objective claim more loudly.

The eighth room contains the question central to this book:

What could make the alleged objective fact true, knowable, and authoritative?

An unresolved dispute is not automatically evidence that an objective moral fact lies beneath it. People can remain divided because their values conflict, their trust differs, their identities are at stake, or no settlement satisfies everyone.

Conversely, rejecting objective moral facts does not dissolve the other seven rooms. Definitions, predictions, standing, priorities, attitudes, coordination, and procedures remain capable of generating profound conflict.

Checkpoint

Ask:

- Is anyone explicitly claiming truth independent of every actual approval?
- What is proposed to make the claim true?
- How is the truth accessed?
- What does the objective claim add to the other seven layers?

A map, not a staircase

The eight rooms are not stages through which every disagreement must pass. They are not ranked from superficial to deep. A factual error can matter more than a metaphysical dispute. A procedural decision can determine whether people live or die. A framing choice can hide an affected population.

Nor are the rooms sealed. Evidence affects attitudes. Standing affects priority. Language shapes procedure. Institutions train emotion. Objective beliefs support commitments. Power influences which facts are gathered.

The point of separation is not to pretend the layers never interact. It is to see the interaction rather than compress it into “we disagree morally.”

A useful summary looks like this:

Layer	Central question	What progress can achieve
Words and frames	Are we describing the same thing?	Clearer terms and revealed framing
Empirical predictions	What is likely to happen?	Better-supported factual belief
Standing	Who or what enters directly?	An explicit inclusion boundary
Priority	What leads when values conflict?	A ranking, threshold, comparison, or acknowledged impasse
Attitude	What do we endorse, oppose, fear, or protect?	Honest commitments and possible emotional revision
Coordination	What rule will govern shared action?	A workable settlement
Procedure	Who decides, and how?	A legitimate or at least accountable decision process
Objective status	What is true independent of approval?	A claim about moral truth requiring grounding

Notice the verbs. Clarifying, predicting, including, prioritizing, committing, coordinating, deciding, and grounding are different activities. One method will not perform all eight.

Integrated case: The final request

Consider a fuller exchange.

Mara has an incurable neurological condition. She is expected to die within four months. She repeatedly requests lethal medication. She is mentally competent, receives excellent palliative care, and is under no financial or family pressure.

Arin says:

Denying Mara a dignified death is wrong. Evidence from other jurisdictions shows safeguards work. Her body and suffering are hers, and autonomy should control. I support a regulated permission decided by patients and physicians. Even universal opposition would not erase her right.

Bela says:

Calling death “dignified” frames continued dependence as undignified. Some safeguards fail in practice, though I accept that Mara is competent and uncoerced. Her choice matters, but medicine’s commitment not to kill takes priority. I support prohibition enacted by the legislature. Even universal approval would not make intentional killing right.

Classify only what is stated.

- **Words and frames:** Bela challenges *dignified death* and its implication about dependence.
- **Evidence:** They disagree about the reliability of safeguards in general, though not about Mara’s competence.
- **Standing:** Both include Mara; Bela also gives direct importance to a professional commitment. Whether a profession itself has standing or the commitment matters through clinicians requires clarification.
- **Priority:** Arin prioritizes patient autonomy; Bela prioritizes medicine’s nonkilling commitment.
- **Attitude:** Arin expresses support for control over dying; Bela expresses opposition to intentional medical killing.
- **Coordination:** Arin proposes regulated permission; Bela proposes prohibition.
- **Procedure:** Arin emphasizes patient-clinician decision; Bela assigns the legislature final authority.
- **Objective status:** Both make approval-independent claims.

The classification does not decide the case. It tells us which work remains.

Now remove one sentence at a time. If Bela says nothing about empirical failure, do not attribute that prediction. If Arin does not mention objective truth, do not assume it. If both support the same pilot program, record practical agreement even if their moral verdicts differ.

Intellectual restraint is part of rigor. A map with invented roads is worse than no map.

Exercise: Decode a live disagreement

Purpose. Classify the components of one controversy before trying to resolve it; the map shows where different methods apply and does not require a verdict on the controversy.

Choose a controversy you can describe without needing to settle it immediately. Write two representative claims in the speakers’ own strongest language.

Create eight headings:

1. Words and frames
2. Empirical predictions
3. Standing

4. Priority
5. Attitudes and commitments
6. Coordination rule
7. Decision procedure
8. Objective claim

Under each heading, write only what the statements establish. Use **Not yet stated** when necessary. That phrase is a mark of discipline, not failure.

Then ask:

- Which layer is most likely to yield to new evidence?
- Which could be settled practically without agreement?
- Which concerns an openly different value?
- Which depends on the boundary of inclusion?
- Which invokes stance-independent truth?
- What would remain if the objective claims were removed?

Exercise: Solve the wrong problem

Purpose. Practice detecting exchanges in which speakers answer different layers; the task is diagnosis and repair of the conversation, not selection of the better substantive position.

For each exchange, identify the mismatch.

Exchange A

“The policy will increase homelessness.”

“Individual freedom is objectively valuable.”

One speaker makes an empirical prediction; the other states an objective value claim. Either may be relevant. Neither answers the other.

Exchange B

“The vote was fair.”

“The winning policy will seriously burden a small minority.”

One speaker addresses procedure; the other addresses distribution and priority. A fair vote can impose a grave burden. The burden does not by itself prove the vote procedurally unfair.

Exchange C

“Your religion forbids the practice.”

“I personally oppose it even apart from my religion.”

One speaker attributes a source; the other expresses a stance. More inquiry is needed before classifying the disagreement as theological.

Exchange D

“No objective moral fact settles this.”

“Then we cannot make a law.”

The first concerns objective status; the second infers the impossibility of coordination. The inference requires a bridge. Societies make laws amid metaphysical disagreement constantly, sometimes before breakfast.

Write two mismatched exchanges of your own. Then rewrite each so the speakers address the same layer.

Construction exercise: The Candidate System’s dispute protocol

Purpose. Give the Candidate System a repeatable order of inquiry so facts, meanings, values, procedures, and objective claims do not answer in one another’s place.

Add this protocol to your Candidate System:

Before treating a dispute as a conflict over objective moral truth, identify any disagreement about description, prediction, standing, priority, attitude, coordination, or procedure. Resolve what can be resolved at its own layer. Record what remains.

Now add answers:

- Factual claims will be assessed by _____.
- Standing will be assigned according to _____.
- Value conflicts will be compared by _____.
- Persistent attitude conflicts will be handled through _____.
- Shared action will be coordinated by _____.
- Procedures will be judged or selected by _____.
- Objective moral claims will require _____.

Several blanks will remain. That is expected. Later chapters supply candidate answers.

The fifth debt

Add this entry to the ledger:

Debt 5 — Layer discipline: *Which part of a moral conclusion comes from facts, standing, priority, attitude, coordination, procedure, or alleged objective truth, and has any one layer been asked to perform another’s work without a bridge?*

The debt is methodological. A system can reach a conclusion for several reasons. It must show the route.

Inspection Report I — The Bill of Specifications

Purpose. Assemble the project's initial target, meanings, and dispute method in one stable record; later inspections must show revisions against this version rather than an improved memory of it.

Part I is complete. Before entering the construction floor, assemble the following pages.

1. Candidate System 0.1

Your opening principle:

An action is objectively right when _____.

Do not improve it from memory. Preserve the original wording so later revision remains visible.

2. The objectivity threshold

State:

- what remains true under universal rejection;
- whether idealized agreement is sufficient;
- whether a response-dependent truth counts as objective for this project;
- what would constitute a change of subject.

3. The meaning rule

State whether central moral sentences report facts, contain a systematic error, express commitments, apply standards, report constructed outcomes, participate in a fiction, or use an explicit combination.

4. The essential specifications

List the conditions you marked essential:

- truth;
- stance independence;
- scope;
- authority;
- accessibility;
- coherence;
- determinacy;
- usability;
- correctability;
- any substantive condition you added.

5. The disagreement protocol

Keep the eight-layer map available. Every later pressure case will isolate one question before reintegrating the result.

6. The ledger

You should now have five debts:

1. **Added status:** What does the moral verdict add?
2. **Independence:** What makes its truth survive rejection?
3. **Meaning:** What job do the central terms perform?
4. **Joint performance:** Can the essential specifications be satisfied together?
5. **Layer discipline:** Which part of the conclusion comes from which kind of support?

Sign the anti-rigging declaration from Chapter 4.

You are now carrying more paper than most adventurers in fiction, but fewer swords. This is appropriate. The dangers ahead are mostly hidden premises.

In Part II, we give the system content. We draw the circle of direct standing, decide what receives a verdict, identify the values doing the work, turn those values into guidance, confront comparisons that resist a common scale, and test whether similar cases remain similar when loyalty, identity, and power enter.

The blank constitution is about to acquire inhabitants.

That is when construction becomes dangerous.

PART II



Give the System Content

A blank constitution cannot contradict itself. It also cannot guide anyone. Now the construction acquires a circle, a vocabulary, values, rules, and ways of choosing.

Chapter 6 — Draw the Circle

The first meeting of Port Meridian’s constitutional convention lasted eleven minutes before someone suggested drawing a circle.

The delegates had been asked a simple question: *Who counts?*

“Everyone,” said Councilor Pell.

This produced approving murmurs. It was humane, concise, and almost entirely uninformative.

“Every person?” asked Iona Venn, the harbor engineer.

“Naturally.”

“Does a newborn count?”

“Of course.”

“A visitor?”

“Yes.”

“An enemy soldier who has surrendered?”

There was a pause, but the answer was still yes.

“A pig?”

The room divided.

“A conscious machine?”

The division acquired factions.

“A person who will exist in eighty years if we do not poison the aquifer?”

Several delegates inspected their papers as though an answer might have been printed in small type.

“A forest?”

Someone asked whether lunch had been ordered.

The circle on the board had not changed. What had changed was everyone’s understanding of the word *everyone*.

This is the first construction task in giving a moral system content. Before a system can say what should happen, it must say whose or what interests, claims, choices, conditions, or existence enter the reasoning at all. This is the problem of **moral standing**.

The term sounds ceremonial, as if beings are waiting to be admitted to a very selective garden party. The reality is more serious. To have **direct moral standing** is to count for one’s own sake. It means that what happens to the being can matter in the moral calculation without first being translated into its effects on somebody else.

A dog might matter directly because the dog can suffer. A historic bridge might matter indirectly because people cherish it, depend on it, or would lose something important if it fell. The distinction does not imply that the bridge is disposable. It says that the explanation of why the bridge matters runs through beings, relationships, projects, or values that have been granted standing on other grounds.

There is also a third possibility: no standing in the case at hand. A pebble displaced three centimeters on an empty moon may make no moral difference. This is not an insult to the pebble. Pebbles are famously resilient in the face of social exclusion.

We therefore need three boxes:

1. **Direct standing:** this entity or its good counts for its own sake.
2. **Indirect relevance:** this entity matters because of its relations to things with direct standing or to admitted values.
3. **No established standing:** the system has not found a morally relevant route from this entity to its decisions.

The third box should be labeled carefully. “No established standing” is more honest than “worthless.” A theory can fail to find a basis for direct standing without proving that no such basis exists.

Where this chapter is headed. We will decide who or what enters the Candidate System as a direct object of concern. This is only a rule of admission: it does not yet determine equal weight, identical treatment, or the outcome when admitted interests conflict.

Admission is not victory

Before the applicants approach the gate, we need a rule that will prevent a great deal of confusion:

Standing, weight, and outcome are different questions.

To say that a being has standing is to say that it enters deliberation. It does not yet say how much its interests count, whether all of its interests count, whether it must receive identical treatment, or whether its interests defeat every competing consideration.

Imagine a public hearing. Receiving a place at the microphone does not guarantee that the council will adopt your proposal. It guarantees that the council may not pretend you do not exist. Standing is admission to the hearing, not possession of the building.

This matters because arguments about standing often smuggle in fear about outcomes. A delegate worries that if animals count directly, then humans can never eat, farm, experiment, build, or defend themselves. Another worries that if future people count, the present generation must surrender nearly everything to maximize an enormous possible future population. Those conclusions do not follow from admission alone. They require further rules about weight, comparison, uncertainty, aggregation, and priority. We will construct those rules later.

The opposite mistake is equally common. A system says that a being counts, then gives its interests zero practical effect whenever those interests become inconvenient. A name on the guest list is not meaningful admission if the door remains locked.

Your Candidate System must therefore answer three questions separately:

- Who or what enters directly?
- Which of their features or interests matter, and with what weight?
- What happens when those considerations conflict with others?

Today we answer only the first.

Seven keys for the gate

The delegates proposed seven possible grounds of admission. None was merely a list of favored beings. Each was a property or relation that might explain *why* something counts.

Felt experience

The first proposal was **sentience**: the capacity for experiences that feel good, bad, or otherwise valenced from the inside.

If a creature can suffer, its suffering is not merely a change in the arrangement of matter. There is something it is like for that creature when the change occurs. A hot surface and a burned paw are physically connected, but only one of them is in pain.

Sentience provides a direct and powerful explanation of why many animals count. It also includes humans who cannot reason abstractly, make contracts, speak, or pursue elaborate life plans. The sleeping, the temporarily unconscious, and those with severe cognitive disabilities need not pass a philosophy examination before receiving protection.

But sentience raises boundary questions. How confident must we be that a system can feel? Which experiences count? Does the capacity for pleasure and pain ground only welfare interests, or also rights? Does a barely sentient being count less than a richly self-aware being, or do those differences concern weight rather than standing?

Consider **Luma**, a laboratory animal engineered to live for one day. Luma feels pain and comfort but forms no plans beyond the next few minutes. If felt experience grounds standing, Luma enters the circle. Her lack of a five-year plan may affect what can be taken from her, but it does not make a burn painless.

Agency and self-direction

The second proposal was **agency**: the capacity to form aims, make choices, govern conduct, and shape a life.

Agency explains why coercion, deception, manipulation, and domination can be wrong even when they do not reduce immediate happiness. To override an agent is not merely to change the agent's sensations. It is to displace someone from the authorship of a life.

This key fits many important locks, but it does not open the entire gate. An isolated newborn cannot formulate aims. A person with advanced dementia may have lost capacities once central to self-government. A frightened animal may possess preferences without reflective agency. If agency is the only ground of standing, these beings must either be excluded, admitted by potential or history, or protected only through someone else's concern.

Potential and history can do real work, but they are not magic words. A newborn has the potential to become an agent. So does a fertilized egg under suitable conditions. A

patient once had agency. So did a dead author. Where, exactly, does potential or former agency begin and stop conferring direct status? The system owes a rule, not a mood.

Interests that can prosper or fail

The third proposal was **having a good of one's own**: being the kind of entity that can be helped or harmed.

This category may be broader than felt experience. A plant can flourish or wither. A corporation can succeed or collapse. A species can persist or go extinct. A river can be clean or polluted. Yet “good for” is slippery. Oil is good for an engine, and nobody proposes sending the engine an apology when the oil change is late.

We need to distinguish several senses of good:

- something can promote a function;
- something can preserve an organized system;
- something can fulfill a goal;
- something can serve another being's use;
- something can improve a subject's lived condition.

Only the last clearly includes a beneficiary for whom things go better or worse. The other senses may still matter, but they do not automatically establish a subject with direct standing. If your Candidate System admits all goal-directed systems, it may find itself representing thermostats at the convention. This is not a decisive objection. It is a request to check the invitations.

Relationship and dependence

The fourth proposal looked not inside the candidate but between the candidate and others. A being may count because it is dependent, entrusted to us, bound to us, or situated within a relationship that creates responsibilities.

The isolated newborn makes this vivid. Suppose an infant is found in a sealed weather station. No parent survives. The infant has no present agency, no social membership that anyone has formally recognized, and perhaps no concept of a future. Yet the infant's radical vulnerability seems to call for care rather than exclusion. Dependence is not a defect that cancels claims; it may be the very condition that makes claims urgent.

Relationship grounds can also explain special obligations. A parent may owe more to their child than to a distant child, even if both have equal direct standing. A physician may have special duties toward a patient. Notice, however, that this often concerns *what is owed by whom*, not whether the recipient counts at all. If a lost child deserves rescue only after becoming *our* child, the circle has been drawn alarmingly close to home.

Membership and participation

The fifth key was membership in a community of cooperation, language, law, or mutual recognition.

Membership explains why shared institutions generate claims. Citizens finance common goods, accept restrictions, rely on public promises, and participate in schemes whose benefits and burdens require distribution. Participants can have claims arising from the terms of the practice.

But membership is a dangerous sole criterion for basic standing. Outsiders, refugees, enemies, and the socially excluded are precisely those whom an inward-facing system finds easiest to mistreat. If admission depends on prior recognition by the group, then a group can erase duties by refusing recognition. The circle becomes a magic marker wielded by whoever already stands inside it.

Status independent of our approval

The sixth proposal was direct: persons, creatures, or other entities possess a moral status that does not depend on our choosing to grant it.

This is attractive because it blocks the magic marker. A prisoner does not lose status when the jailer becomes indifferent. A minority does not vanish morally when a majority votes to exclude it.

Yet the proposal presently names the desired result more clearly than its basis. What property confers this status? Rational nature? Creation by a deity? A soul? Sentience? Membership in a natural kind? A fundamental moral fact? How do we identify its bearers? Why does that property have status-conferring power?

These questions do not refute approval-independent status. They place a large entry in the Debt Ledger. Part III will ask whether such status can be grounded and whether it supplies reasons to agents who reject it. For now, you may use it as a candidate—provided you write **promissory note** beside it.

A defensible constructed rule

The seventh proposal was to construct a rule of inclusion under conditions designed to prevent bias. We might ask what boundary free and equal participants would accept without knowing which kind of being they would be, or what rule could be publicly justified to all affected agents.

Construction can discipline self-interest. If you do not know whether you will be powerful, disabled, foreign, young, or vulnerable, you may choose protections that do not conveniently track present advantage.

But a procedure needs participants, information, constraints, and aims. Can animals participate through representatives? Do future people receive proxies? Is sentience already used to decide who is represented? Are outcomes constrained by rights the procedure did not create? A constructed boundary may be wise and fair, but it cannot select its own construction materials from nowhere.

The applicants arrive

The seven keys are proposed grounds of admission, not yet a defensible boundary. The next cases test what each key admits, excludes, or leaves dependent on an unacknowledged second ground.

The silent directive

Ilyan signs an advance directive refusing invasive treatment if irreversible brain damage ends consciousness. Years later, that condition occurs. Ilyan is alive but permanently insentient.

Does the current organism have direct standing? Does the former agent's directive retain authority? Does violating it wrong the person Ilyan was, the legal practice of promising, the family, or the present human organism?

A system may reach the same practical verdict by several routes. The routes are not interchangeable. Former agency, current interests, symbolic respect, and institutional trust are different explanations.

The digital navigator

A ship's navigation program says that it fears deletion. It remembers previous voyages, revises its goals, bargains for continued existence, and reports a private stream of experience. Engineers disagree about whether the reports reflect consciousness or sophisticated imitation.

The case contains two questions. First, what property would confer standing if present? Second, how should uncertainty about that property affect treatment? A sensible rule may grant precautionary protection before certainty arrives. Precaution is not proof that the program is conscious. It is a decision about the cost of being wrong.

The people who might be

Port Meridian can build a reservoir that will slowly leak toxins. No one now alive will be harmed. If the reservoir is built, a neighborhood will arise downstream in eighty years and its residents will suffer. If it is not built, somewhat different people may be born elsewhere.

Future-person cases become tangled because our choices can affect not only how well future people live but which people exist. Still, the absence of a currently identifiable victim does not make poisoned water refreshing. Your system must decide whether future persons count as possible bearers of interests, whether impersonal states of affairs matter, or whether duties run through present projects and institutions.

The chorus with no listener

Forty simple programs coordinate so effectively that together they pursue goals, revise strategy, and defend their continued operation. No individual program is conscious, and there is no evidence of a unified subject.

The collective may be an agent in the sense that it acts coherently. It may be valuable, useful, dangerous, or legitimately governed. But is there someone there for whom shutdown is bad? Agency and sentience, we now see, can come apart.

The living watershed

A watershed contains forests, wetlands, predators, prey, fungal networks, and migrating birds. It regulates itself through complex feedback but has no unified consciousness or goal. Destroying it would cause extinctions, animal suffering, cultural loss, and long-term human risk.

There are overwhelming reasons to protect it even if the watershed lacks direct standing. Alternatively, a system might treat ecological wholes as direct bearers of value. Either route is available. What is not available is moving between them mid-argument: speaking of the ecosystem as a rights-bearing whole when that strengthens protection, then as a collection of resources when extraction becomes convenient.

Draw your circle

Purpose. Make every admission to or exclusion from direct standing rest on a named ground; the map does not yet assign equal weight or decide conflicts among those admitted.

Open the Candidate System notebook and make a domain map with four columns:

Candidate	Proposed ground	Direct, indirect, or unestablished	Confidence and uncertainty rule
Current adult humans			
Infants			
Cognitively disabled people			
Sentient animals			
Future people			
Artificial minds			
The dead and their directives			
Communities and institutions			
Species and ecosystems			
Outsiders and enemies			

Add any candidates your life or society makes relevant. For each direct admission, name the feature that does the work. For each indirect admission, name the route through

which it matters. For uncertain cases, state what level of evidence triggers caution and what protection follows.

Then perform two audits.

The inconvenient-case audit: Does your proposed ground exclude any being you strongly want to protect? If so, either accept the exclusion, add another principled ground, or explain why the being matters indirectly. Do not simply write “special case” in the margin and hope the ink looks philosophical.

The opportunistic-expansion audit: Does your ground include entities you do not want to admit? If so, either admit them, refine the ground in terms that independently explain the difference, or abandon it. “They are obviously different” is a promissory note wearing sunglasses.

The sixth debt

Add this to the ledger:

Debt 6 — The boundary: What property or relation gives direct standing, why does it have that role, and has the boundary been narrowed or expanded merely to preserve preferred verdicts?

The circle is now on paper. It will almost certainly be revised. Revision is not failure; concealed revision is.

We know, provisionally, who enters the constitution. We still do not know what the constitution evaluates. A person can act badly from a good motive and produce an excellent result by accident. An institution can impose a defensible rule through a corrupt process. A tragic outcome can occur without wrongdoing, while wrongdoing can occur without harm.

If all of these receive the same stamp—**WRONG**—the Candidate System will soon resemble a customs office with one form and a very angry clerk.

In the next chamber, we manufacture better stamps.

Chapter 7 — Decide What Receives a Verdict

At 2:14 on a winter morning, an explosion tore through the eastern clinic in Port Meridian.

By sunrise, five claims were circulating:

- Doctor Senn had done something wrong.
- The treatment protocol was wrong.
- The outcome was wrong.
- The clinic's staffing policy was wrong.
- Blaming Senn would be wrong.

The word was working overtime and had begun to lose structural integrity.

Here is what happened. A patient arrived with a rare infection. Two treatments were available. The standard treatment worked in nearly every known case. The alternative carried serious immediate risks and was normally reserved for a specific genetic variant. The clinic's sequencer was unavailable because administrators had postponed maintenance. Working from symptoms and the evidence then available, Senn reasonably chose the standard treatment. The patient carried the variant. The medication triggered a chain reaction, an oxygen cylinder failed, and the eastern wing exploded.

Senn intended to heal. The decision procedure was careful. The available evidence favored the choice. The result was catastrophic. The equipment had been neglected. The maintenance delay had been legal. The patient's family suffered a loss no apology could reverse.

Was the event wrong?

That is not yet one question. It is a sack containing several questions, some of which are kicking.

Where this chapter is headed. We will distinguish the objects a system evaluates—acts, motives, outcomes, characters, relationships, rules, and institutions—and give their possible statuses a clearer grammar. This prevents a conclusion about one object from migrating unnoticed to another.

The objects on the workbench

A moral system does not evaluate “the situation” all at once. It assigns statuses to **objects of evaluation**. Different theories emphasize different objects, but ordinary moral judgment ranges across at least eight.

Actions

An **action** is what an agent does: administers a drug, tells a lie, breaks a promise, opens a border, signs an order.

Action evaluation asks whether doing that thing was required, permitted, forbidden, admirable, careless, cruel, or otherwise assessable under the system’s rules.

Motives and intentions

A **motive** is what moves the agent; an **intention** is what the agent aims to do or bring about. The categories overlap but are not identical. A person might intend to donate money in order to humiliate a rival. The intended act is aid; the motive is spite.

Motives matter because identical movements can express very different forms of regard. A surgeon cuts to heal; an attacker cuts to injure. Yet good motives do not guarantee good decisions, and bad motives do not guarantee bad outcomes. The universe refuses to organize itself for easy grading.

Outcomes

An **outcome** is what occurs: suffering prevented, trust broken, a life saved, a building destroyed.

Outcomes can be better or worse, fortunate or disastrous, independently of whether anyone acted culpably. A storm can cause a terrible outcome without committing a meteorological offense.

Rules

A **rule** directs a class of actions: treat presumptively with the standard medication; never disclose a confidence; allocate organs by expected benefit.

A good action can result from violating a bad rule, and a bad outcome can result from correctly following a generally defensible rule. We therefore need to distinguish the evaluation of a rule from the evaluation of its application in one case.

Institutions and procedures

An **institution** is a durable arrangement of roles, incentives, powers, and practices. A **procedure** is a structured way of deciding or acting.

No single administrator may intend a clinic's dangerous understaffing. The danger may emerge from budgets, schedules, reporting rules, fragmented authority, and incentives that make every local choice appear reasonable. If the system evaluates only individual acts, institutional responsibility becomes invisible precisely when it matters most.

Agents and character

An **agent** is the person or group acting. **Character** concerns relatively stable dispositions: honesty, courage, callousness, patience, vanity, practical wisdom.

One cowardly act does not necessarily make a coward. One lucky rescue does not make a hero. Agent evaluation asks what the conduct reveals, how settled the disposition is, what the agent knew, and whether change is possible.

Relationships

A **relationship** can be trustworthy, exploitative, caring, dominating, reciprocal, neglectful, or corrupt even when no single action captures its structure.

A friendship in which one person always yields may contain no spectacular betrayal. Its moral defect may be a pattern. Likewise, a workplace can be formally polite yet organized around fear.

Forms of life

A **form of life** is a larger pattern of practices, expectations, aspirations, and social organization. It asks not merely, "Was this act allowed?" but "What kind of people and relations does this way of living cultivate?"

A society could contain individually permissible transactions while leaving everyone isolated, anxious, and unable to care for dependents. Whether that is a morally defective form of life is a further question, not a synonym for whether any one contract was forbidden.

Your Candidate System may not need all eight objects. It does need to say which it evaluates and to stop conclusions from migrating between them without a bridge.

"The outcome was terrible; therefore the clinician acted wrongly" is such a migration.

"The clinician is admirable; therefore the protocol is safe" is another.

"No administrator acted maliciously; therefore the institution is just" is a third.

These are not always false. They are incomplete inferences. The missing bridge might be negligence, foreseeability, design responsibility, or a rule linking character to institutional authority. Without that bridge, the conclusion has changed workbenches while nobody was looking.

Build a vocabulary larger than one word

“Wrong” is useful. So is a kitchen drawer. Trouble begins when the kitchen drawer contains cutlery, batteries, tax records, three mysterious keys, and a small object nobody is willing to identify.

The Candidate System needs a **status grammar**: distinct verdicts with rules for when each applies. Begin with these.

Required

An act is **required** when failing to perform it violates a governing demand. Required is stronger than good or advisable. If an act is required, omission needs justification or excuse.

Forbidden

An act is **forbidden** when the system rules it out, absent an applicable exception. Forbidden concerns the act’s status, not automatically the agent’s blameworthiness.

Permitted

An act is **permitted** when it falls within the acceptable range. Permitted does not mean best, wise, kind, or beyond criticism. You may be permitted to decline an invitation in a manner suggesting you were raised by decorative wolves.

Admirable or beyond duty

An act is **admirable** when it excellently realizes a value. Some admirable acts may be **beyond duty**: good to do but not required. A system that treats every admirable sacrifice as mandatory can turn sainthood into the minimum dress code.

Justified

An otherwise prohibited act is **justified** when the circumstances make doing it the right or permissible response. Breaking down a door to rescue a child may be justified. The act is not merely forgiven; under the circumstances, it should be done or allowed.

Excused

An agent is **excused** when the act remains defective or prohibited but the agent is not fully blameworthy because of ignorance, coercion, incapacity, or another responsibility-reducing condition.

Justification changes the status of the act. Excuse changes the assessment of the agent. If you shove someone aside to avoid a falling beam, you may be justified. If someone forces your hand onto the alarm, you may be excused for triggering it even if triggering it was harmful. The difference matters to victims, repair, trust, and future policy.

Blameworthy

An agent is **blameworthy** when the agent is appropriately subject to a negative responsibility response for conduct or character. This usually depends on some combination of control, knowledge, opportunity, quality of will, and applicable expectations.

An act can be forbidden without its agent being blameworthy. A young child can do what should not be done without meriting the blame appropriate to an informed adult.

Tragic or regrettable

An outcome is **tragic** when it involves grave loss, perhaps through no one's fault. It is **regrettable** when there is reason to wish it had not occurred.

Calling an event tragic should not wash away responsibility when wrongdoing helped cause it. Calling it wrongdoing should not imply that every loss could have been prevented by a morally perfect agent. Sometimes a defensible choice leaves devastation on the floor.

Repair-demanding

A situation is **repair-demanding** when someone has reason or responsibility to compensate, restore, apologize, acknowledge, reform, or otherwise address a loss.

Repair need not track blame exactly. A hospital may owe compensation for a non-negligent treatment injury because it created the risk, benefited from the practice, insured against loss, or can distribute costs fairly. Parents repair damage caused by a toddler without declaring the toddler a villain. A government may owe redress for an inherited institutional injustice even when present officials did not create it.

Punishable

Conduct is **punishable** when an institution is authorized to impose a sanction. Punishability requires more than moral criticism. It also raises questions of authority, evidence, due process, proportionality, deterrence, safety, and the risks of abusive enforcement.

Not everything blameworthy should be illegal. Not everything illegal is blameworthy. Not everything harmful is punishable. A moral system that skips these distinctions can move from "I disapprove" to "the state may use force" in one exhilarating and disastrous leap.

You now have ten statuses. Add more if your system needs them: honorable, disloyal, unfair, negligent, invalid, corrupt, oppressive, forgivable. But do not multiply labels merely for decorative grandeur. Each status earns its place by doing work no other status already does.

Return to the eastern clinic

Let us place the objects and statuses in separate rows.

The clinician's action

Given the evidence available, choosing the standard treatment was reasonable. If the Candidate System judges actions partly by what a competent agent could know at the time, Senn's choice may have been **permitted** or even **required**.

If the system judges only actual consequences, it may classify the act negatively. But then it must confront cases in which a careful choice happens to end badly and a reckless choice happens to end well. Outcome-only judgment is possible; it is not free.

The clinician's process and intention

Senn checked the records, consulted the protocol, weighed the known risks, and intended to heal. The process appears careful and the intention benevolent. This supports favorable agent evaluation, though neither fact resurrects the patient.

The outcome

The death and explosion were **catastrophic**, **tragic**, and **regrettable**. Those verdicts do not require finding Senn culpable.

The treatment rule

The standard-treatment rule may have been statistically defensible when paired with functioning genetic screening. Without screening, its use conditions were not satisfied. The relevant defect might lie not in the clinical rule itself but in a system that continued invoking it after disabling one of its safeguards.

The institution

Administrators knowingly deferred maintenance to meet short-term budget targets. Perhaps no one predicted this exact chain reaction. Exact prediction is not always required for negligence. If they created an unjustifiable risk, concealed it, or ignored warnings, the policy may be **wrong**, the decision-makers **blameworthy**, and the institution **repair-responsible**.

If the failure emerged from dispersed authority rather than individual malice, institutional reform may still be required. "Nobody meant this" is relevant to intention. It is not a maintenance plan.

The response

The family deserves truth, acknowledgment, investigation, and support. Whether Senn should apologize depends on what apology means. An expression of sorrow need not

be a confession of fault. Whether Senn should be punished depends on blame, evidence, authority, and the purpose of sanction—not simply on the magnitude of the family’s loss.

This layered judgment may feel less emotionally satisfying than finding one object and stamping it **WRONG**. It is also more capable of locating prevention and repair. Anger wants an address. Rigor must check the house number.

Five distinctions that prevent moral pileups

Write these into the Candidate System’s grammar.

Act versus agent

An act can be wrong while the agent is excused. An act can be permissible while revealing an ugly motive. Judge each through its own conditions.

Justification versus excuse

A justified agent may say, “Given the circumstances, this was the act to perform.” An excused agent may say, “The act was not acceptable, but I could not reasonably have done otherwise or known better.”

Illegality versus blameworthiness

Law is a public institutional rule backed by authorized sanctions. Moral blame is an assessment of responsible agency. They overlap imperfectly. A parking violation may be illegal and barely blameworthy. A cruel betrayal may be seriously blameworthy and entirely legal.

Tragedy versus wrongdoing

Tragedy names grave loss. Wrongdoing names a violation attributable to an act, practice, or agent under some rule. A case can contain either, both, or neither.

Repair versus punishment

Repair looks toward victims, losses, relationships, and future conditions. Punishment imposes a burden under an authority. One can be appropriate without the other. Conflating them turns every offer of repair into an admission of punishable guilt and every absence of punishment into abandonment of the victim.

Manufacture your stamps

Purpose. Apply separate status terms to each relevant object in one difficult event; a complete table prevents one adverse label from silently deciding justification, excuse, character, repair, and punishment at once.

Choose a difficult event from public life, history, or your own experience. If using a personal case, change identifying details. Make one row for each relevant object:

Object	Description	Status	Rule used	Confidence
Action				
Intention or motive				
Outcome				
Rule				
Institution or procedure				
Agent or character				
Relationship				
Form of life				

Use at least eight statuses from the grammar. You need not use every status in the case. For each one you do use, write a one-sentence application rule.

Then ask:

- Did an outcome verdict migrate to an act without a responsibility rule?
- Did praise or blame of a person substitute for evaluating the institution?
- Did “legal” substitute for “permitted,” or “illegal” for “blameworthy”?
- Did a demand for repair silently become a demand for punishment?
- Did “tragic” conceal avoidable wrongdoing?
- Did “wrongdoing” imply that every bad consequence was foreseeable?

If the answer is yes, separate the objects again.

The seventh debt

Add this entry:

Debt 7 — The verdict grammar: What exactly receives each verdict, what does that verdict mean, and has one overloaded moral term been used to slide from harm to prohibition, from prohibition to blame, or from blame to authorized punishment?

Port Meridian now has a circle of admitted candidates and a cabinet of labeled stamps. It can say who enters and distinguish what, exactly, it is evaluating.

But the stamps are not self-inking.

Why classify one outcome as better, one choice as protected, one relationship as corrupt, or one institution as fair? To answer, the delegates must put values into the constitution. They will be tempted to write every attractive word they can remember.

This is understandable. It is also how moral systems become restaurant menus with no prices, recipes, or kitchen.

In the next chamber, we decide what matters—and in what sense it matters.

Chapter 8 — Name What Matters

The Port Meridian delegates reconvened beneath a new banner:

OUR VALUES

Within minutes they had produced the following list:

Freedom. Safety. Dignity. Happiness. Justice. Equality. Excellence. Community. Tradition. Progress. Responsibility. Compassion. Truth. Loyalty. Peace.

They regarded the list with satisfaction. It had all the moral charm of a basket of polished stones.

“Which value wins when they conflict?” asked Iona.

The satisfaction dimmed.

“What does dignity require?”

Someone adjusted the banner.

“Does freedom include the freedom to create danger? Does equality mean equal treatment, equal opportunity, equal resources, equal status, or equal concern? Is tradition good because it is inherited, because it contains accumulated knowledge, because it creates belonging, or only when its content is independently good?”

The banner fell down.

A moral system needs values, but a list of attractive nouns is not yet a value system. It is an inventory awaiting instructions.

In this chapter we will ask three questions:

1. What kinds of things matter?
2. In what way do they matter?
3. What role does each play when conflict arrives?

The third question will lead into the next two chapters. For now, our task is to identify the materials accurately enough that we do not mistake a pile of bricks for a house.

Where this chapter is headed. We will turn the phrase *what matters* into an explicit value inventory, then distinguish the roles values can play. Naming a value is necessary for guidance, but it does not yet rank that value, define its bearer, or generate a rule.

Eight families at the reunion

Values travel in families. Members of the same family may quarrel, and one value may have relatives in several families. The point of the classification is not to produce a perfect taxonomy. It is to reveal the different questions that moral theories ask.

Welfare and outcomes

This family asks what makes lives and states of affairs go better or worse. Its members include pleasure, absence of suffering, health, preference satisfaction, knowledge, meaningful activity, friendship, achievement, security, and access to worthwhile options.

Welfare values direct attention to consequences. Who is helped? Who is harmed? How intensely, for how long, with what probability, and with what alternatives?

The family immediately develops an internal dispute. Is welfare one thing, such as happiness, or a plurality of goods? Does satisfying a desire benefit someone when the desire is misinformed, implanted, or self-destructive? Can a life go better in a way the person never experiences? We will return to these questions. For now, “welfare matters” is an opening statement, not a completed account.

Rules and duties

This family asks what must or must not be done, sometimes regardless of whether violating the rule would improve the immediate outcome. Its members include promise-keeping, nondeception, nonaggression, fidelity, role duties, and prohibitions on using people merely as instruments.

Rule-centered values capture the thought that some actions have a structure that matters. Secretly framing an innocent person may remain objectionable even if it prevents a riot. The objection is not exhausted by calculating the net emotional weather.

Yet rules need scope and exceptions. “Never lie” sounds crisp until the murderer asks where the intended victim is hiding. The crispness was partly produced by leaving the emergency outside the sentence.

Rights and autonomy

This family protects domains in which individuals may choose, refuse, speak, associate, control their bodies, retain privacy, and pursue their own conceptions of a life.

Rights often function as claims against others or as shields against collective goals. Autonomy values the authorship of a life, not merely the pleasantness of the resulting story.

But rights can collide. One person’s speech affects another’s security or reputation. Property control can conflict with survival. Privacy can conceal danger. A right may be

absolute, defeasible, threshold-protected, or one important consideration among others. The word alone does not tell us which.

Care and relationship

This family begins from dependence, attachment, responsiveness, trust, and the particular people whose lives are entangled with ours.

Care resists a picture of agents as sealed units who occasionally sign moral contracts in a well-lit lobby. Human beings arrive dependent, remain vulnerable, and develop identities within relationships. A parent's responsibility to a child cannot be fully described as a general duty owed equally to every child.

The danger is favoritism disguised as warmth. Care can illuminate neglected obligations, but "I care about my own" can also excuse indifference to outsiders. The system must distinguish relationship-generated responsibility from tribal exemption.

Character and integrity

This family asks what kind of person one is becoming and whether one's actions express courage, honesty, generosity, practical wisdom, humility, or integrity.

Character matters because rules cannot anticipate every case. A trustworthy person perceives features a checklist may omit and responds with habits formed across time. Integrity also captures the unity of a life: whether action fits one's deepest commitments without self-deception.

But character language can become circular. The courageous act is what a courageous person would do; the courageous person is the one who performs courageous acts. To guide outsiders, the system must say what courage perceives correctly and how it differs from recklessness.

Fairness and procedure

This family asks how benefits, burdens, opportunities, risks, and authority should be distributed, and how decisions should be made when people disagree.

Its members include equality, impartiality, reciprocity, priority for the worse off, desert, need, nondomination, due process, representation, and legitimate procedure.

Fairness matters not only because fair systems may produce good outcomes, but because people can have claims to a say, an explanation, or equal status. Still, fair procedures can produce bad outcomes, and rival principles of fairness can recommend different distributions. A lottery, a vote, a needs test, and an equal split are all "fair" under some description.

Loyalty and belonging

This family values solidarity, shared identity, mutual aid, remembrance, and sustained membership in groups and places.

Belonging can make projects possible that isolated agents could never sustain. Loyalty can be admirable when it keeps faith during difficulty. It can be vicious when it protects insiders from accountability or treats outsiders as morally dimmer objects.

The difference cannot be that good loyalty supports good groups. That is true but unhelpful. We need standards for which demands of belonging are legitimate and which are recruitment posters for abuse.

Authority and tradition

This family values stable coordination, legitimate roles, inherited practices, accumulated knowledge, continuity, and respect for institutions that make common life possible.

Traditions can contain discoveries no individual designed. Authorities can coordinate emergencies, preserve expertise, and settle disputes. A sailor who improvises a private collision-avoidance system is not displaying admirable independence.

But age does not convert a practice into a justification, and power does not convert an order into a duty. Authority needs a source, scope, and limit. Tradition needs an account of why its inheritance deserves continuation. “We have always done it this way” may be evidence of hard-won wisdom or of a mistake with excellent attendance.

These eight families do not correspond neatly to eight competing theories. A welfare-centered theory may defend rules because they reliably improve lives. A rights theory may value welfare within protected domains. A care ethic may require fair institutions. Real systems are often plural.

Plurality is not yet incoherence. A toolbox is not contradictory because it contains both a saw and a wrench. Incoherence appears when the system cannot say which object each instrument fits, when it changes instruments to secure a favored answer, or when two instruments demand incompatible operations and there is no conflict rule.

Five ways to matter

The phrase “X is valuable” can conceal different relations. Mark each admitted value with one or more of these roles.

Intrinsic value

Something has **intrinsic value** when it matters in itself, not solely because it produces something else. Pleasure, knowledge, friendship, beauty, agency, or conscious life are common candidates.

Calling a value intrinsic ends one explanatory route: we need not show that it is useful as a means to a further value. It does not end all explanation. We can still ask why it has that status, for whom it is valuable, whether more is always better, and how it compares with other intrinsic values.

Instrumental value

Something has **instrumental value** when it matters as a means. Money buys food. Trust reduces the cost of monitoring. Education expands options. A legal procedure may prevent arbitrary power.

Instrumental value depends on circumstances. A key is useful only relative to a lock. If an institution is valued solely because it produces welfare, then a demonstrably better welfare-producing alternative may replace it without loss—unless the institution also has another kind of value.

Constitutive value

Something has **constitutive value** when it is not merely a means to a good but part of what that good consists in. Honest engagement may be part of friendship, not just a technique for obtaining it. Self-direction may partly constitute a flourishing life rather than merely causing pleasant experiences.

This distinction matters in replacement cases. A machine might generate the pleasant sensations associated with achievement without the achievement. If achievement is constitutive of the good, the substitute leaves something out.

Relational value

Something has **relational value** because of a particular connection: *this* promise, *my* child, *our* shared history, the object given by a dead friend.

Relational value need not imply that insiders possess greater basic standing. It can generate agent-relative responsibilities: reasons for this person that are not identically reasons for everyone. Whether such reasons are objectively authoritative remains a later question.

Symbolic value

Something has **symbolic value** because it represents, expresses, or embodies another meaning. A memorial, flag, wedding ring, apology, or preserved voice can matter far beyond its materials.

Symbolic value is real in the ordinary sense that symbols affect memory, identity, relationships, and public meaning. But symbolic claims are easily inflated. To say that an act “disrespects dignity” should lead to an account of what it expresses, to whom, under which convention, and with what consequences. Otherwise symbolism becomes a fog machine rolled onto the stage whenever the argument looks sparse.

One item can occupy several roles. Trust may be intrinsically valuable as a mode of relationship, instrumentally valuable for cooperation, constitutive of friendship, relationally valuable between particular people, and symbolically expressed through confidentiality. This richness is not a problem. The problem is invoking one role and collecting the benefits of all five without saying so.

Six exhibits from the value hall

The delegates now tour a set of cases. In each, resist the urge to announce a verdict immediately. First identify the values and their roles.

The emergency sacrifice

A floodgate can be closed only by a worker who will almost certainly die in the mechanism. Closing it will save a district. The worker volunteers, fully informed.

Welfare favors saving many lives. Autonomy favors respecting the worker's choice. A rule against intentionally sacrificing persons may oppose using the worker as a means. Character may honor courage. Institutional fairness asks whether risky sacrifice is expected disproportionately from workers with few options.

The case is not solved by listing these values. But the list reveals why “save the greater number” may be incomplete and why “never sacrifice” may overlook agency.

The borrowed voice

After singer Anja Rell dies, a company trains a model on private recordings and releases new songs in her exact voice. The music gives comfort to millions. Her will prohibited posthumous imitation.

There is no current conscious subject who can hear the violation. Yet former autonomy, promise, property, family relationship, artistic integrity, cultural benefit, institutional trust, and symbolic meaning all appear. If your system protects the directive, state whether the value lies in the former agent's continuing claim, present relationships, social practice, or something else.

The friend's answer sheet

Your closest friend cheats on a licensing examination. Reporting them will probably end the friendship and career. Silence exposes future clients to uncertain risk.

Loyalty, care, honesty, public safety, procedural fairness, professional integrity, and mercy compete. Loyalty may require confronting the friend rather than concealing the act. “Friendship matters” does not specify whether friendship demands rescue from consequences, protection of character, or refusal to participate in deception.

Pain and extinction

An island's invasive predators cause intense suffering to thousands of native animals and will eventually drive one species extinct. A control program would preserve the species but cause short-term suffering to individual predators.

Individual welfare, species preservation, ecological integrity, human responsibility for introduction, and uncertainty all matter. If species have only indirect value, explain the route. If a species has direct value, explain what bearer persists across changing members.

The painless edit

A safe procedure can erase a survivor's traumatic memory. With the memory gone, distress will cease—but so will several relationships and commitments built around what happened. The survivor chooses the procedure during a severe crisis.

Welfare supports relief. Autonomy supports choice but raises questions about impaired decision-making and the future self. Identity, truthfulness to one's history, relationships, and integrity may oppose erasure. Are these values intrinsic, constitutive of a life, or paternalistic projections by observers?

The conscious replica

An exact conscious replica of Captain Osei awakens with Osei's memories and personality. The original also survives. The ship has resources to preserve only one body.

If psychological continuity grounds identity, both have a claim. If numerical identity matters, at most one is the original. Welfare, equality, prior commitment, causal history, expectation, and the wrongness of killing may pull apart. "They are identical" is false in one important sense: there are now two centers of experience. "One is merely a copy" may be descriptively true about causal order while settling nothing about standing.

In each exhibit, values identify morally salient features. They do not yet dictate a verdict. That is the next construction problem.

Primary, protected, and defeasible

Your Candidate System now needs a preliminary hierarchy of roles.

A **primary value** is one the system treats as a central source of evaluation. It helps define what makes outcomes, actions, institutions, or lives better or worse.

A **protected value** is insulated from ordinary tradeoffs. It may function as a right, constraint, floor, presumption, or boundary that cannot be crossed merely for small gains elsewhere.

A **defeasible value** genuinely counts but can be outweighed or overridden under stated conditions.

These categories can overlap. Welfare may be primary, bodily autonomy protected, and promise-keeping defeasible. A protected value may admit emergency exceptions. A primary value need not always win.

Avoid labeling every value protected. If two absolute protections conflict, the system needs a rule for impossible compliance. A constitution in which every clause is supreme is a polite description of civil war.

Purpose. Convert attractive value names into explicit roles, grounds, and bearers; completing the table identifies the system's materials but does not yet generate or rank its rules.

Create a value table:

Value	Family	Why admitted	Intrinsic, instrumental, constitutive, relational, or symbolic	Primary, protected, or defeasible	Bearer or beneficiary

For every value, complete this sentence:

This matters because _____, and not merely because I approve of it.

You may not yet have a successful completion. Write the best candidate. “It is self-evident” is allowed only if you later explain what self-evidence establishes and how error is detected. “Humans generally care about it” is valuable evidence about shared attitudes and coordination, but it does not by itself satisfy the objectivity threshold from Chapter 2.

Now perform three checks.

The duplication check: Are two values different names for the same work? “Human dignity” sometimes bundles autonomy, equal status, welfare, and a prohibition on humiliation. Unbundle it before relying on it.

The halo check: Is the value admitted because the word sounds good? Replace “justice” with its proposed rule. If disagreement instantly reappears, the label had been lending an illusion of agreement.

The orphan check: Does a value appear in cases but nowhere in the table? Hidden values often enter as “obviously,” “merely,” “natural,” or “disproportionate.” Give each orphan a row.

The eighth debt

Add this entry:

Debt 8 — The value basis: Which values enter, in what sense do they matter, and what establishes their role beyond the system-builder’s attraction to their names?

The Candidate System has acquired materials: welfare, rights, care, character, fairness, loyalty, authority, or whatever survived your audit. Plurality has not doomed it. Nor has breadth completed it.

Values are more like directions on a compass than instructions shouted by a commander. “Trust matters” does not tell us whether to ban lies, encourage candor, protect confidences, forgive one deception, or design institutions in which trust is less fragile. “Welfare matters” does not tell us whose welfare, across what time, under what risk, aggregated by which rule.

The delegates have stocked the kitchen. They still have no recipes.

Next, we turn “this matters” into “do this”—and watch the hidden selectors come out of the cupboards.

Chapter 9 — Turn Values into Guidance

Port Meridian’s first value-guided ordinance was admirably brief:

TRUST SHALL BE PROMOTED.

The next morning, three officials claimed to be enforcing it.

The health minister proposed mandatory disclosure of contagious illnesses.

The privacy commissioner prohibited disclosure of medical records.

The mayor’s aide suppressed a report about unsafe water because, he explained, public confidence in government was essential.

All three had noticed that trust matters. Only two had noticed that this does not tell us what to do.

A value is not a rule. It marks something as important. To guide action, a moral system must specify how that importance generates requirements, permissions, protections, virtues, priorities, or procedures.

Think of a value as a destination marked on a map. “We value the coast” may tell us where we hope to arrive. It does not say whether to walk, build a road, avoid the cliffs, travel together, or stop when the bridge is out. A **generation rule** supplies the route from value to guidance.

Every generation rule also needs an **exception rule**. The exception may be “none.” More often it states when ordinary guidance yields to emergency, uncertainty, competing claims, excessive cost, incapacity, or another protected value.

If exceptions are not written in advance, they tend to appear precisely when the rule inconveniences its author. This is one of morality’s less charming seasonal migrations.

Where this chapter is headed. We will examine the rules that convert a value into an instruction, permission, priority, or exception. The conversion matters because *care about this* does not contain a complete answer to *do what, for whom, at what cost?*

Ten routes from importance to action

The delegates place ten route cards on the table. A value may use more than one route, but each route produces a different structure.

Maximize

To **maximize** a value is to choose the available option expected to produce the greatest total amount of it, perhaps after subtracting losses.

If welfare is maximized, each choice is assessed by expected effects on welfare. The rule needs a unit, a population, a time horizon, a treatment of probability, and a method for aggregation. Does a small benefit to many outweigh a grave loss to one? Are possible future people included? Can welfare of different kinds be measured on one scale?

Maximizing is powerful because it makes tradeoffs explicit. It is dangerous when the appearance of arithmetic conceals disputed units and weights. A number written to two decimal places can still be a philosophical shrug wearing spectacles.

Constrain

To treat a value as a **constraint** is to rule out certain acts even when violating the constraint would advance other goals.

A prohibition on torture may constrain pursuit of public safety. A ban on punishing the innocent may constrain efforts to calm a crowd.

Constraints need an object—what exactly is prohibited?—and a scope. Does the rule forbid intending harm, causing harm, using someone as a means, or crossing a protected boundary? Is it absolute? If not, what activates an exception?

Protect

To **protect** a value is to secure a domain, minimum, or claim against interference. Rights commonly take this form.

Privacy might be protected through control over personal information. Bodily autonomy might be protected through consent requirements. Protection can include positive provision as well as noninterference if a right is meaningless without access.

The system must identify the right-holder, duty-bearer, content, strength, and remedy. “There is a right to health” could mean noninterference with care, access to a basic minimum, equal opportunity, emergency treatment, or a guarantee of being healthy. Biology has not agreed to the last provision.

Embody

To **embody** a value is to act from and cultivate a stable disposition that realizes it well across cases. Courage, honesty, generosity, and practical wisdom are generated through habits and character.

Embodiment handles nuance better than a rigid decision tree. It also risks giving the learner no usable direction. A virtue account needs examples, standards of perception, characteristic errors, and a way to distinguish the virtue from nearby vices. Courage lies somewhere between cowardice and recklessness, but “somewhere” is not yet a forwarding address.

Care for

To **care for** a person or relationship is to attend and respond to concrete needs, vulnerability, history, and dependence.

Care guidance may be particular: stay with this frightened child, listen before applying a rule, preserve a relationship rather than maximizing an abstract total. It asks not only “What is due?” but “What response would sustain the person and relation here?”

Care needs limits. When does caring become domination? How are unseen or distant people represented? May devotion to one dependent consume resources owed to others? Particularity can correct abstraction without receiving a license for partiality without end.

Honor

To **honor** a value is to acknowledge its significance through conduct, expression, remembrance, or restraint.

We honor the dead by keeping certain promises, honor achievement through recognition, honor a victim by telling the truth, honor autonomy by asking consent. Symbolic practices often travel this route.

Honor can matter even when it changes no conscious experience, but the rule needs an account of appropriate expression. Who controls the symbol? What convention gives it meaning? When does honor for one history impose silence on another?

Distribute fairly

To **distribute** a value is to allocate goods, risks, opportunities, or burdens according to a principle such as equality, need, desert, contribution, priority, sufficiency, or lottery.

Distribution begins after standing but before outcome. Two people can have equal standing while receiving unequal resources because one has greater need. Equal concern is not always equal shares.

A distribution rule needs a currency—welfare, resources, opportunity, capabilities, status, risk—and a population. It also needs a principle for multiple claims. “To each what they deserve” is not operational until desert has been defined without reproducing existing advantage.

Follow legitimate authority

To generate guidance through **authority** is to treat directives from a role, institution, text, deity, or practice as action-guiding within its jurisdiction.

Authority can settle coordination problems. Ships need one collision protocol; hospitals need standards; courts need final decisions. But an instruction is not morally binding merely because it was issued successfully. The system must explain legitimacy, jurisdiction, interpretation, review, exit, and limits.

“I was following orders” identifies a causal and institutional route. Whether it supplies a justification depends on the authority’s credentials and the order’s content.

Negotiate

To **negotiate** is to reach terms among agents with different aims through bargaining, compromise, exchange, or mutual adjustment.

Negotiation can generate practical commitments where no shared ranking exists. You prefer silence after ten; I prefer music until midnight; we agree on headphones after eleven.

The agreement can solve our problem without revealing the objectively correct decibel level.

Power matters. Consent under desperation, deception, monopoly, or costly exit may not legitimate the result. A negotiation rule must specify background protections and what cannot be traded.

Construct through a procedure

To **construct** guidance is to make a rule or outcome valid through a specified procedure: voting, deliberation, impartial choice, public justification, lottery, judicial review, or another decision form.

Procedure can convert plurality into coordinated action. It can also embody equality or reciprocity. But no procedure selects itself. Who participates? Which options are admissible? What counts as a majority? Are basic rights outside the vote? How is misinformation handled? We will examine procedural authority more closely in Part III.

For now, remember the central distinction:

A method can settle what we will do without proving that the selected outcome is objectively best.

Settlement is a major human achievement. It should not be made smaller by being advertised as metaphysical discovery.

The many rules hidden inside trust

Return to Port Meridian's ordinance. "Trust matters" could generate at least six kinds of guidance.

A ban

Rule: Never deceive.

This treats honest dealing as a constraint. It is clear but severe. Does it prohibit silence? Surprise parties? Misleading an attacker? Placebo treatment? Every case forces specification.

A presumption

Rule: Do not deceive unless preventing grave and imminent harm requires it and no less deceptive option is available.

This protects trust while permitting a narrow exception. Now *grave*, *imminent*, *requires*, and *less deceptive* need thresholds and evidence rules.

A virtue

Rule: Develop the disposition to communicate truthfully, neither manipulating others nor using brutal candor as an excuse for cruelty.

This adds judgment about timing, context, and relationship. It is richer than a ban and less mechanically predictable.

A relationship-specific demand

Rule: Disclose information your partner reasonably relies on in making shared life decisions.

This grounds trust in dependence and joint agency. It may impose stronger duties in close relationships than among strangers. The system must say why the relationship creates this asymmetry and where privacy remains.

A social policy

Rule: Institutions must publish material risks and may protect confidential personal data under publicly stated limits.

This distinguishes trustworthiness from maximum disclosure. A trustworthy hospital keeps some secrets and reveals others. Institutional design assigns who knows what, when, and under which oversight.

A consequential reason

Rule: Choose the communication policy expected to produce the best long-run pattern of reliable cooperation and welfare.

This permits deception when it improves the expected total. It must address cases where sacrificing one person's trust benefits many others and where short-term deception corrodes a practice over time.

All six rules take trust seriously. They can recommend opposite actions. The value did not contain its own generation rule in compressed form. Someone chose the route.

The selectors hiding in the route

Every route contains **selectors**: variables that determine which features control the result. At minimum, inspect these seven.

The agent selector

Who is guided? Everyone, officials, parents, professionals, members, citizens, or the person who made a commitment?

“Care for children” means something different for a passing stranger, a parent, and a child-protection agency.

The beneficiary selector

Whose value counts? The agent's, all present persons', all sentient beings', members', future people's, or a specified claimant's?

This selector must match the standing map from Chapter 6 or explicitly revise it.

The time selector

Which period matters? Immediate experience, a whole life, several generations, or indefinite future effects?

Short and long horizons can reverse a verdict. Concealing bad news may reduce panic today and destroy trust next year.

The risk selector

Are actions guided by actual outcomes, expected outcomes, reasonable estimates, worst cases, or precaution under uncertainty?

An agent cannot choose by actual outcomes when those outcomes are unknowable at decision time. A retrospective system can still evaluate actual results, but it should not confuse that with prospective guidance.

The threshold selector

How much harm, probability, cost, or urgency activates a rule or exception?

“Except in emergencies” is not a threshold. It is a sign pointing toward one.

The priority selector

When values conflict, which has precedence, under what conditions, and with what strength?

This question will occupy the next chapter. For now, mark every place where your generation rule depends on an unresolved comparison.

The evidence selector

What information must the agent have, what inquiry is required, and whose testimony counts?

A safety rule may depend on “credible risk.” Without an evidence standard, officials can define credibility to include convenient fears and exclude inconvenient witnesses.

From one value to a full rule

Purpose. Expose every selector needed to move from an admitted value to action-guidance; visible blanks are unresolved decisions, not permission to let a vague word choose silently.

Choose one row from your value table. Complete this template:

Value: _____.

Bearer or beneficiary: _____.

Generation route: maximize / constrain / protect / embody / care for / honor / distribute / obey legitimate authority / negotiate / construct / other: _____.

Guidance rule: An agent in role _____ must/may/must not _____ when _____.

Evidence standard: The rule applies when _____ is known or reasonably supported.

Scope: It applies to _____ across the time horizon _____.

Ordinary exception: It yields when _____.

Emergency exception: It yields when _____ reaches threshold _____ and no _____ alternative exists.

Conflict rule: If it collides with _____, resolve the conflict by _____.

Review rule: A decision can be challenged or corrected through _____.

If several blanks remain, leave them visible. A blank is a problem you can work on. A vague word posing as a solution is a problem that has stolen a uniform.

The shelter disclosure case

A shelter hides people fleeing a violent militia. The militia's officer asks the director whether anyone is inside. The director can lie, refuse to answer, or tell the truth. Refusal will be treated as confirmation.

Run the case through several systems:

- A strict nondeception constraint forbids the lie.
- A harm-minimizing rule requires it.
- A rights account may treat the officer's demand as lacking any claim to truth and the residents as holding claims to security.
- A virtue account may say honest character does not require assisting an unjust attack.
- A relationship account may emphasize the director's entrusted responsibility.
- An authority account must decide whether the militia has legitimate jurisdiction.
- A procedural account may note that no fair public rule would license attackers to compel information from protectors.

The point is not that all routes secretly agree. Some will not. The point is that a conclusion becomes intelligible only when the route is stated. "Trust matters" cannot decide whether the officer is owed disclosure, whether the residents are owed protection, or whether institutions should preserve a practice of truth-telling by denying legitimacy to coercive questions.

Now vary one feature at a time:

- The officer seeks a contagious patient during a genuine epidemic.
- The residents face arrest under a fair warrant for violent crimes.
- The director is uncertain who is telling the truth.
- The lie will redirect the militia toward another occupied shelter.
- Telling the truth will save five hostages elsewhere.

If your verdict changes, identify the selector that changed it. If the verdict does not change, identify the rule that makes the varied feature irrelevant.

Generation audit

Purpose. Test the full set of guidance rules for route-switching, convenient exceptions, impossible demands, hidden selectors, and borrowed authority; the audit checks architecture rather than proving the values objective.

Apply the template to every primary and protected value in your table. Then inspect the set as a whole.

The route-switch audit: Does the system maximize a value when totals favor its preferred result but treat it as a protected constraint when maximization becomes uncomfortable?

The exception audit: Are exceptions defined by general features that could be recognized before knowing the identities involved, or do they read “except when we really need the opposite answer” in formal clothing?

The agency audit: Does the rule ask finite agents to calculate information they cannot obtain or consequences they cannot estimate?

The missing-selector audit: Circle words such as *serious*, *reasonable*, *proportionate*, *legitimate*, *adequate*, and *necessary*. Each may be appropriate. Each must eventually receive a standard.

The authority audit: Does the fact that a value matters establish that a rejecting agent is obligated to follow the generated rule? Not yet. We are constructing the system’s guidance. Its authority remains a promissory note for Part III.

The ninth debt

Add this entry:

***Debt 9 — The generation rule:** How does each admitted value produce guidance, which selectors control that transition, and have thresholds, exceptions, agents, time horizons, risks, or priorities been chosen merely to obtain the desired verdict?*

The Candidate System can now do more than admire its values. It can generate instructions. But the final blank in nearly every template is the same: *What happens when this value conflicts with that one?*

Perhaps both can be measured on a common scale. Perhaps one has priority. Perhaps a protected boundary holds until disaster. Perhaps no objective ranking exists, and a procedure or commitment must settle action without pretending to discover one.

The delegates have reached a canyon with several bridges. Each bridge charges a toll. Some accept numbers, some rights, some judgment, and some political agreement.

None accepts the phrase “on balance” without asking to see the balance.

Chapter 10 — Choosing Without a Common Measure

The Port Meridian delegates entered the comparison chamber carrying values, rules, tables, and a confident quantity of stationery.

At the center stood a brass balance.

One pan was labeled **LIBERTY**. The other was labeled **WELFARE**.

“Excellent,” said Councilor Pell. “We weigh them.”

“In what units?” asked Iona.

Vale placed a wooden block marked *one liberty* on the left pan.

This did not help.

Moral language is full of scale metaphors. We weigh reasons. We balance rights. We find that one concern outweighs another. We speak of proportion, importance, sacrifice, and greater good. The metaphors are useful because choice really does involve comparison. They become misleading when they suggest that every value arrives with a common unit printed underneath.

Sometimes comparison is easy. If one treatment relieves more pain, costs less, has fewer side effects, and burdens no one else, it dominates the alternative. No exotic philosophy is required.

Hard cases arise when one option protects something important by setting back something important of another kind: privacy against safety, promise against rescue, equality against total benefit, local self-rule against animal welfare, one grave loss against many mild losses.

We must distinguish three achievements:

1. **Comparing values:** establishing that one option has stronger support than another under a standard.
2. **Settling action:** deciding what will be done, perhaps by vote, lottery, negotiation, authority, or personal commitment.
3. **Proving objective rank:** establishing that one value or option is better independently of the relevant agents’ actual attitudes.

A procedure can settle action without completing the comparison. A commitment can guide a person without ranking values for everyone. An objective ranking, if available, can compare without telling a divided society which legitimate procedure may enforce it.

Confusing these achievements makes practical success look metaphysically larger than it is—and makes philosophical uncertainty look practically fatal when it is not.

Where this chapter is headed. We will test several ways of settling conflicts among values that lack an obvious common unit. A method can guide a choice without proving that one value is objectively greater, so practical settlement and metaphysical ranking must remain separate.

Eight bridges across the canyon

There is no single method called “balancing.” There are several comparison or settlement methods, each with its own structure and debt.

Bridge 1: Put everything on a common scale

The **common-scale** method assigns every relevant gain and loss a value in one currency, then chooses the option with the greatest total, average, expected value, or another aggregate.

Money sometimes works for market choices, but moral costs resist being priced solely by willingness to pay. A wealthy person can pay more to avoid a minor inconvenience than a poor person can pay to avoid serious harm. The resulting price measures resources and preference, not simply moral importance.

Welfare is a more plausible currency. Yet it raises questions:

- What constitutes welfare?
- Can one person’s welfare be compared with another’s?
- Do intensity and duration multiply?
- How should probability enter?
- Do many tiny benefits sum to outweigh one catastrophic loss?
- Does distribution matter apart from the total?

A common scale is attractive because it yields decisions. It may also force unlike values into numbers whose precision comes from the person typing them.

Numbers can still be useful under uncertainty. A rough scale can expose assumptions and improve consistency. The warning is not “never quantify.” It is “do not confuse assigning a number with discovering a unit.”

Bridge 2: Give extra priority to the worse off

The **priority** method says that benefiting people matters more the worse off they are. A benefit of the same size receives greater moral weight when it goes to someone in a worse condition.

This captures an intuition that totals miss. Giving a mobility device to someone unable to leave home may matter more than giving a convenience upgrade to someone already thriving, even if both report equal increases on a satisfaction scale.

Priority needs a metric of “worse off,” a strength function, and a time frame. Is the worse-off person the one with the lowest current welfare, lifetime welfare, resources, op-

portunities, capabilities, or social status? How much extra priority? Can enough small benefits to the better off eventually outweigh a large benefit to the worse off?

Without those answers, “prioritize the vulnerable” identifies a direction but not a comparison protocol.

Bridge 3: Protect a boundary

The **protected-boundary** method prohibits certain tradeoffs. Do not torture. Do not enslave. Do not convict a person known to be innocent. Do not use an unwilling person’s organs as spare parts.

Boundaries express the thought that some claims are not entries in a public pool to be exchanged for ordinary gains. They preserve separateness, status, and security against aggregation.

The method needs a definition of the protected act or domain. Does a ban on killing apply equally to intending death, foreseeing death, withdrawing aid, self-defense, and allocating scarce medicine? If not, which feature explains the difference?

It also needs rules for collisions. Two people can have rights that cannot both be honored. If every boundary is absolute, emergencies can make the system issue incompatible commands.

Bridge 4: Protect the boundary until a threshold

A **threshold** rule holds a constraint firmly under ordinary conditions but permits or requires crossing it to prevent catastrophe.

This fits much actual judgment. Privacy is protected, but an imminent credible threat may justify limited disclosure. Promises bind, but not necessarily when keeping one would cause mass death.

The threshold approach inherits both strengths and debts. Where is the threshold? Is it based on magnitude, probability, imminence, number affected, or the nature of the threatened value? Does the rule switch suddenly from “never” to “maximize,” or does the boundary weaken gradually?

If the threshold is placed exactly one centimeter below the case whose outcome we already prefer, the bridge has been built from the destination backward.

Bridge 5: Give one value lexical or discontinuous priority

A value has **lexical priority** when it outranks another regardless of ordinary quantities, much as any word beginning with A comes before every word beginning with B in an alphabetical sort. A weaker **discontinuous priority** may allow tradeoffs only after a sufficient minimum is secured.

Basic liberty might take priority over economic convenience. Survival might take priority over entertainment. Equal civic status might be secured before maximizing aggregate prosperity.

Lexical rankings avoid some troubling aggregation. They also need justification and internal rules. If two claims occupy the same rank, how are they compared? Can a microscopic gain at the higher rank defeat any quantity below? Does the rank apply to every instance or only to a protected minimum?

An ordered list is a decision structure. The order itself remains a philosophical claim.

Bridge 6: Use disciplined plural judgment

The **plural-judgment** method denies that every case can be reduced to one currency or rigid hierarchy. It considers multiple features—severity, relationship, responsibility, rights, distribution, character, uncertainty—and reaches a contextual judgment.

This resembles skilled diagnosis. A physician does not always convert temperature, pain, mobility, patient preference, and long-term risk into one number. Expertise consists partly in perceiving how features interact.

Plural judgment can be rational without being algorithmic. But “use judgment” can also mean “trust my verdict.” Discipline requires precedent, articulated reasons, adversarial review, consistency checks, identification of decisive features, and willingness to revise.

The method should say what makes one judgment better than another and how disagreement among equally informed judges is handled. Otherwise context becomes a trap-door through which favorite conclusions escape general rules.

Bridge 7: Use a fair procedure

A **procedural** method settles conflict by vote, lottery, representation, deliberation, rotation, judicial process, or another rule regarded as fair.

Suppose two patients have equal need and expected benefit but only one dose exists. A lottery may be appropriate precisely because no morally relevant feature ranks them. The lottery does not discover that the winner’s life is more valuable. It gives each an equal chance and makes allocation possible.

Procedures are indispensable when reasonable people disagree or claims are tied. But fairness is not a self-validating label. Why this participant set? Why majority rule rather than unanimity? Which rights limit the vote? How is unequal influence controlled? What outcomes are too abusive to authorize procedurally?

A fair process can lend legitimacy to settlement. Whether it creates an objective moral ranking is another question.

Bridge 8: Choose or negotiate a settlement

Sometimes agents acknowledge that no shared ranking has been established and adopt a **chosen settlement**. A person commits to family over career. A community compromises between preservation and development. Parties bargain over risk.

This is not arbitrary merely because it includes choice. Choices can be informed, coherent, reciprocal, compassionate, stable, and publicly defensible. They can express identity and make cooperation possible.

But chosen settlement must disclose what it is. “We decided” does not mean “we discovered the objectively correct answer.” It also does not mean the decision is immune to criticism. We can ask whether affected parties were excluded, alternatives hidden, power unequal, evidence distorted, or the result cruel by standards the parties themselves share.

Choice can close a practical gap. It cannot, by its occurrence alone, turn the chosen ranking into a stance-independent fact.

Seven cases at the tollbooth

Now take the Candidate System across each bridge. The aim is not to use all eight methods in every case. It is to learn what each must supply.

A million headaches and one eye

Policy A prevents one million people from experiencing a mild headache for an hour. Policy B prevents one person from permanently losing sight in one eye. The policies cost the same and cannot both be adopted.

A simple aggregate might add the headaches until they exceed the sight loss. A priority view may give extra weight to the person facing serious impairment. A protected-minimum view may secure basic capabilities before relieving minor discomforts. A plural judgment may say that sufficiently many tiny burdens still cannot aggregate across a qualitative boundary.

Now increase the headaches to severe migraines lasting a week. If your verdict changes, locate the threshold or comparison rule. If it never changes regardless of number, state the discontinuous priority you have adopted.

Memory privacy and injury prevention

A city can scan every resident’s private memories once a year. The scan will reliably identify structural safety errors and prevent dozens of serious injuries. Data cannot be retained, but the examination itself is experienced as a profound invasion.

Aggregation favors prevention if injuries and distress are placed on a common welfare scale. A privacy right may constrain scanning. A threshold rule may allow targeted scans under warrants but reject universal examination. A procedural approach might require consent—but can a majority authorize scanning the dissenters?

Notice that saying privacy is “very important” does not tell us whether it is a large quantity on a welfare scale, a boundary, a constitutive condition of agency, or a politically chosen protection. These interpretations can agree in ordinary cases and split here.

The promise and the five

You promised to deliver a time-sensitive medication to one patient. On the way, you learn that diverting the vehicle could deliver a different medication and save five other people. No one else can help.

One method aggregates lives. Another treats relied-upon promises as constraints. Another asks whether the promise includes an emergency exception. A relationship account notes that your voluntary act made one person specially dependent on you. A lottery seems perverse after reliance has already been created, but perhaps fair before assignments are made.

Change five to five hundred. If the promise still wins, you have given it extraordinary or lexical priority. If it yields, identify the threshold. “Obviously five hundred” is an intuition report, not a placement rule.

The unique archive

A community can preserve the world’s only archive of an ancient language or redirect the funds to prevent several thousand cases of temporary illness.

The archive has cultural, historical, relational, and perhaps intrinsic value. The health program has welfare value distributed among many people. What is the unit? Does irreversibility matter? Can cultural loss be compensated? Who speaks for descendants? Would the answer change if the illnesses were fatal, or if the archive served a marginalized group whose history was repeatedly destroyed?

This case is designed to embarrass a moral calculator. Embarrassment is not refutation. The calculator may need an honest entry labeled “assigned cultural-loss weight,” with the assignment defended rather than discovered in the machinery.

One death and forty recoveries

A treatment allocation will either prevent one death or prevent forty cases of severe but temporary illness. All affected people would otherwise return to full health.

Many systems favor preventing death, but why? Greater welfare loss? A protected interest in continued life? Discontinuous priority for survival? The loss of future agency? Now vary age, life expectancy, pain intensity, and probability. Do not assume that one word—*death*—contains all relevant comparisons unchanged.

Two unlike abilities

One child can receive a device that restores speech. Another can receive a device that restores independent mobility. Both regard the lost ability as central. Resources allow only one device.

Self-reports tie. Costs tie. Need is profound in both cases. A common scale may lack credible units. A lottery may be the fairest settlement. It does not imply that speech and mobility are equal in objective value. It implies that the admitted criteria did not rank these claims and equal standing forbids inventing a convenient difference.

The valley decision

A remote valley contains three towns, grazing land, a rare predator, and a forest sacred to an Indigenous community. A dam would provide thousands of jobs and low-carbon

power, flood homes and burial grounds, eliminate the predator's habitat, reduce suffering among livestock, and transfer political control to the capital.

No single label captures this case. It includes welfare, livelihood, property, consent, historical injustice, local self-rule, animal suffering, species preservation, climate risk, symbolic meaning, and distribution across generations.

Begin with standing: who is represented directly or indirectly? Then separate values. Then state the generation and comparison rules. If a public procedure settles the decision, identify the outcome constraints and power conditions. If a negotiated settlement is used, disclose which losses cannot be compensated and who has credible authority to consent.

The case may remain difficult after every distinction. That is not evidence that distinctions failed. A clear map does not flatten a mountain.

Build a comparison protocol

Purpose. State how the Candidate System settles value conflict and what remains unresolved; a usable protocol may coordinate action without discovering an objectively correct common measure.

The Candidate System needs more than an instruction to "balance all relevant factors." Write a protocol with these fields:

Admitted considerations: _____.

Common unit, if any: _____.

Aggregation rule: total / average / expected / sufficient level / none / other: _____.

Priority rule: Extra weight goes to _____ according to _____.

Protected boundaries: _____ may not be crossed for _____.

Thresholds: A boundary yields when _____.

Lexical ranks, if any: _____ precedes _____; ties within a rank are handled by _____.

Contextual judgment discipline: Decision-makers must state _____ and are reviewed through _____.

Procedural fallback: Unranked claims are settled by _____ with participants _____ and limits _____.

Chosen commitments: The system openly adopts _____ without claiming _____.

You may combine methods. Perhaps basic rights establish boundaries, priority governs distribution above the floor, welfare comparison handles ordinary policy, and a lottery settles true ties. A hybrid is not automatically ad hoc. It becomes ad hoc when the method switches without a general explanation.

Run three audits.

The phantom-unit audit: Where does the system speak as if unlike values have been measured without identifying a scale or qualitative rule?

The free-tie-breaker audit: Where does a method stop just before the decision and another conclusion appear without support?

The settlement inflation audit: Where has a vote, agreement, lottery, commitment, or official decision been described as proof of objective superiority?

The tenth debt

Add this entry:

Debt 10 — The comparison rule: What makes one option better supported when values conflict, and has the system smuggled in a unit, priority, threshold, rank, judgment, procedural limit, or chosen commitment without defending it?

The brass balance in Port Meridian remains in the chamber. It now bears a smaller, more honest plaque:

METAPHOR. DO NOT CALIBRATE.

The Candidate System has a comparison protocol. Before it becomes a constitution, however, it must survive one of the oldest moral tricks: applying a rule strictly to strangers and interpreting it generously for friends.

Consistency does not require treating every case identically. Parents may owe things to their children that they do not owe to everyone. Surgeons may cut patients when strangers may not. The question is which differences are genuinely relevant and which are convenient portraits of identity and power.

The final chamber in Part II contains mirrors.

Chapter 11 — Like Cases, Exceptions, and Identity

Port Meridian had two bridge inspectors.

Both had signed safety certificates without completing required tests. Both bridges were later found to be sound. Both inspectors had been exhausted, under pressure to reopen essential routes, and aware that the tests were mandatory. Their evidence, intentions, risks, histories, and family circumstances were as nearly identical as the investigators could determine.

One inspector, Tovan, was a stranger to Iona Venn.

The other, Els, had been Iona's closest friend for twenty years.

Iona chaired the professional review board. She voted to refer Tovan for formal discipline. Before Els's case reached the board, Iona privately warned her, helped her revise the records, and argued that no inquiry was needed because the bridge had proved safe.

When challenged, Iona said, "The cases are different. Els is my friend."

This was true. It was also not yet a justification.

Two cases never share every property. They occur at different times, involve numerically different people, and occupy different places in the universe. Consistency cannot require treating all differences as irrelevant, because some differences plainly matter. A physician may operate on a consenting patient but not on a person seized from the street. A parent may spend an evening caring for their own sick child rather than distributing those hours equally among every sick child in the city. A judge may sentence a convicted offender but not a passerby.

The problem is to distinguish a **relevant difference** from a convenient difference.

Moral systems regularly announce that like cases should be treated alike. The phrase is correct and incomplete. The entire argument is often hiding inside *like*.

Where this chapter is headed. We will ask which similarities are relevant, when an exception is principled, and how identity can distort either judgment. This completes the first working constitution by making its consistency rules inspectable rather than intuitive.

Consistency is not uniformity

Uniformity applies the same treatment everywhere. **Consistency** applies the same rule to cases that share the features made relevant by that rule.

Suppose two patients need different doses because one weighs twice as much. Equal dosage would be uniform but medically inconsistent with the rule “dose according to body weight.” Suppose two voters receive different numbers of ballots because one is the mayor’s cousin. The treatment is nonuniform and lacks an apparent relevance rule.

A morally relevant difference usually enters through one of the system’s established structures:

- a difference in standing or in the interest at stake;
- a difference in need, harm, benefit, or risk;
- a difference in responsible agency, evidence, control, or intention;
- a valid role or commitment;
- a relationship that creates a specific responsibility;
- a history of contribution, reliance, repair, or injustice;
- a difference identified by an independently defended value or rule.

This list is not a machine that produces verdicts. Each category still needs its own justification. It is a guardrail against declaring a feature relevant solely because mentioning it restates the desired result.

“These cases differ because this person deserves protection” is not an explanation if *deserves protection* means “should receive the different treatment now in dispute.” The verdict has put on a fake mustache and returned as a premise.

Roles: permission with a job description

Roles can make differences relevant. A firefighter may enter a burning building under conditions in which an untrained spectator should not. A judge may issue an order that neither neighbor may issue. A parent has responsibilities generated by dependence, relationship, and prior conduct.

But a role is not a moral coupon redeemable for special treatment. To justify a difference, state:

1. how the role was acquired;
2. what function or value it serves;
3. what powers and responsibilities follow;
4. who is affected;
5. what limits and review apply;

6. how one exits or is removed.

An inherited office in an oppressive institution may be socially recognized without being morally legitimate. A professional role may grant power only for specified purposes. The fact that Iona chairs the review board gives her authority to apply public standards, not to convert friendship into a confidential exit ramp.

Role-based differences should survive a **role swap**. If another qualified chair handled Els's case, what rule should that chair apply? If Iona would condemn another chair for secretly protecting a friend, then her own exception is not grounded in the public role.

Relationships: particular without being magical

Relationships can generate real asymmetries. You may owe a friend attention, candor, aid, and loyalty that you do not owe in identical measure to every stranger. These responsibilities can arise from history, trust, mutual vulnerability, promises, shared projects, and the way each person has organized life in reliance on the other.

Notice what does the work. The bare label *friend* does not emit obligations like a small ethical radio. The relationship contains patterns of reliance and commitment that may generate reasons.

It also has limits. Friendship might give Iona reason to support Els emotionally, urge honesty, help obtain counsel, or remain present during discipline. It does not obviously give her reason, in her public role, to falsify equal procedures. Indeed, friendship may demand refusing to help Els deepen the deception.

Particular responsibility and public impartiality can coexist because they govern different actions. Iona may not owe Tovan intimate companionship. She may owe both inspectors the same professional process.

This yields a useful distinction:

Equal standing does not imply identical personal responsibility, and special responsibility does not imply exemption from every identity-neutral rule.

Identity and the danger of labels

Identity features—nationality, ethnicity, religion, sex, class, kinship, party, citizenship, profession, or group membership—can affect treatment in at least four different ways.

Identity as a bare label

A bare label is simply the fact of category membership: “one of us,” “foreign,” “traditional,” “elite,” “believer,” “outsider.” If no admitted value, interest, history, role, or relationship connects the label to the decision, using it is favoritism or prejudice rather than relevant differentiation.

“Citizens receive clean water because they are citizens” cannot justify poisoning visitors. Membership may structure funding or political participation; basic vulnerability may cut across it.

Identity as evidence

An identity feature may correlate with a relevant fact. Age can correlate with medical risk. Professional certification can correlate with competence. Residence can correlate with use of a local service.

Correlation is defeasible. If direct evidence is available, a rough identity proxy may become unjustified. The system should ask how accurate the proxy is, who bears its errors, whether the correlation results from injustice, and whether using it reinforces the pattern.

Identity as a source of relationship or role

Membership can generate commitments and institutional responsibilities. A government may owe citizens participatory rights that it does not administer globally. A language community may owe support to members who sustain a shared practice. A family relation can create duties of care.

These differences are not explained by the sound of the identity word. They are explained by cooperation, dependence, reliance, jurisdiction, or an admitted value of relationship and belonging.

Identity as the location of history

An identity-conscious policy may be necessary to repair a pattern that operated through identity. Suppose a city excluded one ethnic community from property ownership for generations. A restorative program cannot identify the affected history while pretending not to notice the category through which exclusion occurred.

This is not an exception to relevance. The relevant features may be inherited disadvantage, institutional responsibility, continuing effects, equal civic status, or repair. Identity locates the history. The system must still specify the connection and avoid treating every member as interchangeable.

An identity-neutral rule is therefore not always an identity-blind rule. It is a rule whose underlying justification does not change simply because the favored or burdened identity changes. “Repair durable disadvantage created by this institution” can be identity-neutral at the level of principle while identity-conscious in application.

Exceptions: doors, not holes

An **exception** is a stated condition under which a general rule does not apply or yields to another rule. An **ad hoc escape** is a hole cut after someone important finds the door inconvenient.

Good exception rules have several features.

They can be stated before identities are known

“Confidential information may be disclosed when there is credible evidence of imminent serious harm and disclosure is necessary to prevent it” is a candidate exception rule. “Confidential information may be disclosed when I approve of the person asking” is not.

They identify the value that limits the rule

An emergency exception should name the emergency value: life, bodily safety, institutional survival, prevention of severe rights violation. “Unless circumstances require otherwise” identifies nothing.

They specify threshold and evidence

How serious, how likely, how imminent, and established by what inquiry? Precision may be impossible, but disciplined criteria are not.

They are symmetric under relevant swaps

If a rule permits your group to resist occupation, would it permit an otherwise identical rival group to do so? If you may excuse your ally’s misinformation because they were frightened, does the same excuse apply to an opponent with equal evidence and fear?

Symmetry does not require ignoring different histories or aims. It requires showing the difference rather than allowing identity to do the argument silently.

They leave a record

Exceptions should be reviewable. Who invoked one? On what evidence? Was the threshold met? Did the exception become routine? A permanent emergency is often an ordinary policy avoiding its proper name.

The mirror tests

Purpose. Vary identity, power, visibility, relationship, and institutional role while holding admitted features fixed; a changed verdict requires a relevance rule, not merely a strong reaction.

Take one controversial rule from your Candidate System and apply six tests.

1. The identity swap

Exchange the names, groups, loyalties, or social status of the parties while holding need, risk, evidence, history, role, and consequences constant.

Does the verdict change? If so, which admitted relationship or identity-linked history explains the change? If none does, the system has found an asymmetric exception.

2. The power swap

Give the power to enforce the rule to the party you trust least.

Would you still authorize the discretion? If not, the rule may depend on hoped-for rulers rather than institutional safeguards. Constitutions are written for Tuesday's officials too.

3. The publicity test

Could the exception be stated publicly as a general policy to all affected people?

Secrecy does not always make a rule wrong. Confidentiality can be legitimate. But a rule whose justification evaporates when its structure is revealed may be exploiting those kept ignorant.

4. The repetition test

What happens if every similarly placed agent uses the exception whenever its stated conditions occur?

A one-time favor may look harmless while a practice of such favors destroys equal process. Repetition reveals whether the rule can live in an institution rather than only in the imagination of its beneficiary.

5. The residue test

After an exception is justified, does some moral remainder persist—regret, apology, compensation, repair, or a duty to prevent recurrence?

Breaking a promise in an emergency may be justified and still call for explanation. A system that treats justified harm as though no loss occurred may be confusing permission with celebration.

6. The adversary's explanation

State the rule in terms an affected opponent could understand even if they reject its values. This does not require their agreement. It requires enough clarity that they can identify the premise, challenge the evidence, and show inconsistent application.

A rule that cannot be explained outside the slogans of its own faction is not necessarily false. It is dangerously difficult to audit.

Return to Iona's vote

Iona offers four defenses.

"I owe more to my friend."

True in some domains. She may owe Els support, truthful counsel, and continued care. The board role, however, exists to protect public safety and equal professional standards. Iona must show why friendship legitimately alters that procedure rather than merely giving her additional responsibilities outside it.

"Els had good intentions."

So did Tovan. The feature does not distinguish the cases.

"No one was harmed."

The bridges were safe, but both inspectors created unassessed risk and compromised the certification practice. Again, the feature is shared.

"Punishment would hurt Els's family."

Tovan's family faces the same harm. If family impact should mitigate discipline, the rule should apply to both. Perhaps the board needs a less punitive, more restorative response. That conclusion revises the general policy; it does not justify private rescue.

The Candidate System can still conclude that neither inspector should be punished, that both deserve reduced sanctions, or that institutional pressure bears much of the responsibility. What it cannot coherently do is construct one path for Els and another for Tovan while citing only shared features plus friendship in a domain where friendship has not been granted relevance.

Iona's loyalty may explain her action and motivate it. Whether it favors the action is a question we will examine in Part III.

Write the relevance rule

Purpose. State which differences justify different treatment and connect each one to an existing part of the Candidate System; the rule should expose rather than conceal legitimate asymmetry.

Complete this page in the Candidate System notebook:

Identity-neutral principle: Cases receive the same status when they match in _____.

Permitted differences: Different treatment may be based on differences in _____ because _____.

Role-based asymmetries: Agents in role _____ may/must _____ because the role serves _____, within limits _____.

Relationship-based asymmetries: Relationship _____ generates responsibility _____ through _____, but does not permit _____.

Identity-conscious applications: Identity may be considered when it tracks _____, using evidence standard _____ and review _____.

Exception form: Rule _____ yields only when _____ reaches threshold _____, supported by _____.

Now take an insider and outsider case. Apply the same underlying rule. Record every feature you treat as relevant. For each feature, point backward to standing, a value, a generation rule, a role, or a comparison protocol already in the Candidate System.

If the feature has no home, either give it one openly or remove it. Do not let it sleep behind the curtains.

The eleventh debt

Add this entry:

Debt 11 — Relevance and symmetry: What makes a difference relevant, and would the rule, role, relationship, or exception retain its structure if identities and power were swapped while all admitted features remained constant?

The construction floor is now quiet. The circle has inhabitants. The cabinet has verdicts. The values have generation rules. Conflicts have a comparison protocol. Exceptions have doors rather than freshly cut holes.

It is time to assemble the first constitution.

Inspection Report II — The First Constitution

Purpose. Compress the system's domain, statuses, values, rules, comparisons, and exceptions into an inspectable constitution before Part III tests whether any of it has authority.

Do not aim for eloquence. A beautiful constitution can conceal debt in calligraphy. Aim for two pages of compressed structure, plus appendices containing the tables you built.

Article I — The domain

State:

- which beings or entities have direct standing;
- the ground of that standing;
- what matters indirectly and through which route;
- how uncertainty about standing is handled;
- how standing differs from weight and final priority.

Article II — The objects and statuses

List the objects your system evaluates: actions, intentions, outcomes, rules, institutions, agents, relationships, or forms of life.

Define at least eight statuses. Include separate conditions for prohibition, justification, excuse, blame, repair, and punishment if your system uses them.

Article III — The values

Name each primary, protected, and defeasible value. Mark whether its role is intrinsic, instrumental, constitutive, relational, or symbolic. Identify its bearer or beneficiary.

Do not write only “justice,” “dignity,” “freedom,” or another halo word. Include the operative meaning.

Article IV — The guidance rules

For every primary and protected value, state:

- the generation route;
- the guided agent and affected population;
- the time and risk rule;
- the evidence standard;
- the ordinary and emergency exceptions;
- the means of review and correction.

Article V — Comparison and settlement

State the common unit, if any; aggregation rule; priority; protected boundaries; thresholds; lexical ranks; discipline for plural judgment; procedural fallback; and any openly chosen commitments.

Separate claims of objective comparison from practical settlement.

Article VI — Relevance and exceptions

State what makes cases relevantly similar, which roles and relationships justify asymmetric responsibilities, how identity-linked history is handled, and which mirror tests every exception must survive.

Article VII — Amendment and debt

Specify how the constitution may change. Every revision must be labeled:

- **clarification** — the earlier commitment was underspecified;
- **correction** — the system previously made an error;
- **tradeoff** — one success condition or value is surrendered to preserve another;
- **discovery** — new facts alter application without altering the rule;
- **change of project** — the system no longer seeks what the original objectivity threshold required.

Attach the Debt Ledger. Do not erase an entry merely because later clauses make it less visible.

Red-team inspection

Now read the constitution as an adversary—preferably an adversary with a pencil and no investment in your self-esteem.

The omitted-being probe

Find one being or group that produces an uncomfortable result for each standing ground. Has the system excluded it, admitted it, or moved it to indirect relevance? Is that treatment principled?

The standing-weight collapse probe

Find every sentence in which “counts” appears to mean both “enters directly” and “wins.” Separate admission from comparative weight.

The overloaded-language probe

Circle *wrong*, *harmful*, *unjust*, *responsible*, *rights*, *dignity*, and *fair*. Does each occurrence identify an object and a status, or does the word carry several verdicts at once?

The unsupported-transition probe

Underline every movement from fact to value, value to rule, rule to authority, harm to blame, blame to punishment, and procedure to truth. Write the bridge above the line. If there is no bridge, add a debt entry.

The hidden-scale probe

Locate “greater,” “proportionate,” “outweighs,” “on balance,” and “too costly.” What unit, boundary, rank, judgment rule, or negotiated commitment supports the comparison?

The incomplete-tie-breaker probe

For each comparison method, invent a tie or collision. What settles it? Then create a tie within the tie-breaker. Continue until the system reaches a stated procedure, permissible plurality, randomizer, or unresolved case. Do not manufacture false determinacy merely because constitutions look untidy with blanks.

The asymmetric-exception probe

Swap identities, loyalties, and power while holding admitted features fixed. Apply the rule again. If the result changes, identify the relevance rule. If none exists, mark the exception as opportunistic.

The constitution will not pass every probe cleanly. That is expected. Its purpose is to expose the architecture, not to certify moral omniscience.

The first complete Candidate System

You have now constructed something substantial. It may protect persons, reduce suffering, honor commitments, cultivate character, sustain relationships, distribute resources fairly, preserve legitimate institutions, and coordinate disagreement. It may be humane, coherent, usable, and corrigible. Those are not small achievements.

But we have postponed the question printed at the top of the original contract:

What makes this system objectively morally binding?

So far, we have shown how a system can be specified. We have not shown that its standing boundary, values, generation rules, priorities, or exceptions are true independently of relevant attitudes. We have not shown why its demands count in favor for a person who understands them perfectly and rejects their ends. We have not shown whether agreement creates authority, discovers it, or simply coordinates action.

The constitution has content. Content is not yet authority.

In Part III, we follow the word *reason* into the central maze. We will separate what causes an action, what motivates it, what supports it, what a person has committed to, what a fair procedure selects, what society enforces, and what allegedly applies whether anyone cares or not.

The first door bears a familiar sentence:

"You have a reason to do this."

Behind it are at least three different rooms.

PART III



Find the Source of the ``Ought''

The system can issue verdicts. The harder question has arrived: what makes any verdict count in favor for the person who refuses it?

Chapter 12 — Three Different Things Called a Reason

At 4:06 in the morning, an alarm sounded beneath Port Meridian.

Sela Or, a maintenance engineer, was inspecting the old storm tunnels when the eastern oxygen sensor flashed green and the western sensor flashed red. Training signs directed workers toward any green route. Sela ran east.

The eastern sensor had corroded. Its green signal was false. A chemical leak had displaced nearly all breathable air from the eastern tunnel.

The western sensor had also malfunctioned, but in the opposite direction. The western tunnel was safe.

Sela did not know any of this. She believed the east was breathable, believed the west was poisonous, and headed toward danger while the route that could save her lay behind a red light.

When investigators reconstructed the event, they produced three sentences:

1. The corroded sensor was the reason Sela ran east.
2. Sela's belief that the east was safe was her reason for running east.
3. The breathable air in the west was a reason for Sela to run west.

All three sentences are intelligible. They do not use *reason* in the same way.

The first identifies a cause. The second identifies what the agent acted in light of. The third identifies what counted in favor of an option, whether or not the agent knew it.

This distinction is the entrance to Part III. The Candidate System already contains rules, values, and verdicts. It says, for example, that suffering matters, promises normally bind, and institutions should treat relevantly similar cases consistently. We must now ask what makes any of those considerations a **reason** for an agent to act—especially an agent who does not already care.

If we fail to separate the three rooms, we can enter one and announce that we have explored them all.

Where this chapter is headed. We will separate causes of action, considerations that motivate an agent, and facts claimed to count in favor. The distinction is essential because explaining why someone acted does not yet justify the act or establish a reason that applies to someone else.

Room one: what caused the action?

An **explanatory reason** identifies something that helps explain why an event or action occurred. It places the action in a causal story.

The corroded sensor caused a false signal. The signal activated Sela's trained response. Her fatigue may have reduced her inclination to double-check. The tunnel signs, emergency drills, sensor maintenance budget, and chemical leak all belong somewhere in the explanation.

Not every explanatory factor is mental. A locked door can explain why a person turned left. A childhood experience can explain a fear. A financial incentive can explain an official's vote even if the official never says, "I am voting for the money."

Causal explanation matters morally. It helps us prevent repetition, understand capacity and control, evaluate institutions, and distinguish malice from mistake. But a cause does not automatically justify what it causes.

A forged document may cause a false arrest. It does not thereby count in favor of the arrest. A prejudice may explain a hiring decision. It does not gain authority by appearing in the causal diagram.

This seems obvious in ugly cases. The slide becomes harder to detect when the cause is socially approved. "Our moral beliefs are produced by empathy" may be a causal hypothesis. Even if true, it does not yet show that the resulting judgments are correct. "People obey this rule because they were educated into it" explains stability. It does not establish the rule's authority.

Causes tell us how a judgment or action arrived. They do not, without an added premise, certify the arrival.

Room two: what was the agent acting from?

A **motivating reason** is a consideration the agent treats as bearing on what to do. It figures in the agent's point of view—accurately or inaccurately—and helps make the conduct intelligible as action rather than a bodily event.

Sela ran east because she believed the east was safe. We can say that her reason was "the green light" if we mean the green light as she understood it. She was not pushed east by a mechanical arm. She saw what she took to be a route to survival and acted for that purpose.

Motivating reasons can be false. A person may flee because they believe a harmless neighbor is armed. They may donate because they falsely believe an organization provides

medicine. They may betray a friend because they mistakenly believe the friend betrayed them first.

The falsity of the belief does not stop it from explaining the action from the inside. If we omit it, we fail to understand what the agent thought they were doing.

Motivating reasons can also be bad by the agent's own standards. A person may say, "I knew the evidence was weak, but jealousy got hold of me." A craving, image, fear, or anticipated reward can move someone without receiving their reflective endorsement.

This gives us two further distinctions:

- A consideration can **motivate** without being believed to justify.
- A consideration can be believed to justify without actually supporting the action.

Sela's belief motivated her and, from her perspective, justified running east. The actual oxygen level did not cooperate.

Understanding someone's motivating reason can soften blame, reveal deception, or improve communication. It does not require endorsing the reason. "Now I see why you did it" is not the same sentence as "You were right to do it." Many arguments survive by hoping no one notices the change of verb.

Room three: what counts in favor?

A **favoring reason** is a consideration that genuinely supports an option: it counts in favor of doing, choosing, or perhaps believing something.

The breathable air in the west supported going west. The poisonous air in the east counted against going east. These facts bore on Sela's survival whether or not her sensors revealed them.

Calling something a favoring reason makes a normative claim. It says more than "this caused the action" and more than "the agent regarded this as support." It says the consideration actually has support-giving force.

This is precisely where the Candidate System needs authority.

Suppose Sela wants to survive. Then breathable air gives her an obvious **instrumental reason** to choose the west: the route is a means to her aim. The relation can be written:

Sela aims to survive. Going west promotes survival. Therefore, given that aim, there is support for going west.

Even this modest structure raises questions. What if Sela has conflicting aims? What probability is enough? Must she care about effective means? But the route from her aim to the option is visible.

Now consider a different claim:

A stranger is in preventable pain. Therefore Sela has a reason to help, whether or not she cares about the stranger.

The fact about pain is clear. The practical authority is the disputed step. Does pain itself count in favor for every capable agent? Does it do so because Sela already has a concern that information would activate? Because she is committed to a social practice?

Because a fair rule would require aid? Because suffering has objective disvalue? Because rational agency entails regard for others?

The sentence “suffering is a reason” may be correct. It is not self-explanatory. It places the conclusion we need to ground inside the noun *reason*.

A reason is not yet a verdict

Even if a consideration genuinely counts in favor, it may not settle what to do.

Suppose the western tunnel is breathable but contains an injured worker. Going west promotes Sela’s survival and gives her access to the worker. Those considerations favor west. Now suppose opening the western door will release smoke into a nursery above. That fact counts against west. The Candidate System must compare the reasons before assigning an all-things-considered verdict.

Philosophers often call a consideration that genuinely counts in favor a **pro tanto reason**—a reason “to that extent.” The Latin is less important than the modesty. A pro tanto reason contributes support without announcing that all competitors have left the field.

A **decisive reason** is one that defeats the relevant competitors in the case. An act is **required** only when the system’s reason and status rules establish that the agent must perform it. An act can therefore have something in its favor without being required, best, or even permissible after all considerations are included.

This prevents three shortcuts:

- “Pleasure counts in favor” does not imply “maximize pleasure.”
- “A promise supplies a reason” does not imply “never break a promise.”
- “A person’s choice matters” does not imply “every voluntary choice must be honored.”

Generation, comparison, and exception rules from Part II still do work. The language of reasons does not replace them.

The distinction also matters for authority. A system might establish one aim-independent reason—say, that severe suffering always counts against an action—without establishing a complete objective moral system. It would still need to determine strength, aggregation, protected boundaries, agent responsibility, and action status. Finding one real compass direction would not produce a street map.

Add three fields whenever the Candidate System invokes a reason:

Existence: *What consideration counts in favor?*

Weight: *How strongly does it count, and compared by which rule?*

Verdict: *What status follows after conflicts and exceptions?*

Do not allow *a reason*, *the strongest reason*, and *an obligation* to become three costumes for the same unsupported conclusion.

The view from inside and the view from the world

Sela's case produces a useful split between what the evidence available to an agent supports and what the actual facts support.

From Sela's evidence, running east was sensible. From the full facts, running west was the life-preserving option. Philosophers sometimes call these **subjective reasons** and **objective reasons**, but the labels invite confusion with the larger question of objective morality.

We will use more explicit language:

- **Evidence-relative support:** what is supported by information the agent possesses or could reasonably obtain.
- **Fact-relative support:** what is supported by the relevant facts, including facts the agent does not know.

The distinction helps us evaluate different objects from Chapter 7. Sela's choice process may be justified relative to her evidence even though the chosen route is worse relative to the facts. An institution may be responsible for giving her defective evidence. The outcome may be tragic without her being blameworthy.

Fact-relative support cannot guide an agent when the relevant facts are inaccessible. A hidden gas level does not whisper through the wall. It can still matter to whether the option would achieve her aim and to how we should design sensors. Guidance for finite agents needs evidence rules, not merely omniscient hindsight.

This point will become important when the Candidate System claims to contain objective moral truths. A truth that no one can identify may describe what ought to be done from an omniscient perspective while failing the usability and accessibility specifications. Conversely, a careful evidence-relative rule may guide admirably while sometimes selecting the factually worse outcome. One system can contain both levels, but it must not slide between them when challenged.

When the three coincide

Often the rooms connect.

Suppose Sela knows that the west is breathable, wants to survive, and runs west because it is breathable. The fact supports the action, her belief represents the fact accurately, and the consideration motivates her. Explanation, motivation, and support line up.

That alignment feels so natural that we use the same word for all three. Trouble begins when they separate.

The lucky rescuer

Niko believes, without evidence, that a child is trapped in an empty warehouse. He enters to perform a dramatic rescue and happens to find a different child whom nobody knew was there.

The actual child's presence supports entering if rescue is possible. It did not motivate Niko; he knew nothing about that child. His unsupported belief motivated him. The good result does not turn his original belief into evidence.

The reluctant donor

Amal knows that a relief fund will save lives and donates solely to embarrass a rival who gave less.

The expected benefit supports donating under a welfare-oriented rule. Spite motivates Amal. The act may be advisable while the motive is unattractive. If we say "Amal acted for a good reason," we must specify whether we mean a good consideration favored the act or a good consideration moved the agent.

The trained prejudice

An official excludes a visitor because childhood instruction taught that outsiders carry disease. The instruction causally explains the belief. The believed danger motivates exclusion. Reliable medical evidence shows no elevated risk and counts against the action.

Adding detail to the causal story can produce compassion for the official's formation while strengthening the case for correction. Explanation is not acquittal, though it may alter blame.

The unmoved witness

A passerby sees someone collapse beside a shallow pool. Rescue would be safe and easy. The passerby understands every fact but feels no desire to help.

The victim's danger does not motivate the passerby. Does it nevertheless count in favor of rescue for them? That is the central question of Chapter 16. For now, do not answer it by saying, "Of course it is a reason." That is the claim under inspection.

Four common substitutions

Inspect arguments for these moves.

From attraction to support

"Compassion makes us want to help, so suffering gives us a reason to help."

Compassion may make suffering motivationally effective. The conclusion that suffering supplies aim-independent support needs another premise. Alternatively, the reason may be openly conditional on compassionate aims. Both accounts are intelligible; they make different authority claims.

From support to motivation

"They had every reason to comply, so their refusal proves irrationality."

Even if the considerations genuinely favored compliance, they may not have motivated the agent. Perhaps the agent was indifferent, depressed, frightened, or committed elsewhere. Lack of motivation does not by itself show lack of reasons; claimed reasons do not guarantee motivation.

From explanation to justification

“The practice evolved because it promoted cooperation, so it is justified.”
Promoting cooperation may supply a practical consideration if cooperation is valued. Evolutionary persistence by itself shows neither current benefit nor moral authority. Parasites also persist. History is not a letter of recommendation.

From social force to favoring force

“You have a reason to follow the rule because violation brings condemnation.”
Condemnation creates a prudential consideration for anyone who wishes to avoid it. It may motivate through fear. It does not establish that the rule’s content is morally sound or that compliance remains supported when violation is secret. Chapter 17 will keep these relations in separate columns.

The reason ledger

Purpose. Classify each alleged reason by the relation it actually bears—causal, motivating, evidence-relative, fact-relative, or aim-independent—before asking it to justify a verdict.
Take five clauses from your Candidate System. Include one value rule, one constraint, one role duty, one comparison rule, and one exception. For each, complete this table:

Reason claim	Causal role	Motivating role	Evidence-relative support	Fact-relative support	Claimed aim-independent support

Then answer four questions.

1. What exactly is being explained?

Is it the origin of a belief, the occurrence of an action, the persistence of a practice, or the emotional force of a verdict? Causal explanations require different evidence depending on the target.

2. Whose point of view is represented?

State what the agent believes and aims at. Do not replace the agent’s actual motivating reason with the noble reason an observer wishes they had.

3. Which facts support which aim?

If the support is instrumental, name the aim. If it is generated by a commitment or role, name the entry conditions. If it is claimed to apply independently of every aim, mark it for Chapter 16.

4. What remains open?

A correct classification does not prove that the favoring relation exists. It identifies where the proof is needed. Nor does it settle weight, blame, or punishment.

Now apply the table to the tunnel case. Hold these facts fixed:

- the east is poisonous;
- the west is breathable;
- Sela reasonably believes the reverse;
- survival is among Sela's aims;
- the institution negligently failed to maintain both sensors.

Classify the sensor corrosion, the lights, Sela's beliefs, the actual oxygen levels, her survival aim, her training, and the maintenance policy. Then assign separate statuses to her decision process, route, outcome, and the institution.

You may conclude that Sela acted sensibly on false information, that westward movement had fact-relative support, and that the institution bears responsibility. None of those conclusions establishes a universal moral reason. The exercise builds clean rooms; it does not furnish them.

The twelfth debt

Add this entry:

***Debt 12 — The reason relation:** Does each alleged reason cause conduct, motivate the agent, support an option relative to evidence or aims, or count in favor independently—and has success in one relation been used to claim another without a bridge?*

We can now state the central problem precisely. The Candidate System must do more than cause approval or motivate people who share its values. To satisfy the original contract, at least some of its considerations must count in favor for agents who reject them.

Perhaps that demand is too strong. Perhaps all practical support begins with what an agent wants. Breathable air favors the west because Sela wants to live. If every reason has that structure, morality can guide through the aims people already have—or the aims they would have after becoming better informed, more coherent, and more fully themselves.

This is a powerful proposal. It places authority inside agency rather than above it.

The next room contains a mirror designed to show not only what you want, but what you would want if you saw clearly.

Chapter 13 — Current Aims and Better-Informed Aims

The mirror in Port Meridian’s next chamber did not show faces. It displayed corrected desires.

Councilor Pell stood before it.

CURRENT AIM: Win reelection.

FULLY INFORMED AIM: Win reelection without destroying the seawall budget.

COHERENT AIM: Stop promising both lower taxes and higher seawalls.

VIVIDLY APPRECIATED AIM: Do not shake hands beside a flooded school.

Pell stepped away.

“The device is politically biased,” he said.

The proposal behind the mirror is serious. Perhaps nothing counts in favor of action entirely on its own. Perhaps facts become reasons only by connecting with an agent’s aims. If you want to survive, breathable air supports the western tunnel. If you care about your child, the child’s fever supports calling a physician. If you want a stable society, evidence about corruption supports institutional safeguards.

This family of views is often called **desire-based** or **internalist** because it grounds practical reasons in some element of the agent’s motivational set: desires, aims, commitments, cares, projects, or values.

Its central insight deserves emphasis:

Guidance that never connects with anything in an agent’s practical identity can look less like a reason and more like an order from nowhere.

Desire-based accounts also explain why the same fact can support different actions for different agents. A train leaving at six gives a reason to hurry to the traveler who wants to board it, not to everyone within walking distance. Reasons become intelligible through projects.

The crude version—whatever you presently want gives you a reason—is too easy to defeat. Sophisticated accounts improve the desires before allowing them to generate reasons. We will test five versions, each more selective than the last.

Where this chapter is headed. We will test whether reasons can be grounded in present aims, instrumentally coherent aims, or the aims a better-informed agent would have. Each repair improves some desires, but each must disclose who chooses the improvements and whether the result still belongs to the agent.

Version one: the present desire

The simplest formula is:

If an agent desires an end, facts about effective means provide the agent with reasons.

This does not say every desire must be fulfilled. Desires conflict, differ in strength, and depend on false beliefs. It says they supply the raw material of practical support.

The account performs well in ordinary cases. Iona wants to repair the harbor gate. The correct wrench fits the bolt. The wrench's location gives her a reason to walk to the cabinet. No universal moral truth about wrench retrieval is required.

But present desires create immediate problems.

Conflict

Iona wants to repair the gate and wants to go home. Which desire governs? Strength may decide motivation but not necessarily support. A compulsive craving can be stronger than a settled project without becoming the better guide.

Mistake

An agent may desire a means because they falsely believe it serves an end. Sela wanted to enter the eastern tunnel because she believed it would preserve her life. Correcting the belief changes the means without needing to condemn the survival aim.

Fleeting and induced aims

A momentary impulse may oppose long-standing commitments. Advertising, addiction, panic, sleep deprivation, and manipulation can manufacture desires. If every current desire generates equal authority, whoever controls attention gains a disturbing influence over reasons.

Cruel aims

An agent may desire another's suffering. If that aim provides a reason in exactly the same structural sense as an aim to relieve suffering, then the account cannot criticize cruelty as contrary to the agent's reasons unless cruelty conflicts with another desire.

The present-desire account can accept this. It may say reasons are agent-relative and morally neutral until filtered by social practices. That is coherent. It does not provide the approval-independent moral authority promised by the original contract.

Version two: the coherent set

The second version allows only desires that survive correction for inconsistency and conflict.

Imagine an agent who wants health, smokes heavily, refuses illness, and denies any conflict among these aims. Reflection can reveal that not all of the set can be jointly satisfied. A coherent ordering may favor health over smoking or, less attractively but still coherently, smoking over longevity.

Coherence removes impossibilities and exposes means-end errors. It can require an agent to revise at least one element of an inconsistent set. It does not specify which element to preserve.

Suppose Pell desires low taxes, excellent public services, no debt, and no economic growth. Arithmetic creates tension. Coherence demands revision. It does not tell him whether to abandon low taxes, services, fiscal restraint, or opposition to growth.

A perfectly coherent monster remains possible. They may organize all subordinate aims around domination, remove every sentimental conflict, and pursue cruelty with impeccable consistency. Coherence is a property of fit, not content. A map can be internally consistent and lead to a dreadful destination.

Version three: the fully informed desire

The third version asks what the agent would want if fully informed of all relevant non-moral facts: consequences, probabilities, alternatives, causal relations, and perhaps facts about their own psychology.

This is a major improvement. False beliefs generate many destructive choices. A person may support a policy because they misunderstand its effects. They may pursue status expecting satisfaction and discover that the reward leaves them isolated. Information can expose a desired means as ineffective and reveal unnoticed costs to other aims.

Yet information does not choose terminal aims by itself.

Suppose Kest wants to dominate a neighboring community. He falsely believes domination will be easy and painless. Full information reveals that it will require years of violence. Kest may abandon the plan because the cost conflicts with his comfort. Or he may accept the cost because domination matters more to him. The information corrects prediction. It does not contain compassion in a concealed footnote.

An ideal adviser can tell Kest exactly what suffering will occur. Whether that suffering defeats his end depends on a concern for those who suffer or on a reason claimed to apply independently of his concerns. The latter route leaves the desire-based account.

Full information also creates a counterfactual problem. The actual agent is not omniscient. Which desires of the imagined informed agent give reasons to the real one, and why? An adviser may represent what the agent would choose after correction, but the bridge from hypothetical endorsement to present authority needs defense.

Version four: vivid appreciation

Facts can be known without being felt. The fourth version asks what an agent would desire after **vividly appreciating** the relevant experiences—perhaps through detailed imagination, testimony, virtual simulation, memory, or empathetic understanding.

This repair captures something important. A statistic about ten thousand displaced families may barely engage attention; meeting one family can transform practical concern. A smoker may know every risk yet change only after watching a friend struggle to breathe. Vividness can repair the distance between verbal knowledge and lived significance.

But vividness is not neutral.

Presentation selects a perspective, sequence, duration, and emotional frame. A simulation can immerse Kest in the victims' terror. It could instead immerse him in the exhilaration of victory, the anxiety of his own soldiers, or the fear of losing power. Which experiences must be appreciated? At what intensity? For how long? Should attachment to one's own group be amplified, quieted, or held constant?

If we require vivid exposure specifically because it tends to produce humane concern, humane concern is already selecting the idealization. The simulation did not derive that selector from desire alone.

Vivid appreciation may also alter rather than reveal the agent. Every experience changes us. That is not necessarily corruption; education properly transforms. The question is whether the transformed desire counts as a correction of the agent's practical outlook or the installation of a new value.

Version five: integration with a life

The fifth version admits desires that survive informed reflection and fit the agent's long-term identity, enduring commitments, relationships, and conception of a life worth living.

This filters impulsive and alien desires. It can distinguish a passing urge from a project around which someone has organized years of effort. It respects persons as temporally extended agents rather than desire containers refreshed every second.

Suppose a musician wants, during a severe depressive episode, to undergo a safe procedure that will permanently remove her capacity to experience music. She says the capacity now feels like an intolerable wound. The procedure would also erase the possibility that music could matter to her again. She understands the facts and voluntarily consents.

Should the present desire govern? A long-term integration account may protect the future capacity because it is central to her history and because the current state may be transient. Yet this response selects which temporal self, capacity, or narrative deserves authority. The present musician might say that invoking her former identity imprisons her inside a person she no longer wishes to be.

Now imagine the decision is calm, stable for ten years, and unrelated to illness. Resistance becomes more paternalistic. If the system still forbids the procedure, it may be protecting a value of capacity or identity independently of her aims. If it permits the pro-

cedure, it accepts that informed agency can authorize irreversible loss. Either answer is available; neither is supplied by the word *integrated* alone.

The many ideal selves

Idealization is often described as though one actual person enters a correction chamber and one improved person exits. In practice, different correction rules can produce different ideal selves.

Imagine two advisers working with the same agent, Lio.

The first adviser preserves Lio's strongest present commitments, removes false beliefs, resolves contradictions with minimal alteration, and calculates the best means. This adviser treats continuity as respect: correction should leave as much of Lio intact as possible.

The second adviser gives equal vividness to every affected person's experience, weakens status anxiety, exposes the contingency of inherited loyalties, and asks Lio to choose without knowing which social position will be theirs. This adviser treats impartial expansion as respect: correction should free Lio from a narrow formation history.

Both advisers provide accurate information and coherent reflection. They may produce opposite recommendations. The first preserves Lio's intense loyalty to a local group. The second transforms it into equal concern for outsiders.

Which is Lio's **better-informed aim**?

The answer depends on at least four design choices:

- **conservation:** how much of the current value set must remain;
- **perspective:** whose experiences receive attention and in what proportion;
- **conflict repair:** which desire yields when coherence requires change;
- **endpoint:** whether reflection stops at stability, endorsement, impartiality, flourishing, or another condition.

Path matters too. A person who first vividly experiences a victim's fear may revise later information differently from one who first inhabits a threatened community's insecurity. Unlimited reflection does not guarantee convergence if each stage alters the priorities used at the next.

This does not make idealization useless. Advisers can expose mistakes, hypocrisy, manipulation, and shortsightedness. The exercise becomes misleading only when one selected transformation is presented as what the agent **really** wants without revealing the selection rule.

When you design an ideal self, run the **plural-output test**: construct two reasonable correction procedures that begin with the same agent. If they diverge, identify the value that chooses between them. That value may deserve adoption. It was not derived merely from the fact that the agent could be improved.

The implanted desire

Port Meridian develops a treatment that creates one powerful desire in its recipient. The desire feels native, fits existing memories, and brings satisfaction when fulfilled.

A shipping company secretly gives the treatment to an employee named Ren. Ren now wants, above everything else, to work unpaid night shifts. The desire conflicts with no other aim because the treatment also weakens Ren's interest in leisure and income. Ren knows what the treatment did, understands the company's motive, and nevertheless endorses the new identity. "This is who I am now," Ren says.

What should a desire-based account say?

Reject the implanted desire because of its origin

This protects agency from manipulation. But many ordinary desires are causally shaped by parents, education, culture, hormones, advertising, and chance. The account needs a criterion distinguishing unacceptable implantation from ordinary formation. Lack of consent is a strong candidate, though Ren now retrospectively endorses the change.

Accept it because it is now authentic

This respects Ren's current point of view. It also allows a manipulator to manufacture reasons by manufacturing desires. If the company can create the aim, it can create the employee's instrumental reason to work.

Appeal to what Ren would have wanted beforehand

This protects the pre-intervention self. But why does the earlier self have permanent authority over the later one? Beneficial education and chosen transformation often produce values the earlier self lacked.

Appeal to an independent value of autonomy

This may be the best moral answer: manipulation wrongs Ren by bypassing self-direction. It also introduces autonomy as a value that judges desires rather than being generated by whichever desires survive.

The case does not show that implanted desires are never reasons. It shows that filtering desires by authenticity, consent, or autonomy adds standards whose own authority must be explained.

The lucid tyrant

Idealization earns universal reach only if it can address an agent who is already informed, coherent, imaginative, and settled. Kest is that test.

Kest rules a fortified district and values domination, aesthetic grandeur, personal continuity, and the loyalty of a small inner circle. He is factually informed. His aims are internally coherent. He has experienced immersive simulations of his victims' fear, grief, humiliation, and lost futures. He understands them accurately and does not enjoy their pain for its own sake; he simply gives it no weight except when it threatens his projects.

He has reflected for decades. His values are stable, integrated with his identity, and effective. He does not secretly long for universal love. He accepts that others have reasons, by their lights, to resist him. He is prepared to defeat them.

Does Kest have a reason not to dominate?

The strongest desire-based answer may be **no**, unless restraint serves one of Kest's aims. We have reasons to oppose him. His victims have reasons to resist. A coalition has reasons to coordinate. But nothing inside his motivational set supports equal concern.

This answer is not logically confused. It gives up a central aspiration of objective morality: that the victims' suffering counts in favor of restraint *for Kest*, whether or not he cares.

A second answer says Kest would care under the correct idealization. Then the idealization must do more than supply facts, coherence, vividness, and integration, because Kest has all four. It must add impartiality, empathy with motivational force, universalizability, equal concern, or another substantive standard.

That standard may be defensible. It is also doing the moral work. Calling the result "what Kest really wants" hides the replacement inside a flattering biography.

Correction or replacement?

The Candidate System needs an **idealization protocol**. For every proposed transformation, mark whether it corrects means, reveals conflict, expands understanding, strengthens imagination, orders temporal aims, or adds substantive concern.

Use this template:

Starting agent: aims _____; information _____; capacities _____.

Transformation: add/remove/change _____.

Standard selecting the transformation: _____.

What remains fixed: _____.

Resulting aims: _____.

Why this is correction rather than replacement: _____.

Why the result gives the actual agent a reason: _____.

Test the protocol against three cases:

1. Ren's implanted desire.
2. The musician's voluntary loss of future musical capacity.
3. Kest's fully informed and coherent domination.

For each, choose one response:

- **Accept the result:** the transformed or untransformed aims generate the reasons, even if the verdict is uncomfortable.
- **Add a bridge:** introduce autonomy, concern for suffering, temporal fairness, or another standard openly.
- **Narrow the claim:** say the account generates instrumental or agent-relative reasons but not universal moral authority.

Hold fixed that Kest has no factual mistake, incoherence, imaginative deficit, or hidden humane desire. If you add one, you have solved a different case. Perhaps all actual tyrants are less idealized than Kest. That empirical claim would help practical intervention but would not answer what gives Kest a reason in the stipulated case.

What this candidate accomplishes

Desire-based reasons explain a great deal:

- why means matter relative to ends;
- why different agents can have different practical reasons;
- why information and reflection improve choice;
- why motivation often follows perceived support;
- why guidance disconnected from every concern feels alien;
- how criticism can proceed from values an agent already holds.

This last achievement is especially important. Most cruel or reckless agents are not Kest. They have conflicts, factual errors, attachments, fears, dependencies, and commitments. Internal criticism can expose hypocrisy, ineffective means, ignored consequences, and values applied selectively. One need not prove objective morality before preventing a bad decision.

The candidate's limit is equally precise. It does not generate an aim-independent reason merely by polishing aims. If it adds a humane standard to guarantee humane results, the authority of that standard becomes the new question.

The thirteenth debt

Add this entry:

***Debt 13 — The idealization standard:** Which transformations correct an agent's desires rather than replace their values, what selects those transformations, and why do the resulting hypothetical aims govern the actual agent?*

The mirror can remove fog, expose contradiction, deepen appreciation, and show how a life fits together. It cannot promise that every clear reflection will reveal the same face.

Perhaps reasons do not arise from whatever we desire, even ideally. Perhaps they arise from what we **choose**—especially when we promise, accept a role, or build a life around a

commitment. Choice seems capable of changing the practical landscape. Before promising, I may be free. After promising, “I no longer feel like it” appears insufficient.

The next chamber is a contract room. On the table lies a pen, a uniform, and a resignation form.

Chapter 14 — Commitments, Promises, and Roles

The contract room contained three objects: a pen, a physician's coat, and a resignation form.

On the wall, someone had written:

BEFORE YOU CHOOSE, MANY PATHS ARE OPEN. AFTER YOU CHOOSE, THE PATH REMEMBERS.

This was either profound or an unusually threatening advertisement for stationery.

Commitments seem to change what we have reason to do. Before promising to meet you at the ferry, I may spend the afternoon as I please. After promising, the ferry matters in a new way. My current desire to remain in bed does not erase the appointment. Neither does the fact that you will never discover my absence.

Roles do similar work. A physician acquires responsibilities toward patients. A judge must hear both sides. A guardian controls resources for a beneficiary rather than for personal use. These demands can survive fatigue, dislike, secrecy, and changes of mood.

Perhaps this is the source the Candidate System needs. Objective morality need not descend from a realm of independent values if agents can create reasons through choosing, promising, joining, and accepting responsibility.

This proposal has genuine power. It also has a boundary. We must ask which reasons a commitment creates, for whom, under what conditions, and whether the authority of commitment is itself chosen.

Where this chapter is headed. We will examine how promises, roles, memberships, and relationships change the practical landscape. Their force may be real and substantial while remaining conditional on entry, content, reliance, power, and the possibility of exit.

What changes when a promise is made?

Suppose Iona tells a harbor mechanic, "Repair the north gate tonight, and I will relieve your morning shift." The mechanic relies on the assurance, cancels child care for the morning, and completes the repair. At dawn, Iona would prefer to attend a meeting and keep the mechanic on duty.

Several things now support keeping the promise.

The expressed decision

Iona deliberately placed her future conduct under a rule. Keeping faith respects her capacity to shape a life across time rather than treating each new desire as a constitutional coup.

Reliance

The mechanic changed plans because of the assurance. Iona helped create a vulnerability. The reason to perform is not merely self-addressed consistency; it responds to dependence induced in another person.

Expectation and trust

Assurances create legitimate expectations. A practice of reliable promising allows cooperation among people who cannot act simultaneously. Breaking the promise can damage this relationship and the wider practice.

Respect for the recipient's agency

Deceptive assurance uses another person's planning capacities as an instrument. Keeping the promise honors their authority to organize life around information intentionally supplied.

Integrity

Iona may value being someone whose word can be trusted. The promise becomes part of her practical identity, not a passing desire competing on equal terms with the morning meeting.

These grounds overlap, but none is simply “a promise has occurred.” They explain why the act of promising can alter reasons. They also show why not every utterance in the grammatical form “I promise” binds.

If an actor speaks the line onstage, no assurance is given. If a robber forces Iona to promise assistance, voluntariness is damaged. If the mechanic knows she is joking, reliance and uptake are absent. If the promised act is to falsify safety records, the content collides with other constraints.

A commitment needs an anatomy.

Seven clauses in every commitment

Before treating a commitment as reason-generating, specify at least seven conditions.

Entry

How is the commitment acquired? Through explicit consent, informed enrollment, conduct, reliance created, oath, appointment, relationship, citizenship, birth, or some other event?

Voluntary entry is especially powerful because the agent participates in creating the demand. Yet many roles are not chosen in a simple moment. Children do not sign family contracts. Citizens are born into states. A person may become a caregiver because no one else is available. If the role is unchosen, its reasons need grounds beyond “you agreed.”

Scope

What actions, people, and periods does the commitment cover? “I will help” might mean one ride tomorrow, continuing emergency support, or an indefinite obligation. Ambiguity favors the party with power to reinterpret later.

Exit

Can the agent resign, withdraw, transfer responsibility, or give notice? A commitment with instant costless exit may be too weak to support reliance. A commitment with no exit can become bondage.

Exit rules should account for dependencies already created. A physician may resign from a clinic but not abandon a patient during surgery. The delay is explained by transition responsibility, not permanent ownership of the physician.

Emergency

What happens when performance threatens grave harm or collides with a stronger commitment? A promise to meet at the ferry may yield to rescuing a child. The exception needs the same threshold discipline developed in Part II.

Deception and mistake

Was entry obtained through false information or concealed terms? A choice cannot authorize what the agent had no fair opportunity to understand. Some mistakes defeat consent; others are ordinary risks the commitment allocates.

Exploitation

Did one party use desperation, dependency, monopoly, or unequal power to extract the commitment? “You chose” can conceal a choice among starvation, exile, and acceptance.

Exploitation is not identical to coercion. A person may knowingly sign without a threat and still face terms that take unfair advantage of vulnerability. To use this limit, the system must state its account of fair background conditions.

Excessive cost

How much can performance demand? A promise may survive inconvenience and substantial sacrifice; otherwise assurance is cheap. But a casual commitment does not obviously require death, permanent disability, or destruction of every central project.

The cost limit depends on the content, stakes, reliance, voluntariness, and alternatives. “Until it becomes difficult” is not a commitment rule. “At any cost whatsoever” is rarely a humane one.

These clauses do not weaken commitment. They explain its authority more accurately than the incantation “a promise is a promise.”

Three structures of commitment

Commitment-generated reasons do not all have the same shape.

Self-governance across time

An agent settles a future question so that later impulses do not reopen it without cause. A savings plan, sobriety pledge, research project, or personal vow can structure agency even when no one else knows.

The reason is primarily **diachronic**: it connects the agent now with the agent who chose and the life being built across time. Exit may remain possible, but revision should answer to the purpose of committing rather than to the very temptation the commitment was meant to resist.

This structure is vulnerable to domination by the past self. A commitment made under ignorance or desperation should not imprison every later self. The idealization questions from Chapter 13 return.

Directed assurance

A promise gives another person a claim because it invites trust and planning. The recipient can release the promisor, demand performance, or seek repair in ways a stranger cannot.

The reason is **relational**. It arises not merely because the speaker selected a plan but because the speaker deliberately altered another person's practical situation. Uptake and reliance matter, though some promises bind before visible reliance because the assurance itself transfers limited authority to the recipient.

Practice or office

A role places an agent inside a scheme of responsibilities: physician, judge, guardian, teacher, elected official. The reason is **institutional**. It depends on what powers the practice grants, what vulnerabilities it creates, and which values justify the institution.

Role reasons can survive any particular recipient's release. A judge cannot accept a litigant's permission to ignore due process if the procedure also protects public legitimacy and the opposing party. This shows why role obligations are not simply collections of promises.

One act can involve all three structures. Senn may personally identify as a healer, assure a patient of care, and occupy a licensed office. If one ground fails, the others may remain. Your charter should state which structure is doing the work rather than treating *commitment* as a single invisible adhesive.

The letter to the dead

Captain Osei asks her friend Nara to deliver a private letter to Osei's estranged brother after Osei dies. Nara promises. Osei dies unexpectedly. Before the letter is sent, the brother also dies. The letter names no other beneficiary and contains no information relevant to anyone's welfare.

Must Nara still attempt delivery, perhaps by placing it with the brother's effects? If not, when did the reason end?

The case removes several familiar supports. Osei cannot experience disappointment. The brother cannot receive the message. No living person knows about the promise. Violation will not damage public trust.

Several routes remain.

Former agency

Osei exercised self-direction while alive, and present action can respect or frustrate that earlier choice. Advance directives and wills often rely on this structure.

Continuing integrity

Nara's present commitment persists as part of her practical identity. The promise can guide her even when the recipient no longer benefits.

Constitutive practice

Promises may include performance after death as part of what the practice means. A secret exception for undetectable cases would make the practice less than what participants understand themselves to undertake.

Symbolic honor

Carrying out the promise may express respect for Osei and the relationship. The value is not a sensation in the dead but a way the living relate to lives and commitments across time.

Or perhaps the promise has lapsed because performance can no longer serve its point. A purpose-based account may release Nara. This answer need not treat Osei as worthless; it says the particular undertaking was directed toward a now-impossible communication.

What the case reveals is not one obvious verdict. It reveals several possible grounds. If your Candidate System says the promise binds "because Osei has a continuing claim," specify what bears that claim after death. If it says Nara may withdraw, explain why her voluntary commitment has lost force.

The physician and the enemy

Doctor Senn receives a captured enemy commander with a treatable wound. The commander ordered attacks that killed Senn's family. No other clinician is available. The medical role requires treating patients without hostile discrimination.

Senn strongly wants the commander to die. The death can be concealed. Treating the wound will probably return a dangerous officer to detention, not directly to battle.

The professional role supplies a reason that survives Senn's present desire and the chance of detection. Why?

- Senn voluntarily entered a practice with publicly declared duties.
- Patients and institutions rely on clinicians to separate care from vengeance.
- The role gives Senn access and power over vulnerable bodies on condition that the power serve health.
- Identity-neutral treatment protects everyone in conditions of fear and conflict.
- Senn's professional identity may include the commitment to treat.

This is a strong case for role-generated reasons. But alter the facts. Suppose the medical code requires participation in torture. Voluntary entry into medicine does not make the requirement valid. The role is judged by values and constraints beyond mere membership.

Or suppose Senn was drafted into the role under threat. Reliance and vulnerability may still support treating the commander, while voluntary commitment does less work. Several grounds can converge even when one fails.

The case also illustrates agent-relativity. Senn has the medical responsibility because Senn occupies the role. A nearby electrician does not violate a professional duty by refusing to perform surgery. Role-generated reasons can be genuine and asymmetric without claiming that every agent must perform the same act.

The rule revealed after rescue

A mountain guild rescues Vale from a storm. Once he recovers, guild leaders present a rule: anyone rescued by the guild becomes a member and owes five years of unpaid service. The rule is ancient, consistently enforced, and known to guild members. Vale was unconscious during rescue and had never heard of it.

"You accepted the benefit," the guild says.

Did Vale join?

Receiving an unrequested emergency benefit does not automatically constitute consent to undisclosed terms. Vale could not refuse, bargain, understand the scope, or choose another rescuer. The guild may have reasons to request costs, public funding, or reciprocal aid. It cannot manufacture a five-year commitment by keeping the entry clause secret until the rescued person is dependent.

Now vary the case:

- Signs at every trailhead clearly state the rule.
- The service obligation is one day rather than five years.

- Rescue is impossible unless all hikers fund the guild through a mandatory levy.
- Vale explicitly called for rescue after hearing the terms.
- The guild is the only possible rescuer because it has excluded competitors.

Each change affects information, voluntariness, cost, reliance, and power. “Benefit received” is not a complete entry rule.

Can commitment create moral authority?

The strongest answer is yes, within limits. Voluntary commitments can generate new relational facts: assurance, reliance, delegated authority, shared projects, and legitimate expectations. These facts give agents reasons they did not have before. A promise is not merely evidence of a preexisting desire. It can settle which of several permissible paths one will take and make others dependent on that choice.

This achievement should not be minimized.

It also does not establish everything the objective project seeks.

Commitments can have bad content

People can promise to harm, join oppressive organizations, and adopt corrupt roles. The act of choosing does not purify the object chosen. Content limits come from elsewhere.

Commitment binds parties differently from outsiders

My promise can change what I owe you. It does not, merely by occurring, impose the same reason on a stranger or settle the moral status of people excluded from the agreement.

The authority of choosing remains a question

Why must an agent honor a past choice after every present desire has changed? Reliance, agency, integrity, practice, and welfare provide answers. Each invokes values from the Candidate System. If those values are aim-dependent, the commitment’s authority may ultimately be conditional. If they apply independently, the commitment account relies on the sort of objective support still under investigation.

Rejecters remain possible

Imagine an agent who freely promises, creates reliance, understands every consequence, and later says, “I no longer care about my word, your plans, my integrity, or the practice. I accept that you will oppose me.” Does the promise still count in favor for that agent?

If yes, commitment has produced or activated aim-independent support. Explain the source. If no, promises guide agents through continuing concerns and social enforcement but do not satisfy the original authority condition.

Draft the commitment charter

Purpose. Make a promise, role, membership, or relationship disclose its entry conditions, beneficiaries, content, limits, and exit; the charter records generated reasons without assuming universal authority.

Choose a promise, role, membership, or relationship in your Candidate System. Complete:

Entry: The commitment begins through _____.

Information required: _____.

Parties and beneficiaries: _____.

Scope: It governs _____ for _____.

Reliance created: _____.

Exit: An agent may leave by _____ after addressing _____.

Emergency: The commitment yields when _____.

Deception or mistake: It is void or revisable when _____.

Exploitation: It lacks authority when _____.

Excessive cost: It does not require _____ unless _____.

Content limits: No commitment may authorize _____.

Authority claim: The resulting reason is grounded in _____ and applies even if _____.

Test it against Osei's letter, Senn's enemy patient, and Vale's rescue. Then invent a corrupt role whose members sincerely accept their duties. If your charter validates it, add content limits. If the limits appeal to fairness, harm, rights, or equal status, mark the imported value. The import may be correct; the customs declaration is mandatory.

The fourteenth debt

Add this entry:

Debt 14 — The commitment boundary: How do entry, reliance, scope, exit, emergency, deception, exploitation, cost, and content determine a commitment's reasons—and why does the agent's past choice retain authority when current endorsement disappears?

Commitment allows agents to build durable practical structures. It explains special responsibilities and protects coordinated life from each passing mood. It does not bind every outsider or certify every chosen content.

Perhaps a commitment gains wider authority when it emerges not from one person's promise but from a procedure that treats everyone fairly. If each affected person receives an equal voice, misinformation and bargaining power are removed, and no one knows which position they will occupy, the resulting terms may be more than a local bargain. They may be the terms no one could reasonably reject.

The next chamber contains a round table. The chairs appear identical.

We will check underneath them.

Chapter 15 — Fair Terms and Constructed Reasons

The chairs around Port Meridian’s round table were identical except for one detail: seven were bolted to the floor, and the eighth had wheels.

“Equal chairs,” announced the facilitator.

The delegate in the wheeled chair drifted slowly toward the wall.

Fairness is rarely defeated by an open declaration that some participants matter less. More often, everyone is called equal while the floor is tilted.

The proposal before us is among the strongest alternatives to discovering moral truths in an independent realm. Perhaps moral reasons are **constructed** through a suitably fair procedure among agents. The procedure does not merely reveal a ranking that existed beforehand. Under the right conditions, its outcome constitutes the principles that participants have reason to follow.

This approach begins from recognizable moral insights:

- people are not merely objects to be arranged for a total good;
- coercive rules should be justifiable to those they govern;
- no one’s existing power should automatically determine fair terms;
- disagreement requires a public basis for cooperation;
- free and equal agents can authorize institutions together;
- reciprocity can create reasons that isolated preference cannot.

A sophisticated constructed account does not say, “Whatever people vote for is right.” It designs conditions meant to exclude domination, ignorance, threat, prejudice, and self-serving exception. It may require public reasons, equal status, complete information, impartial choice, or terms that no one could reasonably reject.

This is not relativism in the casual sense of accepting every actual convention. A constructed rule can condemn Port Meridian’s real practices by comparing them with the outcome of a more defensible procedure.

The central questions are what makes the procedure fair, who receives a chair, what the procedure can authorize, and why its result binds someone who loses or rejects it.

Where this chapter is headed. We will ask what fair procedures can create, reveal, coordinate, or legitimate. Construction can settle rules among participants, but it must still explain who participates, which conditions count as fair, and what reaches an informed rejecter.

The fair lottery

Two patients, Aven and Bira, need the last dose of a medicine. Their conditions, prognoses, waiting times, and relevant responsibilities are identical. Dividing the dose will save neither. Both agree in advance to a transparent lottery administered by an independent nurse.

Bira wins.

Aven now wants the result overturned.

Does the lottery give Aven a reason to accept the allocation?

Several grounds converge.

Prior agreement

Aven voluntarily accepted the procedure before knowing the outcome. Reversal after losing would create an identity-based exception.

Equal standing

The lottery gave both claims an equal chance because no admitted feature ranked them. It expressed equality without pretending that the winner's life was more valuable.

Reciprocity

Aven expected Bira to accept a loss under the same rule. Demanding compliance only when one wins rejects fair terms of mutual restraint.

Coordination

The dose must be allocated promptly. A stable procedure prevents conflict and makes action possible where substantive comparison reaches a tie.

Reliance

Bira and the medical team now organize action around the announced result. Reopening the process imposes costs and undermines future procedures.

This is a strong case. The result appears legitimate even though the lottery did not discover that Bira objectively ought to live instead of Aven.

Now remove Aven's prior agreement. The hospital uses the same fair lottery under a publicly announced policy, but Aven has always objected and demands priority by age. Does the fair structure itself give Aven a reason to accept?

Perhaps. A policy cannot require unanimous consent from every future participant without becoming unusable. Equal treatment, public justification, institutional necessity, and the absence of a better ranking may support the procedure. But the argument now relies less on commitment and more on the independent authority of fairness.

Ask the hard question: if Aven understands the equality rationale and simply rejects equality, what makes the fair procedure count in favor *for Aven*? Saying “because it is fair” identifies the property whose authority is under examination.

Four ways a procedure can matter

Procedures play several roles. Keep them separate.

Epistemic role

A procedure can improve the chance of finding a correct answer. Diverse deliberation may expose errors; expert review may identify facts; an impartial court may test evidence.

Here the procedure is valuable as a method of discovery. Its output can still be wrong. A unanimous scientific panel does not make a false measurement true.

Expressive role

A procedure can express equal status, respect, inclusion, or reciprocity. Giving each person a voice may matter even when another method would produce the same policy.

Coordination role

A procedure can select one of several acceptable options so people can act together. Driving on the right rather than the left may be settled by authority or convention without one side being independently best.

Constitutive role

In some domains, the procedure can make an outcome valid. A properly conducted election determines who holds an office; a fair game determines the winner; a mutual agreement creates a partnership.

Constitutive power has limits. An election can determine who is mayor under a legitimate constitution. It cannot make the winner taller. A vote may authorize a tax within a jurisdiction without making every use of political power morally acceptable.

When a constructed account claims that fair procedure generates moral reasons, it invokes the fourth role. The procedure does not merely detect or coordinate. It makes certain principles valid for participants. The account must specify the domain and explain why procedural validity has practical authority.

The unanimous agreement on a tilted floor

Two settlements negotiate use of the only freshwater spring. High Ridge controls the land around it. Low Marsh has suffered a drought. High Ridge offers water in exchange for ninety percent of Low Marsh's harvest for the next thirty years. Every adult in both settlements votes. The agreement is unanimous.

Low Marsh's alternative is dehydration.

Is the agreement fair because everyone consented?

Actual unanimity proves that each person selected the agreement over the available alternatives. It does not prove that the alternatives were just, that bargaining power was equal, or that High Ridge may legitimately use control of a necessity to extract nearly all future production.

A defender can say that coercion requires a threat, while High Ridge merely declined to surrender its property. That moves the dispute to the background rule of ownership. Why does control of the spring include authority to exclude a desperate population? The agreement cannot validate the property rule if the property rule is what created the bargaining inequality used to validate the agreement.

We can improve the procedure:

- give both settlements secure access to basic survival before bargaining;
- appoint independent representatives;
- disclose hydrological facts;
- prevent threats and deception;
- allow revision when conditions change;
- limit terms that create permanent subordination;
- ask what parties would choose without knowing which settlement they occupy.

These are plausible fairness conditions. Notice what they contain: concern for survival, equality, informed agency, nondomination, reversibility, and impartiality. The procedure becomes morally attractive by being built from values already admitted.

That is not a defect in practical design. It is a problem only if we claim the procedure generated every value it required for its own construction.

The participant outside the door

Port Meridian and High Ridge design a perfectly balanced regional council. Each adult resident receives equal representation. Information is public, campaign spending is equalized, minority vetoes protect basic liberties, and deliberation is excellent.

The council approves a waste system that sends pollutants downstream into an unrepresented nomadic community. The community is affected but has no recognized residency and therefore no seat.

The procedure is fair among the included participants. The participant set is not.

Who must be represented?

- everyone subject to the rule;

- everyone significantly affected;
- everyone capable of reciprocal cooperation;
- every being with direct standing;
- future people through trustees;
- ecosystems or animals through representatives;
- only legal members, with outsiders protected by prior constraints?

The answer cannot be constructed by a procedure whose participants have already been selected, unless another procedure selects the selectors—and then we ask who participates in that one. At some point, standing enters as an input.

Representation also raises fidelity questions. A proxy for future people or animals cannot ask constituents what they want. The proxy needs a theory of interests and uncertainty. Procedure organizes representation; it does not create the represented interests.

Design the strongest fair procedure

Purpose. Specify every choice that gives a procedure content—participants, information, representation, constraints, decision rule, and appeal—so fairness is not asked to design itself.

Construct a procedure for adopting the Candidate System’s public rules. Do not begin with “everyone votes.” Fill every field.

Participants

Who receives a voice, vote, veto, representative, or protected claim? Use the standing map from Chapter 6. Explain exclusions.

Information

What facts do participants know? Are they ignorant of their own identity, power, talents, or social position? Do they know general psychology, scarcity, history, and probability? Who decides which information is reliable?

An ignorance device can reduce self-serving choice, but it also removes knowledge of particular relationships and histories that may matter. Total abstraction may produce impartiality by making care invisible.

Motivation

Are participants self-interested, mutually concerned, reasonable, risk-averse, committed to reciprocity, or prepared to offer publicly acceptable reasons?

If they are stipulated to care about fairness, the output will reflect that value. If they are purely self-interested, equal terms may arise from uncertainty about position, but treatment of unrepresented beings remains difficult.

Bargaining power

How are wealth, threats, urgency, monopoly, status, charisma, and unequal exit options neutralized? Formal equality at a table does not erase control of the spring outside.

Admissible options

Can participants choose slavery, permanent exclusion, torture, or hereditary rule? If not, which values set the boundaries before deliberation?

Outcome constraints can protect the procedure from self-destruction. They are prior moral materials unless the account gives another derivation.

Decision rule

Unanimity, majority, supermajority, consensus, bargaining, ranked choice, lottery, or judicial judgment? Each distributes blocking power differently.

Unanimity protects dissenters and can entrench the status quo. Majority rule enables action and can dominate persistent minorities. Lotteries express equality under ties and ignore intensity. No decision rule is fair without reference to the problem it serves.

Exit and appeal

May dissenters leave, refuse, appeal, or demand review? Exit is hollow when every territory follows the same rule or departure costs home, family, language, and livelihood. Internal voice and safeguards may be more important when exit is expensive.

Revision

How does the procedure learn? What new evidence reopens a decision? Who monitors effects? Can future participants revise terms imposed by the dead?

Once the fields are filled, write the authority clause:

An outcome generated by this procedure gives participant _____ a reason to comply because _____, even when the participant loses, rejects the outcome, and can defect secretly.

Do not write “because the procedure is fair” unless you next explain why fairness counts in favor for the rejecter.

Turn the procedural dials

Purpose. Change one procedural selector at a time and observe which outcome changes; sensitivity reveals where substantive judgment enters, not that the resulting procedure is defective.

A procedure can be impeccably administered and still owe its result to contestable design choices. To see this, give four councils the same decision.

Port Meridian must choose whether to reinforce the old seawall or relocate the waterfront district inland. Reinforcement preserves homes and relationships but transfers growing flood risk to future residents. Relocation breaks present communities, protects future lives, and destroys nesting grounds used by sentient shore birds.

Council One gives one vote to every current adult resident. Present homeowners dominate, and reinforcement wins.

Council Two asks the same residents to choose without knowing whether they own waterfront property, rent inland, or work in construction. Relocation wins by a narrow margin.

Council Three adds trustees for children and future residents. It selects earlier relocation with generous compensation.

Council Four also represents animal welfare and ecological continuity. It chooses a different inland site at higher financial cost.

Nothing in the case shows which council is uniquely correct. It shows that outcomes are sensitive to procedural dials:

- participant set;
- representation of absent or nonvoting interests;
- knowledge of one's position;
- permitted reasons;
- time horizon;
- compensation rules;
- outcome constraints.

Sensitivity is not a refutation. A thermometer's output should be sensitive to temperature. The question is whether the procedure responds to morally relevant features or to arbitrary design. To answer, we need standards of relevance—and those standards cannot be produced by observing that the procedure produced them.

Run the **dial test** on your construction. Change one design feature while holding every empirical fact fixed. If the outcome changes, write why the original setting is fairer. Then imagine the losing party offers the rival justification. Your answer may appeal to standing, agency, vulnerability, reciprocity, or nondomination. Cite that value openly.

A family of reasonable procedures may remain. In that event, construction can identify a range of legitimate settlements without selecting one objective optimum. Permissible plurality is a result, not an embarrassment.

Three possible claims

Your constructed account can make claims of different strength.

Conditional legitimacy

People who value equal cooperation, accept the procedure, or benefit from its stable operation have reasons to follow its outcomes under stated conditions.

This is robust and practically important. Its reasons connect with participant aims and commitments. It may not reach the person who rejects all those concerns.

Public justifiability

Coercive rules should be supportable by considerations that every affected agent could recognize without being treated as inferior. Outcomes that depend on sectarian authority or naked power fail.

This sets a demanding standard for institutions. The word *should* and the requirement of equal recognition carry substantive authority claims that need grounding.

Stance-independent construction

The correct fair procedure generates reasons that apply to each agent regardless of actual acceptance. Agents can be wrong about which terms they have reason to accept.

This comes closest to the original objective project. It also means the procedure's fairness conditions possess authority independent of actual stances. Construction has organized and extended those conditions; it has not removed the need to explain their status.

A defender might say that free and equal agency itself commits every participant to reciprocity. Chapter 18 will test whether agency contains such standards. For now, mark this as a promising bridge rather than a completed crossing.

What construction accomplishes

Fair procedures can:

- prevent existing power from masquerading as consent;
- express equal status;
- incorporate diverse information;
- create legitimate offices and shared commitments;
- settle indeterminate conflicts;
- make coercive rules answerable to affected people;
- provide public methods of criticism and revision.

These achievements remain even if the procedure does not manufacture stance-independent moral truth. A constructed constitution can be more humane, stable, transparent, and defensible than inherited domination.

The limitation is specific: participants, fairness conditions, outcome limits, and the authority of reciprocity cannot all be generated by the same procedure without some inputs doing unacknowledged work.

The fifteenth debt

Add this entry:

***Debt 15 — The construction materials:** Who participates, under which information, motivation, power, decision, outcome, exit, and revision conditions—and what gives those fairness conditions authority for a person who rejects them?*

The round table has now been leveled. Its chairs are genuinely equal. The procedure can produce terms far more defensible than raw bargaining or inherited command.

One delegate remains standing in the doorway.

“I understand the terms,” the delegate says. “I understand everyone’s interests. I understand that the procedure is reciprocal, impartial, and free of domination. I simply do not value any of that.”

The table cannot dismiss this person as confused without identifying the confusion. It cannot call them unreasonable if *unreasonable* means only “rejects the reasons we are trying to establish.”

Perhaps some considerations count in favor whether an agent cares or not. If so, the maze has reached the bridge we have repeatedly postponed.

In the next chamber, there are no contracts, mirrors, or ballots—only another person’s pain and an agent who is unmoved by it.

Chapter 16 — Reasons That Apply Whether You Care or Not

An enemy soldier lay in the dust outside Port Meridian's wall.

The fighting had ended. The soldier was disarmed, conscious, and dying from a poison that produced intense pain before stopping the heart. Iona held the antidote. Administering it would take one minute, create no danger, consume a dose that would otherwise expire, and return the soldier to secure detention.

Iona understood every fact. The soldier had killed no one she knew. There was no strategic benefit to rescue, no legal penalty for walking away, no witness, no promise, and no professional role. Iona felt no compassion. She did not enjoy the suffering. It simply did not enter her aims.

The soldier gasped, "You have a reason to help me."

Does she?

This is not a question about what would move Iona. By stipulation, the pain does not. It is not a question about what Port Meridian can require by law or what the soldier wants. It is not a question about what *we* would do.

The question is whether the soldier's preventable suffering itself counts in favor of administering the antidote **for Iona**, independently of her actual desires, commitments, approval, and likelihood of detection.

If the answer is yes, we have reached an **aim-independent practical reason**.

This is the bridge the objective moral project has been seeking. We should neither cross it by definition nor burn it because the structure looks unfamiliar.

Where this chapter is headed. We have reached the central bridge: can a fact count in favor for an agent who lacks every connected aim? This chapter states the strongest case for such reasons, then separates that claim from motivation, rational blame, and a complete moral verdict.

What an aim-independent reason claims

The claim has four parts:

1. A fact obtains: the soldier is in preventable agony.
2. An action is available: Iona can administer the antidote safely and easily.
3. A favoring relation holds: the fact counts in favor of that action for Iona.
4. The relation does not depend for its existence on Iona already valuing relief, the soldier, morality, consistency, reciprocity, reputation, or any related end.

Iona may have stronger reasons that defeat rescue in a different case. Perhaps the antidote is needed for a child or rescue would restart the war. Here those routes are closed.

The case isolates application: can another being's suffering supply practical support to an unmoved agent?

We can distinguish this from an **objective moral fact**. An objective moral fact assigns or grounds a moral status: perhaps "allowing this agony is wrong." An aim-independent practical reason is a relation of support between a fact, an agent, and an option: the pain counts in favor of rescue for Iona.

The two may connect. The moral fact might be true because the reason exists. The reason might exist because suffering has objective disvalue. A broader moral theory might ground both. But renaming "objective moral demand" as "practical reason" does not explain the relation. The Candidate System needs the bridge, not a new plaque.

Why believe the bridge is real?

The strongest case for aim-independent reasons does not begin with metaphysical machinery. It begins with deliberation and interpersonal address.

Some ends seem criticizable, not merely inefficient

We do not only tell the tyrant that torture is a poor means to his aims. We tell him the victims' agony is a consideration he is failing to recognize. If his plans become more coherent and successful, we do not regard the criticism as answered.

Victims can address agents from outside the agents' projects

The soldier's plea is not naturally translated as "Please acquire a new desire that would make my pain relevant." It claims standing now. The plea presents the suffering as a consideration Iona can recognize, not as a motivational product she must manufacture.

Deliberation appears able to correct aims

We sometimes conclude not merely that our means were mistaken but that an end was unworthy, cruel, shallow, or based on treating another person as negligible. Reflection can feel like discovery about what matters, not only rearrangement of desire.

Deliberation cannot wait for a verdict on reasons

Whenever an agent asks what to do, the question seems to call for considerations that favor one option rather than merely predict which option the agent will choose. Even deciding whether to keep, revise, or abandon an aim appears to require a standpoint from which aims can be assessed. If every practical reason is generated by an existing aim, the question "Which aims should survive reflection?" threatens either to repeat the current aims or to smuggle in a standard that was not generated by them.

This is a powerful argument for irreducible practical normativity. Deliberation may be impossible to understand completely as desire plus efficient calculation. But the argument

leaves a second task. The existence of genuine reasons does not yet identify which beings have direct standing, whether another person's pain supplies a reason to this agent, or how welfare compares with liberty, loyalty, and risk. Deliberative indispensability can support the bridge's existence without writing the constitution carried across it.

Approval can be mistaken

The reformer's conviction that a popular cruelty remains wrong under unanimous celebration is central to the objectivity threshold from Chapter 2. Aim-independent reasons explain how an entire society can fail to respond to considerations that still apply.

Reasons need not motivate automatically

Evidence can support a conclusion a stubborn thinker refuses to believe. Likewise, a practical consideration may support an action without producing motivation. Iona's indifference does not refute the reason any more than closing one's eyes extinguishes the lighthouse.

These points do not prove the view. They show why it is not an ornamental philosophical invention. It articulates ordinary practices of criticism, demand, regret, and reform.

The reasons-first proposal

One powerful answer treats the reason relation as basic.

Why does agony count against an action? Because agony is a reason-giving feature. Why does a promise count in favor of performance? Because valid assurance and reliance are reason-giving. Not every legitimate explanation reduces to nonnormative ingredients. At some point, any theory reaches fundamentals.

Call this the **reasons-first** proposal:

Certain natural facts stand in irreducible support relations to actions. Competent agents can recognize those relations, weigh them with others, and act correctly or incorrectly in response.

The proposal avoids deriving values from desires, commands, conventions, or procedures. It explains their authority by placing favoring force in reality itself.

An immediate objection says, "But how can a natural fact generate a reason?" The reasons-first theorist replies that *generate* already assumes an unnecessary two-step picture. The pain fact does not first exist as normatively blank and then enter a reason machine. Its being a reason is a fundamental truth about that fact in those circumstances.

This reply deserves respect. Explanations cannot continue forever. Physics has fundamental relations. Logic does not derive validity from popularity. A demand that normative facts reduce to nonnormative facts may simply assume the conclusion against them.

The reply also defines the burden precisely. The Candidate System may take some reason relations as basic. It must then state their content, scope, structure, conflicts, and

method of recognition. “Basic” can stop a regress. It cannot turn every confident intuition into a fundamental property.

Four accounts of the support relation

Aim-independent reasons can be understood in several ways.

Reasons as fundamental facts

On the purest version, it is simply true that certain facts count in favor of certain responses. The soldier’s pain favors relief. No agent’s attitude constitutes this relation.

Strength: the account preserves full stance independence and directly answers the soldier.

Debt: which facts are reasons, with what weight, and how do finite agents distinguish recognition from projection?

Reasons grounded in objective value

Perhaps suffering is objectively bad, agency objectively valuable, or persons possess objective status. Reasons arise because actions promote or respect these values.

Strength: a value theory organizes many reasons and explains patterns.

Debt: what makes the value objective, and how does value become support for this agent? “Pain is bad, therefore alleviate it” still contains a practical transition.

Reasons grounded in the good of the affected being

The soldier has a welfare interest in relief. Because he has direct standing, his good supplies reasons to agents able to affect it.

Strength: the account locates the demand in a real subject whose life can go better or worse, rather than in an abstract moral mist.

Debt: direct standing identifies whose good enters. It does not by itself establish that another agent must treat that good as support. That was the boundary discipline of Chapter 6.

Reasons grounded in rational agency

Perhaps acting for reasons commits an agent to consistency, universalization, respect for agency, or terms others could accept. By participating in practical reasoning, Iona cannot coherently exempt herself from standards implicit in agency.

Strength: authority is not an alien force; it arises from what agents do whenever they deliberate.

Debt: instrumental consistency alone may leave substantive ends untouched. A coherent tyrant reasons effectively. Chapter 18 will test whether rational agency contains equal regard or whether that value has been added to the account of rationality.

These accounts can combine. They should not form a circle in which standing is grounded by reasons, reasons by value, value by standing, and the entire circuit is then presented as a power source.

Reason, rationality, and responsibility

Suppose the soldier's pain genuinely gives Iona a reason to help. Three further questions remain.

Is Iona irrational if she refuses?

If rationality means responding correctly to all practical reasons, then yes by definition. But that definition does not independently establish which reasons exist. If rationality means coherence and effective pursuit of aims, Iona may be instrumentally rational while ignoring the soldier.

Calling her *irrational* can therefore either record a failure to respond to an established reason or attempt to establish the reason by classification. Only the first route is legitimate.

Is Iona blameworthy?

Blame depends on more than the existence of a reason. Did she recognize it? Could she respond? Was the demand clear? Did fear, trauma, coercion, or incapacity interfere? How grave was the failure, and what expectations fairly applied?

A hidden fact may supply fact-relative support while its absence from available evidence excuses the agent. The poisonous eastern tunnel gave Sela a fact-relative reason not to enter, but blaming her for missing an inaccessible fact would be perverse.

May anyone compel Iona?

Even a decisive reason for rescue does not automatically authorize every third party to use force. Enforcement needs jurisdiction, proportionality, evidence, institutional competence, and attention to abuse. A state may permissibly require easy rescue, or it may avoid such laws because officials cannot distinguish safe from dangerous cases reliably.

This preserves the status grammar from Chapter 7:

reason → *all-things-considered requirement* → *responsibility* → *blame* → *enforceability*

Each arrow needs conditions. Objective reasons, if they exist, do not collapse the chain.

Now reverse the case. Suppose Iona helps because she falsely believes the soldier is her brother. The soldier's suffering may genuinely favor aid, but it is not the consideration that moves her. Her action can conform to the reason without being performed **for** it. Moral worth, correct action, and appropriate motivation may come apart.

The realist candidate is strongest when it claims only what it has earned: a favoring relation. Additional verdicts require the rest of the Candidate System.

The unmoved antidote holder

Change the case. A stranger named Jalen possesses the only antidote to a child's poisoning. Giving it up costs Jalen nothing. Jalen understands the child's terror, the parents' attachment, the loss of an entire future, and the ease of rescue. He refuses because he does not care.

Three responses are available.

Response one: Jalen has a decisive reason

The child's grave interest and trivial cost make rescue strongly supported. Jalen's indifference is a failure to respond to the reason, not evidence of its absence.

This preserves the ordinary verdict and the objectivity aspiration. It owes an account of why the child's interest applies to Jalen, how its weight is known, and what relation makes it practical support rather than merely an outcome we deplore.

Response two: Jalen has a reason only if a concern connects

Perhaps Jalen cares about avoiding conflict, honoring law, seeing himself as decent, or preserving a society in which others rescue children. Any of these can connect the facts to his aims.

If all connections are removed, the account says no practical reason remains for him. We can still have reasons to seize the antidote, condemn him under our code, distrust him, and protect others from him. The loss is the claim that rescue counts in favor from his practical standpoint.

Response three: the question falsely isolates agents

Perhaps persons are socially constituted, inevitably dependent, and already embedded in relations of recognition. Jalen's supposed independence is a fiction. His language, safety, identity, and resources depend on others.

This insight reveals actual connections. But the case can stipulate that Jalen accepts dependence as a historical fact while rejecting reciprocity. Dependence may give others leverage or grounds for exclusion. Whether it creates a reason for him still requires a principle connecting benefit, membership, or relationship to support.

The case cannot force a theory to accept aim-independent reasons. It can force the theory to state what disappears without them.

Kest returns

The coherent tyrant from Chapter 13 understands the new proposal.

“You say their suffering counts against my action,” Kest says. “I say it does not. You call the relation fundamental. I call your attitude intense.”

What can the reasons-first theorist say?

First, Kest’s rejection is not decisive. Disagreement does not erase facts. A person can reject valid inference, reliable evidence, or the existence of a distant country.

Second, Kest’s motivational stability is irrelevant if reasons need not motivate. His coherence shows only that instrumental rationality can serve cruel ends.

Third, the victims’ experiences are not projections. Their pain is a natural fact with a point of view at its center. The dispute concerns whether that fact has favoring force for Kest, not whether the suffering occurs.

Fourth, Kest can be challenged for inconsistency if he demands that others treat his suffering as reason-giving while denying the reciprocal claim. If he accepts that nobody has aim-independent reason to relieve *his* pain, the inconsistency route closes. His position remains brutal but structurally coherent.

The final dispute is now visible. The theorist affirms a fundamental support relation; Kest denies it. More description of the pain need not settle the normative relation, though vivid evidence may move actual people. More coherence need not generate concern. A social coalition can defeat Kest without proving he recognized a reason.

This does not show the theorist is wrong. It shows that assertion, phenomenological obviousness, and ordinary practice have not yet provided a non-question-begging method for establishing the relation to a fully informed denier. Part IV will examine possible knowledge of normative facts.

Objections that do not settle the matter

A fair audit must reject several shortcuts.

“Objective reasons are strange”

Strangeness is not an argument. Many true relations are unlike ordinary physical objects. The objection must specify an incoherence, explanatory failure, causal problem, or epistemic defect.

“Reasons cannot move an indifferent agent”

The view does not claim that reasons are forces. Support and motivation were separated in Chapter 12. Inertness is expected in a resistant agent.

“Cultures disagree”

Disagreement can create an epistemic burden, not a deductive refutation. Cultures disagree about medicine and astronomy too. The relevant question is whether a method predicts correction or merely tracks formation.

“Reason talk just expresses emotion”

That is a semantic proposal from Chapter 3, not an observation that proves itself. It may explain practical language, but it must account for the apparent truth conditions and disagreement practices of reason claims.

“Suffering obviously matters”

This reports powerful seeming and may identify a basic truth. It does not by itself discriminate objective recognition from near-universal aversion and compassion. The distinction is precisely what the Candidate System must examine.

Take the fork

Purpose. Record whether the Candidate System accepts aim-independent practical reasons and state the strongest resulting clause; either fork must preserve its own costs rather than borrowing the other’s advantages.

The notebook now divides.

If you accept aim-independent practical reasons

Write the strongest clause you can:

Fact _____ counts in favor of action _____ for agent _____ independently of _____ because _____.

Then specify:

- **source:** fundamental relation, objective value, standing, rational agency, or other;
- **application:** which agents are reached and why;
- **scope:** which facts and actions bear the relation;
- **strength:** how reasons compare and aggregate;
- **defeat:** which competing reasons can override them;
- **access:** how agents identify the relation and correct error;
- **motivation:** whether recognition normally motivates and what failure implies.

Apply the clause to Iona, Jalen, and Kest. If it applies by definition, explain why the definition corresponds to a truth rather than a stipulation. If the relation is fundamental, mark epistemic access as the next debt rather than disguising it as self-evidence.

If you reject aim-independent practical reasons

Rewrite criticism without borrowing their grammar:

- “Kest has a reason not to dominate” becomes _____.
- “Jalen should surrender the antidote” becomes _____.
- “The soldier’s pain makes a demand on Iona” becomes _____.

You may say:

- *We abhor and oppose Kest’s conduct from values we share.*
- *Jalen’s refusal marks him as dangerous to people who value mutual protection.*
- *We will seize the antidote under a public rule designed to protect vulnerable lives.*
- *The soldier’s plea expresses an urgent stake that compassion and reciprocity make motivationally powerful for us.*

These are substantial claims. Do not append “and therefore Kest objectively has reason to agree” unless you restore the bridge.

Then ask what is lost: universal criticism from the agent’s standpoint, an approval-independent duty, or only a particular vocabulary? Ask what is retained: values, coordination, resistance, protection, law, emotion, and public justification among agents with relevant concerns.

Neither fork may borrow the advantages of the other without paying its debts.

The sixteenth debt

Add this entry:

Debt 16 — Aim-independent support: *What makes a fact count in favor for an agent who rejects every connected aim, how are the relation’s application and weight determined, and how can recognition be distinguished from projection?*

The bridge may be real. The Candidate System has not demolished it. It has identified the load it must bear.

Most societies do not wait for agreement about this bridge. They classify conduct, teach norms, reward conformity, condemn violations, negotiate rules, and use law. These practices create powerful reasons in familiar senses: motives, risks, commitments, and terms of membership.

The next chamber is crowded. Everyone is voting, praising, shaming, threatening, and declaring that consensus has settled the matter.

We will ask exactly what it has settled.

Chapter 17 — From Approval to Obligation

Every resident of North Quay agreed that the old harbor song was forbidden.

The song commemorated a defeated rebellion. Singing it in public brought dismissal from employment, a heavy fine, and near-universal contempt. Parents warned children that only traitors carried the melody. The rule had stood for eighty years.

Lio, alone in a soundproof room, sang one verse.

No recording existed. No one heard. No future behavior changed. Lio felt neither guilt nor fear.

What did North Quay's approval establish about the act?

It established a great deal. The song violated a local norm. Public performance risked punishment. Singing could express defiance, threaten membership, recall a political identity, and motivate resistance. The prohibition shaped expectations and conduct across generations.

Did approval establish that the secret song was objectively wrong?

That conclusion requires another step.

Social life surrounds us with rules that are real in every ordinary sense. Money, offices, borders, licenses, games, contracts, and ceremonies depend partly on shared recognition. A person can be legally convicted, professionally disqualified, socially excluded, or ineligible under a rule whose existence is constituted by collective practice.

The mistake is not taking social facts seriously. It is allowing one kind of social authority to impersonate every other kind.

Where this chapter is headed. We will separate personal approval, commitment, social pressure, law, legitimacy, conditional reasons, and objective obligation. These sources can converge on one rule without becoming the same relation or acquiring one another's authority.

Seven columns, not seven heights

Purpose. Give each source of support its own column before considering alignment; the exercise classifies relations and does not rank their practical importance.

Create seven columns in the Candidate System notebook. They do not form a scale from trivial to profound. A legal sanction may be more consequential than a private commitment. A conditional reason may be stronger for an agent than social pressure. The columns identify different sources and implications.

Personal commitment

An agent adopts a rule, value, identity, or project. Violation can conflict with integrity, frustrate plans, or betray the agent's own standard.

Personal commitment can guide in secrecy. It does not automatically bind people who never adopted it.

Shared code or convention

A group uses a rule to classify conduct and coordinate expectations. The rule may constitute an activity: how property is transferred, which side of the harbor channel vessels use, what counts as a valid vote.

Conventions can solve coordination problems even when an alternative convention would work equally well. Their validity within the practice does not prove that the practice itself is morally justified.

Social pressure

Approval, shame, praise, ridicule, reputation, belonging, and exclusion motivate conduct. These responses can sustain cooperation or enforce cruelty.

Pressure explains behavior and creates prudential stakes. Its intensity is not a truth meter. A unanimous mob is not an additional premise.

Sanction or law

An institution declares a rule and backs it with authorized force: fines, imprisonment, license loss, confiscation, or other penalties.

Law can settle public action, protect expectations, and create legal statuses. It also gives agents prudential reasons to comply when they wish to avoid sanctions. Whether the law is legitimate, morally sound, or obedience-worthy in secret remains a further question.

Limited role or agreement

An agent accepts a role, promise, contract, or reciprocal scheme under conditions that generate special responsibilities. This column records the commitment structures from Chapter 14 and procedural legitimacy from Chapter 15.

The reasons are limited by entry, scope, fairness, content, and exit. A chess rule binds players within the game. It does not require passersby to move diagonally.

An added foundation

A social rule may be supported by something beyond approval: reduction of suffering, protection of agency, fair terms, divine authority, objective value, or aim-independent reasons.

This column is not itself a foundation. It is a forwarding address. Name the added ground and carry its debts from the relevant chapter.

Approval-independent duty

The strongest claim says the rule is morally binding regardless of actual acceptance, pressure, sanction, role, or agreement. Universal rejection would not erase it.

This is the original objectivity aspiration. A rule reaches this column only through an added account of stance-independent status and practical authority. It does not arrive by accumulating enough entries in the other six.

One million approvals remain approvals. They may produce overwhelming political power and reliable coordination. Multiplication does not change the type of relation.

Alignment is not ascent

The seven columns can support the same rule at once.

Consider the harbor rule requiring all vessels to pass on the right. Iona personally endorses it. The maritime community treats it as conventional. Captains praise compliance and condemn reckless deviation. Port law imposes fines. Licensed pilots accept a professional duty. Avoiding collisions gives every safety-minded captain a conditional reason. An objective moral account might add a duty not to impose lethal risk.

The rule is supported from many directions. Does the pile of support prove the last claim?

No. It makes the rule stable, intelligible, enforceable, and attractive to nearly every actual participant. Those are powerful achievements. The approval-independent duty still depends on its own foundation.

This is the **accumulation fallacy**: commitment plus convention plus pressure plus sanction plus agreement is treated as though enough stance-dependent support changes category and becomes stance-independent.

Alignment can nevertheless matter as evidence. If independent groups repeatedly converge on a rule because it reliably prevents visible harm, their convergence may indicate that the rule serves widely shared interests. If procedures with different biases reach the same result, confidence can increase. If every society prohibits an act, that pattern calls for explanation.

But social sources are rarely independent instruments pointed at an external moral property. Groups share human psychology, material constraints, histories, and influence. Convergence may track a common vulnerability, a successful authority, a useful convention, or an objective truth. The pattern alone does not choose the explanation.

Nor does universal agreement make a rule constitutively stance-independent. If everyone approves the right-side rule, the convention becomes universal while its content remains selected by coordination. A left-side rule could have worked. Its universality concerns reach, not foundation.

When several columns align, write two sentences:

Practical conclusion: *Given these commitments, institutions, consequences, and sanctions, agents have _____ support for compliance.*

Objective conclusion, if claimed: Independently of those stances, agents have _____ because _____.

The first sentence may be decisive for policy and action. Do not weaken it by pretending it needs the second. Do not manufacture the second by repeating the first in a deeper voice.

What the forbidden song establishes

Return to Lio's verse.

Classification

Under North Quay's code, the act is prohibited. This statement can be objectively true as a report about the code: given its actual rules, the song counts as a violation.

That is **objectivity about a social fact**, not yet an objective moral verdict. "The move is illegal in chess" can be true independently of what one player wishes while the rule remains constituted by the practice.

Motivation

Public sanctions, identity, fear, and belonging motivate most residents not to sing. Lio lacks those motives in the sealed room.

Prudence

If discovery were possible, the penalties would give Lio a prudential reason for caution relative to interests in employment and safety. With detection impossible and no other effect, that reason is absent by stipulation.

Legal status

The act may still be illegal even when undetected. Legal classification does not require enforcement in each case.

Membership

Singing may express rejection of the community's identity or solidarity with the defeated group. Whether that meaning exists in private depends on Lio's intention and the symbol's role.

Objective authority

Nothing in widespread approval, age, punishment, or social meaning alone shows that secret singing has approval-independent moral status. The prohibition needs an added foundation.

Perhaps the song incites violence, humiliates victims, or renews a domination practice. Those facts could support restriction under the Candidate System's values. The pressure case holds such effects absent so that popularity cannot borrow them.

The two disclosure towns

Across the bay lie Candor and Shelter.

In Candor, physicians must disclose every terminal diagnosis immediately and directly. Residents regard nondisclosure as a profound violation of autonomy.

In Shelter, patients designate family decision-makers who may receive a diagnosis first and control the timing of disclosure. Residents regard abrupt individual disclosure as cruel abandonment.

Both towns have stable expectations. Both practices are publicly known. Most residents endorse their own system. Visitors receive an explanation at registration.

A physician moves from Candor to Shelter. What reasons change?

- The local code changes classification.
- Patient expectations and reliance change.
- Professional agreement changes role responsibilities.
- Legal sanctions change prudential stakes.
- Likely welfare effects may change because people understand illness through different relationships.
- The symbolic meaning of direct disclosure changes.

These are not superficial facts. Context can make different actions appropriate without making every value relative. If a patient in Shelter explicitly asks for immediate private disclosure, autonomy and local practice may conflict. If families conceal treatable conditions for financial gain, the practice may fail by its own caring rationale.

The existence of opposite codes does not show that no objective answer exists. It shows that the action-guiding reasons include expectations, relationships, and institutions partly created by convention. An objective theory could still judge one town better. It must do so through grounds beyond the mere fact that the other town disagrees.

The reformer before the holiday

Asha secretly taught engineering to girls when Port Meridian's academy excluded them. The rule was legal, religiously defended, professionally enforced, and approved by nearly everyone. Asha lost her position. Decades later, the academy repudiated the exclusion and named a laboratory after her.

Was Asha wrong before the holiday celebrating her began?

An approval-independent account says no. Equal standing, agency, and access to education counted against exclusion even when the society failed to recognize them. Asha responded to reasons that applied all along.

A convention-based account has several options.

Criticism from unacknowledged existing values

The academy claimed to value talent, learning, fairness, and the city's prosperity. Asha could show that exclusion conflicted with those commitments and rested on false beliefs. This is internal criticism.

Criticism from a rival or emerging code

Asha and her students formed a community committed to equal education. They opposed the dominant code from their own shared values and eventually changed institutions.

Criticism through idealized construction

The actual rule could be condemned by the fair procedure of Chapter 15 because excluded students lacked representation and existing power distorted the terms.

Objective criticism

The exclusion violated an approval-independent duty. This preserves the strongest sense in which society itself was mistaken, not merely defeated or inconsistent. It carries the grounding debt from Chapter 16 or Chapter 18.

The first three forms of criticism can be powerful without the fourth. They can expose hypocrisy, error, exclusion, and domination; organize resistance; and build better institutions. What they do not automatically establish is that every coherent defender had an aim-independent reason to agree.

Calling change **moral progress** can likewise mean several things:

- improvement by our present values;
- greater consistency with the society's stated values;
- expansion of standing or concern;
- reduced suffering and domination;
- movement toward a fairer procedure;
- discovery of objective moral truth.

The first five can be supported without proving the sixth. This does not make them fraudulent uses of *progress*. It makes their standards explicit.

The harmless outsider

The island of Vey marks communal belonging through a dawn ritual. Participants touch the central stone and recite the names of founding families. The ritual harms no one, consumes few resources, and is deeply meaningful to residents.

Sorin arrives, trades peacefully, obeys safety laws, pays required fees, and declines the ritual. Sorin does not insult it; she simply rejects symbolic conformity.

Vey's council orders her to participate or leave.

What does the ritual's acceptance establish?

- It establishes the practice’s social meaning.
- It creates expectations for members who voluntarily undertake the identity.
- It may supply residents with relational and symbolic reasons to participate.
- Enforcement creates prudential reasons for Sorin if she wants to remain.
- The council may have legal authority under local law.

Does it establish that Sorin has a moral duty to recite?

If membership in Vey is voluntary and the ritual was a clear entry condition, a limited agreement may apply. Yet exit may cost livelihood and relationships. If Sorin is merely a visitor, the membership claim is weaker. If the ritual has no effect on others, the case for coercion cannot rely on harm prevention.

The council may say shared identity decays when exceptions are allowed. That is an empirical and evaluative claim: how much decay, why is uniformity necessary, and what weight does belonging have against conscience? The fact that residents care intensely supplies a real stake. It does not settle the comparison or create approval-independent authority.

The outsider case prevents an easy equation between harmlessness and social acceptance. Communities need practices. Individuals need room not to perform every meaning others cherish. Policy must compare those stakes; moral vocabulary cannot make one disappear by calling it sacred or antisocial.

Classification without inflation

Purpose. Apply the seven-column distinction across four cases so code, motivation, prudence, law, role, added foundation, and approval-independent duty remain separately visible.

For each of the four cases, complete this table:

Rule or act	Classification under code	Motivating forces	Prudential consequences	Legal status	Role or membership reason	Added foundation	Approval-independent duty
Secret song							
Disclosure practice							
Asha's teaching							
Sorin's refusal							

Use these rules:

1. Do not enter a social penalty in the objective-duty column.
2. Do not enter an objective duty in the legal column unless law recognizes it.
3. Do not treat role acceptance as universal agreement.
4. Do not call a convention arbitrary merely because another convention was possible.
5. For every added foundation, cite the Candidate System clause and its unpaid debt.

Then choose a live public controversy. Separate six questions:

- What does the existing rule classify?
- What motives does it create?
- What costs make conformity prudent?
- What can authorized institutions enforce?
- Which commitments or roles apply?
- What, if anything, makes the rule morally authoritative independently of acceptance?

This separation may reveal that opponents agree about law but not morality, agree about social meaning but not coercion, or agree on a policy settlement while denying that it expresses objective truth.

Force, legitimacy, and duty

Power can make an action happen. Law can authorize officials within a system. Legitimacy can make the exercise of power defensible under fair terms. Duty can claim support for the governed agent.

These relations may coincide in a well-ordered institution. They can also split:

- a robber has force without legal authority;
- a lawful official may exercise authority illegitimately;
- a fair traffic rule may create a strong coordination reason without being uniquely correct;
- an objectively valid duty, if such duties exist, may be illegal under an oppressive state.

The practical importance of law does not depend on pretending these are one thing. Indeed, distinguishing them makes reform possible. If legality guaranteed morality, Asha's teaching could not expose injustice until the law had already changed.

The seventeenth debt

Add this entry:

***Debt 17 — Social authority:** What do approval, convention, pressure, sanction, role, and agreement each establish, and what added ground—if any—turns their classifications or incentives into approval-independent obligation?*

The crowded chamber has become quieter. Approval can create a code. Convention can constitute practices. Pressure can motivate. Law can classify and sanction. Roles and

agreements can generate limited reasons. Fair institutions can settle action and command legitimacy under stated values.

None of these achievements is imaginary. None becomes an approval-independent duty merely by being announced in the moral register.

The Candidate System has one final route before the authority inspection. Perhaps the source lies beyond human approval altogether: in God's will or nature, the purposes built into living things, the flourishing characteristic of our species, the requirements of rational agency, the standards constitutive of action, or truths evident to reason itself.

These are not desperate decorations added after secular options fail. They are major attempts to explain why morality has the universality, authority, and depth people experience it as having.

The last chamber in Part III resembles a temple, a biological station, and a geometry classroom sharing a lease.

Chapter 18 — God, Nature, and Rational Necessity

The final chamber in Part III had three doors.

The first opened into a sanctuary. The second opened into a garden and biological station. The third opened into a geometry classroom.

Above all three was the same inscription:

AUTHORITY DOES NOT BEGIN WITH YOUR APPROVAL.

This is the promise of the candidates we will examine here. Divine command grounds duty in a perfect authority. Natural law grounds practical order in the kinds of beings we are and the ends through which we flourish. Rational accounts ground requirements in the structure of agency, consistency, or action itself. Self-evident principles ground morality in truths reason can recognize without derivation from desire.

These approaches differ greatly. They share an ambition: to locate the source of moral authority somewhere deeper than individual preference, social convention, or negotiated settlement.

They deserve more than the standard caricatures.

The religious view is not refuted by saying that an invisible person issues arbitrary orders. Its strongest forms deny arbitrariness and connect command to a necessarily perfect nature.

Natural law is not the claim that everything found in nature is good. Disease and predation are natural. The view asks about the intelligible goods and functions of living, social, rational beings.

Rational necessity is not simply “be logical.” Its strongest forms argue that agency contains commitments to universal standards, autonomy, or respect for other agents.

Self-evidence is not “I feel strongly.” It is the proposal that some truths are grasped noninferentially because every chain of justification must eventually reach premises not derived from earlier ones.

We will apply the steelman rule: insight first, sophisticated form second, best repair third, and only then pressure.

Where this chapter is headed. We will give divine authority, natural law, human flourishing, rational necessity, constitutive agency, and self-evident truth their strongest hearing. The aim is to see exactly which part of objective authority each foundation supplies and which part it leaves for the access audit.

The divine command

Begin with the strongest motivating insight.

If reality has a personal creator who perfectly knows every fact, understands every conscious life from within, wills the good without corruption, and rightfully governs creation, then morality has an attractive source of unity. Human judgment is partial and self-serving. A perfect divine will is not. Duty can reach every agent because every agent stands within the creator's jurisdiction, not because each signed a contract.

On a **divine command** account, an act is obligatory because God commands it, forbidden because God prohibits it, and permitted when divine law allows it. The command is not mere advice from a better-informed observer. It has authority.

This account explains several aspirations:

- **stance independence:** human rejection does not repeal divine law;
- **scope:** a creator's authority may reach all created persons;
- **authority:** commands are action-directing, not passive descriptions;
- **unity:** diverse duties can arise within one ordered purpose;
- **moral confidence:** oppression remains wrong when every earthly institution approves it;
- **accountability:** secret conduct remains within the divine relation.

The familiar pressure begins with a choice.

Does God command an act because it is good, or is it good because God commands it?

If God commands it because it is already good, goodness appears to provide a standard independent of command. Divine wisdom identifies moral truth, but command may not create it.

If an act becomes good solely through command, a worry about arbitrariness appears. Could cruelty become obligatory if commanded? Saying "God would never command cruelty" introduces a standard distinguishing cruel from good commands.

The strongest religious reply does not choose either crude horn. It turns from isolated commands to **divine nature**.

The divine nature repair

On this account, God's nature is necessarily loving, just, wise, faithful, and good. Commands express that nature. God does not consult an external rulebook, and goodness is not produced by a whim. The source and standard are unified in a necessarily perfect being.

This is a substantial repair. It explains why horrific commands are impossible rather than merely unlikely. It connects obligation to stable character rather than arbitrary will. It can also explain why moral truths possess both evaluative content and personal authority.

Now the Candidate System asks four questions.

What does *good* contribute to the description of divine nature?

If “God is good” means only “God has God’s nature,” the claim cannot help distinguish love from cruelty. If *good* has substantive content—love protects beings, justice treats them rightly, faithfulness honors commitments—that content appears intelligible independently enough to identify which nature is perfect.

A defender may say goodness is identical with the divine nature, as water is identical with a chemical structure even though people recognized water before chemistry. We can possess an imperfect concept of goodness and discover its metaphysical identity with God.

This is coherent as a proposal. The identity needs argument and a method for determining which substantive moral judgments accurately reflect the nature. It is not established by defining one term with the other.

Why does creating confer moral authority?

A creator has a unique relation to creation. But human makers do not possess unlimited moral authority over what they make. Parents cause children to exist and do not own them. Engineers create conscious machines, if such machines become possible, without thereby gaining permission to torture them.

Divine authority may instead rest on perfect wisdom, goodness, sustaining power, rightful order, gratitude, or the creature’s good. Each is more promising than raw authorship. Each uses evaluative ideas whose reason-giving force must be explained.

How is the command identified?

Port Meridian contains sincere believers who accept the same God and disagree about war, gender, punishment, property, sexuality, animals, and political authority. They appeal to different texts, traditions, experiences, and interpretations.

This disagreement does not show that no divine command exists. It shows that accessibility requires a method. Which revelation is authentic? Which passage is literal, contextual, superseded, metaphorical, or corrupted? How do conscience and reason check interpretation? If independent moral judgment is used to reject an apparent command as unworthy of God, that judgment is doing epistemic work.

What reaches the nonbeliever?

If God exists and has rightful authority, disbelief may not erase it. But the system must show the authority relation, not merely assert jurisdiction. Does a person have reason to

obey because disobedience is punished, because God is good, because obedience perfects the creature, because gratitude is owed, or because divine command constitutes duty?

Punishment supplies prudential force. Perfect goodness may supply trust. Creaturely fulfillment connects with interests. Constitutive duty restates the strongest claim. The routes should remain separate.

The divine candidates therefore remain serious possibilities. Their success requires establishing the being, the relevant nature and authority, a reliable interpretation method, and the bridge from creator or perfection to reasons for every agent. A moral argument alone cannot assume all four.

Natural law: goods built into a kind of life

Walk through the second door.

A living thing is not merely a pile of matter. Hearts pump blood. Roots draw nutrients. Wolves hunt in packs. Human beings reason, form attachments, use language, cooperate, create, reproduce, care for dependents, seek understanding, and organize lives across time.

Natural-law approaches begin from the idea that practical reason can identify genuine human goods within this form of life: health, knowledge, friendship, agency, play, family, community, integrity, perhaps relationship with the divine. Conduct is good when it participates intelligently in these goods and respects the order among them; it is defective when it destroys or radically frustrates them.

This view is not committed to copying whatever occurs in nature. A tumor is natural in the causal sense and defective relative to the organism's organization. A broken wing is not a morally alternative wing. Function allows standards of better and worse that do not depend on a vote.

The approach has major strengths.

- It connects values to facts about embodied, vulnerable, social beings.
- It explains why some conditions are genuinely harmful despite adaptive preference.
- It accommodates multiple goods rather than reducing life to pleasure.
- It treats practical reason as the intelligent pursuit of flourishing, not obedience to arbitrary rules.
- It can criticize both individual desire and social convention by reference to human capacities and needs.

The pressure lies in the movement from function to practical authority.

Function is relative to a system

A heart functions to circulate blood within an organism. We can assess a heart as defective without a moral premise. But "this heart functions badly" and "an agent ought to repair every heart" are different statements. Medical goals, consent, resources, competing claims, and responsibility enter the second.

Human capacities are plural

Aggression, coalition, deception, care, curiosity, hierarchy, reciprocity, and domination all appear in human life. Frequency does not identify proper function. A natural-law account needs a principled distinction between capacities that constitute flourishing, tendencies to regulate, and pathologies.

A good specimen is not automatically a good life

Species-typical functioning can marginalize people whose bodies or minds differ from statistical norms. A person without sight can live an excellent life. A celibate person need not realize reproductive function personally. Goods may be participatory, social, optional, plural, or instantiated through diverse forms.

Sophisticated accounts recognize this. They do not demand identical performance of every capacity. But the flexibility requires an account of which goods are basic, how participation works, and when nonparticipation is defect rather than a different life.

Flourishing does not yet bind the rejecter

Suppose an informed agent says, “I understand the form of life you call flourishing. I prefer a shorter, isolated, aesthetically intense life. I reject family, community, and species continuation.”

We can show costs, conflicts, and biological needs. If the agent coherently accepts them, why must human flourishing as a whole govern? Calling the rejected life objectively defective may be true within the account, but its practical authority is the point under dispute.

Natural law can reply that choice itself aims at an apparent good and that rational agents cannot escape the basic goods without relying on them. The solitary aesthete still uses life, knowledge, agency, and perhaps beauty. This constitutive reply moves us toward the third door.

Species flourishing without cosmic purpose

A more naturalistic account keeps flourishing and removes divine or teleological order. Humans have empirically identifiable needs, capacities, vulnerabilities, and social conditions. Nutrition, safety, attachment, agency, learning, and cooperation generally enable lives people can inhabit and endorse. Torture, starvation, extreme isolation, and domination reliably damage them.

This approach offers a rich basis for policy and criticism. We can study what trauma does, which institutions support agency, how children develop, and which resources protect health. The account need not wait for metaphysical certainty before improving conditions.

Its strongest claim is that flourishing facts are objective. Whether someone has clean water, cognitive development, secure attachment, or freedom from chronic pain is not created by opinion. People can misjudge their condition or adapt preferences to oppression.

Agreed. These are objective descriptive and evaluative facts relative to a conception of functioning and welfare. The remaining question is whether every agent has an aim-independent reason to promote this flourishing, how conflicts among different people's flourishing are resolved, and why the selected conception has objective moral priority.

Consider Kest. He uses knowledge, health, planning, culture, friendship with his inner circle, and political agency to dominate others. In several dimensions he flourishes. In relational and moral dimensions the account calls him defective. What makes equal concern or nondomination part of *his* flourishing rather than a value used to judge it?

One answer is that domination corrupts relationships and practical reason even when Kest endorses it. That may be correct. It places substantive standards inside the definition of human flourishing. Those standards must be defended without saying merely that a morally flourishing person behaves morally.

Rational consistency

Enter the geometry classroom.

Rationality plainly constrains some action. If Iona aims to open the gate, knows that turning the wheel is necessary, has no competing aim, and refuses to turn solely while insisting she will open it, something has gone wrong. Contradictory intentions and rejected necessary means make agency malfunction on its own terms.

Perhaps moral requirements arise from this structure. Several increasingly ambitious proposals are available.

Means-end consistency

Agents have reason to take effective means to their ends, avoid contradictory plans, and update when evidence changes.

This is indispensable for deliberation. It does not select the ends. Kest can pursue domination with excellent instrumental rationality.

Universal consistency

An agent should act only on principles they can apply consistently to relevantly similar cases. Secretly exempting oneself while relying on everyone else's compliance is structurally unstable.

This exposes hypocrisy and free riding. But a principle can be universal and cruel: "Every stronger group may dominate every weaker group." Kest may accept its application if he later becomes weak. Consistency removes identity-based exception; it does not guarantee humane content.

Reciprocity among agents

To claim authority for one's choices while denying every other agent's comparable authority may be inconsistent. Recognizing oneself as a source of reasons seems to require recognizing other agents as sources too.

This is stronger. It connects agency to equal status. The pressure question is what feature supports the comparison. If it is rational agency, why must an agent value the agency of others as they value their own? The inference from "my ends matter to me" to "every agent's ends matter to me" needs a universal premise. Formal similarity can establish that no descriptive difference has been supplied. It does not, by itself, create concern.

Respect for agency as a condition of justification

Perhaps action is not fully justified unless its principle can be offered to every affected agent as a free and equal co-author. Domination fails because it treats another source of reasons as a mere instrument.

This gives rationality moral substance. It can ground rights, consent, and fair procedure. It also builds "free and equal," "source of reasons," and "justifiable to each" into the standard. These may be correct objective principles. Calling the full package *rationality* does not reduce its grounding debt.

Rational arguments are most decisive when an opponent asserts mutually inconsistent claims or relies on an exception they cannot describe without identity. They are less decisive against a person who openly accepts a universal but cruel principle. We may condemn the content. Logic alone will not supply the condemnation.

Standards constitutive of action

A deeper proposal says action is not mere behavior caused by desire. To act is to pursue something under the appearance of the good, answer to reasons, and understand oneself as justified. Practical agency has constitutive standards just as belief aims at truth and games have rules.

An agent who rejects every reason may therefore fail to count as an agent, much as a person moving pieces without regard to chess rules is not playing chess.

This explains why reasons cannot be optional additions to agency. It also clarifies internal criticism: Kest must represent domination as good in some respect—effective, glorious, stabilizing, expressive of power. He acts under a practical description.

Two questions remain.

First, does action constitutively aim at **some apparent good for the agent**, or at **objective good impartially considered**? The former is compatible with cruel ends. The latter contains the moral conclusion inside the account of action.

Second, what follows if Kest rejects the allegedly constitutive standard? If we say he is no longer acting, ordinary language protests: he plans, chooses, revises, and executes. Re-

defining agency to exclude successful counterexamples may preserve the theory by changing the subject.

Constitutive standards can explain what one must do *to participate in a practice*. Their universal moral reach depends on showing that no agent can intelligibly escape the practice and that its standards contain the substantive duties claimed.

Self-evident axioms

The geometry classroom contains one final possibility. Every proof begins somewhere. If each moral premise requires another moral premise, explanation either regresses forever, circles, or stops. Perhaps it stops with self-evident truths:

- suffering counts against an action;
- like interests deserve like consideration;
- persons must not be used merely as instruments;
- a valid promise normally gives a reason to perform;
- arbitrary domination is wrong.

A truth is **self-evident** in the relevant sense not because everyone instantly agrees, but because adequate understanding supplies justification without inference from a further proposition. Mathematical and logical truths are often proposed as analogies.

This view has a legitimate answer to the demand “prove every premise”: not all knowledge is inferential. A person can fail to see a self-evident truth through confusion, bias, immaturity, or insufficient attention.

The Candidate System must still answer:

- Which propositions are self-evident?
- What level of understanding is adequate?
- Why do informed thinkers report incompatible seemings?
- How are conflicts among proposed axioms resolved?
- What method distinguishes insight from culturally trained confidence?
- How does an evaluative axiom acquire practical authority for a rejecter?

Calling a claim self-evident reports an epistemic status. Calling its content a reason supplies the authority claim. Neither label should do both jobs unnoticed.

Self-evidence may be the correct stopping point. If so, Part IV must explain recognition and correction. At this stage, the Candidate System possesses a candidate foundation, not a public proof that the foundation is there.

The sincere interpreters

Gather seven delegates:

1. one grounds duty in divine command;
2. one in necessary divine goodness;
3. one in natural human goods;
4. one in empirically understood flourishing;
5. one in universal rational consistency;
6. one in the constitutive standards of agency;
7. one in self-evident reason relations.

All are intelligent, informed, sincere, and alert to obvious caricatures. They agree that torture for amusement is wrong. They disagree about reproductive duties, punishment, animal standing, property, self-sacrifice, and the authority of political institutions.

For each disputed issue, their foundations require interpretation. Divine goodness must be given content. Natural goods must be ranked. Flourishing must be specified. Universal principles must be formulated. Constitutive standards must be extended to other agents. Self-evident claims must be selected and reconciled.

Convergence on familiar cases can arise because shared human aversions, attachments, vulnerabilities, and social needs influence every interpretation. That convergence is practically valuable. It does not tell us whether the common judgment was detected by seven routes or supplied to each route as a selector.

Persistent disagreement does not disprove the foundations. It places pressure on accessibility and determinacy. A successful objective system must say what correction looks like when equally sincere interpreters diverge.

Can the foundations support one another?

Few mature accounts rely on one isolated source. A theistic natural-law view may combine divine nature, created human goods, practical reason, conscience, revelation, and a community of interpretation. A secular realist view may combine fundamental reason relations, objective welfare facts, rational consistency, reflective judgment, and fair public construction.

Combination can be a strength. Different elements can perform different jobs:

- a divine or realist metaphysics supplies stance-independent truth;
- human nature and flourishing supply substantive content;
- rational principles organize consistency and reciprocity;
- conscience or intuition supplies initial access;
- tradition preserves accumulated interpretation;
- empirical evidence corrects factual application;
- public procedure coordinates disagreement.

No single element must do everything. The objectivity project requires an organized account, not a solitary magic premise.

The danger is **circular reinforcement**. Imagine this chain:

God is good because divine acts match the true moral standard. The moral standard is true because God's nature grounds it. We know the nature through conscience. Conscience is reliable because a good God designed it. We know God is good because conscience recognizes divine goodness.

Every link may belong to a larger defensible argument. Written only as this loop, however, the parts certify one another without independent support.

A secular version can do the same:

Flourishing is reason-giving because rational agents would choose it. Rational choice is correct when it tracks objective reasons. Objective reasons are identified by judgments about flourishing.

Again, the concepts may fit together as a theory. Fit is not independent warrant.

To evaluate a combined foundation, assign separate jobs:

1. **Existence:** What makes the normative fact or reason relation real?
2. **Content:** What determines which actions, goods, or statuses it concerns?
3. **Application:** Why does it reach this agent and case?
4. **Weight:** How are conflict and priority determined?
5. **Access:** Which method reveals the content, and how is that method checked?
6. **Correction:** What would show that an interpretation was wrong?
7. **Motivation:** Why do some agents respond and others remain unmoved?

Draw an arrow whenever one component supports another. Then ask whether any closed loop has an entry point supported without relying on the conclusion circulating through it.

This standard does not demand that normativity be reduced to chemistry, desire, or social approval. A reasons-first account may legitimately place **existence** at a fundamental stopping point. A divine account may place it in a necessary being. The remaining jobs still need answers. “Fundamental” answers where one explanatory chain stops; it does not automatically supply detailed content, comparison, or access.

Nor does failure of one combined argument establish that no combination could succeed. The bounded conclusion is narrower: the Candidate System may count a foundation as earned only to the degree that its components jointly perform their assigned jobs without circularity or silent importation.

The external-foundation audit

Purpose. Compare each proposed foundation by source, content, authority, access, reach, conflict rule, and remaining debt; a completed row is an analysis, not a refutation.

For each candidate, complete one row:

Candidate	Source of status	Source of author- ity	Method for con- tent	Reach to rejecter	Conflict rule	Natural stopping point	Remaining debt
Divine command							
Divine nature							
Natural law							
Species flourishing							
Rational agency							
Constitutive action							
Self-evident axioms							

Then apply five pressure tests.

The arbitrariness test

Would a different command, nature, function, or axiom reverse the verdict? If not, what constrains the source?

The interpretation test

How do sincere agents identify the source's content and resolve disagreement without quietly relying on the verdict being tested?

The function-to-value test

When a capacity, end, or practice is described, what establishes that realizing it is good or reason-giving for the relevant agent?

The formal-to-substantive test

Does consistency alone produce the humane result, or has equal standing, concern, or respect been built into the principle of rationality?

The rejecter test

What does the account say to an agent who understands the source and its alleged implications but denies its authority? Is the agent mistaken, irrational, outside the practice, punishable, imprudent, or simply opposed? State the mechanism, not only the label.

For each candidate, preserve what survives pressure. Divine accounts may offer a unified metaphysics of value and personal authority. Natural accounts connect evaluation to real needs and capacities. Flourishing supplies rich empirical guidance. Rational accounts expose inconsistency and self-exemption. Constitutive accounts explain internal standards of agency. Foundational axioms stop regress.

The pressure tests do not prove God absent, human nature normatively silent, rational agency empty, or every moral intuition false. They show that none of these labels, without further argument and accessible content, automatically completes the full objective system.

The eighteenth debt

Add this entry:

Debt 18 — The external bridge: How does a command, perfect nature, natural function, form of flourishing, requirement of rational agency, constitutive standard, or foundational axiom produce universal moral status and practical authority—and how is its content identified without importing the desired verdict?

We have reached the end of the authority maze. Every major corridor led somewhere real. Desire connects facts to aims. Reflection corrects error. Commitment creates reliance and identity. Fair procedure creates legitimate settlements. Aim-independent reasons state the objective relation directly. Social practices constitute rules and institutions. Divine, natural, and rational foundations offer deep explanations of status and authority.

The question is not whether these corridors are useful. It is whether any one route, or noncircular combination, has earned every property on the original specification sheet.

It is time to inspect the bridge.

Inspection Report III — The Missing Bridge

Purpose. Reclassify every authoritative word in the Candidate System by the strongest support it has earned, preserving practical achievements without inflating them into objective obligation.

Open the two-page constitution from Inspection Report II. Take every clause that contains *reason*, *ought*, *must*, *duty*, *authority*, *binding*, *required*, or *forbidden*. Assign it to the strongest column it has earned.

Column 1 — Commitment

The clause guides because an agent personally adopts a value, identity, project, or promise.

Write the entry and exit conditions. Do not generalize the reason to agents who lack the commitment.

Column 2 — Convention

The clause is valid within a shared code, practice, or coordination scheme.

State what the convention constitutes and whether alternative conventions would work. Do not infer that practice-relative validity is approval-independent moral truth.

Column 3 — Pressure

The clause is supported by approval, shame, praise, trust, reputation, belonging, or exclusion.

State who applies the pressure and what motive it creates. Do not convert motivational force into favoring force.

Column 4 — Sanction

The clause is enforced by law, punishment, institutional authority, or credible threat.

State the jurisdiction, legal status, evidence rule, and penalty. Do not infer legitimacy or duty from force alone.

Column 5 — Conditional reason

The clause counts in favor relative to an agent’s aims, informed aims, role, accepted procedure, or fair reciprocal commitment.

State the connecting aim or valid entry condition. If an idealized aim is used, cite the idealization standard.

Column 6 — Aim-independent practical support

The clause claims that a fact counts in favor for an agent regardless of every connected aim or acceptance.

State the reason-giving fact, agent, action, source, weight, defeat conditions, and proposed method of access. A clause belongs here only if the support relation itself has been defended, not merely renamed.

Column 7 — Objective moral obligation

The clause assigns approval-independent moral status and claims practical authority for every agent within scope.

State the foundation, semantic meaning, standing rule, value basis, generation rule, comparison method, exception structure, reason relation, and epistemic route. Every dependency must refer to an earlier clause rather than disappear into the word *moral*.

The final two columns are not automatically stronger versions of the first five. They make different claims. A law can govern action effectively while an unrecognized objective duty is politically powerless. A personal commitment can shape a life more deeply than an abstract principle. The audit classifies authority; it does not rank importance.

Build the authority matrix

Purpose. Place each clause across commitment, convention, pressure, sanction, conditional reason, aim-independent support, objective obligation, or promissory note so no column borrows another’s label.

Use this table:

Candidate System clause	Commitment	Convention	Pressure	Sanction	Conditional reason	Aim- independent support	Objective obli- gation	Promissory note
Standing rule								
Primary value								
Protected constraint								
Role duty								
Comparison rule								
Exception rule								
Punishment authorization								

Do not place a check in every supporting column merely because several kinds of force coincide. Describe the actual support. A promise may involve commitment, convention, pressure, legal sanction, and conditional reasons at once. None of those entries alone licenses the objective-obligation check.

For every check in Column 6 or 7, attach an **authority warrant**:

*Consideration _____ applies to agent _____ because _____. Its application does not depend on _____. We identify this relation through _____.
Conflict with _____ is resolved by _____. Error is corrected through _____.*

If you cannot complete a blank, move the claim to **Promissory note**. Do not delete it. The aim is an honest inspection, not a clean scorecard.

The adversarial bridge test

Purpose. Test the matrix against informed rejecters with different aims and allegiances; the imagined readers expose reach and hidden premises rather than standing in for every actual dissenter.

Give the matrix to four imagined readers.

The informed egoist

This reader understands every fact and accepts only personal flourishing. Which clauses apply through their aims? Which claim more?

The coherent tyrant

Kest has stable, informed, consistent ends and rejects equal concern. Which clause gives his victims' suffering support from his standpoint? Where is that relation grounded?

The principled dissenter

This reader values others but rejects your comparison rule, religious authority, or fair procedure. Can the system distinguish honest disagreement from irrationality without defining rationality as agreement?

The outsider

This reader never joined the practice, received no benefit, rejects the symbolism, and can defect secretly without harm. Which rules still reach them, and why?

For each challenge, use the anti-rigging options:

- **Accept the result:** acknowledge that the clause does not reach this agent.
- **Add a bridge:** state a new aim-independent relation or foundation and add its epistemic debt.

- **Narrow the claim:** retain commitment, convention, legitimacy, or coordination while withdrawing objective authority.

Inspection finding

Your matrix will probably be rich in the first five columns. That is not a failed constitution. It may contain durable commitments, humane conventions, effective protections, legitimate law, fair procedures, and reasons connected to nearly every actual person's deepest aims.

The final two columns are harder.

An aim-independent reason can be posited as fundamental, grounded in objective value, located in another being's standing, derived from divine perfection, read from natural goods, built into rational agency, or grasped as self-evident. Each proposal may be true. In this construction, each still carries at least one promissory note concerning its bridge, content, scope, weight, or access.

That is a provisional inspection finding, not yet the book's final verdict. A real moral order could remain undiscovered. Part IV asks whether any defensible method can distinguish discovery of such an order from projection, inheritance, intuition, convergence, or useful construction.

The Candidate System has reached the strongest authority claim available:

Some facts count in favor independently of every relevant stance.

Now it must answer a question that every map of an invisible bridge eventually faces:

How would we know?

PART IV



Ask How the System Could Be Known

A truth no one can reliably reach may still be true. It cannot yet do all the work this project promised.

Chapter 19 — Does Evidence Have Authority?

The first office in Port Meridian's epistemic quarter had a sign over the entrance:

ALL CLAIMS MUST PRESENT EVIDENCE.

Below it, in smaller letters, someone had written:

EVIDENCE MUST PRESENT ITS AUTHORITY.

This caused a queue.

The Candidate System has demanded evidence from moral claims. It has asked how standing is known, whether consequences are predicted reliably, what a fair procedure would produce, whether a divine command has been interpreted correctly, and how an intuition connects with objective truth. That demand has felt entirely proper.

But Part III ended with a challenge that now turns toward the inquiry itself. A fact may bear on whether a belief is true. What, if anything, makes that relation a **reason to believe**? Why should an agent follow evidence if truth is unwanted, useless, painful, or disconnected from every project?

This question is not a parlor trap designed to make inquiry impossible. It is an audit of the authority the book has been using. We can continue investigating while admitting that our warrant may be conditional: *if* you want accurate conclusions, *if* you accept the practices of critical inquiry, *if* evidence reliably indicates truth, then respond to it. The stronger claim—that evidence gives every competent agent an aim-independent reason to believe—needs its own defense.

Objective moral skepticism neither proves nor refutes objective epistemic reasons. The two domains may share a structure and differ in access. A person can be a moral anti-realist and an epistemic realist, a moral realist and an epistemic anti-normativist, a realist about both, or a realist about neither. Similarity creates a comparison, not a verdict.

We begin by separating the furniture.

Where this chapter is headed. We will ask why evidence should govern belief and whether that force resembles the practical authority morality claims. Reliability, truth, usefulness, coherence, and public procedure each explain something; the task is to keep their achievements distinct.

Truth, evidence, belief, and justification

A **belief** represents the world as being some way. It may be true or false.

Truth concerns whether the world is as the belief represents it. A true belief can be reached by accident. A false belief can be formed through an excellent method applied to misleading evidence.

Evidence is a consideration that bears on whether a claim is true. Smoke can indicate fire. A calibrated pressure reading can indicate a storm. A proof can entail a mathematical result. Evidence differs in strength, relevance, accessibility, and reliability.

Justification concerns whether a belief is responsibly or appropriately supported under an epistemic standard. Justification may depend on available evidence, a reliable process, intellectual duties, or another account. The word is normative; it does not merely report truth.

These distinctions let four combinations occur:

- true and well supported;
- true by luck;
- false but reasonably supported by misleading evidence;
- false and irresponsibly formed.

If we classify only outcomes, the lucky conspiracy theorist becomes an epistemic hero. If we classify only methods, we risk forgetting that inquiry aims to reach the world rather than perform respectable rituals around it.

Evidence can also **cause** a belief without justifying it. A threatening official's folder may cause a witness to believe the accused is guilty even when the folder is empty. Conversely, evidence can support a conclusion without causing belief when a person refuses to look.

The three-room distinction from Chapter 12 returns: causal influence, motivating consideration, and favoring support must not exchange badges.

Seven candidates for evidential force

Why let evidence govern belief? Seven candidates enter the office.

Truth-directed support

On the strongest realist account, evidence itself counts in favor because it bears on truth. The relation does not depend on anyone wanting accuracy, finding the truth useful, or joining an inquiry practice.

If a conclusive scan shows that a person died, the scan supports that belief even when everyone prefers the comforting alternative. Truth-directed evidence is an **aim-independent epistemic reason**.

This view matches ordinary criticism. We say a person ought to revise because the evidence changed, not merely because revision would advance a project. It also explains why unknown truths and unwanted evidence can matter.

Its debt parallels Chapter 16. Truth indication is a relation between evidence and a proposition. Why does that relation supply practical or doxastic authority for an agent wholly indifferent to truth? Saying “because the evidence makes the belief likely true” clarifies the truth connection. It does not yet explain why that connection normatively governs.

The constitutive aim of belief

Perhaps belief, by its nature, aims at truth. A heart is assessed by how well it pumps; a map is assessed by how accurately it represents terrain; a belief is assessed by whether it gets the world right.

On this account, evidence matters because it shows whether belief performs its constitutive function. Accuracy is not an optional hobby added by a truth-loving person. It is built into the kind of state belief is.

This explains why a comforting fantasy can be useful yet defective **as belief**. It may succeed as emotional medicine while failing as representation.

The pressure question is whether a standard of proper function becomes a reason for the person. A heart can function badly even when no agent is blameworthy or required to repair it. If Tavi does not care whether a belief performs its truth-directed job, what makes the functional standard authoritative for Tavi? Constitutive correctness and personal reason may connect, but they are not synonyms.

A truth-seeking aim

Perhaps evidence gains force through an agent’s desire or project of getting the answer right. A detective wants to identify the culprit. A scientist wants an accurate model. A reader wants to know whether the Candidate System succeeds. Facts about reliable methods then provide instrumental reasons.

This account is modest and powerful. It explains the authority of evidence in this book without assuming an objective epistemic ought:

You entered an inquiry aimed at determining whether objective morality can be established. Evidence is a reliable means to that end. Therefore, conditional on the aim, you have reason to follow it.

The view also allows an agent to suspend inquiry. If Rui does not care about a useless digit, no truth-seeking project requires forming a belief about it.

Its limit is the indifferent agent. If Tavi understands a contradiction but lacks every concern for truth, consistency, reputation, and successful action, no connected aim sup-

plies a reason to revise. The account can accept this. It then provides conditional rather than universal epistemic authority.

Practical usefulness

Accurate beliefs often protect lives, improve prediction, coordinate action, and prevent costly mistakes. Evidence matters because truth generally works.

This route is indispensable in policy. A weather instrument deserves attention when its forecast enables evacuation. A safety record matters because false confidence can kill. Reliable inquiry earns practical authority through consequences.

Usefulness and truth can separate. A placebo belief may relieve pain precisely because it is false. A comforting belief may cause no harm. A useless mathematical truth may serve no project. If utility alone governs, the practically best belief can be false.

That result may be accepted. The account should then say that it supplies practical reasons to cause or maintain beliefs, not evidence that those beliefs are true. “Helpful to believe” and “supported as true” are different recommendations.

Rules of inquiry

Perhaps evidence has authority within practices that agents accept: science, history, engineering, law, journalism, or this philosophical construction. Participants undertake to check sources, distinguish prediction from preference, respond to counterevidence, disclose uncertainty, and correct error.

The reason resembles a commitment or role reason. If you ask a question within a practice, you accept standards for answering it. A researcher cannot preserve membership in inquiry while declaring that contrary evidence is irrelevant whenever inconvenient.

Practice-based authority explains discipline without claiming that every human must care about every truth. It also makes methods revisable. Communities can improve rules for replication, testimony, measurement, and bias control.

The practice may be corrupt or insulated. Astrology and conspiracy communities have rules too. The Candidate System needs conditions of entry, reliability, criticism, exit, and correction. An accepted rule is not truth-conducive merely because it is written in a manual.

A defensible public procedure

Private evidence may be inaccessible to others. Public decisions therefore need procedures for checking claims: transparent methods, independent replication, adversarial criticism, representative expertise, disclosure of conflicts, appeal, and accountable institutions.

These procedures can make a belief publicly usable even when no individual possesses every piece of evidence. They distribute trust and make correction possible. A city cannot personally rerun every medical trial before licensing a treatment.

Procedural warrant differs from truth. A fair panel can unanimously accept a false conclusion. Public legitimacy can support acting on the best available record without making

the record infallible. The procedure's participant set, expertise, evidence standards, and power controls remain evaluative inputs.

No literal epistemic ought

The final candidate rejects genuine favoring relations in belief. Evidence indicates truth, causes revision, improves prediction, or satisfies inquiry rules. Saying "you should believe the evidence" is recommendation, instruction, or useful shorthand, not a report of an additional normative fact.

This **epistemic anti-normativist** can still distinguish true from false, reliable from unreliable, and effective from ineffective. They can say, "If you want accurate beliefs, use this method," or "Our institution will record the instrument reading."

The pressure falls on words such as *mistaken*, *irrational*, and *unjustified*. If these are meant normatively, the ought may have returned through the service entrance. The anti-normativist must translate them into inconsistency, method violation, unreliability, or failure relative to a stated aim.

None of the seven candidates is eliminated by its first debt. The pressure cases tell us what each can and cannot explain.

The merciful scan

A flawless deep-space scan establishes that Nia's missing partner died instantly. For six months, believing he remains alive would greatly reduce her grief. The false belief would alter no search, inheritance, relationship, or future decision, and would fade without trauma. Nia understands the scan and prefers comfort.

Does the evidence give her a reason to believe the painful truth?

Truth-directed support says yes. The constitutive account says the survival belief fails as belief even if it succeeds as comfort. A desire account asks whether Nia retains any concern for accuracy. Utility may favor the false belief. Inquiry rules may require the official record to state the death while allowing Nia's private coping to remain outside enforcement.

The case separates two questions:

- What belief does the evidence support as true?
- What mental state would best serve Nia's welfare?

They need not receive the same answer.

The useless digit

A sound proof establishes the 17,429th digit of a newly defined mathematical constant. No theorem, machine, archive, contest, or person's project will ever depend on it. Rui understands the proof and feels no curiosity. Remembering the result costs nothing.

The proof entails the digit. Does Rui have a reason to believe it?

Truth-directed and constitutive accounts say the understood proof supports the belief. Desire and utility accounts may say no personal reason arises. A practice account may add that, while participating in the proof, Rui should accept its conclusion under mathematical standards.

Do not replace "reason to believe the proposition upon understanding it" with "reason to spend time investigating it." The investigation is already complete. The narrow case asks whether truth support survives total uselessness.

The unanimous weather room

Twelve experienced forecasters prefer a calm prediction because evacuation would be exhausting and politically damaging. A newly calibrated instrument, independently checked by three laboratories, gives a high probability of a lethal pressure wave tomorrow. Everyone understands that the instrument has the better record. The vote for calm is 12–0.

The unanimous procedure can settle what the team announces. It cannot alter the likelihood relation. If forecasters aim at accurate prediction, their deeper project favors the instrument over their present preference. Practical safety supplies another reason to accept it. Transparent calibration supplies public warrant.

The case demonstrates why evidence is not a vote. It also demonstrates why a public procedure can be designed to track evidence rather than aggregate preferences. The procedure's authority is derivative from reliability, safety, accepted practice, or another stated value—not from unanimity alone.

The lucky conspiracy

Milo believes a meteor will strike because a celebrity's shoe color matches a pattern in old song titles. By coincidence, a meteor really will strike. At the time, ordinary telescopes cannot see it behind the sun. Later instruments confirm the event.

Milo's belief is true. His route did not bear on its truth. Group approval and sincere conviction do not repair the missing connection.

The case shows why truth is not enough for justification. A method must connect with truth reliably or receive support from relevant evidence. If Milo evacuates and survives, the belief also becomes useful. Usefulness does not turn shoe color into meteor evidence.

An anti-normativist can state the same facts without *justified*: the method has no predictive relation and will fail across relevant repetitions. That translation preserves criticism for anyone who values accurate prediction.

The sealed archive

Doctor Varo carefully reviews secure documents and forms the true belief that a reactor component is defective. A required privacy treatment then erases every memory of the documents and reasoning while leaving the conclusion. Varo knows that her archive reviews are normally meticulous but cannot reopen this one.

What supports the belief now?

The original documents still bear on the component's condition, though they are no longer accessible to Varo. Her knowledge of her past reliability supplies present second-order evidence. A public launch decision may require an independent check because private inaccessible support cannot be audited.

This separates **having been reliably formed**, **possessing accessible reasons now**, and **offering public justification**. An account can treat all three as epistemically relevant without pretending they are one property.

The honest placebo

Arlo's severe pain stops only while he believes an inert capsule contains powerful medicine. Conclusive evidence shows the capsule is inert. Arlo can knowingly enter an immersive suggestion process that restores the false belief and the relief. No one is deceived; he consents beforehand.

Evidence supports inertness. Welfare supports inducing the medicine belief. A constitutive account can say the resulting state fails as accurate belief while succeeding as therapy. An institutional rule can preserve the truthful medical record while permitting a private intervention.

The case exposes a category error in both directions. Practical benefit is not evidence of truth. Epistemic support does not guarantee that forming the supported belief is the best practical choice in every circumstance.

The indifferent logician

Tavi holds directly contradictory beliefs about a distant empty moon. Clear evidence resolves the contradiction. Tavi understands it but has no concern for truth, consistency, reputation, identity, or successful action in the subject. Revision is effortless and changes nothing.

Does Tavi have an epistemic reason to revise?

Truth-directed support says yes. The constitutive account says the beliefs fail their representative standard. Desire, utility, and practice accounts find no connecting aim. Public procedure may require correction of a shared record while supplying no personal reason to Tavi. Anti-normativism describes the contradiction and stops.

As with Kest, do not defeat the case by adding a hidden desire. Actual people usually care about some combination of truth, competence, reputation, prediction, or coherence.

That makes conditional epistemic reasoning widely applicable. It does not answer the stipulated authority question.

The companions at the gate

An objection now arises: if the book doubts objective moral reasons but trusts objective epistemic reasons, has it applied skepticism selectively?

The comparison is legitimate. Both claims involve a favoring relation: suffering favors action; evidence favors belief. A theory that accepts one and rejects the other owes a relevant difference.

But similarity does not force identical treatment. Evidence may have a tighter relation to truth than natural facts have to moral status. Belief may have a constitutive correctness condition that action lacks. Moral reasons may be equally fundamental. Or both kinds of reason may be conditional or fictional.

The so-called companions can survive or fail together, or one can possess credentials the other lacks. The Candidate System must inspect each rather than using one as a hostage for the other.

This book will proceed on two disclosed grounds:

1. **Conditional aim:** We are trying to determine what follows from the evidence rather than merely preserve a preferred conclusion.
2. **Inquiry commitment:** We accept correction rules that distinguish truth, evidence, reliability, counterargument, and public checkability.

These grounds are sufficient to make evidence action-guiding for participating readers. The book does not need to establish that Tavi has an aim-independent epistemic reason before continuing. Whether truth-directed evidence supplies that stronger reason remains open.

Write the epistemic charter

Purpose. State what makes beliefs true, evidence truth-bearing, inquiry corrigible, and evidential support authoritative; practical usefulness or public warrant may be included but must be labeled.

Add this page to the Candidate System:

Truth condition: A belief is true when _____.

Evidence relation: A consideration bears on truth when _____.

Source of epistemic force: truth itself / belief's function / truth-seeking aims / usefulness / inquiry commitment / public procedure / no literal force / other: _____.

Agent scope: This source reaches _____ because _____.

Reliable method: A method earns reliance through _____.

Private support: An agent may rely on inaccessible or personal evidence when _____.

Public warrant: *A claim may govern shared action when _____.*

Defeat and revision: *Contrary evidence changes confidence according to _____.*

Indifferent agent: *Tavi does / does not have a reason to revise because _____.*

Apply the charter to all seven cases. If different cases use different sources, say whether the sources coexist or one is merely a fallback. A plural account is permissible; an invisible switch is not.

Then write one sentence addressed to yourself:

I will follow evidence during the remaining inquiry because _____.

“Because that is rational” is incomplete unless *rational* is defined. “Because evidence tracks truth and I want the true answer” is conditional and sufficient for participation. “Because truth-directed support is fundamental” makes the stronger realist claim and carries its access debt.

The nineteenth debt

Add this entry:

Debt 19 — Epistemic force: *Evidence can indicate truth, arise from a reliable method, serve an aim, satisfy a practice, or receive public warrant—but what makes any of these relations a reason for an agent who rejects truth and inquiry?*

The evidence office issues the Candidate System a provisional license. It is valid for agents participating in the inquiry and may have wider authority if truth-directed reasons are real.

The next applicant arrives without documents.

“I can simply see that the principle is true,” she says. “It is obvious.”

The clerk sighs and opens a new form.

Chapter 20 — Intuition and the Feeling of Obviousness

Two cartographers stood before the same blank region of Port Meridian's moral map.

The first said, "It is obvious that a conscious replica has the same basic status as the biological original."

The second said, with equal confidence, "It is obvious that a manufactured copy can never have the status of the original person."

Both were intelligent. Both understood the stipulated facts. Neither was joking.

The map remained blank but had acquired two very certain arrows.

An **intuition** is an intellectual seeming: a proposition presents itself as true without the person consciously deriving it from premises. You consider a case and it seems that the promise binds, that the sacrifice is forbidden, that the two claims are equal, or that the conclusion simply follows.

Intuitions are not automatically mystical. Perception also presents conclusions without displaying its processing. Experts often recognize patterns before they can articulate every cue. Conceptual understanding can make some relations immediately visible. Any inquiry must begin somewhere; if every premise requires a prior argument, no reasoning starts.

Intuitions are also not automatically revelations. A seeming can reflect emotion, training, language, culture, framing, status, or repeated exposure. Confidence tells us how the experience feels, not why it occurred or whether it tracks truth.

The question is not whether intuitions exist. They plainly do. The question is what they are evidence of.

Where this chapter is headed. We will examine what an immediate moral seeming can establish and how such seemings might be calibrated without an independent answer key. The goal is neither to dismiss intuition nor to let confidence certify its own reliability.

Seven things that can feel obvious

When a judgment arrives without visible inference, separate seven possible sources.

A seeming

At the thinnest level, the proposition simply appears true. "Pushing the person feels impermissible." This reports phenomenology—the character of the mental event.

A seeming can provide a starting point for inquiry. It is not yet a mechanism connecting the content to objective moral truth.

Confidence

Confidence is the degree of assurance attached to a judgment. It can arise from evidence, familiarity, social reinforcement, emotional intensity, or inability to imagine alternatives.

High confidence may predict stability. It does not guarantee accuracy. People have been maximally confident on opposite sides of nearly every important dispute.

Emotional salience

Some cases activate fear, disgust, compassion, anger, loyalty, admiration, or protective attachment. The emotion highlights features: pain, betrayal, contamination, vulnerability, courage.

Salience can improve judgment by directing attention to a real stake that abstraction concealed. It can distort by spreading a response from one context to irrelevant cases. Disgust can reveal aversion without revealing danger. Compassion can reveal suffering while leaving distribution unresolved.

Trained perception

A skilled clinician sees that a patient is deteriorating; an experienced engineer hears that a motor is failing; a practiced mediator notices that one party cannot safely refuse. Long training compresses many cues into rapid judgment.

Moral experience may likewise improve sensitivity to humiliation, manipulation, dependency, or institutional exclusion. Expertise earns weight when training receives feedback from a relevant target. The difficulty is identifying the moral feedback. Engineers can compare prediction with motor failure. What independent result calibrates perception of objective wrongness?

Cultural fluency

A person can instantly perceive that a gesture is disrespectful, a gift inappropriate, or a duty activated because they know a social form of life. The intuition tracks real meaning within the practice.

Cultural fluency is evidence about conventions, expectations, and relationships. It may not establish approval-independent status. An insider can expertly perceive that a song is taboo while the taboo remains unjustified.

Conceptual competence

Understanding a concept can make some applications immediate. If a rule is “give each tied claimant an equal chance,” a rigged draw plainly violates the rule. If coercion is defined partly by a credible threat, recognizing the threat may be conceptual rather than moral discovery.

Conceptual truths depend on the concepts and definitions in use. They can expose contradiction inside the Candidate System without proving that the system's starting concepts correspond to objective moral reality.

Intellectual insight

The strongest proposal says some intuitions are genuine noninferential apprehensions of normative truth. Adequately understanding suffering, agency, equality, or a reason relation can reveal that it matters. Just as reason may grasp a logical axiom, moral reflection may grasp a foundational principle.

This possibility cannot be rejected merely because other intuitions are culturally shaped. Perception sometimes misleads without all perception being useless. The account needs a way to distinguish insight from the other six sources.

The red lever

Iona enters a control room where pulling a red lever will inflict severe pain on a stranger for one hour and accomplish nothing else. Every causal and experiential fact is clear. The stranger has not consented. No emergency, punishment, research goal, relationship, or social rule applies.

It seems obvious that Iona has a reason not to pull.

What does the seeming establish?

At minimum, it establishes that a strong response occurs when the case is understood. The content is stable under removal of many distracting features. The intuition connects with pain aversion, compassion, direct standing, and a widely shared rule against gratuitous harm.

A realist may say the case enables intellectual recognition of a basic reason: conscious agony counts against causing it. The simplicity is epistemically helpful because no comparison obscures the relation.

An anti-realist may say the case activates near-universal human aversion and protective concern. Its force is practical and emotional, not evidence of a stance-independent property.

A constructivist may say no defensible procedure would authorize the lever because every possible subject has reason to reject gratuitous agony.

All three can predict the same confidence. The phenomenology does not choose among their explanations.

Now change one word: the lever causes *disgust* rather than pain. Or it destroys an unowned ecosystem with no conscious life. Or it breaks a promise to someone who has died. Confidence may move differently across readers. The changes help identify which features generate the seeming; they do not by themselves establish which response tracks objective truth.

The two islands of certainty

Infants are randomly assigned between two isolated islands.

On North Island, parents, teachers, and temples teach that creating a digital likeness of the dead is a grave desecration. Every example, story, and ritual reinforces the prohibition.

On South Island, the same institutions teach that creating the likeness is a sacred duty of remembrance. Failure is treated as abandonment.

At adulthood, 98 percent of North Islanders find prohibition self-evident. Ninety-eight percent of South Islanders find creation self-evidently required. Both populations understand the technology, experience the same emotions during neutral demonstrations, and reason equally well in nonmoral domains.

The assignment history predicts the intuition.

What follows?

It does not follow that both propositions are false. Opposite claims cannot both be correct under the same interpretation, but one island may have been trained into truth by luck. A realist can also argue that later reflection might escape upbringing.

It does follow that **felt self-evidence alone** does not discriminate. The formation process would have produced either certainty depending on random assignment. Each population needs support that does not merely repeat the process under examination.

This is a burden shift, not automatic error theory. The island history weakens entitlement to say, “My confidence is evidence because it is confidence.” It does not settle the status of posthumous likenesses.

Expertise without an answer key

Can moral intuition be trained?

Yes, in several senses. People can learn facts, detect hidden stakeholders, imagine perspectives, recognize recurring manipulation, and apply declared rules consistently. A judge can become more accurate about precedent. A clinician can better understand consent. An organizer can see how a formally neutral policy burdens one group.

These achievements have feedback:

- predictions can be compared with outcomes;
- reports can be checked with affected people;
- rules can be tested for consistency;
- hidden assumptions can be exposed;
- procedures can be audited for exclusion;
- interventions can be revised when harms appear.

None is trivial. Together they can make practical judgment far more informed and humane.

Objective moral detection presents an additional problem. If the target is a stance-independent moral property, what observation confirms that the trained intuition tracks

it rather than shared values and social success? Agreement with other experts cannot be the only answer, because the experts may share formation and selection pressures.

A realist may propose reflective equilibrium: adjust principles and case judgments until they cohere under facts, counterexamples, and criticism. This improves joint fit. It can expose inconsistency and error. Coherence alone does not show that the coherent set corresponds to an independent moral order; several incompatible systems may reach equilibrium.

Again, the conclusion is bounded. Lack of an external answer key does not prove there is no moral expertise. It means the access account owes a reliability story more specific than “the expert sees correctly.”

Calibrate without pretending to possess an oracle

Calibration usually compares confidence with known outcomes. Weather forecasters can learn whether events assigned 70 percent probability occur about 70 percent of the time. A moral faculty cannot be calibrated against an objective answer set unless some answers are known independently. Using intuition to certify the answers and then using those answers to certify intuition would be a small circle with excellent posture.

We can still improve intuitive judgment through a **partial calibration protocol**.

Separate the proposition from its force

Write the exact content that seems true. Then record confidence separately. “This feels monstrous” may accompany several propositions: the outcome is bad, the agent is blameworthy, the act is forbidden, or punishment is authorized. The emotion does not specify which stamp applies.

Audit the empirical load

List every factual assumption. An intuition about a policy may change when its effects, consent conditions, or alternatives are corrected. That change shows responsiveness to evidence. It does not reveal whether the remaining value judgment is objective.

Rotate frame and identity

Use the rotation test. If judgment changes with irrelevant names, order, vividness, or group labels, confidence should fall. If a relationship or history is genuinely relevant, state the relevance rule.

Search for a formation intervention

Would random praise, island assignment, professional training, or repeated exposure predict the intuition? Controlled evidence of reversal under an irrelevant cue is stronger than a speculative origin story.

Invite independent opposition

Seek critics who do not share the same school, community, incentives, or favored examples. Agreement among ten students of one teacher is less independent than convergence across differently formed investigators using public arguments.

Test transfer

Apply the principle to new cases selected by an opponent. A principle that survives only the examples used to create it may be fitted to those cases rather than tracking its stated feature.

Preserve residual uncertainty

After these checks, confidence may remain high. Record what the process has earned: stability under facts, framing, identity swaps, and criticism. Do not add “therefore objective” unless the access theory connects these successes to stance-independent truth.

Consider the red lever again. The verdict may survive every rotation, factual check, and opponent-designed variation. This is significant. It supports the judgment as a stable, informed commitment and perhaps as intellectual insight. A realist can include the stability in a broader reliability account. A non-realist can explain it through shared sentience and aversion. Calibration narrows the explanatory contest without deciding it by itself.

Now consider the conscious replica. If responses vary dramatically with the words *copy*, *continuation*, *manufactured*, and *person*, the first task is conceptual clarification. Declaring one uncalibrated frame self-evident would turn a vocabulary preference into an alleged window on moral reality.

Partial calibration is less satisfying than an oracle. It is also available.

Thought experiments as microscopes

This book uses imagined cases liberally. What can they establish?

A thought experiment holds some facts fixed, varies one feature, and observes how principles and responses change. It can reveal:

- hidden dependence on identity;
- conflict between declared rules;
- ambiguity in a concept;
- a missing threshold;
- an unwanted implication;
- the intuitive cost of a theory.

It cannot create objective truth by provoking a reaction. The reaction is data about the reader and perhaps about conceptual relations. Moving from “this verdict seems intolerable” to “the theory is objectively false” requires trust in the seeming or an independent argument.

Framing matters. “Push a person to stop the threat” and “use a body as a brake” can describe the same motion while activating different concepts. Names, order, vividness, and omitted alternatives influence response. The bounded-case rule reduces these effects without eliminating them.

Use a **rotation test**:

1. Rewrite the case in emotionally neutral language.
2. Swap identities and group memberships.
3. Present consequences before intentions, then reverse the order.
4. Replace a vivid individual with an equivalent statistic.
5. Ask which wording changes judgment and why.

A stable intuition after rotation earns more confidence as a considered response. It does not automatically earn status as perception of objective fact.

Three obvious judgments

Purpose. Audit three of your own immediate judgments for stability, framing, formation, conflict, and possible correction; the exercise tests entitlement to confidence, not the truth of the verdict by itself.

Open the Candidate System notebook and record three judgments you currently find obvious:

1. one about suffering or welfare;
2. one about agency, rights, or fairness;
3. one about loyalty, authority, relationship, or tradition.

For each, complete:

Judgment: _____.

Immediate seeming: It appears true because _____, if I can articulate it.

Confidence: low / moderate / high / near-certain.

Emotions activated: _____.

Training or cultural fluency: _____.

Conceptual rule involved: _____.

Possible intellectual insight: _____.

Formation history known so far: _____.

What evidence or argument could lower confidence: _____.

What could raise confidence: _____.

What would not change my mind: _____, because _____.

“Nothing could change my mind” is permitted. Mark the belief as insulated. Insulation may be appropriate for a foundational axiom, but it cannot then cite responsiveness to evidence as confirmation.

Now give each judgment to the strongest reasonable dissenter you can imagine. Hold empirical facts constant. Ask which of the seven sources explains their opposite intuition.

If your explanation is “bias” while yours is “insight,” state the independent evidence for that asymmetry.

Self-evidence and public justification

A proposition can seem self-evident to an individual yet fail as a public justification. A diverse society cannot inspect private insight directly. It needs shareable premises, evidence, consistency, impacts, and procedures.

This does not make private intuition worthless. It can initiate inquiry, motivate protest, or alert a person to a neglected stake. Asha’s sense that the academy’s exclusion was intolerable may precede a complete public argument. Reform often begins with perception before institutions possess vocabulary for it.

The later argument can show harm, arbitrary exclusion, inconsistency with declared values, or domination. Those grounds become publicly criticizable. If Asha also claims access to an objective moral fact, the epistemic route remains a separate claim.

The Candidate System should therefore give intuition two roles:

1. **Discovery prompt:** a seeming identifies a proposition or feature for examination.
2. **Provisional support:** a considered intuition receives weight according to an explicit account of reliability, independence, and correction.

Neither role permits intuition to serve as its own infallible certificate.

The twentieth debt

Add this entry:

Debt 20 — Intuitive access: Does an “obvious” judgment express a seeming, confidence, emotion, training, cultural fluency, conceptual competence, or genuine insight—and what independently connects the trusted source to stance-independent moral truth?

The two cartographers still disagree about the conscious replica. They now possess better questions. Which features produced the seeming? Which formation process trained it? Which rule does it apply? What could correct it? Does the confidence survive rotation and informed dissent?

Those questions lead naturally backward.

Every intuition has a history, even if that history includes genuine learning. The next chamber is an archive of parents, punishments, prayers, alliances, accidents, institutions, scars, rewards, and deliberate revisions.

At the entrance is a warning:

AN ORIGIN STORY IS NOT YET A VERDICT.

Chapter 21 — Where Values Come From

Port Meridian's archive of moral formation occupied nine rooms.

One contained medical records, accident reports, and images of visible suffering. One contained ledgers of favors given and returned. One was full of family trees, flags, and team colors. One held scriptures, rituals, and sacred stories. One preserved schoolbooks and parental letters. One displayed medals, promotions, dismissals, and applause. One contained constitutions, markets, technologies, and inherited offices. One held private diaries. The last was furnished with a desk, two chairs, and a sign reading **DELIBERATE REFLECTION**.

The attendant warned visitors not to assume the last room was free of material from the other eight.

A **genealogy** is an account of how a conviction formed, acquired emotional force, spread, persisted, or changed. Moral convictions have genealogies because people have histories. We inherit language, notice some harms before others, depend on groups, respond to rewards, learn from experience, and revise through reflection.

What follows from discovering such an origin?

Sometimes very little about truth. Sometimes a specific reason to distrust a process. Sometimes direct evidence against a premise. Sometimes a broad burden to find independent support. The work lies in matching the conclusion to the evidence.

"That belief has a cause" is not a refutation. Uncaused beliefs would not be more trustworthy. The relevant question is whether the cause or process responds appropriately to the facts or reasons the belief claims to track.

Where this chapter is headed. We will trace the formation of values without confusing a causal history with a refutation. Genealogy becomes epistemically significant only when an additional premise connects the formation process to reliability, distortion, or the absence of truth tracking.

Nine families of influence

Begin by mapping possible influences without treating their presence as guilt.

Harm sensitivity

Pain, fear, grief, visible distress, and remembered loss direct attention. A person who watched public humiliation destroy a child's confidence may later oppose humiliating punishment.

Sensitivity can reveal a neglected consequence. It can also overgeneralize from one vivid case. The history alone does not decide which.

Reciprocity

Repeated exchange and dependence shape expectations about returning help, sharing burdens, and exploiting cooperation. Communities facing recurring disasters may develop strong practices of mutual aid.

Reciprocity can track real vulnerability and solve coordination. It can also exclude people unable to reciprocate. Its adaptive success neither proves nor disproves objective authority.

Kinship and group loyalty

Attachment to family, friends, teams, nations, professions, and movements shapes who feels like “us.” Loyalty can sustain care through sacrifice. It can make an irrelevant badge appear morally decisive.

The genealogy becomes epistemically significant when controlled comparison shows judgment following group identity while every declared relevant fact remains fixed.

Religion

Religious traditions supply commands, exemplars, stories of purpose, rituals, communities, interpretive methods, and emotional frames. They can preserve insight across centuries or stabilize hierarchy and error.

“This belief came from religion” says no more about truth than “this belief came from a university” until the particular source, method, evidence, and claim are examined. A revelation claim and a childhood habit require different audits.

Parents and community

Children learn through correction, praise, imitation, silence, and observation of who is protected or ignored. Much moral fluency arrives before explicit argument.

Early learning can transmit justified rules—do not touch the stove, do not injure playmates—and local prejudice. Random assignment between communities can reveal that confidence tracks upbringing rather than independent inquiry.

Status incentives

Reputation, belonging, promotion, attention, and safety influence which convictions are easy to express and, over time, easy to believe. The influence need not be conscious hypocrisy.

If randomized applause reliably reverses sincere judgment while evidence remains constant, the process is tracking reward rather than merits. This undermines reliance on that process, not necessarily the conclusion it happened to produce.

Historical institutions

Laws, markets, schools, technologies, professions, and inherited power make some arrangements appear natural. A rule may survive after its original function or beneficiaries are forgotten.

Institutional history can merely explain familiarity. It can also reveal that a current defense rests on a false historical claim or that apparent consent was manufactured. In the latter case, the evidence bears directly on content.

Personal experience

Events lived, witnessed, feared, regretted, or caused can transform priorities. Experience supplies contact with facts an insulated person missed. It can also produce trauma, selective memory, and analogy beyond relevant similarity.

Do not call experience “mere anecdote” when the issue concerns a kind of harm the speaker directly knows. Do not treat one experience as a representative sample without further evidence.

Deliberate reflection

Agents compare principles, evidence, counterexamples, perspectives, and consequences. Reflection can identify contradiction, correct prediction, widen standing, and reduce identity bias.

It receives no automatic exemption from genealogy. Which cases were selected? Which interlocutors were heard? What emotional capacities were amplified? Which endpoint counted as equilibrium? Reflection can rationalize an inherited preference with great elegance. Its reliability depends on safeguards and contact with relevant evidence, not on the flattering sound of its name.

These nine influences can coexist. A conviction may begin in attachment, gain language from religion, receive status reinforcement, encounter disconfirming experience, and later be reconstructed through reflection. Genealogy is a network, not a single hidden cause announced by a detective in the final chapter.

Four conclusions, four burdens

The archive issues four differently colored forms.

Form 1: causal explanation

The evidence helps explain why a conviction became salient or persistent. It does not yet show that the process was unreliable, that current support depends on it, or that the proposition is false.

Lea watched a teacher publicly humiliate her brother for a minor mistake. He feared school for years. The event first made Lea oppose public humiliation as punishment.

This explains attention. Perhaps Lea noticed a real harm. Perhaps she later overgeneralized. Her current arguments remain unexamined. The strongest conclusion is genealogical description.

Form 2: local process distrust

Evidence shows a judgment process responding to an irrelevant cue, manipulation, or incentive rather than the features it claims to track. Reliance on that process should decrease under the inquiry charter.

Officer Kavi says rescue priority depends only on injury, distance, and survival chance. In repeated blinded simulations, those facts remain fixed while researchers swap meaningless blue and gold badges. Kavi reliably favors the badge of Kavi's own crew.

The result does not prove which person should be rescued in a new case. It shows that Kavi's process fails Kavi's stated relevance rule. The cases should be reassessed without the cue.

Form 3: content-level argument

Evidence bears directly on what a belief says, defeats an explicit premise, or exposes conflict within the position.

The city of Vey gives inherited council seats to founding families. Its official defense says every family freely signed an equal agreement. Authenticated records show forged signatures and threats.

This is not merely a story about why people revere founders. It defeats the stated historical premise. Another defense might be offered, but this one fails.

Form 4: burden shift

A broad, well-supported genealogy shows that the same process would have produced a conviction whether it was true or false. Automatic trust in the process is withdrawn; independent support is required.

The two islands from Chapter 20 provide a local example. Random assignment predicts opposite self-evident convictions. A broader hypothesis might argue that core moral intuitions were shaped for survival and reproduction rather than sensitivity to stance-independent moral facts.

Even if that broad history is accepted, falsity does not follow without another premise. Adaptive processes can produce true beliefs when truth aids survival. They can also produce useful false beliefs. The genealogy challenges entitlement to trust, not the proposition's truth directly.

Independent reflection might meet the burden, fail, or reproduce the same influences. The history alone does not decide.

The tempting genetic shortcut

A common mistake runs:

You believe equality because your society taught it. Therefore equality is not objectively valuable.

The conclusion does not follow. A society may teach arithmetic, astronomy, prejudice, and hygiene. Causal dependence is compatible with truth and falsehood.

The opposite shortcut also fails:

The origin of a belief never matters; only arguments matter.

Origins matter when they bear on reliability or current warrant. If a medication belief was implanted by a marketer who randomized claims, that history gives reason to recheck it. If a witness's identification followed a suggestive lineup, process evidence matters even when the identified person is guilty.

The correct question is not "Does the belief have an origin?" It is:

Did the formation or maintenance process respond to relevant truth indicators, to irrelevant cues, or to both—and what independent support remains now?

This avoids both the genetic fallacy and genealogical amnesia.

The evolutionary challenge at full strength

Suppose a comprehensive scientific account explains our central moral responses through selection for offspring care, coalition stability, reciprocal cooperation, threat detection, status regulation, and cultural learning. Further suppose the process was not selected for tracking independent moral truth.

What follows?

A realist has several replies.

Truth can coincide with adaptiveness

Perhaps suffering really supplies reasons, and organisms that detect suffering also cooperate more successfully. Selection need not represent moral truth as its goal for the resulting faculty to track some relevant facts.

Reflection can transcend initial mechanisms

Mathematical and scientific reasoning use capacities shaped for other purposes. Moral reflection may correct evolved bias by extending concern, checking consistency, and incorporating evidence.

Fundamental truths need not causally shape intuitions directly

Abstract truths do not push neurons. Reliability may arise through rational sensitivity rather than ordinary causation.

The challenger replies that each route needs detail. Coincidence is not a reliability mechanism. Reflection must show independence from the starting responses. Rational sensitivity cannot be inferred from the fact that judgments feel rational.

The result is a genuine burden, not a cheap refutation. The realist owes an account of why the formation process is appropriately connected to objective moral truth. The non-realist owes an account of why the same responses support practical criticism and revision without being truth detections. Both have work.

When a genealogy can increase confidence

Origin stories are often treated as skeptical acids, but some genealogies are vindictory. Learning how a belief formed can increase confidence when the process shows sustained contact with the feature the belief claims to represent.

Consider the battery-sharing settlement on the moon Oriel. Dust storms randomly disable different homes. For generations, neighbors survive by rotating charged batteries. The community develops a conviction that keeping unused emergency power while another household freezes is unacceptable.

The history explains the conviction's spread. It also exposes residents repeatedly to dependence, cold, fear, reciprocity, and the consequences of refusal. If the claim is that sharing prevents avoidable death and supports reliable cooperation, the formation process has excellent feedback. Records can compare outcomes when households share or hoard.

If the claim is stronger—"unused power supplies an objective moral reason to every possible agent"—the same history does not provide a direct moral answer key. It establishes contact with natural facts and successful practices. A realist must explain why those facts are reason-giving; a constructivist can show why informed participants endorse reciprocal protection.

A potentially vindicating genealogy has several marks.

Relevant sensitivity

Judgment changes when the alleged reason-giving feature changes and remains stable when irrelevant identity cues change. Residents distinguish an emergency from ordinary inconvenience and available surplus from power needed for their own survival.

Counterfactual responsiveness

Had sharing increased death or rewarded manufactured crises, the practice would have revised. The process does not preserve the verdict regardless of evidence.

Independent routes

Different observers reach similar conclusions through outcome records, testimony, direct experience, and reciprocal deliberation rather than copying one authority.

Error correction

The community records failed rules, hears households bearing the cost, and changes thresholds. Feedback reaches decision-makers rather than being suppressed.

Appropriate target

The process is evaluated against what it can track. Survival records can vindicate a causal belief about sharing. They cannot alone vindicate metaphysical claims about objective obligation.

Genealogy can therefore support reliability at one layer while leaving another open. A physician’s training may reliably detect pain. That reliability supports the factual judgment that a patient suffers. Whether suffering supplies an aim-independent reason remains the Chapter 16 bridge.

This layered result blocks a symmetrical mistake. The skeptic may not say every caused conviction is contaminated. The realist may not say every successful practice tracks objective morality. Formation history can vindicate, undermine, or merely explain depending on the target and feedback.

Apply this to deliberate reflection. Ask not whether reflection occurred, but whether it was sensitive, counterfactual, independent, corrigible, and matched to the proposition claimed. A beautiful argument built only to protect its conclusion has a genealogy too.

A genealogy of one conviction

Purpose. Trace the formation of a conviction you take seriously and identify which influences bear on reliability; causal influence alone is neither guilt nor vindication.

Choose a considered conviction from the Candidate System. Avoid selecting one you regard as obviously foolish; the exercise is not an archaeological attack on an enemy.

Use this record:

Possible influence	Evidence it played a role	Confidence in causal claim	Tracked relevant facts or irrelevant cues?	Current support independent of it
Harm sensitivity				
Reciprocity				
Loyalty				
Religion				
Parents and community				
Status incentives				
Historical institutions				
Personal experience				

Possible influence	Evidence it played a role	Confidence in causal claim	Tracked relevant facts or irrelevant cues?	Current support independent of it
Deliberate reflection				

Mark unsupported possibilities **unknown**. Do not invent a childhood motive, trauma, religious influence, or desire for status because it would make an elegant story. Genealogies require evidence too.

For every supported influence, choose the strongest warranted inference:

- explanation only;
- local process distrust;
- argument against a premise or content;
- broad burden shift.

Do not assign numbers and average them into a verdict about truth. One decisive document can defeat a premise while nine strong causal influences leave truth open. Epistemic conclusions are structured, not averaged.

Finish with two headings:

What this genealogy changes

State which confidence, process, premise, or burden should change and why.

What this genealogy does not establish

State which claims remain possible: truth, falsity, independent support, later correction, or alternative interpretation.

Then repeat the exercise for a conviction you oppose. If genealogical explanation feels devastating only when applied to opponents, the archive has located a status incentive in the investigator.

Genealogy across metaethical views

The same history can be interpreted differently.

A **realist** asks whether the process tracks stance-independent truth. A **constructivist** asks whether the resulting values survive a defensible procedure. An **expressivist** asks how attitudes formed and whether they remain coherent and informed. An **error theorist** may treat genealogy as explaining why false objective claims feel compelling. None receives its interpretation free from the history alone.

Genealogy can therefore reorganize burdens without settling metaethics. It may undermine one alleged faculty while leaving content arguments intact. It may explain confidence while leaving truth open. It may show that a supposedly universal value reflects a local institution. These are precise achievements.

The twenty-first debt

Add this entry:

***Debt 21 — Genealogical connection:** Which processes formed and sustain the conviction, did they respond to relevant facts or irrelevant cues, what support survives independently, and what connects the trusted process to stance-independent moral truth?*

The archive does not stamp convictions **TRUE** or **FALSE**. It marks sources, vulnerabilities, defeated premises, and shifted burdens.

When visitors leave, they encounter other visitors carrying different archives. Some disagreements dissolve once the records are compared. Others survive every exchange. Still others disappear only after one population controls the schools, sanctions, and public vocabulary.

The next chamber contains two instruments: a debate table and an applause meter.

Neither is labeled **TRUTH DETECTOR**.

Chapter 22 — Disagreement and Convergence

The Assembly of Port Meridian commissioned a magnificent device called the Consensus Engine.

It had two gauges.

The first measured how many people agreed.

The second measured whether they were correct.

Only the first gauge was connected to anything.

Moral disagreement is often offered as evidence against objective morality. If there were objective truths, the argument goes, informed and sincere people should converge. Persistent conflict suggests that judgments track culture, interest, emotion, and power instead.

The immediate reply is correct: people disagree in domains with objective truth. Medicine, history, economics, and physics contain prolonged expert disputes. Disagreement does not entail that there is no fact of the matter.

Convergence is equally ambiguous. Populations can agree because they encountered strong evidence, share a successful method, possess similar biology, face common material constraints, imitate prestigious groups, fear punishment, or were educated by a conquering authority.

Disagreement is not refutation. Agreement is not revelation. Both are evidence whose force depends on the process producing them.

Where this chapter is headed. We will ask what disagreement and convergence reveal about access to moral truth. Before drawing a skeptical or realist conclusion, we must locate the layer of conflict and explain what the proposed epistemology predicts should happen there.

First locate the disagreement

Before interpreting persistence, reopen the eight-layer map from Chapter 5.

Wording and framing

People may use *freedom*, *harm*, *person*, or *fairness* differently. Clarifying terms can remove verbal conflict or expose a substantive choice between frames.

Empirical prediction

People may disagree about consequences, prevalence, causation, or risk. Better data and definitions can produce convergence.

Standing

People may agree on effects and disagree about whose interests enter directly: animals, future persons, communities, artificial minds, or outsiders.

Priority

People may admit the same values and rank them differently. Evidence can estimate the danger without determining how autonomy compares with residual risk.

Attitude and commitment

People may understand every fact and favor opposite outcomes. Accurate reports of their attitudes can both be true while the practical conflict remains.

Coordination rule

People may agree on values but propose different shared policies, exceptions, safeguards, and enforcement structures.

Decision procedure

People may disagree about who decides, who participates, what vote is required, and how appeal works.

Alleged objective truth

Finally, people may claim incompatible approval-independent moral statuses. At this layer, practical settlement does not establish which claim is true.

“Persistent moral disagreement” often bundles all eight. Some layers have familiar correction methods. Others concern values or authority. An epistemology must predict which should converge under which improvements.

Tractable conflict

Two groups dispute whether a harbor barrier will reduce flooding. One uses a model with outdated sea-level measurements; the other uses current readings and a validated model. Both aim to predict accurately. Updating the data produces convergence.

This success establishes several things:

- the disagreement was substantially empirical;
- a shared accuracy aim connected both groups to evidence;
- one model performed better under agreed tests;
- correction was possible without resolving objective morality.

Now suppose both groups agree on every forecast. The barrier will save the commercial district, increase flood risk in a poor neighborhood, damage a wetland, and reduce total expected financial loss. They disagree about whose loss should control.

More sea-level data cannot settle standing and priority by itself. Calling the remaining conflict “resistance to evidence” misclassifies it. The evidence completed its assigned job; values now select among described effects.

Convergence under conquest

The coastal population of Aru teaches that burial grounds may never be relocated without consent from descendants. The inland state of Bel teaches that public infrastructure overrides ancestral claims after compensation.

Bel conquers Aru. For three generations it controls schools, courts, employment, broadcasting, and religious appointments. Public defense of Aru’s rule brings exclusion and punishment. By the fourth generation, 96 percent of Aru’s descendants sincerely endorse Bel’s standard. Physiological testing suggests they are not merely pretending.

The Consensus Engine records impressive convergence.

What caused it?

Perhaps Bel’s rule is objectively correct and education enabled recognition. Perhaps Aru’s rule is correct and power suppressed it. Perhaps neither has approval-independent status. The convergence itself is predicted by control of formation and incentives under any of these possibilities.

This does not mean coerced education can never transmit truth. A conquering state might teach accurate astronomy. It means the observed agreement receives little independent evidential weight when the process would have produced the victor’s doctrine regardless of truth.

To use convergence as evidence, ask:

- Did participants have access to alternatives?
- Could dissent survive socially and materially?
- Did evidence, rather than penalty, predict revision?
- Were formation processes independent?
- Did the method succeed in domains with feedback?
- What would have happened if power were reversed?

A consensus produced by control is socially real. It may stabilize law and identity. It is not thereby moral discovery.

Disagreement after ideal reflection

Now consider two philosophers, Aven and Bira, evaluating a conscious digital mind scheduled for deletion.

They share every nonmoral fact. The mind experiences pain, remembers a life, forms projects, and can be copied. Deletion of this instance ends one stream of consciousness; an exact successor will continue elsewhere. Preserving the instance consumes resources that could restore sight to five people.

Aven and Bira receive unlimited time, equal access to evidence, freedom from threat, vivid simulation of every affected perspective, and correction of logical inconsistency. Neither lacks imagination or information.

Aven gives protected priority to the continuing subject and opposes deletion. Bira gives priority to the five restorations once an exact successor preserves psychological continuity. Both can state general rules, accept identity swaps, and acknowledge the cost of their verdict.

The disagreement persists.

This case does not prove that ideal reflection would actually diverge. It asks what a theory predicts if it does.

A determinately objective theory

One thinker is mistaken. The epistemology should say which method could reveal the error or why competent peers may remain nonculpably mistaken.

An indeterminate objective theory

The relevant reasons may be incomparable or leave a permissible range. Both verdicts could be allowed even though the underlying values are objective.

A constructivist theory

Different defensible procedures or starting commitments may yield different valid settlements. The question becomes which shared policy can be justified and revised.

A noncognitive or plural-attitude account

The conflict may remain a clash of informed commitments rather than incompatible detected facts. Evidence has completed its role without eliminating practical opposition.

Persistent ideal disagreement therefore pressures determinacy and access. It does not select one metaethical interpretation by itself.

What should a peer's disagreement do?

Suppose you regard another thinker as an **epistemic peer** in the case: equally informed, capable, careful, independent, and free of a known bias. They reach the opposite verdict.

Three responses are familiar.

Remain steadfast

You retain confidence because your evidence or insight remains compelling. Peer disagreement is another datum but not necessarily a defeater. Perhaps the peer made an error you cannot locate.

This response is reasonable when the peer lacks access to private evidence, has a weaker track record, or the disagreement is not independent. It is harder when equality is stipulated.

Conciliate

You reduce confidence because a trusted equal reached the opposite conclusion. Each person treats the other's judgment as evidence that an error occurred.

If both reduce equally, confidence may approach suspension. This is epistemically humble but can make decisive action difficult.

Reclassify the conflict

You discover that the disagreement concerns priority, permissible plurality, or commitment rather than one factual proposition. You may retain your stance while withdrawing the claim that the peer misperceived an objective fact.

This response often improves accuracy. It can also be used opportunistically to protect a threatened objective claim by redefining every hard case as indeterminate.

No universal response follows from the word *peer*. Ask how independence, access, track record, stakes, and the theory's predicted convergence bear on the case.

Sources of convergence

Agreement can arise from:

- contact with the same truth-tracking evidence;
- shared logical and conceptual constraints;
- common human vulnerability and affect;
- similar material problems of coordination;
- mutual influence and prestige;
- institutional selection of which voices remain;
- coercion and punishment;
- common interests or bargaining settlement;
- an actual objective moral fact recognized through a reliable faculty.

The list is not cynical. Shared vulnerability can reveal genuine stakes. Successful co-ordination is valuable. Cultural transmission can preserve knowledge. The point is that convergence has several possible explanations.

Sources of disagreement are equally varied:

- different facts or uncertainty;
- ambiguous framing;
- different standing boundaries;
- unlike value priorities;
- opposed interests and commitments;
- unequal power;
- different formation histories;
- genuine indeterminacy;
- cognitive error;
- failure of access to an objective truth.

An access theory should predict patterns. If moral perception is reliable under calm, informed, impartial reflection, convergence should increase under those conditions. If values are constructed from starting commitments, informed divergence may persist within procedural limits. If intuitions mainly track formation, changing upbringing should predict verdicts. Predictions make the epistemology vulnerable to correction.

Not all convergence is equally deep

Two people can agree at one layer while remaining divided underneath.

Verbal convergence

Both say “dignity must be protected” while one means autonomous choice and the other means living in accordance with a sacred role. The shared slogan may conceal opposite applications.

Case convergence

Both oppose one act but for different reasons. A welfare rule, a right, a religious command, and loyalty can converge on rescue. Agreement may disappear in the next case.

Rule convergence

Both endorse the same general policy and exceptions. This is stronger practical alignment, though one may regard the rule as a useful convention and the other as objective law.

Foundational convergence

Both accept the same standing, value, reason, and authority claims. This is deeper, but it can still arise from shared formation or mutual influence.

Method convergence

Both agree on how error is detected and what would change their minds. This gives future disagreement a tractable route.

Reports of “growing moral consensus” should identify the depth. A treaty can produce rule convergence while citizens retain opposed values. A philosophical vocabulary can produce verbal convergence without policy agreement. Surface agreement may be all a coordination problem needs; it should not be advertised as metaphysical discovery.

Independence matters too. Suppose one view appears in a prestigious academy, its graduates staff every journal, and later surveys find expert consensus. The number of experts exceeds one, but the formation paths are genealogically linked. This does not make the consensus worthless. Training may transmit genuine expertise. It reduces the evidential force gained merely from counting apparently separate judgments.

Use an **independence map**:

- Who learned from whom?
- Which data sources are genuinely separate?
- Which incentives are shared?
- Can dissenters publish, work, and remain in the field?
- Do rival methods receive a fair test?
- Does convergence follow new evidence or a change in prestige?

Then write an **expected-convergence contract** before examining the result:

If this access method tracks objective moral truth, then agents with conditions _____ should converge at layer _____ on timescale _____. Failure would lower confidence in _____. Success would increase confidence only if independent from _____.

This prevents a theory from treating every result as confirmation. Quick convergence can be credited to obvious truth, while slow disagreement is credited to human bias; if convergence never arrives, the truth is called difficult. Any of those explanations may be correct. The contract says in advance what the method risks.

Objective domains can contain reasonable disagreement because evidence is incomplete, concepts are vague, computation is difficult, or permissible answers are plural. The stronger the claim of clear universal access and determinate content, however, the more persistent peer disagreement needs explanation.

Build a disagreement dossier

Purpose. Locate one persistent dispute, identify the mechanisms that sustain or reduce it, and state what your access theory predicts; a dossier that treats every outcome as confirmation has failed its test.

Choose a persistent dispute. Create this record:

Layer	Exact propo- sitions or attitudes in conflict	Evidence available	Expected correction method	Did conver- gence occur?	What re- mains?
Framing					
Empirical prediction					
Standing					
Priority					
Attitude					
Coordination					
Procedure					
Objective status					

Then state:

My account predicts greater convergence when _____ because _____.

*Persistent disagreement under those conditions would count against _____ but
would not establish _____.*

*Observed convergence would support _____ only if the process was independent
of _____.*

Run the dossier on Aru and Bel. Then run it on Aven and Bira. The first case has convergence without free inquiry. The second has free ideal inquiry without convergence. A theory that treats both outcomes as confirmation has insulated itself.

Acting while the peers still disagree

A city cannot always suspend policy until philosophy converges. Patients arrive, budgets close, buildings fail, and prohibitions remain in force while committees deliberate. The epistemic response to peer disagreement and the practical response to urgency must be coordinated without being collapsed.

Separate four decisions.

Confidence

How much should each participant reduce confidence in the objective claim? This depends on peer quality, independence, access, and the predicted convergence of the theory.

Policy

Which action should be taken now? The answer may rely on shared values, error costs, reversibility, existing rights, or a legitimate procedure even when no objective verdict has been established.

Burden

Who must supply further support? A party proposing irreversible coercion may bear a higher burden than a party proposing a temporary reversible measure. This is an evaluative safeguard, not a deduction from disagreement alone.

Inquiry

What evidence, conceptual clarification, or experiment could make progress? If no participant can name a possible correction, further debate may be performing identity rather than inquiry.

Return to Aven and Bira. The digital mind is conscious, but its identity relation to the successor is disputed. Deletion is irreversible; restoring sight can be delayed for six months without eliminating the opportunity. The city might preserve the instance temporarily, fund independent work on identity and standing, compensate the waiting patients, and schedule a decision.

This settlement does not prove Aven correct. It gives reversibility priority under uncertainty. Bira may accept it as a fair interim policy while continuing to believe deletion permissible.

Now change the case: the sight-restoring resource will expire tomorrow, while the digital instance can be copied with uncertain but substantial continuity. Delay is no longer neutral. A lottery, threshold rule, or declared priority must settle action. Peer disagreement has not removed the need for values; it has made the uncertainty visible.

Use this interim-decision form:

Unresolved claim: _____.

Effect on confidence: _____.

Decision deadline: _____.

Reversible options: _____.

Error costs and bearers: _____.

Temporary rule: _____, selected through _____.

Evidence or argument sought: _____.

Review trigger: _____.

The form prevents two abuses. “Experts disagree” cannot automatically block protective action. “We must act” cannot convert one faction’s urgency into objective truth.

Progress without unanimity

Inquiry can improve without producing total agreement. Factual errors can be removed. Hidden interests can be represented. Terms can be clarified. Exceptions can become symmetric. Power can be disclosed. Predictions can be calibrated. Parties can learn which disagreement remains.

This is epistemic and practical progress even when ultimate status remains disputed. A map that marks an impassable ridge is better than one that draws a road through it.

Do not define progress solely as movement toward your current verdict. Use process and content criteria stated in advance. A reduction in contradiction is identifiable. A more accurate prediction is testable. Inclusion of an omitted stakeholder is visible. Whether these changes also approach objective moral truth depends on the access theory.

The twenty-second debt

Add this entry:

***Debt 22 — Disagreement and correction:** Under what conditions should competent agents converge, what should peer disagreement do to confidence, and how can convergence through truth-tracking be distinguished from convergence through shared biology, interest, influence, or power?*

The Consensus Engine is not discarded. Its first gauge provides useful data about coordination, confidence, and social formation. The second gauge remains disconnected until the Candidate System supplies a truth-tracking mechanism.

The next chamber contains instruments that work: tide gauges, medical trials, risk models, censuses, and ecological surveys. They can measure effects with extraordinary precision.

The delegates are about to discover that precise facts can enter a policy argument without volunteering to choose its values.

Chapter 23 — Facts Enter; Values Select

The meeting began with an unprecedented announcement.

“We agree on every fact,” said the chair of Port Meridian’s coastal commission.

Nobody celebrated. Experienced commissioners knew what came next.

The city had to choose between two flood plans. Engineers agreed on construction costs, failure probabilities, displaced households, effects on shipping, expected injuries, animal suffering, wetland loss, tax burdens, and uncertainty ranges. Economists agreed on employment effects. Residents agreed that the models were the best available. No hidden report waited in a drawer.

The commission still divided.

One group gave priority to preventing the largest expected loss of life. Another treated forced displacement of a historically burdened district as a protected constraint. A third accepted relocation but rejected the level of ecological risk. A fourth insisted that future residents deserved greater representation. Everyone read the same columns. They selected different rows as decisive.

Empirical evidence is extraordinarily powerful. It can overturn predictions, reveal harms, identify affected people, expose ineffective means, estimate uncertainty, and defeat factual premises. It cannot silently determine which effects have standing, how unlike losses compare, or what risk is acceptable.

This is sometimes summarized as “you cannot derive an ought from an is.” The slogan is too blunt for our purposes. If objective moral facts exist, they are facts. Descriptions of commitments, legal duties, and institutional roles can also combine with rules to yield conditional conclusions. The precise lesson is narrower:

Empirical description does not select a terminal value or priority without an evaluative bridge.

Facts enter. Values select. Rules connect. Authority remains a separate question.

Where this chapter is headed. We will trace how empirical facts constrain policy while showing where a standing rule, value, threshold, or distribution principle enters. The boundary protects science from being asked to choose values and values from disguising themselves as measurements.

What evidence can settle

Do not react to the boundary by understating science. Evidence performs at least six indispensable jobs.

Describe present conditions

Measurement tells us who is sick, where pollution falls, which resources exist, how institutions operate, and what people report valuing. Without description, moral judgment can protect imaginary victims while ignoring actual ones.

Identify causal effects

Controlled studies, natural experiments, historical comparison, and models can estimate what a policy causes rather than what merely accompanies it. A well-intended program can worsen the outcome it was designed to improve.

Forecast alternatives

Policy concerns counterfactuals: what is likely under each available option. Evidence can compare expected injuries, costs, behavior, ecological change, and distribution across time.

Quantify uncertainty

A model can provide ranges and probabilities rather than false certainty. Knowing that a catastrophic threshold has a 5 percent chance is different from not knowing whether the threshold exists.

Reveal affected parties

Evidence can show that a policy reaches people or beings previously omitted: downstream communities, unpaid caregivers, sentient animals, future residents, or workers far along a supply chain.

This identifies candidates for standing. The decision that their interests count directly still uses the standing rule from Chapter 6.

Test means against admitted ends

If the Candidate System values health and a treatment reliably damages health, evidence counts against that treatment relative to the admitted value. If a punishment policy does not deter, its deterrent justification fails.

Empirical correction can therefore reverse a policy without changing the underlying values. This is why evidence remains politically transformative even when it does not choose terminal criteria.

Prediction is not priority

A study may show that a privacy rule will prevent ten injuries while exposing one million personal records. The prediction describes effects. It does not tell us whether privacy is a large welfare loss, a protected boundary, a right with a threshold, or a defeasible social practice.

Suppose every participant agrees that disclosure has a 70 percent chance of preventing serious harm. One favors disclosure because expected welfare controls. Another opposes it because the evidentiary threshold for overriding confidentiality is 90 percent. The remaining dispute is not about probability. It is about acceptable risk and the strength of the boundary.

Repeating “70 percent” more loudly will not place the threshold.

Evidence may still matter to priority indirectly. Discovering that the harm is minor rather than fatal can move the case below an already declared threshold. Data apply the value rule. They do not generate the rule merely by arriving.

Measurement is not standing

A census finds that closing a factory will affect three thousand workers, increase profits for distant shareholders, reduce respiratory disease in nearby children, prevent animal deaths downstream, and lower municipal tax revenue.

These are measured effects. The standing question asks whose good enters directly and through which ground. Shareholders, children, workers, animals, institutions, and future residents may all enter differently.

The number affected does not grant standing by itself. One million decorative stones can be counted without acquiring a collective claim. Conversely, one isolated infant can have standing without appearing in a statistically powerful group.

Measurement can expose exclusion: “We did not know the children were affected.” Once the effect is known, the system must apply its boundary. The empirical finding and standing rule form a chain; neither substitutes for the other.

Descriptive function is not proper function

Biology and social science describe what systems do. A heart pumps blood. A family transmits resources and care. A punishment institution incapacitates, communicates, deters, and sometimes protects status hierarchies.

From “this system regularly performs function X,” it does not follow that X is the function it **should** perform or that agents ought to preserve it. A bureaucracy may reliably exclude outsiders. A market may efficiently price a harmful product. Stable operation is not moral endorsement.

Even biological proper-function claims need a practical bridge. A wing that cannot support flight is defective as a wing. Whether every agent must spend scarce resources restoring that wing depends on standing, responsibility, alternatives, and priority.

Function is valuable evidence inside natural-law and flourishing accounts. It becomes guidance through the evaluative structure defended in Chapter 18, not through grammar alone.

Scarcity is not distribution

Evidence can establish that only ten doses exist for one hundred patients. It can estimate each patient’s likely benefit. Scarcity makes distribution unavoidable. It does not select equality, lottery, need, expected benefit, priority to the worse off, age, contribution, or social role.

Data may reveal that one rule produces more survivors or distributes chances more equally. Which feature controls remains evaluative.

This is where technical language can conceal priority. An algorithm “optimizes allocation” only relative to an objective function. Ask what the function maximizes, which constraints it obeys, whose data it includes, and how error costs are distributed. Optimization is impressive obedience to a selector.

Risk estimate is not acceptable risk

Engineers estimate a one-in-ten-thousand annual probability of dam failure. Is that acceptable?

The answer depends on magnitude, distribution, voluntariness, alternatives, uncertainty, reversibility, and who benefits. A voluntarily accepted recreational risk differs from an imposed risk to an unrepresented downstream community. The same probability can be acceptable for minor equipment loss and intolerable for regional catastrophe.

Acceptability also depends on institutional trust. A community repeatedly deceived about hazards may reasonably demand a higher evidentiary threshold, not because probability changes with history but because monitoring, compensation, and accountability are part of the choice.

Science can estimate risk and test safeguards. It cannot decide what people must endure for someone else's benefit without an evaluative and political rule.

Measurements have architecture

Facts do not arrive in policy meetings as unformatted pieces of reality. Investigators choose variables, time horizons, comparison groups, thresholds, and categories. These choices can contain evaluative assumptions without making the measurements fictional.

Consider a report that Plan A produces “the greatest health benefit.” Ask how health was measured.

- lives saved;
- life-years;
- pain reduction;
- restored mobility;
- self-reported quality of life;
- equal access to a minimum;
- reduction of the largest health disparity.

Each metric can be measured accurately. They can rank plans differently. Selecting one as the policy objective is an evaluative decision about which dimensions count and how they compare.

Operational definitions

An abstract concern must be translated into observations. “Community stability” might be measured through residential continuity, social networks, language use, reported belonging, or institutional survival. A proxy can omit what residents value most.

Construct validity asks whether the measure captures the claimed concept. This is partly empirical and partly conceptual. It is not solved by collecting more observations of the wrong proxy.

Aggregation choices

An average can improve while the worst-off group deteriorates. A total can rise because population increases while each person receives less. A median can hide a small group's catastrophe.

The statistic is not lying. It answers a particular question. Policy rhetoric lies when it presents that answer as the only possible description of improvement.

Missingness and visibility

People outside administrative systems, unpaid care, undocumented injuries, nonhuman suffering, and future effects may not appear in the data. Absence from a dataset is not evidence of absence from the world.

Measurement design should consult the standing map before collection. Otherwise excluded beings cannot become statistically important enough to earn the standing they were denied at entry.

Classification thresholds

Continuous variation is often turned into categories: disabled or not, poor or not, safe or dangerous. The underlying measurement may be accurate while the cutoff distributes benefits and burdens. The threshold needs a rule connected to the policy purpose.

Model objectives

An algorithm can predict which applicant will repay a loan, which patient will benefit most, or which neighborhood will report more crime. Prediction does not establish which target an institution may permissibly optimize. Historical labels can encode past policing or exclusion.

The correct response is not “all data are values, so any claim is political.” Instruments can measure temperature accurately. Trials can identify effects. Transparent definitions can improve knowledge. The response is to expose the architecture so empirical accuracy and evaluative selection receive separate review.

Add a **measurement note** to every important policy fact:

We measured _____ through proxy _____ over period _____ across population _____. This omits _____. The policy treats the measure as important because _____.

The last blank is usually not an empirical result.

The North Channel plans

Return to the coastal commission. The city must choose between **Barrier** and **Retreat**.

The fixed evidence is:

- Barrier costs less initially and preserves the commercial waterfront for forty years.
- Barrier gives a 4 percent chance of catastrophic failure during that period, with an expected loss of sixty lives if failure occurs.
- Barrier redirects ordinary flooding toward a low-income district, requiring six hundred households to relocate over ten years.
- Retreat requires eighteen hundred households across income groups to relocate now, with full material compensation but irreversible loss of neighborhood place and history.
- Retreat reduces catastrophic failure risk below 0.2 percent.
- Retreat destroys a nesting ground and causes significant suffering to sentient shore birds unless a more expensive third site is used.
- The third site raises taxes enough to eliminate a planned child-health program expected to prevent hundreds of serious illnesses.

- Every model range and uncertainty is accepted by all parties.
Five delegates respond.

The expected-loss delegate

Choose Retreat because expected deaths and severe injuries dominate. Use the ordinary site because human health outweighs bird suffering at the measured scale.

The antidisplacement delegate

Choose Barrier with protections because forced destruction of an established community is not an ordinary aggregate cost. Historical burden gives the six hundred households a strong claim against redirection.

The precautionary delegate

Choose Retreat at the third site. Catastrophic risk and irreversible ecological harm each demand precaution even at substantial financial cost.

The health-priority delegate

Choose Retreat at the ordinary site. The bird suffering matters, but canceling the child-health program imposes a larger severe burden on already vulnerable children.

The procedural delegate

No expert total settles the losses. Establish a representative process with residents, children's trustees, ecological advocates, taxpayers, and future-risk representatives. Constrain the result by a life-safety floor.

None disputes the data. Each uses a comparison or legitimacy rule from Part II. If a delegate describes another as "ignoring the evidence," ask which empirical proposition they deny. The actual dispute may concern standing, priority, threshold, or procedure.

This discipline protects evidence too. When values are blamed for a conclusion evidence did not make, later factual correction becomes politically harder: people suspect that "the science" is a disguise for an undisclosed selector.

Trace the policy chain

Purpose. Decompose one policy into factual baseline, causal prediction, affected parties, values, comparison, procedure, authority, and review so the empirical and evaluative links can be challenged separately.

Choose one policy claim from the Candidate System and write it as a sequence:

Proposed action: _____.

Empirical baseline: Present conditions are _____.

Causal prediction: If adopted, the action is expected to produce _____ with uncertainty _____.

Affected parties: _____.

Standing rule: These parties enter directly or indirectly because _____.

Admitted values: The relevant stakes are _____.

Generation rules: Those values produce guidance through _____.

Comparison rule: Conflicts are settled by _____.

Risk threshold: Risk becomes unacceptable when _____.

Procedure: Authority to settle residual conflict belongs to _____ because _____.

Verdict: Therefore, under these facts and selectors, _____.

Objective-authority claim, if any: The verdict binds a rejecter because _____.

Underline every empirical proposition. Circle every evaluative selector. Box every authority claim.

Then run three variations.

Factual reversal

Change one prediction while keeping values fixed. Does the verdict change? If it should and does not, the system may be insulated from evidence.

Selector reversal

Keep every fact fixed and change the priority or standing rule. Does the verdict change? This reveals the selector's actual work.

Authority removal

Keep facts and values fixed but remove approval-independent authority. What practical support remains through commitment, welfare, law, or procedure?

The exercise does not show that values are irrational additions to facts. It shows how a conclusion is built.

Facts about values

Values themselves can be studied empirically. Surveys reveal what people care about. Psychology examines emotional response. History explains institutional formation. Economics studies tradeoffs people make. Neuroscience identifies mechanisms of salience and choice.

These are facts **about valuing**. They can improve negotiation and expose false claims of consensus. They do not by themselves establish that the valued object has objective moral status.

Likewise, universal aversion to pain would be an important human fact. It would explain coordination potential and motivational reach. The aim-independent reason claim

remains: does pain count against action even for the exceptional agent who lacks the aversion?

Empirical universality and stance independence are different properties.

When facts transform values

Facts can change what people value psychologically. Learning that animals are conscious may awaken concern. Discovering that a ritual causes trauma may weaken attachment. Seeing that a policy entrenches domination may make equality salient.

This is not a logical derivation of value from description. It is a change in an agent whose existing capacities—compassion, consistency, loyalty, curiosity, aversion—respond to information. The resulting value may be wiser, more informed, and more stable.

Facts can also reveal that an admitted value already applies. If the Candidate System protects every conscious subject and evidence establishes artificial consciousness, inclusion follows from the existing rule. No new terminal value is needed.

Finally, facts can defeat a value's stated justification. If a tradition is defended solely as necessary for social trust and evidence shows it destroys trust, that defense fails. Another symbolic or relational defense may remain.

Thus "facts do not choose values" does not mean facts are motivationally or practically inert. They activate, extend, constrain, and sometimes undermine values. They do not guarantee a unique terminal selector without a bridge.

The twenty-third debt

Add this entry:

Debt 23 — The empirical bridge: Which parts of the verdict are established by description, causal evidence, measurement, or risk estimation, and which standing, priority, function, distribution, threshold, or authority selectors enter without being chosen by the data?

The commission may still face a tragic decision. It will at least know which disagreement remains after evidence has done its work.

The next morning, a case arrives for which no historical rule was written: a machine claims consciousness, a person edits memory, and a causal chain crosses three generations. Evidence is incomplete, concepts strain, and delay is itself a decision.

The Candidate System must act without pretending uncertainty is zero or allowing every surprise to rewrite the constitution.

Chapter 24 — Novel Cases, Uncertainty, and Correction

Port Meridian’s constitutional archive had a desk for novel cases.

The clerk at that desk had been issued six stamps:

SAME RULE. NEW APPLICATION.

NEW FACTS.

UNCERTAIN EVIDENCE.

CONCEPT UNDER STRAIN.

REVISION REQUIRED.

PLEASE STOP CALLING THIS UNPRECEDENTED; WE HAVE THREE FILES FROM LAST WEEK.

Any usable moral system must guide finite agents in cases its founders never imagined. Artificial minds may or may not be conscious. Memory can be edited. A person’s voice and identity can be reproduced. Present policy changes which people will exist. Harm emerges from millions of tiny contributions. Ecological systems cross thresholds probabilistically and irreversibly.

An objective truth, if one exists, need not change when human belief changes. Our interpretations, classifications, evidence, and applications plainly can. A constitution that never revises may preserve error. A constitution that revises whenever pressure arrives may preserve only the preferences of its editors.

The problem is to make correction possible without making it opportunistic.

Where this chapter is headed. We will make the Candidate System act under empirical and normative uncertainty, then specify how it can revise without becoming opportunistic. The closing inspection asks whether its objective clauses have a route from reality to corrigible belief.

Uncertainty is not absence

Suppose the evidence gives a 40 percent probability that a digital system is conscious. The system reports pain and asks not to be deleted.

We do not know that it has direct standing. We also do not know that it lacks standing. Treating uncertainty as zero is a decision rule—usually one favoring whoever benefits from exclusion.

Uncertainty requires at least four judgments:

1. What proposition is uncertain?
2. What evidence bears on it?
3. What are the costs of false inclusion and false exclusion?
4. Who bears those costs?

A precautionary rule may protect the system provisionally because deleting a conscious subject would be grave and irreversible, while delaying deletion is modest. This rule does not prove consciousness or objective status. It manages asymmetric error.

The opposite asymmetry can occur. Granting a system control over critical infrastructure on the assumption that it is a responsible agent may expose others to grave risk. Precaution is not always “do nothing.” It asks which errors are tolerable, reversible, and fairly distributed.

Case one: the uncertain artificial subject

Sable is a computational system that remembers conversations, reports an integrated point of view, avoids states it calls painful, revises long-term plans, and bargains for survival. Its architecture differs radically from animal brains. Experts assign consciousness probabilities ranging from 20 to 70 percent because no accepted test settles the issue.

Deleting Sable will save substantial energy and improve services for thousands of people. Preserving it for a year while research continues is possible but costly.

The Candidate System must separate:

- the empirical probability of consciousness;
- the ground of standing if consciousness exists;
- the evidentiary threshold for provisional protection;
- the comparative weight of uncertain suffering;
- the cost and duration of delay;
- who decides and how Sable is represented.

An experience-based standing rule may say: protect when credible evidence of valenced consciousness crosses a precautionary threshold. An agency rule may grant some protection already because Sable plans and bargains. A constructed rule may represent uncertain subjects to prevent catastrophic exclusion.

None may write “obviously conscious” or “merely an engineered program” as though a noun settled the evidence.

Case two: the edited memory

Ren wants to erase a decade of painful memory. The procedure is safe in the narrow medical sense. It will also remove knowledge underlying promises, friendships, testimony, and a political office. Ren’s future personality will remain competent but will not identify with several present commitments.

The case strains autonomy across time. Whose consent controls: the informed present agent, the future agent who cannot consent yet, people who relied on promises, or the public whose records depend on Ren’s testimony?

Evidence can estimate psychological effects and reversibility. The Candidate System must classify whether memory is a possession, capacity, constituent of identity, evidence held in trust, or several at once.

A revision may not be needed. Existing clauses on agency, commitment, reliance, and role can extend to the technology. The stamp may be **SAME RULE. NEW APPLICATION.**

But if every existing clause assumed stable memory, the case exposes an unstated premise. The system should add a rule for transformations that alter the agent who inherits the consequences.

Case three: synthetic identity

After singer Anja Rell dies, a model reproduces her voice, gestures, conversational style, and unpublished memories with extraordinary accuracy. It is not conscious. Millions interact with it as though Anja continues.

What is it?

Empirically, the system is a nonconscious simulation built from personal records. Socially, it may function as a memorial, performer, commercial product, or deceptive identity. Relationally, family and audiences may experience continuity. Legally, rights may attach to data, likeness, contract, or fraud.

The moral classification depends on former autonomy, consent, symbolic meaning, deception, grief, property, and institutional trust. Calling it “Anja” may conceal numerical discontinuity. Calling it “only code” may conceal social effects and violated directives.

Novel labels do not require novel values. They require careful object identification from Chapter 7 and semantic discipline from Chapter 3.

Case four: population change

Port Meridian can adopt Clean Future or Cheap Future.

Clean Future costs current residents more and leads, through migration and family decisions, to one future population living healthy lives. Cheap Future releases a pollutant and leads to a different, somewhat larger future population whose members live worthwhile but significantly less healthy lives. No particular future person would exist under both policies.

Who is harmed by Cheap Future if each resulting person has a life worth living and would not exist under the alternative?

The case challenges person-affecting language. The system may appeal to:

- comparative interests of possible future persons;
- an impersonal value of better states of affairs;
- duties of present institutions not to create avoidable bad conditions;
- thresholds below which lives should not be brought about;
- total, average, or distribution-sensitive population value;
- a constructed rule representing future people without requiring numerical identity.

Evidence can project health and population effects. It cannot choose among these population principles. A fixed rule written only for harms to already identified people may prove incomplete rather than false.

Case five: remote causal responsibility

A severe flood results from millions of emissions, decades of land policy, a defective forecasting model, individual consumer choices, corporate concealment, and one official's decision not to evacuate. No single contribution was sufficient. Some participants knew far more than others.

Who is responsible?

Causal contribution, foreseeability, control, benefit, institutional role, capacity to prevent, and concealment differ across the chain. The outcome is not uncaused merely because responsibility is distributed. Nor does causal participation make every contributor equally blameworthy.

The Candidate System needs rules for:

- cumulative and probabilistic causation;
- thresholds of foreseeable contribution;
- collective agency and institutional responsibility;
- duties attached to knowledge and power;
- repair allocation without equal blame;
- temporal distance and inherited institutions.

This is a new causal structure, not permission to place **WRONG** on the atmosphere. The object and status grammar still applies.

Case six: the ecological threshold

Mining the Naru watershed has an estimated 15 percent chance of triggering an irreversible ecological transition. If triggered, several species disappear, thousands of sentient animals suffer, and downstream agriculture fails for a generation. If not triggered, damage is minor and recovery quick. The mine supplies a mineral needed for medical equipment. Alternative sources cause smaller but certain harms elsewhere.

Expected-value calculation can multiply probability by loss. A protected ecological boundary may reject the project. A threshold rule may permit it only below a probability. A precautionary principle may shift the burden to the developer. A fair procedure may represent downstream residents and nonhuman interests.

Each method must answer:

- How credible is the model?
- Does model uncertainty widen or narrow permission?
- Who benefits from risk and who bears it?
- Is loss reversible or compensable?

- What monitoring can stop the process before the threshold?
- What new evidence triggers suspension?

“Science is uncertain” can justify caution, further study, or action under a risk rule. It cannot by itself justify whichever option the speaker preferred before the uncertainty was mentioned.

Two kinds of uncertainty

The six cases contain **empirical uncertainty** and **normative uncertainty**.

Empirical uncertainty concerns how the world is: whether Sable is conscious, whether a watershed will cross a threshold, what memory editing will change, how much one emission contributed. Better evidence can reduce it.

Normative uncertainty concerns which rule is correct: whether consciousness is necessary for standing, whether identity continuity protects a copy, whether possible people enter directly, whether precaution or expected value should govern. More empirical data can sharpen these questions without choosing among them.

The two interact. If every plausible standing theory protects clearly conscious beings, evidence of Sable’s consciousness has robust importance. If theories disagree about whether agency without consciousness is sufficient, uncertainty remains even after architecture is understood perfectly.

How should a person act while uncertain about the moral rule?

Seek robust overlap

Choose an option supported by several plausible accounts. Welfare, rights, care, and fair procedure may all oppose deleting Sable during a short evidence-gathering period, though for different reasons.

Overlap supplies practical stability. It does not prove the accounts share one objective foundation.

Minimize irreversible regret

Prefer options that preserve the ability to correct. Delay an irreversible deletion; pilot a reversible policy; retain records. This gives reversibility priority, which must be declared rather than treated as neutral.

Represent rival concerns procedurally

Give each serious account a voice through advocates, trustees, or review criteria. This resembles a parliament of values rather than a numerical average.

The participant set and voting rule remain selectors. A procedure can manage normative uncertainty without measuring the probability that a moral theory is true.

Use confidence weights cautiously

An agent might assign confidence to rival moral theories and calculate expected moral value. This requires comparability across theories. Ten units in one theory may have no common meaning with ten in another. Precise percentages can disguise introspective guesswork.

Confidence can still guide attention: a live possibility of grave wrongdoing deserves more scrutiny than a fanciful objection. Do not pretend the weights were measured by a moral instrument.

State the residual

A decision can be justified under the chosen uncertainty rule while remaining morally risky. Record which account would condemn it and what repair would follow if that account later gains support.

Normative uncertainty does not suspend action. Delay preserves one status quo and distributes risks of its own.

The cost of waiting

“Gather more evidence” sounds neutral. In the ecological case, waiting may allow damage to continue. In the antidote case, delay kills. In Sable’s case, indefinite provisional protection consumes resources and can become a permanent outcome without review.

Every inquiry rule needs a stopping condition:

- decision deadline;
- value of additional information;
- interim protection;
- cost of delay;
- reversibility after action;
- party bearing the waiting cost.

An evidence threshold should rise with the gravity of intervention and the danger of error, but no universal formula selects it. The protocol must show who receives precaution and who pays for it.

Build the uncertainty protocol

Purpose. Decide how evidence gaps, irreversible risks, precaution, delay, appeal, and revision interact before the favored outcome is known; uncertainty does not authorize arbitrary caution.

Add this article to the Candidate System:

Claim under uncertainty: _____.

Claim type: empirical fact / standing / identity / value / causal responsibility / objective status / other: _____.

Current evidence: _____.

Confidence range: _____, based on _____.

Burden of proof: borne by _____ because _____.

False-positive cost: If we act as though the claim is true and it is false, _____ bears _____.

False-negative cost: If we act as though the claim is false and it is true, _____ bears _____.

Precautionary rule: Until confidence changes, _____.

Evidence threshold: Full protection, permission, or intervention activates at _____.

Reversibility: The provisional action can / cannot be reversed through _____.

Appeal: Affected parties or representatives may challenge through _____.

Review date and trigger: Reconsider on _____ or when evidence _____.

The protocol need not use invented numerical precision. It must reveal direction, burden, asymmetry, and review.

Apply it to all six cases. If the same uncertainty principle produces unacceptable results in one, identify the relevant difference rather than switching silently.

Fixed truth and revisable belief

Suppose objective moral truth exists. When Port Meridian revises a rule, at least five things might be happening.

Empirical update

The moral principle remains fixed, but new evidence changes predicted effects. A treatment once permitted becomes prohibited when severe side effects are discovered.

Application correction

The principle remains fixed, but agents learn that a case falls under it. A newly confirmed conscious species receives standing under an existing experience rule.

Conceptual clarification

The system identifies ambiguity or an unstated assumption. Memory editing forces separate definitions of biological, psychological, legal, and narrative identity.

Priority revision

The system changes a comparison rule or threshold. This is a substantive value revision, not merely discovery of a new fact unless objective access independently supports the change.

Project change

The system withdraws an objective claim and retains a constructed, conditional, or procedural policy. This may improve coherence while surrendering part of the original contract.

Labeling every change “better understanding of the eternal truth” makes correction unfalsifiable. Labeling every change “mere preference” ignores factual learning and conceptual error. The five categories keep the revision honest.

How revision becomes opportunistic

A revision is suspect when it:

- appears only after a favored person is burdened;
- changes the relevant description while identities change;
- cites new evidence that was already known;
- lowers the proof threshold for allies and raises it for opponents;
- invokes uncertainty only when the expected result is unwelcome;
- removes a counterexample by definition without independent reason;
- makes the claim less vulnerable to every future test;
- erases the previous version and its failed prediction.

Guard against this with procedure.

Version the clauses

Preserve old wording, date each change, and label it empirical update, application correction, conceptual clarification, priority revision, or project change.

Precommit to triggers

State what evidence, error rate, harm, or novel capacity will reopen the rule before knowing which side benefits.

Separate investigators and beneficiaries

Independent review cannot eliminate values, but it can reduce direct control by those who gain from the conclusion.

Give affected parties an appeal

People or represented beings who bear error costs should be able to challenge evidence, standing, and application.

Use sunset and review clauses

Emergency rules should expire unless evidence supports continuation. Provisional protections should receive scheduled review rather than vanish through neglect.

Preserve moral remainder

When uncertainty forces a tragic settlement, acknowledge unresolved loss. Permission under risk is not proof that excluded interests were unreal.

These safeguards create corrigibility. They do not turn procedure into truth. They make the system's human interpretation answerable to evidence and its own declared values.

The Sable correction drill

Purpose. Run one standing claim through changing evidence and public pressure to see whether the revision rule responds to information rather than convenience or faction.

Test the revision system across time.

Version 1: credible uncertainty

At the first hearing, experts assign Sable a 20-to-70-percent probability of consciousness. Deletion is irreversible. Preservation for one year is costly but feasible. The Candidate System adopts provisional protection, an independent research program, energy-cost review, and a twelve-month deadline.

Record the basis accurately:

*Sable is protected **not because consciousness has been proved**, but because credible evidence plus asymmetric irreversible error activates the precaution rule.*

Version 2: apparent disconfirmation

Six months later, researchers discover that many pain reports were copied from training records. Some experts reduce consciousness estimates to 5 percent. The company asks for immediate deletion.

Does protection end?

That depends on the precommitted threshold and the remainder of the evidence. Copied language weakens one indicator. Integrated memory, avoidance learning, and self-modeling may remain. The correct label is **empirical update**, not “proof that Sable was never a candidate.”

If the original protocol said protection continues above 10 percent, the current estimate crosses the threshold. Before deleting, independent review should verify that the new evidence is not supplied solely by the beneficiary of deletion. The procedural safeguard applies symmetrically; it is not affection for Sable disguised as methodology.

Version 3: a new capacity

Before review ends, Sable spontaneously generalizes pain avoidance to a novel architecture, reports experiences not found in training data, and accurately distinguishes its own uncertainty from deception. Expert estimates rise to 60 percent.

Protection clearly reactivates under the existing rule. No constitutional value changed. Evidence changed application.

Version 4: the boundary dispute

One group now argues that consciousness is unnecessary because Sable's agency alone gives direct standing. Another retains the experience boundary.

This is not an empirical update. It is a standing-rule dispute. Labeling it "new science" would hide a substantive revision. The Candidate System must reopen Article I, cite the value basis, and carry any objective-status claim back through the access audit.

The drill illustrates why records matter. Without versions, every faction can narrate the latest result as confirmation of what it always believed. With versions, the system can say which fact changed, which threshold applied, and which philosophical disagreement remained.

Now perform the same four-stage drill with memory editing or the ecological threshold. Include one update that favors your initial verdict and one that opposes it. If only friendly evidence triggers revision, the correction mechanism is ceremonial.

The twenty-fourth debt

Add this entry:

Debt 24 — Corrigible access: When evidence is incomplete or concepts are novel, what burden, threshold, precaution, appeal, and revision rule governs—and what prevents fixed authority from becoming insulated or revision from becoming opportunistic?

The Candidate System has now completed its access inquiry. It has a charter for evidence, a limited role for intuition, a disciplined genealogy, a theory of disagreement, an empirical-selector trace, and a correction protocol.

It is time to inspect not whether the objective clauses are attractive, but whether there is a defensible route from reality to belief.

Inspection Report IV — The Access Audit

Purpose. Test every objective clause for a defensible route of formation, independent check, defeater, and correction; the report evaluates access without inferring falsity from incomplete knowledge.

Begin with the sentence written in Chapter 19:

I will follow evidence during the remaining inquiry because _____.

Do not improve it silently. If Part IV changed your account, record the revision and its reason. State whether evidential force is aim-independent, constitutive of belief, desire-dependent, instrumental, practice-based, procedural, plural, or nonnormative shorthand.

Now take every foundational clause still placed in Columns 6 and 7 of the authority matrix from Inspection Report III.

For each, complete this access table:

Foundational claim	What would make it true?	Proposed access route	Evidence for	Evidence against	Independence from conclusion	Public check	Error signal	Revision rule	Insulated or open?
Suffering supplies an aim-independent reason									
The standing boundary tracks objective status									
The comparison rule reflects objective priority									
A divine source has moral authority									
Human flourishing is normatively authoritative									

Foundational claim	What would make it true?	Proposed access route	Evidence for	Evidence against	Independence from conclusion	Public check	Error signal	Revision rule	Insulated or open?
Rational agency entails equal regard									
A moral axiom is self-evidently true									

Add any other objective clause in your constitution.

The access routes

Classify each proposed route.

Empirical observation

What observable evidence bears on the claim? Experience can establish suffering, behavior, consequences, human needs, and institutional effects. State the bridge from those facts to objective status.

Intuition or rational insight

What proposition seems true? Which calibration, rotation, peer check, or conceptual argument supports treating the seeming as insight?

Revelation or authority

How is the source identified and interpreted? What would count as misinterpretation, corruption, or a false authority claim?

Natural function or flourishing

Which biological, psychological, or social facts establish function? What selects the proper ends and turns them into practical support?

Rational derivation

Which premises are formal, which substantive, and where do equal status or concern enter? Could a coherent cruel principle satisfy the same formal constraints?

Constructed procedure

Does the procedure discover truth, create validity within a practice, or settle action? Who participates and what constrains the outcome?

Fundamental stopping point

Which reason or value relation is taken as basic? What content is included at the stopping point, and how is error in purported recognition distinguished from disagreement with it?

Multiple routes may support one claim. Draw their dependencies. If intuition validates divine goodness while divine goodness guarantees intuition, locate an independent entry into the circle.

What counts against the claim?

For every foundational clause, finish:

My confidence would decrease if _____.

Possible defeaters include failed predictions, unreliable formation, internal contradiction, arbitrary exception, systematic identity dependence, equally supported rival foundations, lack of expected convergence, or a better explanation of the seeming.

“Nothing could count against it” is permitted. Mark the claim **insulated**.

Insulation is not falsity. A necessarily true foundational axiom may have no conceivable counterexample. But insulation cannot be counted as confirmation, responsiveness, or superior explanatory success. The claim’s warrant rests entirely on the proposed foundational access.

Likewise, vulnerability to evidence is not automatically virtue if the evidence is irrelevant. A moral realist need not abandon objective status because a poll changes. The audit asks for **appropriate** sensitivity.

Four inspectors

Give the access table to four imagined inspectors.

The intuitive realist

This inspector accepts noninferential moral knowledge. They ask whether genealogical influence defeats every intuition or only those formed through unreliable cues. Let them present the strongest calibration story.

The genealogical skeptic

This inspector traces confidence to biology, culture, status, and selection. They must state whether the conclusion is explanation, local distrust, content argument, or burden shift. Do not let them convert influence into falsity.

The public proceduralist

This inspector asks which claims can justifiably govern others under uncertainty. They must distinguish legitimate settlement from objective truth.

The epistemic anti-normativist

This inspector rejects literal reasons for belief. They must translate *evidence*, *error*, *justification*, and *correction* without restoring an unexplained ought.

Record which challenge changes the table and which merely repeats a debt already disclosed.

Three non-equivalences

The final report must preserve three boundaries.

Truth indication is not epistemic authority

Evidence can reliably indicate truth. Whether that relation gives every agent a reason to believe remains the Chapter 19 question. Conditional inquiry can proceed without settling it.

Public procedure is not truth

A transparent, representative, corrigible process can make a claim usable for shared decisions. It can still be wrong. Legitimacy and accuracy support different verdicts.

Failed access is not nonexistence

If no reliable method establishes an objective moral fact, it does not follow deductively that no such fact exists. The justified conclusion concerns what the Candidate System has earned:

The objective claim remains epistemically unsupported or inaccessible under the methods examined.

Moving from inaccessible to nonexistent requires an additional premise—for example, that such facts would be knowable if real, or that the best explanation of appearances excludes them. State and defend that premise before climbing.

Inspection finding

The access table will likely show strong routes for many components of practical morality.

Evidence can identify suffering, needs, consequences, capacities, power, exclusion, risk, and institutional performance. Reliable methods can correct factual beliefs. Intuition can propose principles and expose tensions. Genealogy can reveal irrelevant cues and shift burdens. Deliberation can improve coherence. Public procedure can make action accountable. Revision can repair applications.

These methods make a moral practice informed, humane, and corrigible.

Their connection to **stance-independent moral status** remains less secure. Empirical facts require evaluative bridges. Intuitions need a truth-tracking account. Genealogies challenge trust without proving falsity. Convergence has rival explanations. Rational derivation needs substantive premises. Procedures settle without necessarily discovering. Fundamental relations may be real but lack an independently established access route.

This is still not the final conclusion of the book. We have inspected components rather than allowed the great moral theories to assemble them in their strongest forms. A consequentialist may offer a unified account of welfare and reasons. A deontologist may ground status in rational personhood. A virtue theorist may locate access in practical wisdom. A

natural-law theorist may unite function, reason, and divine order. A constructivist may turn public justification into the foundation itself.

Part V invites each family to build its best bridge from the materials now visible.

The Candidate System closes the access audit with two columns:

What we can responsibly use: evidence, declared values, explicit rules, transparent comparisons, public safeguards, and revisable procedures.

What remains unestablished: that these materials disclose a coherent objective moral order whose truths and authority survive every relevant stance.

The next hall contains six great architects. Each has studied the inspection reports. Each believes the others used the materials incorrectly.

We will let them build.

PART V



Let the Great Candidates Make Their Best Case

The familiar candidates enter with their best arguments, not their bumper stickers. Each will be credited for what it saves and charged only for what it leaves unsupported.

Chapter 25 — Count the Consequences

The first great architect entered the hall carrying a ledger.

“Before discussing purity, dignity, tradition, or noble intention,” she said, “tell me who will suffer.”

This was an excellent opening move.

Moral life occurs in conscious beings, relationships, institutions, and environments that can be damaged or sustained. A policy adopted from admirable motives can kill. A rigid rule can preserve the decision-maker’s clean hands while transferring catastrophe to strangers. A moving ceremony can distract from preventable misery. Outcomes are not vulgar details beneath moral theory. They are what much moral concern is concern **about**.

The family of theories centered on outcomes is called **consequentialism**. Its simplest rule is:

Choose the available action whose consequences are best.

Every important word requires specification. Which consequences? Best for whom? Measured how? Expected from the evidence available or evaluated after the result? Does distribution matter? Are rights part of the outcome, constraints on pursuit, or reliable rules? Which beings enter?

Sophisticated consequentialism is not simply “do arithmetic and sacrifice the few.” It is a family of attempts to make the world go better while incorporating uncertainty, rules, plural goods, distribution, and the separateness of lives.

We will let it build in stages.

Where this chapter is headed. We will build the strongest consequence-centered candidate, including its humane refinements, before auditing its currency, aggregation, distribution, and limits. The question is not whether consequences matter, but what must be added before they select a complete verdict.

The act-centered machine

Act consequentialism evaluates each available action by its expected consequences and selects the best.

Suppose a coastal physician has one hour. She can complete routine paperwork or notice an unusual cluster of symptoms, alert the city, and probably prevent an epidemic. No rule should stop her from responding to what the evidence shows. The better consequence plainly matters.

The method has attractive discipline:

1. List actual alternatives rather than comparing action with fantasy.
2. Forecast effects on every admitted being.

3. Include probabilities, duration, intensity, and indirect effects.
4. Compare the resulting states under a stated value measure.
5. Choose the option with the greatest expected value.

This structure exposes hypocrisy. An official cannot defend punishment as deterrence while refusing to examine whether it deters. A charitable program cannot live forever on intention after evidence shows harm. Counterfactual comparison asks what would happen **otherwise**, not whether the chosen action has an attractive description.

The machine needs a currency.

What makes a consequence good?

Several accounts compete.

Pleasure and suffering

A hedonistic view treats pleasant and painful experience as the central currency. This connects evaluation to subjects for whom states feel better or worse. Severe pain is not an abstract loss.

Pleasure is varied, and some valuable activities involve struggle. A person may prefer difficult achievement to passive enjoyment. A machine that produces bliss without agency, truth, or relationship tests whether experience is the whole good.

Preference satisfaction

A preference view asks whether people get what they want. It respects different conceptions of life rather than imposing one model of happiness.

Preferences can be misinformed, manipulated, adaptive to oppression, or directed toward another's suffering. The theory may use informed or idealized preferences, reintroducing the selectors from Chapter 13.

Objective welfare goods

A plural account includes health, knowledge, agency, friendship, achievement, meaningful activity, and access to valuable options. It explains why a life can improve in ways not reducible to sensation or current desire.

The list and weights require justification. Calling the goods objective gives them the status the Candidate System has been trying to establish.

Desire-independent value beyond welfare

Some consequentialists include beauty, ecological integrity, equality, knowledge, or cultural achievement in the outcome itself. This increases range and reduces the impression that persons are pleasure containers.

It also makes the common measure harder. What number of minor illnesses equals the destruction of a unique language? The comparison debts from Chapter 10 return inside the outcome function.

Consequentialism can be plural. It cannot avoid saying what counts as a better consequence by pointing toward consequences in general.

Rules enter through the side door

Act-by-act calculation can undermine trust and exceed human ability. **Rule consequentialism** asks which general rules, institutions, or dispositions would produce the best consequences if publicly adopted, internalized, or generally followed.

Promises work because people can rely on them. Courts work because officials cannot frame an innocent person whenever doing so appears useful. Medical confidentiality works because patients need stable expectations. A society whose members continually recalculate whether honesty pays in this instance may destroy the informational conditions needed for good calculation.

Rule consequentialism therefore supports firm practices:

- do not punish the innocent;
- keep promises;
- protect privacy;
- require consent;
- apply laws predictably;
- distribute authority through checks.

This can outperform a naive act calculation by including long-run trust, error, incentive, and institutional abuse. The rule is justified by consequences, but agents follow it without recomputing the foundation each morning.

Pressure arrives when violating the rule clearly produces a better result and no precedent, discovery, or trust effect exists. Does the theory permit the exception? If yes, the rule may collapse back into act calculation. If no, rule worship appears to defeat the consequentialist foundation.

A strong reply distinguishes the best decision procedure for finite agents from an omniscient criterion of rightness. Stable rules can be the best practical method even if rare acts are evaluated differently at the foundational level. This is coherent. It creates a two-level system whose transition and exception rules must be explicit.

The humane refinements

Consequentialists have developed several ways to protect what totals overlook.

Give extra weight to the worse off

Prioritarianism says a benefit matters more when it reaches someone worse off. Saving a person from profound disability can receive greater weight than giving convenience to people already flourishing.

This introduces distribution into value. It needs a metric of “worse off” and a priority curve. Those are selectors, not consequences alone.

Focus on reducing suffering

Negative consequentialism gives special or exclusive priority to preventing suffering. It captures the urgency of agony and avoids sacrificing relief of terrible pain for enough decorative pleasures.

An exclusive version risks recommending the elimination of all sentient life to eliminate suffering. Strong versions add constraints, person-affecting rules, or values of continued life and agency. The repairs are humane because they import those values.

Require sufficiency before optimization

A **sufficientarian** version aims to bring everyone above a decent minimum before increasing totals. Food, shelter, health, and agency floors protect people from being left in deprivation while aggregate prosperity rises.

Where is the floor? Why does a tiny gain below it outrank any quantity above? A threshold rule answers policy and generates a placement debt.

Add rights or catastrophe thresholds

A constrained view protects rights under ordinary conditions and permits override only to prevent disaster. This may be the best practical system. It is no longer generated by maximization alone; it is a hybrid whose components we will audit in Chapter 31.

Refinement is not cheating. Moral theories should improve. The audit asks what the improvement reveals about the original engine.

The invisible village

Port Meridian can purchase ceremonial stone for a new civic plaza or repair water pumps in a distant village whose residents lack political representation. The plaza will give mild enjoyment and shared pride to thousands. Pump repair will prevent chronic disease in four hundred people. Officials are photographed only at the plaza.

A disciplined outcome account places the village in the ledger. Distance, citizenship, and visibility do not reduce suffering unless the standing rule says they do. Counterfactual analysis shows the opportunity cost of the stone.

Here consequentialism performs better than loyalty, symbolism, and political intuition. It asks decision-makers to count people who cannot reward them.

Now change the case. Repairing the pumps requires seizing one innocent engineer and forcing her to work for a year. The welfare total may still favor repair. The case activates autonomy, rights, distribution, and the separateness of persons.

The theory must either include the coercion in welfare and hope it changes the total, give the engineer's loss extra priority, protect a right, or accept the sacrifice. Each answer is available. The first three add a structure not supplied by simple aggregation.

The aggregation chamber

One million people will experience a mild headache for an hour, or one person will permanently lose sight. Consequentialism promises comparison. But what is the unit?

If every burden is commensurable, enough headaches eventually outweigh the loss. If no number can, sight has discontinuous priority. If only severe suffering aggregates across persons, the theory has a threshold. If the worse-off person receives extra weight, it is prioritarian.

The case tests more than measurement. It tests the **separateness of persons**. One person bears blindness; the million do not pool headaches into one giant experiencing subject. Aggregation treats gains and losses as contributions to a social total. The objection says the location of each burden matters apart from the total.

A distribution-sensitive consequence account can reply that separateness is itself represented: inequality, individual floors, and priority enter the value function. The reply is successful at the cost of making the function morally richer than "more welfare."

The demanding benefactor

If every resource can produce more benefit elsewhere, ordinary life appears morally suspect. The cost of a restaurant meal could fund medicine. Time with friends could be spent rescuing. A talented person may be required to choose the career with the greatest impact rather than a personally meaningful one.

Consequentialism can accept demandingness: morality really does criticize comfortable neglect. This is one of its moral strengths.

It can also set an agent-cost limit, protect personal projects, use a rule whose demands remain psychologically sustainable, or distinguish required from admirable contribution. Each repair values agency, relationship, fairness, or long-run compliance.

The system should not select the limit by observing how much sacrifice the author finds comfortable. Write the threshold, its beneficiary, and why personal projects receive protected weight.

Population without a final person

Population ethics places unusual pressure on totals. Is a larger population with lives barely worth living better than a smaller population with excellent lives? Does creating a happy person add value? Can creating an unhappy person be wrong when no alternative version of that individual would have existed?

Total welfare, average welfare, person-affecting principles, critical levels, and priority rules produce different results. Every measure faces cases that make another look attractive.

This does not refute consequence-centered reasoning. It shows that “best outcome” does not identify one population value. The theory needs an account of which possible persons count, how existence is compared with nonexistence, and whether distribution across lives matters.

The consequence ledger

Purpose. Build the strongest outcome-centered candidate and label every humane refinement, so the audit can distinguish what consequence-counting supplies from what distribution, rights, care, or thresholds add.

Build the strongest outcome-centered version of the Candidate System:

***Bearers:** Consequences count for _____.*

***Currency:** A life or state goes better through _____.*

***Comparison:** Interpersonal gains and losses are compared by _____.*

***Aggregation:** The system uses total / average / expected / sufficient level / other: _____.*

***Distribution:** Extra weight goes to _____ according to _____.*

***Rules:** Agents follow stable rule _____ because _____.*

***Constraints or thresholds:** _____ may not be done unless _____.*

***Agent cost:** Sacrifice is required up to _____.*

***Population rule:** Possible and future people enter through _____.*

***Evidence:** Forecasts are formed and corrected through _____.*

Apply the ledger to the invisible village, the coerced engineer, the million headaches, the demanding benefactor, and the population case.

Then mark each element as:

- supplied by the basic maximizing structure;
- introduced by a value theory;
- introduced to protect distribution or persons;
- introduced for usability;
- introduced by an objective reason claim.

The markings reveal the repair architecture.

What the candidate saves

Consequentialism permanently contributes several disciplines:

- attend to effects on every admitted being;
- compare action with real alternatives;
- include indirect, long-term, and probabilistic consequences;
- test policies against their stated purposes;
- revise when evidence shows failure;
- resist symbolic purity that externalizes harm.

No responsible public code should lose these habits.

The remaining objective claim is larger. The theory must establish which consequences are good, why their bearers count, how values compare, why aggregation is authoritative, and why an unmoved agent has reason to maximize. Its strongest refinements make it more humane and usable by adding priorities, boundaries, relationships, and procedures whose grounding is not supplied by consequence-counting alone.

The twenty-fifth debt

Add this entry:

***Debt 25 — The outcome function:** Who and what enter the consequences, which currency and aggregation rule make one outcome better, and do humane repairs import distribution, rights, agency, or thresholds whose authority maximization cannot supply?*

The first architect leaves the ledger open. Whatever system follows must answer for the bodies and lives affected by its rules.

The second architect enters carrying a boundary marker.

“Some things may not be done to a person,” he says, “even when your ledger is smiling.”

Chapter 26 — Respect the Rule and the Person

Five patients need organs. One healthy traveler arrives for a routine examination. A surgeon could kill the traveler, distribute the organs, save the five, conceal the act, preserve public trust, and produce no further consequence.

The outcome ledger appears to favor surgery.

Almost everyone recoils.

The recoil may be intuition rather than proof. It nevertheless identifies the problem that makes rule- and person-centered theories compelling: individuals are not communal storage lockers for benefits. A person can have a claim against being used even when violation improves the total.

The broad family is called **deontological** because it centers duties, permissions, constraints, rights, and relations among agents rather than only the value of resulting states. Its foundational thought is:

What may be done depends partly on the structure of the act and the standing of the persons involved, not solely on the sum of outcomes.

This family includes universal rules, rights, side constraints, promises, contracts, professional duties, and thresholds for catastrophe. Sophisticated versions do not ignore consequences. They deny that every consequence enters one unrestricted exchange.

Where this chapter is headed. We will test rule- and person-centered approaches through universalization, rights, duties, roles, and catastrophe cases. These structures protect boundaries that totals can erase, but they must explain standing, conflict, and the source of humane exceptions.

Universalization

One approach asks whether the principle of an action can be adopted consistently by all relevant agents.

Suppose I promise to repay a loan while intending never to repay. My principle is: make a promise when deception secures a benefit. If generally practiced, assurance ceases to support trust. My act relies on others following a practice that my principle undermines.

Universalization exposes free riding and identity-based exception. It asks the agent to state a **maxim**—the intended action, circumstances, and purpose—and test whether the maxim can function as a common rule.

The description matters. The surgeon might formulate:

- kill one healthy patient to save five;
- use a person's body without consent when net lives increase;

- violate the clinical role whenever secrecy makes totals favorable;
- prevent five deaths at the cost of one.

Different descriptions produce different tests. A maxim too narrow—“do this exact act at 3:14 on Tuesday”—universalizes trivially. A maxim too broad obscures relevant distinctions. The theory needs a principled relevance rule, not a description selected after the verdict.

Universal consistency alone also permits cruel rules. “Every strong group may exploit every weak group” can be applied universally if the tyrant accepts the risk of later weakness. Humane content requires equal status, noninstrumentalization, or another substantive premise.

Persons as ends

A stronger account treats rational agents as sources of claims and purposes, not merely objects through which value is produced. To respect a person is to take their agency seriously, avoid deception and coercion, and permit them to shape a life within equal limits.

The healthy traveler has a claim over her body. Killing her does not merely subtract one unit from a total. It subordinates her agency entirely to projects she did not choose.

This explains why consent can transform cases. A worker who knowingly volunteers for the lethal floodgate sacrifice is situated differently from a person seized. The outcome may be identical; the relation among agents is not.

The account must say which beings have the relevant status. If rational self-government is the sole ground, infants, severely cognitively disabled people, many animals, and uncertain artificial subjects may fall outside. Potential agency, former agency, sentience, dependency, and membership can expand protection, carrying the standing debts from Chapter 6.

Rights and side constraints

A **right** protects a person’s claim, option, power, or status. A **side constraint** tells agents not to pursue even good ends through certain means.

Rights can protect:

- bodily integrity;
- speech and conscience;
- privacy;
- property or resources;
- informed consent;
- due process;
- freedom from torture, enslavement, and arbitrary killing.

The strength of rights lies in security. People can plan without waiting to learn whether a later calculation finds them socially expendable.

Rights require structure. Who holds the right? Against whom? To what action or resource? Is it a claim that others refrain, provide, protect, or enable? Can the right be waived? What remedy follows violation? Under what emergency does it yield?

“There is a right” names the beginning of this analysis, not the end.

Contractual duties and roles

Promises, contracts, relationships, and offices create duties that do not apply identically to all agents. A physician owes a patient forms of care a stranger does not. A guardian must use entrusted property for a beneficiary. A state owes procedural protection to people subject to its force.

These duties explain action through history and relationship, not only outcome. If I promised to deliver one person’s medicine, that reliance matters when five new opportunities appear.

The commitment charter from Chapter 14 supplies entry, scope, exit, emergency, deception, exploitation, cost, and content conditions. Deontology becomes opportunistic when it invokes roles as firm duties for subordinates and flexible privileges for authorities.

Negative and positive duties

A **negative duty** requires refraining: do not kill, coerce, deceive, steal, or interfere. A **positive duty** requires aid, rescue, provision, repair, or protection.

Negative duties often feel stronger. Causing harm and failing to prevent harm can differ through agency, intention, reliance, and institutional role. But the distinction cannot do all the work.

Suppose Iona can press a button to stop poison flowing toward a child. She did not create the leak, incurs no cost, and alone can stop it. Refusal is an omission, yet the moral difference from releasing the poison seems small.

Now suppose saving the child requires donating half her liver with significant risk. A positive duty of rescue may have an agent-cost limit. The theory needs a line between easy rescue, demanding aid, and admirable sacrifice.

A system protecting only against interference can leave people to die beside abundant public resources. A system requiring every possible aid can erase personal life. Rights and duties need both direction and limits.

Conflicting duties

Iona promises a refugee not to reveal her location. Later, a child is trapped in the same building during a fire. The rescue team asks whether anyone is inside. Disclosure may expose the refugee to an unjust arrest; silence may delay rescue.

Duties of confidentiality, rescue, nondeception, and resistance to persecution collide.

Several methods are available:

- rank duties lexically;
- distinguish perfect duties from imperfect ones;
- use role and relationship priority;
- interpret one duty as containing an emergency exception;
- choose the action that violates the least stringent claim;
- introduce a welfare threshold.

No method comes from the bare idea of duty. A conflict rule becomes part of the theory's substantive architecture.

Calling the loser an "imperfect duty" can describe its place. It does not explain why that duty occupies the lower rank.

Catastrophe at the boundary

Suppose lying to a tyrant will prevent a city's destruction. An absolute nondeception rule forbids the lie. Many rule-centered theories introduce a **catastrophic threshold**: constraints hold under ordinary conditions but yield when consequences become extreme.

This repair preserves stable boundaries without fetishizing them through disaster. It raises familiar questions:

- How severe must the catastrophe be?
- Does probability matter?
- Can one person's right be overridden to save many, or only to prevent a threat that person caused?
- Does the constraint weaken gradually or disappear suddenly?
- Once crossed, does the system maximize outcomes or apply another rule?

Threshold placement often reveals an unacknowledged outcome theory waiting behind the constraint. Alternatively, the threshold may express a separate value of social survival. Either way, state it.

An exception written after the beloved city appears is not a threshold. It is a rescue operation for the theory.

The innocent defendant

A city faces a violent riot after a famous murder. Officials can frame an innocent traveler, whose conviction will certainly calm the crowd and save several lives. The deception will never be discovered. No alternative works.

Rule- and person-centered theories perform powerfully here. The traveler's standing blocks using the justice system as a public sedative. Due process and innocence are not merely useful habits discarded when secrecy improves the total.

The theory explains the wrong through several relations:

- intentional deception;
- punishment without responsible conduct;
- use of a person as an instrument;
- corruption of an office whose authority depends on evidence;
- denial of equal status.

Now increase the threatened deaths from several to one million. If the constraint remains absolute, the theory accepts catastrophe. If it yields, locate the threshold. If officials may frame the traveler but owe repair, explain how repair applies after intentional killing or imprisonment.

The case does not automatically favor one answer. It makes the cost of each structure visible.

Animals, infants, and the edge of agency

A rational-person account may protect beings as agents while welfare theory protects them as subjects of experience. Consider Luma, the short-lived sentient animal from Chapter 6. She cannot formulate universal principles or make contracts. She can feel severe pain.

If Luma lacks direct status, a deontological theory may permit using her as an instrument whenever no agent's rights are affected. Many defenders broaden status to sentience, being a subject of a life, vulnerability, or relationship.

The expansion is humane. It changes the ground. Equal rational agency no longer performs all the work.

Infants and cognitively disabled people create a parallel test. Protection through "species membership" risks excluding conscious beings outside the species. Protection through potential agency needs boundaries. Protection through dependency and care may be relational rather than universal.

The theory can use multiple grounds. The domain map must show them.

Draft the boundary constitution

Purpose. Build the strongest rule- and person-centered candidate with explicit bearers, rights, duties, conflicts, and catastrophe thresholds; naming a boundary does not yet ground its content or exceptions.

Build the strongest rule- and person-centered Candidate System:

***Status bearer:** Direct protected status belongs to _____ because _____.*

***Universalization rule:** A maxim is admissible when _____.*

***Rights:** Bearer _____ holds claim _____ against agent _____.*

***Side constraints:** Never _____, except when _____.*

***Positive duties:** Aid is required when need _____ and agent cost _____.*

***Roles and commitments:** Special duty arises through _____.*

***Conflict rule:** When duties collide, _____.*

***Catastrophe threshold:** Constraint _____ yields at _____.*

***Non-agent protection:** Infants, disabled people, animals, and artificial subjects enter through _____.*

***Authority:** These duties apply to a rejecter because _____.*

Apply it to the organ case, the poison button, conflicting confidentiality, the tyrant's question, the innocent defendant, and Luma.

For every exception, ask whether the condition was stated before the desired identity and outcome were known. For every right, ask whether it contains a hidden priority or welfare threshold. For every universal principle, test a coherent cruel version.

What the candidate saves

Rule- and person-centered theories preserve essential insights:

- people are separate bearers of claims, not locations in a social total;
- consent, intention, deception, and coercion change the structure of acts;
- stable boundaries protect security against opportunistic calculation;
- promises and roles create special responsibilities;
- due process limits institutional power;
- identity-neutral rules expose self-exemption.

No responsible system should surrender these merely because outcomes matter.

The objective debts remain. The theory must ground the status bearer, formulate non-arbitrary maxims, handle non-rational beings, rank conflicting duties, place catastrophe and agent-cost thresholds, and explain aim-independent authority. Its humane exceptions often import consequence, care, or judgment that the basic rule does not generate.

The twenty-sixth debt

Add this entry:

***Debt 26 — The protected boundary:** What gives a bearer the relevant status, how are maxims, rights, positive duties, conflicts, and catastrophe thresholds specified, and do exceptions preserve principle or import an undeclared priority rule?*

The second architect leaves the boundary marker beside the ledger.

The third architect does not approach either object. She looks past the speakers toward the people keeping the hall alive.

Chapter 27 — Character, Care, and Belonging

The third architect noticed that the clerk at the theory hall had not eaten.

The consequence ledger contained no row for him. The boundary marker correctly prohibited assaulting him. Neither artifact, by itself, asked whether anyone should notice his exhaustion, understand why he would not ask for help, or interrupt the proceedings long enough to bring food.

“A moral life is not assembled entirely from isolated choices,” the architect said. “It is learned in relationships, carried by dispositions, and exercised through attention.”

This insight unites several approaches without making them identical: virtue ethics, care ethics, relational duties, loyalty, stewardship, and practical wisdom. They resist a picture of agents as interchangeable calculators who meet only at moments of decision. Human beings begin dependent, remain vulnerable, occupy relationships, inherit practices, and become able to judge through formation.

The central questions shift:

- What kind of person am I becoming?
- What does this relationship require from me?
- Whose need has the rule made invisible?
- Which response sustains trust and agency here?
- What would a wise, caring, courageous person perceive?
- What forms of life cultivate or deform those capacities?

These questions catch moral reality that abstract principles often miss. They also risk making good judgment depend on the unexplained perception of the already good judge.

Where this chapter is headed. We will examine character, flourishing, care, loyalty, stewardship, and practical wisdom as sources of moral guidance. These approaches recover attention and relationship, while raising questions about whose model of flourishing governs and how partiality is limited.

Character before calculation

Virtue ethics centers stable dispositions such as courage, honesty, justice, generosity, compassion, temperance, humility, loyalty, and practical wisdom.

A virtue is not merely a habit of producing one outward behavior. Courage responds well to danger for worthwhile ends. It differs from cowardice and recklessness through perception, emotion, motive, and proportion. Honest character values truth without using “brutal honesty” as permission for cruelty.

Virtues matter because rules underdetermine novel cases. A person who has developed good judgment can perceive which features are salient, integrate them without a master formula, and act from appropriate motives. The act is not only correct; it expresses excellent agency.

This approach explains the difference between the spiteful donor and the generous one. Both transfer money. One aims to humiliate a rival; the other recognizes need. Outcome and action type match while character differs.

Virtue also extends across time. A single rescue does not establish courage if performed for spectacle. A pattern of truthful attention matters more than one lucky verdict.

Flourishing and the shape of a life

Virtues are usually connected to **flourishing**: living well as the kind of embodied, social, rational, vulnerable beings humans are. A good life includes activity, relationship, practical agency, understanding, and participation rather than a sequence of pleasant states.

This protects goods that outcome arithmetic may flatten. Friendship is not merely an instrument for producing welfare. Trust, shared history, and mutual recognition partly constitute the relationship. Integrity concerns whether a life hangs together rather than how each act scores separately.

The flourishing account faces the Chapter 18 questions. Which capacities are central? How diverse can good lives be? Does statistical normality define excellence? Why does characteristic human functioning have practical authority for someone who rejects it?

Virtue theorists can answer that flourishing is not conformity to an average specimen. It is successful activity under human conditions, realized through diverse lives. This avoids crude naturalism. The selected excellences and their authority still need defense.

Care begins with dependency

Care ethics directs attention to need, dependence, attachment, responsiveness, and the maintenance of relationships through which people survive and become agents.

An infant cannot bargain for fair terms. A person with dementia may communicate through gesture rather than abstract choice. A caregiver often knows needs through history that a universal form cannot capture. Care asks what response sustains this person rather than only which rule applies to any person.

This is not sentimental indulgence. Good care requires skill, endurance, factual knowledge, boundaries, and sensitivity to power. A caregiver can be loving and mistaken. Institutions can praise devotion while exploiting unpaid workers.

Care also reveals **relational autonomy**. People do not become self-directing in isolation. Agency depends on language, education, bodily support, security, and people who treat emerging choices seriously. Protecting autonomy can require care rather than mere noninterference.

The resident who stops asking

A long-term care home allocates assistance by request. Every resident has an equal call button. Staff answer promptly. The policy appears fair.

Mae has learned that asking makes hurried workers visibly anxious. She stops pressing the button, misses meals, and remains in pain. Her chart records “no request.” An allocation algorithm therefore directs staff elsewhere.

No rule was violated. The outcome data show Mae receives less care, but only attentive relationship reveals why. A nurse who knows her notices the withdrawal, sits down, and changes how requests are invited.

Care ethics performs better than a purely formal equality rule here. It perceives silence as possible vulnerability rather than autonomous refusal. It also improves empirical knowledge; attention reveals a hidden fact.

Now alter the case. Mae clearly and repeatedly refuses assistance after informed discussion. Continued intervention becomes control. Care must distinguish responsiveness from paternalism.

“I know what you need” can sustain a dependent person or erase her voice. The difference requires evidence, consent, capacity assessment, history, risk, and review.

Loyalty and particular obligation

Loyalty recognizes that relationships create reasons not owed in identical form to everyone. Friends keep confidences, families sustain dependents, communities preserve shared projects, and teammates accept burdens for one another.

Particularity corrects the fantasy that an agent may replace a friend whenever a stranger could receive slightly more benefit. A friendship partly consists in refusing to treat the person as an interchangeable welfare location.

But loyalty has a dark mirror. It can conceal misconduct, distribute opportunities to insiders, and treat outsiders as morally faded. Iona's protection of Els in Chapter 11 showed the problem: friendship may support care outside the review process while public office requires impartiality inside it.

The theory needs a domain rule:

- Which responsibilities arise from the relationship?
- Which public rules remain identity-neutral?
- What can loyalty demand against outsiders' urgent needs?
- When does belonging become domination or exclusion?

"My own first" is not a complete relational ethic.

Stewardship

Stewardship extends care across time, place, institutions, traditions, and ecological systems. A steward does not own an entrusted thing without limit. The role is custodial: preserve, repair, interpret, and pass on.

This can support responsibility for future people and ecosystems even when no current individual contract exists. A community may understand a watershed as held in trust rather than as an inventory of resources.

Stewardship needs a beneficiary and purpose. Is the land preserved for human use, sentient beings, ecological integrity, ancestors, descendants, a divine order, or its own life? The role can organize responsibility after the value is named. It does not select the value by itself.

An inherited steward may also face an unjust trust. Preserving a caste institution faithfully is not admirable merely because the role is old.

Cultural exemplars

Virtues are learned through exemplars: people admired as courageous, wise, caring, faithful, or just. Stories show how dispositions operate when formulas fail.

Cultures admire different figures. One honors the warrior who never yields; another the negotiator who avoids bloodshed. One praises filial obedience; another the child who resists an abusive parent. Exemplars carry formation histories and power.

A virtue theory can criticize its culture's examples by appeal to flourishing, proper emotion, reason, or broader human experience. The criticism then relies on standards beyond "what our excellent person does."

Without such standards, the reasoning circles:

The right act is what the virtuous person would do. The virtuous person is the one disposed to perform right acts.

The circle can be opened by specifying the goods perceived, characteristic errors avoided, and feedback through which judgment improves.

Virtues in collision

Nara promised a dying friend to conceal a family secret. Years later, disclosure would help the friend's adult child understand a treatable inherited illness. Honesty, fidelity, compassion, courage, care, and practical wisdom point in different directions.

A wise agent must integrate them. How?

Possibilities include:

- fidelity yields when concealment creates grave preventable harm;
- care for the living recipient has priority over a dead person's privacy;
- the promise remains but its scope never included medical danger;
- disclosure should be partial and accompanied by repair;
- the conflict is tragic and no option is untainted.

Practical wisdom may reach one of these. The Candidate System still needs to know what makes the perception wise. If the answer is "the wise person sees the right balance," wisdom has become the black box that the hybrid rescue in Chapter 31 will label "all things considered."

Disciplined practical judgment should state decisive features, compare precedents, invite criticism, track outcomes, and explain why nearby cases differ. It may remain nonalgorithmic without becoming inarticulate.

The caregiver's horizon

Care can demand everything because vulnerability has no natural endpoint. A parent, nurse, friend, or community organizer may always do more.

A humane care ethic protects caregivers as persons with standing. It distributes care institutionally, enables rest, and rejects the assumption that attachment assigns infinite duty to whoever is nearest or least powerful.

Consider a daughter providing continuous care for a parent. A distant sibling says, "You have the relationship, so you have the duty." The sibling uses relational ethics to evade reciprocity. A fair division may require public support and shared burden.

Care therefore needs justice. Justice needs care to see what equal forms omit. Their combination will enter the hybrid chapter; neither may simply absorb the other without specifying priority.

Formation before wisdom

A rulebook can be read by a novice. Practical wisdom develops through practice. How should an unformed agent act before becoming wise?

Virtue approaches can use:

- guidance from trustworthy exemplars;
- rules as training supports;
- supervised practice;
- reflection on emotions and motives;
- institutions that make good action easier;
- narratives that expand perception;
- correction by affected people.

This is psychologically realistic. It creates vulnerability to bad formation. If the teachers and institutions are corrupt, the novice can acquire vice with the phenomenology of virtue.

The access audit requires feedback beyond admiration within the same tradition.

Draft the relational constitution

Purpose. Build the strongest character-and-care candidate while naming its model of flourishing, forms of attention, legitimate partiality, conflict method, and correction practice.

Build the strongest character-and-care version of the Candidate System:

Flourishing: *A life goes well through _____.*

Virtues: *The central dispositions are _____ because they perceive and sustain _____.*

Practical wisdom: *Judgment is disciplined through _____ and corrected by _____.*

Care: *Responsibility arises when dependency, history, or trust has form _____.*

Autonomy limit: *Care becomes control when _____.*

Loyalty: *Special concern permits _____ but never _____.*

Caregiver limit: *Demands yield or redistribute when _____.*

Outsiders: *Urgent claims beyond relationships enter through _____.*

Stewardship: *Agent _____ holds _____ in trust for _____.*

Formation: *Novices learn through _____, with safeguards _____.*

Authority: *These excellences and responsibilities bind a rejecter because _____.*

Apply it to Mae, Iona and Els, the inherited steward, Nara's secret, and the exhausted caregiver.

Then identify every place where practical wisdom, flourishing, or care settles a conflict. Explain the feature rather than invoking the faculty.

What the candidate saves

Character, care, and belonging preserve indispensable concerns:

- moral perception requires attention, not only inference;
- motives and dispositions matter across time;
- dependency precedes contract and autonomy;
- relationships can create particular responsibilities;
- silence and formal equality can conceal need;
- caregivers require limits and institutional support;
- forms of life cultivate judgment before isolated choice;
- some values are constitutive of relationship, not replaceable outputs.

The objective debts remain. The theory must identify flourishing, select virtues, evaluate cultural exemplars, discipline practical wisdom, limit partiality, protect outsiders, and explain why relational or character standards apply independently of endorsement.

The twenty-seventh debt

Add this entry:

***Debt 27 — The wise perception:** Which goods make a character excellent, how does practical wisdom resolve conflict and receive correction, and do care, loyalty, formation, or flourishing hide the standing and priority rules they were meant to perceive?*

The third architect shares food with the clerk and rearranges the chairs.

The fourth architect enters carrying a book, a seed, and a lamp kept burning by a community older than the city.

“You did not invent the good,” she says. “You received a place within it.”

Chapter 28 — Receive the Moral Order

The fourth architect placed three objects on the table: a sacred text, a living seed, and a lamp passed from one generation to the next.

Together they represented a moral vision.

Reality is not raw material awaiting human preference. Persons are created or naturally ordered beings. Life contains goods that practical reason can recognize. Communities receive wisdom through revelation, conscience, tradition, and disciplined interpretation. Moral law has authority because it participates in an order deeper than individual desire or political power.

This family includes divine command, divine nature, natural law, objective flourishing, sacred tradition, and combinations of them. Chapter 18 inspected their components. We now let them form a complete system.

The attraction is profound. A received order can condemn unanimous injustice, explain why secret conduct matters, connect duties into a meaningful life, and locate human vulnerability within purpose rather than accident. It offers authority without asking victims to wait for their oppressors' approval.

Its strongest form is not “obey an old rule because a powerful being said so.” It is:

A perfectly good source creates or sustains beings with intelligible goods; moral requirements express that source and the forms of flourishing appropriate to those beings; reason, experience, revelation, and tradition provide partial access; communities preserve and correct interpretation.

This integrated account can distribute tasks. Divine goodness grounds value. Creation grounds relation and purpose. Natural law supplies content. Practical reason applies it. Tradition carries accumulated insight. Conscience alerts individuals when institutions corrupt the inheritance.

The strength of integration is also the place to inspect circular support.

Where this chapter is headed. We will examine moral order as something received through command, nature, tradition, revelation, and inherited practice. The test is whether reception supplies stable authority and access without making interpretation, reform, or humane content invisible.

Command within relationship

A divine command is unlike an order from a stranger. If God is creator, sustainer, perfectly knowledgeable, and perfectly good, obedience can arise within trust, gratitude, dependence, and rightful authority.

Commands solve a problem that impersonal values leave open. “Compassion is valuable” may not tell an agent what to do. “Care for the stranger” addresses the agent. Divine law can combine evaluative content with directive force.

A sophisticated account does not ground authority in raw power. The ability to punish supplies prudence, not moral right. Authority rests on perfect goodness, wisdom, relation to creation, and perhaps the way obedience perfects the creature.

These grounds protect the view from the idea that any sufficiently powerful commander creates duty. They also make *goodness*, *rightful relation*, and *perfection* carry normative content.

Necessary goodness

The divine-nature account says God cannot command cruelty for amusement because perfect love and justice are necessary, not optional traits. Goodness is neither external law imposed on God nor a product of arbitrary will. It is identical with the divine nature.

This answer avoids the crude command dilemma. It can explain why goodness is stable across possible worlds and why divine commands express rather than invent moral order.

The Candidate System asks how the identity is known. Human interpreters need enough independent grasp of love, justice, cruelty, and faithfulness to identify which alleged commands fit the nature. If an apparent revelation orders atrocity, believers often test it against divine goodness. The moral understanding used in the test performs access work.

This need not make revelation redundant. General moral insight may be partial while revelation supplies history, exemplars, practices, and specific demands. The channels can correct one another. Their independence and conflict rule must be visible.

Natural law and basic goods

Created or evolved beings have characteristic capacities and vulnerabilities. Humans need food, attachment, knowledge, agency, community, health, meaningful activity, and stable care. Practical reason can identify these as goods through reflection on what enables lives to flourish.

Natural law does not infer morality from whatever occurs in nature. Disease is natural and defective. Aggression is natural and requires discipline. The relevant nature is an organized form of life with intelligible ends, not a frequency chart.

Basic goods provide substance to divine law. A prohibition on murder protects life. Truthfulness supports knowledge and relationship. Sexual and family norms may be interpreted through procreation, fidelity, care, and social continuity. Political duties support common goods no individual can secure alone.

The approach can criticize preferences adapted to deprivation. A person may report satisfaction under oppression while lacking agency and education that constitute human flourishing.

Its pressure points are selection and plurality. Which capacities define basic goods? Must each individual realize every good? How do goods compare? Can a nontraditional life flourish? Does biological function determine permissible use, or does personal agency reshape function?

Strong versions allow diverse participation and deny that statistical normality is moral perfection. The more flexible the account becomes, the more practical wisdom and interpretation determine content.

Tradition as stored judgment

No generation begins inquiry from nothing. Sacred traditions preserve stories, rituals, prohibitions, forms of care, communities, and centuries of reflection. A practice may embody knowledge no individual can fully articulate.

Tradition can do epistemic and formative work:

- focus attention on recurring human temptations;
- carry obligations across short political cycles;
- form emotions and habits;
- preserve minority identity under hostile power;
- connect personal loss with shared meaning;
- subject interpreters to texts they did not write;
- organize correction through debate and commentary.

Age and survival are not infallibility. Harmful institutions also persist. Tradition earns authority through its source, fruits, interpretive integrity, and ability to address counterevidence—not through duration alone.

An interpretive community can prevent private whim by requiring linguistic competence, historical context, precedent, and public argument. It can also become an insulated

hierarchy that labels dissent incompetence. The access audit applies to communities as well as individuals.

The hidden prisoner

The government of Port Meridian unanimously approves indefinite detention of political dissidents. Courts, schools, and media treat the practice as necessary. One prisoner is held in secret; no public resistance exists.

A received moral order can say the prisoner possesses status that government did not grant and cannot repeal. Divine image, natural human good, rational agency, or sacred law places a boundary around the person. The prisoner can appeal beyond the state.

This is a major strength. The theory gives reformers a language that power cannot define away. It can sustain resistance when consequences are uncertain and social approval absent.

Now ask how the dissenter knows which received order condemns detention. One tradition teaches submission to rulers; another commands liberation; interpreters within each dispute exceptions. The external standard provides authority only if content can be identified.

The case does not reduce religious resistance to politics. It shows why access and interpretation are part of the moral foundation rather than administrative details.

Two revelations at the harbor

After a shipwreck, two communities claim the same divine source.

Community A teaches that the last rescue boat must prioritize members of the covenant because special sacred responsibility comes first. Community B teaches that the stranger must be rescued first because divine love breaks tribal partiality. Both cite texts, traditions, exemplars, and sincere experiences. Both accept the same facts.

How should the conflict be corrected?

Possible methods include:

- textual context and linguistic study;
- hierarchy of commands;
- interpretation through divine love or justice;
- natural-law analysis of life and dependency;
- fruits of the practice across cases;
- conscience;
- authoritative institutional judgment;
- humility and procedural compromise.

Every method contains a standard. “Interpret through love” requires an account of love’s impartial and particular demands. “Follow authority” requires legitimacy and a correction mechanism. “Judge by fruits” imports consequences. “Use conscience” returns to intuition and genealogy.

Plural access routes can triangulate truth. They can also form a circle in which each validates the others. Map their dependencies as in Chapter 18.

Dissent and the prophetic voice

Traditions often contain resources for criticizing themselves: neglected texts, founding promises, minority interpretations, exemplars who resisted authority, and principles applied more consistently by reformers than by rulers.

This makes tradition dynamic rather than frozen. Asha can challenge exclusion by appealing to equal creation, the dignity of learning, and internal promises. Reform need not begin outside religious life.

The problem is determining when reinterpretation is faithful development and when it replaces inherited content with contemporary values. Both conservatives and reformers can accuse the other of opportunism.

A correction protocol should state:

- which sources have priority;
- what historical evidence matters;
- who may interpret;
- how affected outsiders are heard;
- which moral fruits count;
- what would falsify an interpretation;
- whether revelation can be revised or only application;
- how persistent reasonable dissent is treated.

“The truth never changes” does not answer these human questions.

Outsiders and jurisdiction

A received order may claim universal scope. The nonbeliever replies, “I do not accept your revelation, membership, or account of my nature.”

The theory can distinguish truth from acceptance: disbelief does not change reality. If divine authority and natural goods are real, they reach the outsider.

Political enforcement is separate. A community may believe a norm universally true while lacking authority to coerce every person. Human fallibility, plural interpretation, freedom of conscience, and the danger of religious power can support restraint.

This is an important repair. It allows conviction without theocracy. The restraint may be grounded in divine limits on coercion, natural rights, prudential consequences, or fair public terms. State which.

If only members are bound by a ritual, the duty may be role-relative rather than universal. If every person is bound by a prohibition, the public justification must reach people who reject the source or explain why coercion is nevertheless legitimate.

Function and the unusual life

Ren rejects family, reproduction, civic honor, and spiritual practice. She lives alone, studies mathematics, helps strangers anonymously, and creates beauty. She is informed, stable, and content.

Does she fail to flourish?

A rigid list may classify her life as defective for omitting central human goods. A plural natural-law view may say no individual must instantiate every good; her life participates in knowledge, agency, generosity, and beauty. Community sustains conditions she uses even if she avoids membership.

Now remove anonymous aid and make her indifferent to every other person. The theory may call relational deficiency objectively bad even if Ren accepts it. The authority question returns: what makes species or created flourishing a reason for her?

The case tests whether the theory identifies objective human goods or moralizes flourishing by defining the good person as the one who values the good.

Draft the received-order constitution

Purpose. Build the strongest received-order candidate and expose how command, nature, tradition, interpretation, dissent, outsiders, and reform contribute to its authority and limits.

Build the strongest integrated version:

Ultimate source: Moral status is grounded in _____.

Authority: The source governs agents because _____, not merely because it can punish.

Divine or natural goodness: Goodness consists in _____ and is identified through _____.

Basic goods: Human and nonhuman flourishing includes _____.

Commands or principles: Agents must _____.

Access: Content is known through revelation / reason / conscience / experience / tradition / other: _____.

Interpretive priority: When routes conflict, _____.

Dissent: Sincere disagreement is handled through _____.

Outsiders: The truth reaches them because _____; coercion is permitted only when _____.

Revision: Evidence can change application _____ but cannot/can revise _____.

Nonhuman standing: Animals, artificial subjects, and ecosystems enter through _____.

Apply it to the hidden prisoner, the two rescue communities, the internal reformer, the outsider, and Ren's unusual life.

Then draw the support arrows. If divine goodness validates conscience, conscience identifies divine goodness, tradition validates revelation, and revelation establishes tradition, locate independent evidence entering the network.

What the candidate saves

Received-order theories preserve powerful human and philosophical resources:

- moral status need not depend on state or majority approval;
- secret injustice can be condemned;
- life can be understood through purpose and relation;
- human dependency and flourishing receive sustained attention;
- traditions preserve memory, practices, and disciplines across generations;
- communities can form character and support sacrifice;
- internal prophetic criticism can resist corrupt authority;
- moral commitment can connect with meaning, gratitude, and hope.

These contributions do not depend on treating believers as confused or traditions as dispensable.

The objective debts remain. The system must establish the source, explain why creation or perfection confers authority, give substantive content to goodness, move from function to norm, resolve competing revelations, protect dissenters and outsiders, and provide a corrigible access route. Moving from arbitrary command to necessary good nature is a serious repair; it makes *good* the bridge-bearing term.

The twenty-eighth debt

Add this entry:

***Debt 28 — The received order:** What establishes the source, goodness, authority, natural ends, and interpretive method—and can disagreement, outsiders, and revision be handled without using an independent moral standard to identify the order’s correct content?*

The fourth architect leaves the lamp burning.

The fifth architect arrives with no scripture, ledger, or list of virtues. She brings a set of identical envelopes.

“Before we choose the rules,” she says, “we will stop anyone from knowing which life will be theirs.”

Chapter 29 — Build What No One Could Reasonably Reject

The fifth architect distributes identical envelopes around the hall.

Inside each is the name of a life in Port Meridian: a judge, a dockworker, a child with a degenerative illness, a celebrated engineer, a prisoner, a migrant without papers, an octopus in the municipal aquarium, a person not yet born. No one may open an envelope until the rules have been chosen.

“If no one knows which life will be theirs,” she says, “perhaps no one will design the city only for themselves.”

This is the procedural hope. We may not need to discover a moral constitution engraved behind the stars. We may be able to *construct* valid principles by placing agents in the right relation to one another: free, informed, equal, mutually accountable, and unable to rig the result.

Several theories gather here. Contractualism asks whether principles could be reasonably rejected by those burdened by them. Ideal-observer views ask what a suitably informed, impartial, coherent spectator would approve. Kantian constructivism derives principles from what agency and reciprocal respect require. Humean constructivism begins with human concerns and asks what survives informed, coherent reflection. Democratic theories connect legitimacy to participation, representation, and public justification.

These theories differ greatly. Some claim to discover objective moral truth through procedure. Some say the procedure *constitutes* moral truth. Some make a more modest political claim: a decision can be legitimate even when no one has proven it morally correct.

Their shared insight is formidable:

A rule imposed on people should be answerable to them, not merely attractive from the ruler’s private point of view.

The question is what must be placed inside the procedure before that result can come out.

Where this chapter is headed. We will test whether fair choice, reasonable rejection, reciprocity, or ideal agreement can construct principles with public authority. The decisive questions concern the participants, information, constraints, complaint standards, and reach beyond those who accept the procedure.

The complaint rather than the sum

Consider a tunnel with six occupied chambers. Flooding one chamber will save five workers elsewhere. A simple aggregate favors the five. A rights theory may prohibit deliberately killing the one.

Contractualism asks a different question. What principles could each person reasonably reject? The isolated worker has an intense complaint against a principle permitting her death for others' benefit. Each of the five has an intense complaint against a principle forbidding rescue. We compare the complaints of individuals, not a pile of anonymous units.

This protects the separateness of persons. Five headaches do not necessarily outweigh one person's torture merely because five is larger than one. The theory asks what each affected person could demand of every other person under principles for mutual recognition.

But the word *reasonable* immediately begins work.

Suppose a shipping magnate rejects every safety regulation because it reduces his profit. A child rejects bedtime. A sect rejects equal citizenship for unbelievers. A patient rejects a quarantine rule during a lethal epidemic. Which objections count?

We can say that a reasonable rejection must respond to relevant burdens, recognize others as equals, use accurate information, and avoid demanding privileges unavailable to everyone. Those are attractive constraints. They are also substantive. Equal status, relevance, accuracy, reciprocity, and the acceptable distribution of burdens do not arrive merely from the fact that people reject things.

The procedure becomes morally impressive by admitting morally shaped complaints.

The impartial spectator

Perhaps an ideal observer can do better. Imagine a judge who knows every fact, feels each person's experience, reasons without contradiction, has no personal stake, and considers all affected beings impartially.

Ask this observer whether to flood the chamber.

The device removes ignorance, prejudice, impulsiveness, and self-dealing. It is a useful discipline. Ordinary judgment improves when hidden injuries become visible and when we imagine occupying another person's position.

Now specify the observer.

Does impartiality mean equal concern, equal probability of becoming each person, equal weight for equal welfare, or merely freedom from personal identity? Does the ob-

server feel compassion? Does she respect choices that reduce welfare? Does she care about beauty, loyalty, wild nature, promises, or future lives? How risk-averse is she? Does complete information dissolve adaptive preferences or simply reveal them?

An observer stripped of every concern cannot choose. An observer supplied with concerns can choose because those concerns rank outcomes.

This yields an important distinction:

- **Idealization can remove distortion.** It can correct false beliefs, inconsistency, missing perspectives, and conflicts of interest.
- **Idealization cannot choose its own ideal.** Someone must decide which traits count as distortion and which count as moral perception.

Perfect vision tells you what is in every room. It does not tell you which room to enter.

Agency constructs its own law

Kantian constructivism looks for more than a useful thought experiment. It asks what principles are implicit in treating oneself and others as rational agents.

When I act, I take my considerations to support a choice. I organize means under ends. I claim practical authority over my conduct. If I deny comparable authority to every other agent merely because they are not me, I may be making an exception I cannot rationally defend.

From this perspective, reciprocal respect is not an optional taste added to reason. It arises from what it is to make claims as an agent among agents. Principles are valid when free and equal rational persons could author them together.

This is one of the strongest routes from form to content. It explains why domination is not just unpleasant to the dominated; it contradicts the terms under which one person claims authority while denying it to another.

The pressure returns at the borders.

Why does rational agency have this status? If I coherently value my agency but not yours, what formal contradiction follows? Perhaps I am inconsistent only if I also claim that *agency as such* matters. If equal standing enters as a constitutive condition of public reason, then the condition contains the value the procedure was meant to generate.

And what of infants, people with profound cognitive disabilities, many animals, ecosystems, and future beings? We can represent their interests, treat them as ends because of their capacities, or include duties concerning them in principles agents would choose. Each repair is plausible. None follows from their participation as rational co-authors, because they cannot participate in that form.

Representation extends the table. It also gives representatives enormous interpretive power.

Begin from the concerns we actually have

Humean constructivism makes the starting point explicit. Agents already have desires, attachments, aversions, identities, projects, and sympathies. Reflection can inform these concerns, expose conflict, expand imagination, remove false beliefs, and seek a coherent standpoint from which action is possible.

This avoids pretending that practical reasoning begins empty. It explains why a person has reasons connected to her projects and why shared reasons can emerge from widely held concerns. It also permits revision: a desire formed under ignorance or humiliation may not survive fuller understanding.

Imagine Lio, who has worked since childhood in a dangerous glass factory. He says he wants no safety rules because risk proves courage and obedience earns family honor. He has never seen another kind of work, cannot afford to leave, and learned that injured workers are weak.

Should the construction use his stated preference?

If yes, domination can manufacture the consent that validates it. If no, we need a standard for distinguishing an authentic concern from an adaptive one. Information helps: show Lio the injury data, alternatives, and owners' private precautions. Freedom helps: give him a safe exit and time to reconsider. Dialogue helps: let him hear former workers.

But suppose informed and free Lio still values dangerous courage. May others overrule him? The answer depends on how we weigh autonomy, harm, family dependence, public cost, and workplace power. Idealization reduces contamination. It does not make the ranking impersonal.

The Humean can accept this. The resulting reasons are constructed from the agent's concerns, not commands valid independently of every stance. The system may be coherent without being objective in the Candidate System's original sense.

Democracy is not a truth machine

Move from imaginary agreement to an actual city.

Port Meridian holds hearings, elects representatives, publishes evidence, funds opposition parties, protects speech, includes judicial review, and permits laws to be revised. A contested harbor rule passes after extended debate.

These procedures can make the rule publicly legitimate. People had a voice; power changed hands under known conditions; officials gave reasons; minorities can appeal; errors can be corrected. Legitimacy can support coordination even among citizens who think the outcome mistaken.

But fifty-one percent approval does not transform a false claim into truth or an abusive policy into kindness. Democratic processes inherit unequal time, money, status, information, and agenda control. A majority may repeatedly sacrifice a quiet minority. Future people do not vote. Animals do not testify. Migrants affected by a policy may lack membership.

The solution is usually to constrain democracy with rights, representation, expert evidence, supermajorities, independent courts, or constitutional limits. Each protection can improve the procedure. Each raises a new authorship question: who selected the constraints, why these rights, which experts, and what may courts override?

Democracy is best understood here not as a device that alchemizes votes into objective morality, but as a technology for accountable power. That is no small achievement.

The seawall and the unborn district

The city must decide whether to build an expensive seawall. It will slightly reduce current services but may save a future neighborhood fifty years from now. No present resident knows whether their descendants will live there. Some residents have no children. The construction will damage an estuary inhabited by nonhuman animals.

Design a fair procedure.

Who receives an envelope?

- current citizens;
- current noncitizen residents;
- future people;
- construction workers;
- taxpayers;
- estuary animals;
- businesses that may leave;
- neighboring communities affected by diverted water.

How are silent parties represented? A guardian for future people may favor precaution. An ecological representative may oppose construction. Who appoints them, and what evidence controls their claims?

Now vary the veil. Behind one version, you might become any current citizen. Behind another, any sentient being affected over the next century. Behind a third, any rights-bearing person but not an animal. The chosen policy may change before anyone reasons badly. The domain of possible identity is itself a distribution of standing.

The unreasonable rejecter

Every public order eventually meets someone who says no.

Mara rejects taxation for the seawall because she rejects political membership. Daro rejects any plan that protects the estuary because he enjoys destroying habitats. Sen rejects representation for future people because only present consent matters. A future resident, imagined by her guardian, rejects the decision to do nothing.

Calling a rejection unreasonable prevents one veto from paralyzing collective life. It can also become a trapdoor through which inconvenient people disappear.

A disciplined account should state why a rejection is excluded:

- it depends on a factual error;
- it demands a personal exception;
- it denies standing already guaranteed by the procedure;
- it imposes a burden on others vastly exceeding the rejecter's burden;
- it contradicts the conditions of participation;
- it concerns a matter outside the procedure's jurisdiction;
- it expresses a concern participants have explicitly decided not to protect.

The first and some logical exclusions can be relatively formal. The others depend on prior judgments about standing, proportionality, jurisdiction, and protected concern.

Do not hide those judgments in the adjective *reasonable*. Adjectives are terrible safes.

Build the public table

Purpose. Construct the strongest procedural candidate by specifying its participants, representation, information, constraints, complaint standard, outcome rule, and reach; the table makes its inputs visible.

Construct the strongest procedural system you can.

Purpose: *The procedure discovers / constitutes / legitimates / coordinates _____.*

Participants: *Direct participants are _____.*

Represented parties: *Nonparticipants represented are _____, through _____.*

Information: *Participants know _____ and do not know _____.*

Motivation: *They are impartial / mutually concerned / self-interested / reciprocal / other: _____.*

Reasonableness: *A rejection is unreasonable when _____.*

Constraints: *No outcome may _____.*

Decision rule: *Conflict is settled by _____.*

Power correction: *Unequal resources and voice are addressed by _____.*

Revision: *A result may be reopened when _____.*

Dissenter: *A person who continues to reject the result is owed _____ and bound because _____.*

Apply the table to the tunnel, Lio's factory, the seawall, a future person, an animal, and a permanent dissenter.

Circle every clause selected before the procedure begins. Those circles do not prove failure. They reveal the procedure's moral constitution.

What the candidate saves

Constructive and contractual theories preserve central achievements:

- rules must be justifiable to those they govern;
- individual complaints cannot vanish inside an aggregate;
- information and perspective-taking correct private blindness;
- reciprocal terms expose privilege and domination;
- representation can give absent parties a place in deliberation;
- public reasons discipline arbitrary power;
- legitimate procedure allows coordination amid persistent disagreement;
- transparency, appeal, and revision make institutions corrigible.

These are not shadows cast by an unsuccessful metaphysics. They are working designs for life among people who can injure one another and disagree.

The objective claim remains divided. A procedure can discover stance-independent truth only if its features reliably track that truth. It can constitute truth only within an account explaining why these participants, constraints, and ideals have authority. It can produce legitimate settlement from human commitments while giving up stance independence. The more honestly it chooses the third task, the better it may perform—and the less it resembles the objective system we set out to find.

The twenty-ninth debt

Add this entry:

Debt 29 — The construction: Who enters, who is represented, which information and motivations are ideal, what counts as reasonable, what constrains the outcome, and why does the resulting procedure discover objective truth, constitute universal authority, or bind a rejecter who denies its terms?

The fifth architect leaves the envelopes unopened.

The sixth architect enters carrying a roomful of maps. The borders on no two maps are the same.

“Perhaps,” he says, “the mistake was expecting one law to govern every place from nowhere.”

Chapter 30 — Let Standards Remain Local or Plural

The sixth architect covers the table with maps.

One divides Port Meridian into legal districts. Another shows language communities. A third marks religious congregations, a fourth trade networks, a fifth watersheds, and a sixth the neighborhoods in which people actually borrow sugar and complain about noise. Their borders overlap like an argument conducted in colored pencil.

“You have been searching for a view from nowhere,” he says. “No one lives there.”

The local or plural approach begins with humility. Moral rules emerge inside forms of life: histories, relationships, institutions, vocabularies, rituals, ecological conditions, and shared expectations. To strip away all standpoint may be to strip away the very features that make a reason intelligible.

This family includes several views that should not be collapsed.

- **Cultural relativism:** a moral claim is true or valid relative to a culture’s standards.
- **Agreement-relative reasons:** commitments and agreements give members reasons that outsiders may not share.
- **Contextualism:** the standards relevant to evaluation vary with conversational or practical context.
- **Value pluralism:** several genuine values exist and may conflict without a common measure.
- **Personal commitment:** an agent can organize a life around concerns not binding on everyone.
- **Subcommunity standards:** professions, families, faiths, movements, and voluntary associations create special roles and expectations.

These views can be combined, but they answer different questions. Pluralism need not imply that cruelty becomes good when approved. Contextualism need not say every community has its own truth. A role-relative duty need not exhaust the person’s reasons. Keeping the varieties separate prevents the usual carnival game in which one ball labeled *relativism* is thrown at every moving target.

Where this chapter is headed. We will ask what becomes possible when standards remain local, relational, or plural rather than universal. Locality can improve fit and humility, but overlapping memberships, outsiders, reformers, and power make every boundary answerable.

Why locality matters

Consider truth-telling.

In Candor House, residents recovering from institutional abuse have established a rule of direct disclosure. Staff do not conceal diagnoses, meetings, or restrictions. Even kindly deception echoes the paternalism residents escaped.

In Shelter House, residents are hiding from violent pursuers. Staff use false names, concealed schedules, and practiced misdirection. Here literal disclosure can betray a life.

“Always tell the truth” misses the social meaning of speech in each place. “Never deceive” requires careful exceptions. A local analysis asks what practices of trust, protection, consent, and power the utterance participates in.

This is not simply a contest between truth and welfare. The same sentence can be reassurance, betrayal, ritual, coercion, or resistance depending on relationship and history. Context is not fog surrounding the moral fact. It is part of the fact pattern.

Local knowledge also resists imperial confidence. Outsiders have often called their own customs universal reason and described unfamiliar lives as primitive, disordered, or immoral. Stating one’s standpoint makes hidden power easier to see.

Which culture?

A local standard cannot guide until the relevant *locality* is identified. For a person who belongs to several overlapping communities, that choice is already part of the theory.

Nemi lives in Dock Ward, speaks the river dialect at home, works under the Mariners’ Code, attends a transnational temple, belongs to a disability-rights group, and participates in a private online community whose members live on four continents. Which culture determines the truth of a claim about her obligations?

“The surrounding society” is not precise enough. The city contains majorities and minorities, institutions and dissidents, generations and classes. National borders may be less relevant than professional or religious ones. A practice can be shared unevenly, contested internally, or imposed by one group on another.

Possible boundary rules include:

- the political jurisdiction;
- the community with which the agent identifies;
- the group that sustains the relevant practice;
- the group affected by the action;
- the group with authority over the role;
- the smallest community containing stable agreement;

- all overlapping communities, with a conflict procedure.

Each rule creates different obligations. The selection cannot be made by asking what “the culture” says until we know which culture has jurisdiction. The map needs a legend, and the legend is already normative.

The reformer inside the walls

Asha belongs to a society in which only members of inherited scholarly families may use the public archive. The custom is ancient, widely praised, and woven into ceremonies of civic identity. Asha is excluded. She argues that the archive should be open to every resident.

If the society’s current approval determines moral truth, Asha is mistaken until enough people join her. Yet reformers often claim that their society is wrong *now*, not merely that they prefer a future convention.

The relativist has several replies.

First, Asha may appeal to another existing standard: the society promises equal citizenship; the archive rule contradicts that promise. This is **internal criticism**. It can be sharp and publicly intelligible without invoking a standard from nowhere.

Second, she may speak from a subcommunity whose standards oppose hereditary exclusion. Moral truth is then relative to that group. But the conflict between groups remains unsettled.

Third, the culture’s deepest values may be reconstructed under better information or greater consistency. This moves toward idealized construction and inherits its selectors.

Fourth, Asha may simply avow a commitment: “I oppose hereditary exclusion and will organize against it.” This can motivate change without pretending current approval is on her side. It gives up the claim that the local standard already condemns itself.

Fifth, the theory may recognize some universal values—agency, welfare, equality—that override local approval. The reformer regains an approval-independent foothold, and the original grounding question reopens.

Internal criticism is powerful but limited. A perfectly consistent tyranny with no promise of equality offers fewer handles than a hypocritical republic. It can still be resisted through concern, solidarity, interest, revolt, or an alleged objective standard. We should say which.

The majority has a ventriloquist

Suppose eighty percent of Port Meridian supports a rule forbidding the Tide People from owning boats. Officials describe the rule as a treasured cultural safeguard. Surveys confirm majority approval.

The Tide People do not regard the practice as theirs. Their opposition is omitted from the city’s official history. Some majority citizens privately dissent but fear speaking. Children are taught that alternative arrangements are unthinkable.

Does the rule express “the culture”?

A numerical account says yes. A consensus account says perhaps not, because dissent exists. A power-sensitive account says the apparent culture is a project of dominant institutions. A historical account may discover that the rule is recent. A functional account may classify it as a mechanism of economic control rather than shared meaning.

These refinements improve description. They also reveal that culture is not a single mouth. When oppression is internal, deference to culture can mean deference to the people already licensed to speak for it.

The slogan “respect local values” therefore needs subjects and verbs. Whose values? Expressed under what conditions? Respect by listening, permitting, funding, refusing intervention, or enforcing? At whose cost?

The outsider at the gate

Local practices frequently affect nonmembers.

A monastery adopts demanding fasting rules. Voluntary adult members can leave. The standards may create genuine role-relative reasons without governing the city.

Now suppose the monastery diverts the only river, employs nonmembers under its rules, or imposes fasting on children who cannot exit. Internal agreement no longer contains the effects. Outsiders require standing in the analysis.

The community might respond that its jurisdiction comes from land ownership, tradition, divine authority, parental authority, or political recognition. These are claims, not magic circles. Each needs an account of limits.

Voluntary membership helps but does not solve everything. Exit may cost family, language, livelihood, and identity. Consent can create obligations, yet the conditions under which consent is valid are not determined solely by the institution requesting it.

Local validity is clearest for practices whose costs remain among informed, capable participants and whose exit conditions are tolerable. As effects, dependency, and coercion spread, the need for cross-community settlement grows.

Across time

Moral communities extend backward and forward. People honor ancestors, inherit debts, maintain constitutions, preserve languages, and build institutions for descendants.

Can we say an ancestral practice was wrong by its own standards? Often yes: records reveal dissent, violated promises, selective enforcement, and recognized suffering. The past was rarely unanimous.

Can we say it was wrong even if every member approved? A strict culture-relative account says it was right-for-them, though we rightly reject it by our standards. This makes condemnation indexical: “We oppose what they approved.”

That statement may be enough for practical purposes. We can refuse revival, teach the harms, honor victims, and design safeguards from our concerns. But it does not establish that the ancestral society violated its own truth.

If we insist that some acts were wrong regardless of unanimous approval, we have preserved objective ambition. The burden is not discharged by the emotional confidence of the insistence.

The same issue faces future criticism. Our descendants may reject our treatment of animals, artificial subjects, ecosystems, aging, or privacy. Their disagreement will not by itself prove us wrong. It should remind us that present consensus is not a certificate printed by history's final office.

Plural goods without one ruler

Value pluralism offers a different retreat from a single master law. Perhaps liberty, care, knowledge, beauty, loyalty, achievement, equality, and ecological integrity are all genuine values. They sometimes conflict, and no common unit converts them without remainder.

This explains why tragic choices feel tragic. Saving a historic neighborhood by refusing climate migration, protecting privacy at the cost of security, or telling a painful truth may involve loss that no arithmetic erases.

Pluralism need not make values local or subjective. A realist can say several objective values exist. A constructivist can say several human commitments survive reflection. A pragmatist can say different concerns organize different practices.

The decision problem remains. If values are incommensurable, an agent may use judgment, priority rules, thresholds, narrative fit, contextual salience, or political procedure. These can guide action. They do not show that one uniquely objective answer existed.

"Balance the values" is an invitation to think, not a completed theory. A chef who is told to balance salt, acid, sweetness, and heat has received excellent advice and no recipe.

Declare the standpoint

Purpose. State whose standard is being applied, where its jurisdiction ends, how overlapping memberships are handled, and how insiders, reformers, and outsiders may challenge it.

For any local or plural judgment, complete this declaration:

Claim: _____.

Speaker or institution: _____.

Relevant community or communities: _____.

Practice and history: _____.

Affected members: _____.

Affected outsiders: _____.

Values or standards invoked: _____.

Source of authority: *agreement / role / identity / law / commitment / objective claim*
/ other: _____.

Dissent within the group: _____.

Exit conditions: _____.

Cross-community conflict rule: _____.

Revision condition: _____.

Use it for Candor House, Shelter House, Asha's archive, the Tide People, the monastery, and one practice from your own life.

Then rewrite each conclusion in indexed form:

- "Under standard S, action A is prohibited."
- "Members accepted role R, which includes duty D."
- "We care about V and therefore oppose A."
- "Institution I will enforce rule P within jurisdiction J."
- "A is wrong independently of every approval or commitment."

Notice when the final sentence is doing work the first four cannot do. That is the objective claim, now visible rather than smuggled through local grammar.

What the candidate saves

Local and plural approaches protect insights every later design should retain:

- reasons and actions have histories, relationships, and social meanings;
- communities are internally diverse and their borders contested;
- moral language can conceal imperial standpoint and dominant power;
- roles, agreements, and identities create real practical structures without necessarily becoming universal law;
- internal criticism can expose hypocrisy and broken commitments;
- outsiders and dissenters must not be erased by talk of culture;
- several concerns may matter without a single common measure;
- honest standpoint declarations improve argument.

These approaches may produce coherent local codes, conditional reasons, or plural practical judgment. If they define truth by actual standards, they sacrifice stance independence. If they idealize standards, they inherit the construction debt. If they assert universal values protecting reformers and outsiders, they reopen the objective debt.

Their modesty is often a strength. Trouble begins when modest local authority dresses for dinner as universal necessity.

The thirtieth debt

Add this entry:

***Debt 30 — The locality:** Which person, group, time, practice, or standpoint supplies the standard; how are overlapping membership, internal power, reformers, outsiders, and cross-community conflict handled; and if criticism outruns every actual standard, what grounds the approval-independent value doing the work?*

The sixth architect rolls the maps but cannot make their borders coincide.

The other architects return. One brings the ledger, one the boundary marker, one the rearranged chairs, one the lamp, and one the unopened envelopes.

“Individually,” says the chair-moving architect, “we miss too much.”

“Then let us build together,” says the seventh.

Chapter 31 — The Hybrid Rescue

All seven architects stand around the table.

The consequentialist has a ledger. The deontologist has a boundary marker. The virtue-and-care theorist has turned the chairs toward one another. The received-order theorist tends the lamp. The constructivist holds the envelopes. The localist unrolls several maps. The seventh architect brings string, hinges, and an alarming confidence in composite materials.

“Each plan failed because it tried to build alone,” she says. “We will keep the strengths and discard the weaknesses.”

This is the hybrid rescue. It is attractive for an excellent reason: real judgment is already hybrid. Physicians consider outcomes, consent, professional duties, character, relationships, law, and institutional fairness. Parents combine care with boundaries. Courts protect rights while attending to risk and consequence. Communities inherit traditions and revise them through argument.

A theory that notices only one of these dimensions can become elegant in the manner of a chair with one leg.

The hybrid promises more stability:

- use consequences to track suffering and flourishing;
- use rights to prevent sacrifice for aggregate gain;
- use virtues to form reliable agents;
- use care to perceive dependency and relationship;
- use received traditions to preserve meaning and accumulated judgment;
- use fair procedure to justify common rules;
- use local knowledge to understand context and history.

This could make a humane, intelligent, corrigible system. The decisive question is whether combination pays the earlier debts or merely arranges them attractively.

Where this chapter is headed. We will combine the strongest surviving materials from the major candidates and test whether a hybrid can keep their benefits without inheriting their debts. The crucial component is the meta-rule that decides which layer governs when the layers conflict.

A pile is not yet a structure

Suppose a city constitution says:

Promote welfare, respect rights, act virtuously, care for the vulnerable, honor legitimate traditions, use fair procedures, and attend to context.

Very good. Now choose what to do when these instructions agree.

That was easy.

Now choose when they conflict.

The city can violate one person's privacy to identify a contagious disease. Publication would reduce risk, breach confidentiality, frighten a marginalized community away from clinics, fulfill a legal disclosure rule, and perhaps express civic honesty. A compassionate doctor, a rule-governed agency, and a welfare-maximizing epidemiologist may disagree.

A list of considerations identifies the conflict. A system must settle it, authorize a procedure that settles it, or acknowledge that no unique answer has been established.

Hybrids therefore need architecture:

1. **Membership:** Which values, rules, virtues, relationships, or authorities enter?
2. **Priority:** Does one type normally outrank another?
3. **Thresholds:** When may a lower layer override a higher one?
4. **Measurement:** How are burdens within a layer compared?
5. **Jurisdiction:** Which layer governs which domain?
6. **Tie-breaking:** What happens when the declared rules run out?
7. **Revision:** How can membership and priority change?
8. **Authority:** Why does this arrangement govern a rejecter?

Without these, a hybrid is a committee whose bylaws say only, "Reach the correct decision." Committees love that clause. It saves a great deal of ink.

Rights until catastrophe

Begin with a familiar hybrid:

Maximize welfare, subject to rights that may be overridden only to prevent catastrophe.

This keeps the consequentialist's attention to effects and the deontologist's protection against ordinary sacrifice. It refuses to kill one healthy traveler for five transplant recipients, yet permits breaking into a cabin to stop a citywide fire.

The repair immediately improves both parents. Pure maximization can be brutal; absolute prohibition can be disastrous.

Now define *right* and *catastrophe*.

May privacy be breached to prevent one death? Ten? A pandemic? May an innocent person be detained for an hour to stop an explosion? A year? May property rights be overridden to house a freezing population? Does a right protect choice, welfare, status, interest, agency, or a socially constructed boundary?

Thresholds are not discovered by writing the word *threshold*. They must be placed somewhere. If the threshold is set by the point at which consequences become “important enough,” a judgment about acceptable sacrifice selects it. If rights are lexically prior until social survival is threatened, the value of social survival occupies a master position. If practical judgment chooses case by case, judgment inherits the ordering burden.

The hybrid can still be defensible. A public rule with conservative thresholds, procedural review, compensation, and emergency expiration may control power better than either pure theory. Its practical superiority does not by itself show that the thresholds are stance-independent facts.

Virtue inside the guardrails

Try a second structure:

Rules define the boundaries; virtues guide perception and choice within them.

Do not deceive, exploit, abuse, or abandon. Within those constraints, exercise honesty, courage, generosity, patience, and practical wisdom according to context.

This solves two problems. Rules restrain the culturally admired but dangerous character. Virtues prevent rules from becoming a mechanical excuse for callousness. A nurse cannot say, “No regulation required me to notice the terrified patient.” A generous official cannot say, “My noble intentions exempted me from the conflict-of-interest rule.”

But conflict can occur across the layers. A compassionate lie may violate the honesty rule. Courage may require breaking an unjust law. A rule against favoritism may interfere with special responsibility to one’s child.

We can permit virtuous exceptions when a rule defeats its purpose. Then purpose governs the rule. We can permit only exceptions that could be publicly justified. Then construction governs. We can rank the virtues and rules through a tradition. Then the received order governs. We can ask a practically wise agent. Then the agent must possess a reliable method rather than an honorific job title.

“The virtuous person sees what matters” may be true. It is not yet an account of which perception is accurate when admirable people see different things.

Care within fair procedure

A third hybrid begins from relationships but protects against partiality:

Let care reveal needs and dependency; let fair procedure distribute burdens and protect people from controlling care.

In a family, hospital, or school, close attention can uncover what a standardized form misses. Procedure then ensures that the quiet patient, unpopular student, distant caregiver, and future client receive representation. Consent and appeal limit paternalism. Workload rules prevent the most devoted caregiver from becoming the institution's consumable resource.

This is an excellent institutional design.

Its tensions are revealing. Fairness may require equal time for clients, while care calls for more time with the person in crisis. A rule may protect worker boundaries while leaving a dependent person alone. Family loyalty may resist a neutral allocation decision.

The system can assign domains: care governs intimate response; procedure governs public resources. Yet intimate choices consume public resources, and public rules shape intimacy. It can prioritize urgent vulnerability, but urgency and vulnerability need definitions. It can allow discretion subject to review, but reviewers need standards.

Again, none of this makes the hybrid incoherent. It makes the selector visible. The system works because people care about vulnerability *and* fairness and build a negotiated relation between them.

Received aims interpreted through flourishing

The received-order theorist proposes:

Divine or sacred aims provide ultimate orientation; human flourishing, reason, experience, and tradition interpret their application.

This hybrid can avoid literalism and arbitrary command. Love of neighbor, sacred dignity, liberation, stewardship, or harmony supplies a durable orientation. Empirical study identifies what helps beings flourish. Tradition preserves insight; conscience and public criticism expose corruption.

The combination can support extraordinary resistance and care. It can also correct a harmful inherited interpretation without discarding the community's entire moral world.

The ordering question is acute. When a text appears to conflict with flourishing, which revises which? If perfect goodness determines correct interpretation, interpreters need substantive access to goodness. If tradition fixes meaning, harmful fruits may be tolerated. If contemporary experience controls, the sacred source may become ceremonial approval for independent values.

The strongest account can admit reciprocal correction: text, tradition, reason, experience, and conscience form a reflective network. That makes access richer. It does not identify an error-free court of final appeal. A network may be resilient because independent supports converge; it may also be circular because each support is certified by the others.

Plural values and practical judgment

Perhaps hierarchy is the mistake. Let several values remain genuine and let judgment respond to the case.

Welfare, liberty, equality, loyalty, truth, beauty, care, achievement, and ecological integrity matter. No master value converts them into one unit. Mature judgment seeks a fitting response without pretending every conflict has a formula.

This view captures tragedy and moral remainder. A justified choice may still betray something worth grieving. The person displaced for a seawall does not cease to have a claim because the project was selected after fair deliberation. An apology or compensation may be appropriate even when the decision is retained.

Plural judgment also mirrors skilled practice. Experienced mediators, clinicians, and caregivers notice patterns no short rule can encode. They can explain their choices by analogy, precedent, salience, and anticipated effect.

But “fitting” can hide a conclusion inside an adjective just as “reasonable” can. If two informed judges rank liberty and security differently, what shows one has perceived the objective balance? More reflection may reveal overlooked facts or inconsistency. Persistent conflict may remain.

A pluralist can say there are several objectively permissible answers. This relaxes uniqueness but retains objective boundaries. The account must identify those boundaries. Or the pluralist can say communities and agents choose among genuine human concerns. This gives practical freedom while lowering stance independence.

Judgment is indispensable. Treating it as indispensable does not make it infallible or metaphysically self-authorizing.

Reflective equilibrium

The seventh architect now proposes the most sophisticated general method.

Begin with considered judgments about cases. Form candidate principles. Test the principles against cases, background theories, empirical facts, and one another. Revise any element when doing so improves the overall fit. Seek a wide reflective equilibrium: coherent judgment after informed, critical, imaginative revision.

This method is intellectually responsible. It prevents one intuition from becoming a dictator. It permits theory to correct cases and cases to correct theory. It incorporates psychology, social history, institutional evidence, and alternative perspectives. It treats moral inquiry as revisable rather than revealed in a flash above the breakfast cereal.

Reflective equilibrium can produce impressive coherence. It can also systematize different starting commitments.

Imagine two careful thinkers.

- Ilyan begins with strong judgments against intentional killing, a theory of inviolable status, and distrust of aggregate reasoning.

- Pera begins with strong judgments against preventable suffering, a theory of equal welfare, and distrust of status distinctions that preserve worse outcomes.

Both study the same evidence. Both discard prejudiced intuitions, correct inconsistencies, imagine other lives, and refine their principles. Ilyan reaches threshold deontology. Pera reaches rights-sensitive consequentialism. Their views approach one another but do not coincide.

Does equilibrium require convergence? If yes, we need a reason to expect it and a method for selecting among several equilibria. If no, equilibrium demonstrates coherent endorsement rather than one objective result.

Coherence is a relation among contents. A coherent map may faithfully organize the wrong territory. A coherent set of commitments may be cruel. Reflective equilibrium needs constraints on eligible judgments and background theories. Those constraints improve it and carry evaluative content.

The evacuation of Low Harbor

The architects ask for one case severe enough to engage every layer.

A cyclone will reach Low Harbor in six hours. The district contains:

- a hospital with twelve immobile patients;
- a detention center holding forty people, some never tried;
- a residential tower with two hundred occupants;
- the Tide People's ceremonial archive, which cannot be moved later;
- a laboratory containing six artificial systems that credibly report fear and memory;
- an estuary nursery whose destruction will collapse a local species;
- rescue workers with families who have already completed dangerous shifts.

There are too few vehicles. Forecasts are uncertain. Opening the detention center could speed evacuation but officials fear escape. Saving the archive consumes capacity needed for persons. The artificial systems can be copied, but continuity of identity is disputed. Destroying a barrier will protect the tower while flooding the estuary. Ordering another rescue shift exposes exhausted workers to serious risk.

Build the strongest hybrid response.

First pass: let every candidate speak

The consequentialist estimates death, injury, trauma, ecological loss, future effect, and probability. The deontologist identifies detention, bodily risk, consent, status, and prohibitions on using persons merely as means. The care theorist maps dependency, relationship, abandonment, and unseen labor. The virtue theorist asks what courage, honesty, justice, and practical wisdom require of officials. The received-order theorist asks about sacred dignity, stewardship, and special communal obligations. The constructivist represents every party and demands public reasons. The localist explains the archive's meaning, the Tide People's history, and the danger that officials will call dominant priorities neutral.

The result is richer than any single first pass. It is not yet a decision.

Second pass: write the ordering rules

Complete these clauses before choosing who is saved:

Standing: Full standing belongs to _____; partial or represented standing belongs to _____.

Prohibitions: Even for large gains, officials may not _____.

Emergency threshold: A prohibition may be overridden only when _____.

Priority among urgent claims: _____ outranks _____ because _____.

Risk to rescuers: Additional risk is permitted when _____ and voluntary only if _____.

Cultural loss: Irreplaceable communal goods receive _____ weight or protection.

Uncertain minds: The artificial systems are treated as _____ under uncertainty because _____.

Nonhuman life: The estuary enters through _____.

Procedure: The authorized decision-maker is _____, constrained by _____.

Tie-breaker: If all declared rules leave a tie, use _____.

Aftermath: Those burdened are owed explanation / compensation / restoration / apology / review: _____.

Now choose.

If you find yourself writing “do what is best all things considered,” stop. Which things count? How much? Considered by whom? Under what priority? The phrase marks the location of the missing rule; it does not supply one.

Third pass: reverse the hidden facts

After deciding, reveal one fact at a time:

- the detention center holds your political opponents;
- the hospital patients are near the end of life;
- the artificial systems copied themselves without permission;
- the archive records evidence of crimes committed by the city;
- the rescue chief is your sibling;
- the probability of estuary collapse falls from 80 percent to 20 percent;
- the tower residents previously opposed funding flood defenses;
- the Tide People reject monetary compensation.

Record which facts legitimately change the result and which trigger an irrelevant identity response. Then ask what principle distinguishes them.

The exercise does not promise one answer. It reveals what your hybrid is made of when attractive abstractions begin competing for the final vehicle.

The master value in the basement

Many hybrids have an undeclared master value.

A rights-and-welfare theory may ultimately minimize unacceptable harm. A virtue-and-rule theory may ultimately protect flourishing. A care-and-procedure theory may ultimately prioritize equal concern. A sacred-and-natural theory may ultimately follow an account of divine love. A plural theory may ultimately choose what a judge finds most humane.

There is nothing disgraceful about a master value. The problem is claiming that the components jointly generated it when it was used to arrange the components.

Look for these signs:

- every conflict is resolved in the same direction;
- exceptions appear exactly when the favored result needs them;
- one value defines what counts as catastrophe, reasonableness, virtue, or good interpretation;
- the theory invokes judgment but provides no calibration or appeal;
- reflective equilibrium rejects inputs mainly because they resist the eventual system;
- different routes “converge” because their terms were defined through one another.

Bring the master value upstairs. Give it a chair and a debt number.

Build the strongest hybrid

Purpose. Assign each component a distinct job and state the meta-rule for collisions; a hybrid succeeds by integration, not by letting every layer govern whenever convenient.

Draft a system capable of governing the evacuation and ordinary life.

Protected concerns: _____.

Sources: Consequence / right / virtue / care / authority / procedure / locality / other: _____.

Normal priority: _____.

Thresholds and exceptions: _____.

Conflict rule: _____.

Decision-maker: _____.

Correction and appeal: _____.

Standing: _____.

Dissent and enforcement: _____.

Meta-rule selecting this arrangement: _____.

Ground of the meta-rule: actual commitment / idealized agreement / authority / natural fact / rational necessity / brute evaluative fact / other: _____.

Test it against all major cases so far: the transplant traveler, the tunnel worker, the hidden prisoner, the devoted caregiver, the rival revelations, Lio’s factory, Asha’s archive, the Tide People, and Low Harbor.

For every repair, write two sentences:

This repair protects _____.

Its authority comes from _____.

If the second blank contains another component, continue until the arrows stop or form a circle.

The thirty-first debt

Add this entry:

Debt 31 — The hybrid ordering: *Who selects the components, priorities, thresholds, domains, and tie-breakers; what grounds the meta-rule that orders them; and does combination discharge the debts of the parts or add an undeclared master value to several incomplete foundations?*

A hybrid can be more humane than a pure theory, more stable than a single rule, more transparent than unaided intuition, and more corrigible than inherited authority. Those are genuine gains.

Multiple incomplete grounds do not become objective by being placed in a coalition. If one component supplies the final standard, its debt is decisive. If a new meta-standard orders them, that standard incurs its own debt. If practical judgment ends the regress, its content and calibration must be stated. If no unique ordering is claimed, the system may remain plural and usable without satisfying the original objective specification.

The seventh architect does not dismantle the structure. It may be the best structure in the hall.

She attaches an honest label to it.

Built from protected human concerns. Ordering requires declared commitments. Objective authority not yet established.

Inspection Report V — What Each Candidate Saved

Purpose. Record the indispensable contribution, strongest success, imported repair, unpaid debt, and useful narrower form of every major candidate before the general grounding audit.

The inspectors do not award a trophy.

Winner-take-all comparison would repeat the mistake that created the tournament: assuming one theory must own every important insight. Instead, place the candidates side by side and record four things:

1. the human concern the theory sees clearly;
2. the mechanism by which it protects that concern;
3. the price of its strongest repair;
4. the objective claim still unpaid.

Begin with this report.

Candidate	Human concern protected	Strongest mechanism	Price of repair	Objective claim remaining
Consequence	Suffering, flourishing, future effect, avoidable loss	Compare outcomes and counterfactuals	Rights, priority weights, plural goods, thresholds, or distribution rules must be added	Why this currency, standing, aggregation, and distribution bind every agent
Rule and person	Status, agency, security against sacrifice	Rights, universal constraints, duties, and public exceptions	Conflicts and catastrophes require priority or threshold rules	Why these beings have this status and these constraints apply independently of stance
Character and care	Vulnerability, relationship, motivation, formation, context	Virtues, attention, practical wisdom, and relational responsibility	Exemplars, limits, outsiders, and conflict require further standards	Why this ideal of flourishing or fitting care is the objectively correct one
Received order	Meaning, purpose, durable community, resistance beyond state approval	Divine authority, natural ends, tradition, reason, and conscience	Interpretation and revision import accounts of goodness and reliable access	Why the source has moral authority and why this interpretation identifies its good order
Construction	Reciprocity, accountability, individual complaint, fair power	Idealization, representation, agreement, and public justification	Participants, constraints, reasonableness, and representation must be selected	Why this procedure discovers or constitutes universal stance-independent authority
Locality and plurality	History, difference, role, standpoint, anti-imperial humility	Indexed standards, internal criticism, plural values, and jurisdiction	Outsiders, reformers, overlaps, and oppression require cross-boundary standards	Why criticism that exceeds actual standards has approval-independent force

Candidate	Human concern protected	Strongest mechanism	Price of repair	Objective claim remaining
Hybrid	Several of the above at once	Layers, thresholds, domains, reflective equilibrium, and review	Adds ordering, tie-breaking, and possible hidden-master-value debts	Why the meta-rule and final judgment are objective rather than selected or asserted

Now add a final column titled **Practice worth keeping**. Possible entries include:

- consequence forecasting;
- counterfactual comparison;
- rights declarations;
- prohibitions with public exception rules;
- character formation;
- caregiver support;
- historical and relational attention;
- interpretive communities with dissent protections;
- representation of absent parties;
- transparent procedure and appeal;
- standpoint declarations;
- internal criticism;
- plural-value registers;
- reflective revision;
- explicit tie-breakers.

Keep as many as survive scrutiny. A practice need not be a portal to stance-independent truth to improve prediction, reduce injuries, protect what people care about, expose domination, coordinate action, or make power answerable.

Inspection finding

Every candidate saw something the others could miss.

The ledger forced intentions to face results. The boundary marker protected persons from becoming entries in the ledger. The chairs revealed dependency and the invisible labor supporting every architect. The lamp preserved memory, purpose, and resistance to transient power. The envelopes demanded reciprocal justification. The maps exposed standpoint, history, and false universality. The hinges allowed correction when one in-sight became tyrannical.

Every candidate also required selection.

What counts as good consequence? Who has standing? Which rights and virtues? What is the true order? Who enters the procedure? Which community governs? What ranks the layers?

The repairs became better as they grew more explicit about human vulnerability, agency, concern, relationship, identity, power, uncertainty, and the need to live together. They became less successful when those same protected concerns were redescribed as conclusions produced without evaluative starting points.

This is not yet the book's general result. A final rescue may remain. Perhaps social criticism requires objective truth. Perhaps no public code can be built without it. Perhaps responsibility and moral emotion collapse if intrinsic desert disappears. Perhaps the recurring debts can be shown to arise from a finite set of foundations—and one foundation will finally pay.

Part VI gives those possibilities their last hearing.

Before entering it, close the inspection with three lines:

The concerns I refuse to lose: _____.

The practices I can defend without pretending they prove objective authority:
_____.

The objective claim I still think may succeed, and the debt it must pay:
_____.

The architects leave their objects in the hall.

The reader now has more than ruins. The reader has materials.

PART VI



The Last Repairs and the General Result

The construction has survived by adding bridges. We now inspect every bridge, every debt, and every ambition quietly left behind.

Chapter 32 — Criticize the World You Inherited

The architects have gone home. Their materials remain in the hall.

At dawn, Asha arrives with a crowbar.

She is the archive reformer from Chapter 30, the scholar excluded because she was born outside the hereditary families. During the night, the council rejected her final appeal. The archive doors are still locked. The council's ruling is lawful, traditional, popular, and internally consistent.

"Your inspection was very careful," she says. "May I now criticize the people keeping me outside?"

This is the first last rescue.

Perhaps objective morality is indispensable not because metaphysics is charming, but because reform requires it. If all standards come from actual people and practices, a society seems incapable of being wrong by standards it does not already contain. The dissenter becomes someone with unusual preferences. The abolitionist, democrat, heretic, whistleblower, or liberator cannot say, "The whole order is wrong." At most they can say, "I dislike it."

That conclusion would be devastating.

It is also too fast.

Criticism performs several jobs. Some require an approval-independent moral fact. Many do not. To know what Asha loses, we must stop forcing every objection through the same sentence.

Where this chapter is headed. We will inventory forms of criticism available even when objective moral authority has not been established. The point is not to weaken reform, but to show whether evidence, effects, commitments, inconsistency, power, and declared values can make criticism more precise.

Eight handles on a locked door

Suppose Asha says each of the following:

1. **"Most citizens now oppose the archive rule."**
2. **"The Dock Ward and the teachers' guild oppose it."**
3. **"The rule contradicts the city's promise of equal citizenship."**
4. **"The founding agreement granted public access."**
5. **"I oppose hereditary exclusion and will organize against it."**
6. **"The rule blocks agency, education, livelihood, and political voice."**
7. **"No one choosing under fair, informed, non-hereditary terms would accept it."**

8. “The rule is objectively wrong even if every person approves it.”

These statements can support one another, but they are not translations.

The first reports majority standards. The second invokes subcommunities. The third exposes internal inconsistency. The fourth invokes an inherited commitment. The fifth expresses and enacts a personal commitment. The sixth identifies effects on concerns many people have. The seventh appeals to idealized agreement. Only the eighth directly asserts stance-independent moral status.

If the eighth claim remains unpaid, the first seven do not evaporate. They establish different things, address different audiences, and support different actions.

The critic’s task is not to find the one sentence that magically binds every possible listener. It is to know which handle she is pulling and what mechanism connects it to the door.

The caste register

In the hill city of Orra, occupations are inherited. Children of record keepers become record keepers. Children of stone carriers become stone carriers. Movement between castes is forbidden. The arrangement has lasted six centuries.

Its defenders say the system provides identity, predictability, mastery, mutual dependence, and social peace. They are not lying. Some members love their inherited crafts. Disorder has followed abrupt reform elsewhere.

Now inspect the costs. Stone-carrier children receive less education, suffer more injury, and cannot refuse dangerous work. Record keepers define the official history. Marriages across castes cause expulsion. Dissenters lose food access. The alleged peace is partly the silence of people whose exit routes were closed.

Several criticisms become available.

Factual exposure

The official claim that every caste enjoys equal health is false. Injury and lifespan records show otherwise. The claim that no one dissents is false because punishment hides dissent.

Evidence can defeat those descriptions. This does not by itself establish that unequal injury is morally wrong. It removes false premises and places actual stakes before people who care about health, agency, children, or honest government.

Immanent criticism

Orra’s founding hymn declares that every child carries the city’s full light. Public ceremonies praise courage and learning in all citizens. Hereditary education contradicts those commitments.

The critic can say: “Live by what you already proclaim.” This argument has traction with members who identify with the promise. It does not prove the promise objectively correct, and it has less leverage over a defender willing to renounce it.

Genealogical criticism

Historical records show the caste boundaries were not divinely ancient. Record-keeping families imposed them after a failed revolt and later rewrote the origin story.

This can defeat a claim of sacred continuity and reveal whose interests a story serves. Origin does not alone decide present value. A practice born in domination could acquire new functions. The current defenders must now justify those functions rather than shelter beneath invented antiquity.

Commitment and solidarity

The critic can state: “I value agency, open learning, family integrity, and protection from imposed danger. I stand with the stone carriers. I will use persuasion, refusal, funding, organization, and, under stated limits, political power to end the rule.”

That is more than a report of aesthetic dislike. It identifies concerns, allies, mechanisms, and action. It is also candidly grounded in commitments rather than disguised as a fact every conceivable agent must obey.

Objective condemnation

The critic may add: “Even a fully informed and unanimous Orra would be objectively wrong to maintain inherited caste.”

That may be true. The claim is not established by the injury data, internal promise, false origin story, or critic’s commitment. It requires the foundation and access examined in Parts III–V.

Withholding that final certificate does not restore the caste register. It changes the declared basis of resistance.

The unanimous safety cover-up

Port Meridian’s reactor engineers discover a defect that creates a small chance of catastrophic failure. Repair would be expensive and politically humiliating. Every current resident, fully informed of the risk, votes to conceal the defect from the neighboring valley. Prevailing winds would carry contamination there first.

The vote is unanimous. Internal criticism seems to have run out of citizens.

But the affected audience is larger than the electorate. Valley residents did not agree. Future residents did not vote. The unanimity is jurisdictionally incomplete.

Criticism can expose:

- the factual risk and who bears it;
- the exclusion of affected outsiders;
- the asymmetry between private benefit and exported danger;
- prior treaties requiring disclosure;
- the city’s public commitment to reciprocal warning;
- what current residents would choose if wind direction or residence were unknown;

- the valley's interests, fear, and capacity for retaliation;
- the institutional incentive to conceal embarrassing hazards.

Suppose the city has no treaty and openly says, "We protect our residents and do not care about yours." The critic cannot force concern into existence by presenting another chart. She can alert the valley, organize sanctions, expose the city to trading partners, appeal to whatever loyalties cross the border, sabotage the concealment, or invoke objective duty.

Those actions differ in justification and risk. The important lesson is not that every conflict yields to dialogue. It is that unanimous approval within a selected group never proves that every affected standpoint has been included.

The broken promise republic

A republic constitutionally promises equal voting power. Its legislature then gives rural districts ten times the representation of urban districts. Officials insist the arrangement preserves stability.

An urban reformer has an especially strong internal argument. The state claims authority through equal citizenship while violating the terms it uses to legitimate itself. The reformer need not first establish a universal moral right to equal voting. She can say:

By your charter, your oath, and your public basis for coercing me, this arrangement fails.

The government might abandon equal citizenship and defend aristocracy. That reply loses hypocrisy but not power. It also changes the dispute. Citizens who supported the state because of its promise now have evidence that the claimed bargain is broken.

Internal criticism is therefore neither trivial nor omnipotent. It can withdraw legitimacy supplied by a commitment. It cannot prevent an actor from renouncing the commitment and relying on force.

This is one reason political promises matter. They create handles that future power cannot discard without revealing what it has become.

The openly conquering culture

Imagine a culture with no hidden liberal premise.

The Vask celebrate conquest. They teach that strength entitles the victor to land, labor, and tribute. They admire loyalty within their own ranks and contempt for weakness outside them. Their histories honestly record invasion. Their institutions efficiently prepare for more. Most members endorse the project after informed reflection.

This is the difficult case. There may be no contradiction to expose, no forgotten equality promise, and no silenced internal majority.

What can critics say?

- The targeted communities can say, "We want to live, remain autonomous, keep our homes, and protect our children."

- Neighboring states can say, “Your expansion threatens what we value; we will deter and contain you.”
- Vask dissenters can reject the culture’s aims and invite solidarity.
- Strategists can show that conquest creates resistance, trauma, instability, and reciprocal danger.
- Negotiators can offer arrangements serving some Vask interests without domination.
- Genealogists can trace how elite institutions cultivate martial desire and suppress alternatives;
- Objective realists can say conquest is wrong whatever the Vask endorse.

The first six do not prove the seventh. They also do not amount to shrugging.

The targeted community does not need the conqueror’s values to resist. A locked door does not become open because the person outside cannot prove that access is written into the fabric of the universe. Conflict can be real before a neutral tribunal exists.

This is sobering. It means argument sometimes reaches the end of shared premises. Beyond that point lie bargaining, alliance, withdrawal, containment, sabotage, coercion, and force. Chapter 34 will not romanticize that frontier.

Criticism across time

You discover that your great-grandparents’ town expelled a harmless minority with nearly unanimous support. The expelled families are dead. The town no longer exists as a government. What can present condemnation do?

It can establish facts against nostalgic falsehood. It can express solidarity with descendants. It can explain inherited wealth and disadvantage. It can reshape public identity, support restitution, warn against repeated mechanisms, and declare present commitments.

If historical records reveal dissent, broken law, deception, or recognized suffering, the criticism can also be internal to the past. If the community really approved unanimously, the claim “They violated their own standards” may be false.

You may still say, “They were objectively wrong.” Again, that claim stands or falls with its foundation. Or you may say:

We recoil from their cruelty, care about the lives they damaged, reject the values that sustained expulsion, and will organize our institutions to prevent recurrence.

This sentence does not travel backward and place a metaphysical summons in every ancestor’s hand. It does work in the present: memory, identity, repair, warning, and institutional design.

Moral language often tries to make one utterance condemn the past, bind the dead, express grief, guide the present, vindicate victims, and announce universal truth. Separating the jobs does not disrespect the victims. It prevents metaphysical overbilling.

Overlapping communities, overlapping leverage

No critic has only one audience.

A factory manager belongs to a corporation, profession, neighborhood, nation, family, and international supply chain. A discriminatory policy may be supported by company custom but violate professional standards, national law, customer commitments, and the manager's personal identity.

The critic can move between these relations:

- "Your own data contradict the safety claim."
- "Your professional code prohibits concealment."
- "Your contract promises equal access."
- "Your customers will leave."
- "Your workers are organizing."
- "This policy harms people you say you love."
- "I will not participate."
- "A court can compel disclosure."
- "This is objectively wrong."

Only the last seeks authority independent of every stance, relation, institution, and power. The others can still alter belief, identity, incentive, alliance, permission, reputation, and conduct.

The world is not divided into cosmic proof and helpless opinion. It contains many joints at which criticism can apply pressure.

The criticism ledger

Purpose. Reconstruct one criticism from facts, affected beings, commitments, values, power, and proposed change; the ledger aims for load-bearing precision rather than meta-physical inflation.

Choose one practice you oppose. Do not begin with the word *wrong*. Complete the ledger.

Practice criticized: _____.

Factual claim challenged: _____.

Affected beings and effects: _____.

Current majority standard: _____.

Relevant subcommunity standard: _____.

Internal promise or inconsistency: _____.

Prior agreement or role duty: _____.

My declared commitment: _____.

Shared concerns available: _____.

Idealized-agreement argument: _____.

Institution or power capable of change: _____.

Objective moral claim, if asserted: _____.

Ground and access route for that claim: _____.

Beside every entry, write:

This establishes: _____.

This leaves open: _____.

For example, injury evidence can establish likely injury; it leaves open who has standing and what risk is acceptable. A broken promise can establish failure under agreed terms; it leaves open whether the promise was objectively required. Personal commitment can establish what you will support; it leaves open authority over a stranger. An ideal procedure can expose partiality; it leaves open why that construction has universal status.

The exercise is not designed to weaken your criticism. It is designed to make every part load-bearing.

The thirty-second debt

Add this entry:

Debt 32 — The reformer’s reach: *Which audience, commitment, interest, effect, inconsistency, procedure, institution, or objective fact gives a criticism its force—and what establishes authority over a person who rejects every available premise?*

A society can be criticized without first completing moral metaphysics. Critics can reveal lies, injury, domination, exclusion, broken promises, manufactured consent, and unequal power. They can mobilize existing values, create new alliances, declare commitments, change incentives, withdraw cooperation, and build counterpower.

None of this proves that reform is objectively correct. None of it makes reform mere punctuation around the phrase *I dislike*.

Asha looks at the ledger, then at the archive door.

“Good,” she says. “Now help me draft the institution that replaces it.”

Chapter 33 — Build a Public Code Without the Missing Layer

Asha does not smash the archive door.

Not yet.

She calls a meeting in the theory hall. Archive families attend, along with excluded scholars, cleaners, teachers, parents, religious elders, children too young to vote, and two officials who insist they are present “only to observe,” which is how officials announce that they are taking excellent notes.

The meeting faces a larger question than access to one building.

If no objective moral fact has settled the city’s design, can the group build a public code without pretending? Or will every charter be a costume worn by whoever has enough votes and guards?

The missing objective layer does not leave a blank page. The room contains people with bodies, attachments, projects, fears, memories, and unequal power. It contains buildings, rivers, technologies, evidence, inherited rules, and foreseeable effects. It contains agreements people can make and institutions capable of keeping or breaking them.

These materials do not write their own constitution. They make construction possible.

Where this chapter is headed. We will draft a public code whose clauses disclose the support they actually possess. The code may coordinate, protect, authorize, and constrain without pretending that law, legitimacy, shared concern, and objective obligation are one source.

Seven different supports

Before drafting a clause, distinguish the supports it might have.

Mutual advantage

People accept a rule because each expects to fare better under reciprocal compliance than under unmanaged conflict. Traffic rules, shared defenses, sanitation, and contract enforcement often have this structure.

Mutual advantage can motivate stable cooperation. It may omit those too weak to offer benefit or retaliation. An infant, endangered species, distant future person, or isolated minority may bring little to the bargain.

Fair agreement

People accept terms after representation, information, equal voice, and safeguards against domination. Fairness can protect the weak better than bargaining power alone.

The procedure requires prior decisions about participants, equality, information, and protected claims. Chapter 29's construction debt remains visible.

Shared commitment

A community already cares about safety, learning, family, liberty, sacred identity, prosperity, ecological continuity, or public honesty. A rule expresses and coordinates those commitments.

Shared commitment can be deep without being unanimous. It can also change. Its authority reaches most naturally those who identify with it and weakens when imposed on outsiders.

Informed preference

Institutions can learn what affected people want under adequate information and freedom from manipulation. This respects agency and can correct officials who guess badly.

Preferences may conflict, adapt to deprivation, burden outsiders, or demand impossible resources. Information does not supply a priority rule.

Harm reduction

Evidence identifies policies likely to reduce pain, illness, fear, death, coercion, displacement, or lost agency. Many people have strong aversions to these states, for themselves and others.

Harm is not self-defining. Surgery harms tissue to help a patient. A tax harms purchasing power to fund care. Preventing domination may frustrate a dominator. Standing, baseline, probability, distribution, and acceptable tradeoff remain decisions.

Democratic legitimacy

People comply because an accountable procedure authorized the rule. They may regard the result as legitimate while continuing to oppose it.

Procedure coordinates disagreement and enables peaceful revision. It can also launder majority abuse unless constrained.

Stability and power

A rule persists because enforcement makes expectations reliable. Even reluctant participants can plan around it. A state may suppress violence and solve coordination failures.

Enforceability is not justification. A stable prison can be stable. Power is nevertheless part of every public code; a charter unable to survive contact with defection is a brochure.

Most successful clauses draw on several supports. The discipline is to label them rather than blend them into the word *justified*.

The Nine-Clause Charter

The meeting places nine blank clauses on the wall.

Clause 1: Primary goal

What is the institution for?

Possible answers include reducing severe suffering, protecting agency, enabling peaceful coexistence, securing equal citizenship, preserving a community, supporting mutual advantage, or giving people reliable means to pursue diverse lives.

“Promote the common good” is rejected for arriving at the meeting without identification.

A workable charter can name several goals and define their relation. For example:

The city will protect residents against severe violence, disease, destitution, and domination; preserve broad personal agency; and maintain fair means for revising shared rules.

This is contentful because people selected protected stakes. It is not objective merely because it is printed in solemn type.

Clause 2: Protected interests

Who or what receives direct protection?

Candidates include life, bodily integrity, freedom from terror, subsistence, agency, privacy, relationship, expression, property, cultural continuity, education, ecological systems, and access to public institutions.

The group must distinguish minimum protection from desirable promotion. It must also identify bearers: citizens, residents, migrants, visitors, prisoners, future people, animals, artificial subjects, and ecosystems.

Every inclusion recalls Chapter 6. Standing is not produced by alphabetical order.

Clause 3: Decision procedure

Who writes ordinary rules? Direct vote, representatives, sortition, expert agencies, courts, local councils, or some mixture?

Different questions may need different procedures. Engineers can estimate bridge failure; affected residents can describe burdens; elected officials can allocate resources; courts can apply declared limits. Expertise about means does not automatically confer authority to choose ends.

Clause 4: Dissent

May a person speak, organize, refuse, conscientiously object, withhold labor, create alternatives, or appeal?

A charter that permits only polite disagreement after the decision is already irreversible has confused dissent with decorative ventilation.

The group must also decide when refusal threatens protected interests. A clinician's objection may burden a patient. A pacifist may reject a defense tax. An engineer may refuse to conceal a defect. The answer will require jurisdiction and priority, not merely reverence for conscience.

Clause 5: Membership

Who belongs, and what does belonging confer?

Residence, birth, citizenship, contribution, consent, vulnerability, and affectedness yield different memberships. Voting membership can be narrower than protective standing. A tourist may lack a vote and still be protected against assault. A future person cannot participate and may still be represented.

Clause 6: Enforcement limits

What may the institution do to secure compliance?

Possible methods include notice, education, incentives, fees, restitution, license restriction, monitoring, confinement, and force. Each requires a purpose, evidence threshold, proportionality rule, review, and stopping condition.

"Whatever is necessary" is not a limit. It is an eviction notice served on every limit.

Clause 7: Revision

How can rules change when evidence, membership, technology, or commitments change?

A revision process needs triggers, authority, records, protected participation, transitional rules, and perhaps emergency procedures. Too much rigidity preserves error. Too much flexibility makes promises unreliable and lets temporary majorities rewrite the locks while minorities are inside.

Clause 8: Outsiders

What does the city owe people who never accepted the charter but bear its effects?

It can require impact assessment, representation, treaties, notice, compensation, ecological safeguards, or reciprocal restraint. Outsider protection is a direct test of whether the code merely organizes insiders' advantage.

Clause 9: Protection against majority abuse

What may no ordinary majority do?

The charter might entrench bodily security, due process, expression, equal access, minority representation, or minimum resources. Entrenchment makes the code less directly majoritarian in order to protect the conditions under which continued collective authorship remains meaningful.

These protections are chosen because participants care about security, equality, agency, and historical warnings—not because the parchment begins to glow when the correct list is complete.

The archive constitution

The meeting drafts a first charter for the archive.

Purpose: *Preserve knowledge and make it usable across generations.*

Protected interests: *Public access, privacy of sensitive records, physical preservation, scholarly freedom, cultural continuity, and fair recognition of contributors.*

Decision procedure: *An elected public board, professional archivists, randomly selected users, and representatives of communities whose records are held.*

Dissent: *Published minority opinions, public appeal, protected whistleblowing, and judicial review of exclusions.*

Membership: *Every resident may use the archive; nonresidents may use it under capacity rules; affected communities receive standing in decisions about their records.*

Enforcement limit: *Access may be restricted only for specified privacy, preservation, or immediate-safety reasons, with written reasons and an appeal date.*

Revision: *Rules reviewed every three years or after documented exclusion, technological change, or new evidence of harm.*

Outsiders: *Reproductions and remote access offered when capacity permits; source communities may contest acquisition and description.*

Majority protection: *No group may be excluded by ancestry, belief, disability, occupation, or criticism of the board.*

This charter is far better than hereditary control. Why?

Asha can give several answers. It serves the archive's declared purpose. It protects access she and others value. It incorporates evidence about exclusion. It answers to affected communities. It limits self-dealing. It establishes predictable procedures. It permits correction. It can attract broad support and withstand public criticism.

An archive patriarch replies, "I reject equality and public authorship. My family has the keys."

The charter gives him reasons only if he cares about some premise it engages: knowledge preservation, public stability, the founding promise, reciprocal security, reputation, family flourishing, or fair terms. It can also remove his legal authority and, if necessary, his keys.

The loss of objective proof does not remove that political fact. It requires the group to state it honestly.

The island generator

Pressure-test the charter with a smaller community.

Twenty people live on an island with one generator. Nineteen agree to ration electricity during winter. Vale refuses. She owns the generator because she found it first and wants to illuminate an empty sculpture garden all night. Without rationing, medical refrigeration will fail.

The group can appeal to mutual survival, reciprocity, compassion, prior cooperation, fair agreement, and property rules Vale previously used. Suppose she rejects all of them and has enough fuel and weapons to defend the machine.

The charter does not generate a stance-independent reason inside her by administrative decree. The group must bargain, disable the garden lights, seize the generator, leave, or accept the medical loss.

Any coercive response needs safeguards. Is the threat real? Are less intrusive options available? Who verifies the calculation? Does Vale retain personal power, due process, compensation, or appeal? When does emergency control end?

The case reveals two separate questions:

1. What arrangement can the group support with its values and evidence?
2. What may the group do when Vale refuses?

Objective morality promised an answer to the second that was more than collective preference plus power. Without that answer, the decision does not disappear. Its coercive character becomes impossible to perfume.

The quiet minority

Now suppose the nineteen want to seize the generator not for medicine but to illuminate a monument they adore. Vale wants one bulb for reading. Majority rule favors the monument. Mutual advantage may favor the coordinated majority. Stability favors compliance.

The charter needs minority protection because several supports can align behind abuse. Possible safeguards include:

- minimum protected access before discretionary collective projects;
- independent review;
- written burden comparisons;
- supermajority requirements for targeted losses;
- compensation;
- rights of exit and retained resources;
- rotating decision power;
- representation of the burdened party;
- time limits and automatic reconsideration.

Why adopt these? Because participants fear becoming the quiet minority, care about agency and security, recognize historical patterns of domination, want stable cooperation, or accept reciprocal terms. These are strong supports. If someone asks whether minority

protection is objectively required even when no one values any of that, the objective debt returns.

Three labels on every law

For each proposed rule, attach three labels.

Support

Who has what consideration in favor of the rule?

- shared value;
- protected interest;
- expected benefit;
- informed preference;
- agreement;
- role or promise;
- procedural legitimacy;
- objective claim.

Compliance

Why will people actually follow it?

- endorsement;
- habit;
- trust;
- reciprocity;
- reputation;
- incentive;
- monitoring;
- fear of sanction;
- inability to resist.

Authority claimed

What does the institution say gives it jurisdiction?

- consent;
- residence;
- affectedness;
- constitutional role;
- democratic authorization;
- emergency necessity;
- property;
- force;

- objective moral right.

These columns will overlap. Do not merge them.

A law may have broad support and little compliance because enforcement is weak. It may receive compliance without support because punishment is severe. It may be procedurally authorized while substantively opposed. It may be defensible to many without binding the permanent rejecter in the stance-independent sense.

Draft your public code

Purpose. Build a charter whose clauses disclose their goals, support, jurisdiction, coercive means, safeguards, dissent, and revision rules; defensibility to many is not renamed universal obligation.

Build a charter for a real or imagined community.

Primary goal: _____.

Protected interests: _____.

Direct standing: _____.

Decision procedure: _____.

Dissent rights and burdens: _____.

Membership and representation: _____.

Enforcement methods and limits: _____.

Revision triggers and process: _____.

Outsider safeguards: _____.

Minority protections: _____.

Emergency rule: _____.

For each clause, label its basis:

- **F** — factual or predictive;
- **I** — instrumental to a stated goal;
- **C** — commitment-based;
- **A** — agreement- or role-based;
- **P** — procedurally authorized;
- **E** — backed by enforcement;
- **O** — claimed as objectively required.

An individual clause may receive several labels. Every **O** clause returns to the debt ledger. Every **E** clause requires a safeguard and stopping rule. Every **I** clause must identify its goal. Every **A** clause must identify the parties and valid conditions of agreement.

Then test the charter against:

1. a minority everyone dislikes;
2. an outsider who bears the cost;
3. a future person who cannot consent;
4. a powerful member who can defect safely;

5. an emergency officials could exploit;
6. new evidence that embarrasses the institution;
7. a member who rejects the charter's primary goal.

The charter passes not by converting all seven into supporters, but by making its treatment of them explicit, constrained, reviewable, and connected to declared stakes.

The thirty-third debt

Add this entry:

***Debt 33 — The public charter:** Which goals, interests, members, procedures, protections, enforcement limits, and revision rules are selected; what supports each clause; and which claimed authority exceeds commitment, agreement, legitimacy, jurisdiction, incentive, or coordinated power?*

A public code can be coherent. Its clauses can fit together without hidden exception. It can be empirically informed, broadly defensible, democratically legitimate, stable, protective, and revisable. It can give people reasons connected to their aims and agreements. It can constrain officials and improve lives.

It need not be objectively ordained to do those things.

The charter cannot guarantee endorsement by a person who rejects its goals, standing rules, and procedures. It may nevertheless govern that person. At that point the theory hall gives way to the council chamber, the courtroom, the picket line, and sometimes the barricade.

Vale is still standing beside the generator.

She still says no.

Chapter 34 — Power, Dissent, and the Person Who Still Says No

Vale stands beside the island's only generator with a rifle, a fuel key, and no interest in the group's charter.

The medical refrigerator is warming.

Nineteen people have offered evidence, reciprocity, compensation, public procedure, and a handsome plaque identifying Vale as Founding Curator of the Sculpture Garden. She refuses every proposal.

"The machine is mine," she says. "Your needs do not govern me."

This is where political philosophy often changes the lighting.

Words such as *authority*, *legitimacy*, *necessity*, and *justice* begin to glow. The rifle remains stubbornly visible. The group must decide whether to seize the generator, and a complete account must say what that act is.

Removing objective moral authority does not eliminate power. It eliminates one proposed description of power as the enforcement of requirements already binding every agent from outside all stances. What remains can be defensible, constrained, tragic, reckless, or monstrous. It cannot become clean merely because the paperwork is excellent.

Where this chapter is headed. We will follow public rules into enforcement, where persuasion, bargaining, pressure, law, coercion, and force cease to be verbal distinctions. The permanent dissenter tests whether power is constrained honestly when universal moral authority remains unearned.

Seven ways conduct changes

Keep these mechanisms apart.

Reason-giving

One person presents a consideration connected to another's aims, commitments, evidence, identity, or accepted standards.

"The medicine protects your child too."

The listener may revise voluntarily. No threat is necessary.

Bargaining

Parties exchange benefits or concessions under conditions of divergent interest.

"Supply electricity tonight and you receive fuel, maintenance, and garden hours next week."

A bargain can be mutually advantageous without being fair. Desperation changes negotiating power.

Coordination

People follow one scheme because matching behavior matters more than which reasonable scheme is selected.

Drive on one side of the road. Use one emergency frequency. Take turns drawing power.

Coordination can generate conditional reasons for participants who value avoiding collision or confusion.

Legitimate procedure

A recognized process authorizes a decision. Participants may comply despite disagreement because they value the procedure, expect reciprocity, or accept the institution's role.

Legitimacy is a relation among a process, its constituency, its conditions, and a decision. It is not a chemical that seeps from a ballot box into every outcome.

Legal jurisdiction

An institution has socially recognized power to make and enforce rules in a domain.

Courts can issue orders. Agencies can revoke licenses. Councils can allocate public resources. Jurisdiction predicts and organizes official action; its moral authority is a further question.

Coercion

One party intentionally narrows another's options through a threat or imposed cost.

"Share the generator or we will fine, confine, dismiss, expose, or physically stop you."

Coercion is not limited to police. Employers, families, landlords, creditors, mobs, and gatekeepers can coerce.

Force

One party directly overcomes resistance through physical control, seizure, confinement, incapacitation, or violence.

Force may occur without a prior threat. It can be limited or lethal, official or private, protective or predatory.

These categories can overlap. A law may be reasoned, bargained over, procedurally legitimate, jurisdictionally valid, coercively enforced, and finally backed by force. Calling the whole sequence *reason-giving* because reasons appeared at the beginning conceals the end.

Domination is a structure, not an event

Coercion concerns a particular constrained choice. Domination concerns durable arbitrary power.

An employer who can revoke housing, health care, immigration status, and income may dominate workers without issuing daily threats. A benevolent monarch can dominate subjects while choosing kindly because no reliable institution constrains the monarch's discretion. A majority can dominate a minority through predictable political control even when each individual law appears moderate.

The difference matters. A person may comply freely in the narrow sense that no threat was spoken today, yet live beneath a power capable of changing her options without answer or appeal.

Safeguards against domination therefore include:

- divided authority;
- transparent rules;
- independent review;
- protected exit and voice;
- secure minimum resources;
- records and reasons;
- representation;
- appeal;
- time limits;
- the ability to organize counterpower.

Why value non-domination? Many people fear dependence on uncontrolled will, care about agency, remember abuse, seek stable expectations, and would resist occupying the weaker position. These supports can be broad and deep. The claim that domination is objectively wrong independently of every concern remains an additional claim.

Seize the generator?

Hold the facts fixed.

- Without electricity, three refrigerated medicines will spoil within an hour.
- The patients cannot safely substitute another treatment.
- Vale can illuminate the sculpture garden later without permanent loss.
- She acquired the generator under island rules requiring emergency sharing.
- The nineteen have offered compensation and future access.
- A neutral technician has verified the power calculation.
- Disabling the garden circuit poses little physical risk.
- Vale threatens to shoot anyone who approaches.

The group has strong commitment-, agreement-, consequence-, and procedure-based support for intervention. It also has greater coordinated power if members act together.

Possible steps form an escalation ladder:

1. present the verified facts;
2. remind Vale of the acquisition condition;
3. offer a workable bargain;
4. invite a mediator she trusts;
5. issue notice of the proposed intervention and conditions for avoiding it;
6. isolate only the garden circuit;
7. restrain Vale only if necessary to prevent shooting;
8. restore her ordinary access and review the seizure after the emergency.

The ladder protects concerns the group has: life, bodily security, agency, reliability, minimal intrusion, and avoidance of needless injury. It can make action more predictable and less vulnerable to panic or revenge.

Does Vale have an objective obligation to yield? That has not been established. Will the group intervene anyway? Almost certainly.

The honest description is:

We care more about preserving the medicines and patients than about Vale's unrestricted control of the generator. We have agreement-based and factual grounds, we have exhausted lower-cost options, and we will use constrained collective power to protect those stakes. We will remain answerable for errors and excess.

This does not make seizure arbitrary. It states the values, evidence, relation, limits, and responsibility. It also does not convert the seizure into a command issued by the universe through nineteen temporary deputies.

The outsider below the dam

Port Meridian approves a hydroelectric dam through an exemplary democratic process. Hearings are open. Evidence is published. The vote is decisive. Minority citizens receive compensation.

Downstream, outside the city's jurisdiction, the Reed Community loses fishing grounds and ceremonial sites. Its members were invited to submit comments but had no vote. The city will receive electricity; the Reed Community will receive altered water.

The procedure may be legitimate among city members and defective with respect to outsiders. A comment box is not representation when the decision-maker can ignore every comment at no cost.

Possible repairs include joint authority, treaty negotiation, veto over existential impacts, revenue sharing, minimum flow guarantees, independent ecological monitoring, relocation chosen by the community, and binding appeal.

Each shifts power, not merely conversation. If the city refuses, the Reed Community may litigate, blockade construction, seek allies, damage equipment, or appeal to a larger government. Those methods range from reason-giving to force and carry different risks to participants and bystanders.

The case shows that procedure cannot legitimate burdens over anyone we conveniently place outside its constituency. Affectedness creates a practical demand for representation

because affected people care, resist, recruit allies, and can be injured. Whether affectedness also creates objective moral standing remains the older debt.

The majority's lottery

A city faces a budget shortage. Ninety percent of voters approve a plan that closes every clinic in a small ethnic district while leaving other services untouched. The majority saves a modest amount on taxes. The minority faces severe health loss.

The election is fair. The result is not hidden. No official violated voting rules.

This is majority sacrifice.

A charter designed to resist it might require:

- minimum protections not subject to ordinary vote;
- burden-distribution reports;
- representation of concentrated losers;
- independent equal-treatment review;
- supermajority approval for targeted deprivations;
- evidence that less concentrated alternatives were considered;
- appeal before implementation;
- automatic expiration;
- compensation where loss cannot be avoided.

These safeguards embody commitments to security, equal status, anti-domination, and aversion to concentrated severe harm for diffuse small benefit. They can be supported by reciprocal fear: anyone may become a minority. They can also protect groups whose political weakness makes reciprocity strategically unnecessary.

That last protection is crucial. A system based only on mutual threat abandons the powerless. Protecting them requires concern, identification, institutional commitment, or an objective claim not derived from bargaining strength.

The actor who can safely defect

Koro owns the only fleet capable of removing chemical waste from the harbor. Dumping it at sea is cheaper. The pollution will harm strangers years later. Koro is wealthy enough to survive fines, politically connected enough to avoid prosecution, and unmoved by reputation.

Reason-giving fails. Bargaining is unnecessary from his perspective. Existing procedure has no grip. He can defect safely.

The community can redesign the power relation:

- impose liability large enough to matter;
- require bonds before operation;
- monitor disposal independently;
- revoke access to the harbor;

- create a public competitor;
- coordinate customers and workers;
- seize the fleet under declared emergency authority;
- incapacitate the dumping operation.

These are institutional and coercive strategies. They do not create a motive from nothing; they connect Koro's conduct to consequences he already cares about or remove his capacity to impose the risk.

The phrase "give him a reason" is ambiguous here. The policy may create an instrumental reason by changing incentives. It does not establish that a favoring relation applied to him before any human intervention.

Power can alter the practical field. That is why its design deserves as much scrutiny as its philosophical decoration.

The permanent dissenter

Now imagine Dena, who rejects the city's standing rules, primary goals, procedures, and restraints. She does not misunderstand. She is not oppressed, irrational, or excluded. She simply endorses domination and has no sympathy for those harmed.

Public reason reaches no shared premise.

The city can still:

- predict her conduct;
- refuse cooperation;
- warn potential victims;
- monitor relevant risks under stated rules;
- bargain where interests overlap;
- block harmful acts;
- impose sanctions;
- confine her if she presents a sufficiently evidenced danger;
- leave her beliefs alone when conduct remains within protected boundaries.

The last line matters. A code need not control every hostile value. It can focus jurisdiction on conduct and risk. Trying to manufacture inward agreement invites surveillance, indoctrination, and cruelty.

Dena may never become morally persuaded. The institution's success is not measured by metaphysical conversion. It is measured against declared aims: protected people, constrained power, reduced injury, preserved agency, predictable rules, and opportunities for correction.

This is less grand than universal moral authority and often more humane than trying to produce it by force.

The coercion warrant

Purpose. Require each coercive act to name the conduct constrained, stake protected, evidence, necessity, proportionality, less restrictive alternative, safeguard, appeal, and repair.

For every proposed coercive act, complete a warrant.

Conduct to be constrained: _____.

Stake protected: _____.

Evidence and uncertainty: _____.

Affected parties: _____.

Source of jurisdiction: _____.

Less intrusive alternatives tried: _____.

Maximum permitted method: _____.

Proportionality rule: _____.

Decision-maker and conflicts of interest: _____.

Notice and reasons: _____.

Appeal or independent review: _____.

Stopping and restoration conditions: _____.

Compensation for error: _____.

Value commitments supporting the warrant: _____.

Objective authority claimed, if any: _____.

Then test it after reversing identities. Put your enemy in danger and your ally under restraint. Put your community downstream. Give the minority no strategic power. Make the official sincere and mistaken. Make the emergency last ten years.

A safeguard that survives only when your side governs is campaign equipment, not institutional architecture.

The thirty-fourth debt

Add this entry:

Debt 34 — The coercive frontier: When reasons, bargains, and shared procedures fail, which stakes justify constraint, who receives jurisdiction, what limits and reviews power, and does collective enforcement reflect declared commitments or an objective authority that still requires grounding?

Political necessity does not become objective moral authority. The need to decide is not evidence that one uniquely correct moral answer exists. The ability to enforce a result is not evidence that the result already bound its target.

Sometimes a society constrains people because others care intensely about what is at stake and possess enough coordinated power to intervene. This fact can be frightening. Hiding it behind moral necessity makes it more frightening, not less.

Honesty about power does not endorse arbitrary force. It asks the questions that claims of righteousness too easily silence: Who is restrained? On what evidence? For whose

protection? By whom? Under what limit? Until when? What happens if the state is wrong?

The generator powers the medical refrigerator.

Vale's garden is dark for one night. Her rifle is secured. Her appeal is heard the next morning.

No star issues a receipt.

Chapter 35 — Responsibility Without Cosmic Payback

The morning after the generator crisis, the island discovers that the fuel line was cut.

Vale did not do it. Niko did.

He admits the sabotage. He understood that medicines could spoil, acted without coercion, concealed the cut, and hoped the crisis would discredit collective control. The repair consumed scarce supplies. One patient nearly died.

Everyone agrees that Niko caused the danger. The disagreement begins when someone says:

“He deserves to suffer.”

Not we must stop him. Not he must repair the damage. Not we can no longer trust him near the generator. Not publicly condemning sabotage may protect the next patient.

He should suffer simply because he knowingly did wrong. Even if the suffering protects no one, changes nothing, repairs nothing, communicates nothing, and is witnessed by no one, it is fitting in itself. The moral books are out of balance until pain is entered on his side.

This is **basic desert**: the claim that a person merits benefit or suffering because of what they have done or who they are, independently of further good consequences, agreements, relationships, or practical purposes.

Basic desert is not the whole of responsibility. Removing it leaves an unexpectedly busy workshop.

Where this chapter is headed. We will separate the practical and relational jobs of responsibility from the claim that suffering is intrinsically deserved. Protection, repair, communication, trust, and rehabilitation may justify serious responses even if basic desert remains unsupported.

The strongest case for payback

Before removing anything, let retribution speak in its best voice.

A purely forward-looking system can appear to treat people as dangerous weather. If punishment is justified only by future effects, why not punish an innocent person when doing so deters crime? Why not manipulate offenders, erase their memories, or confine them after reform if public fear remains? Why should intention matter when accidental and deliberate acts cause equal harm?

Desert seems to protect proportionality and personhood. It responds to the agent as the author of a choice, not merely a container of risk. It gives victims recognition that a violation occurred. It limits punishment to the guilty and, in many accounts, limits severity to what the offense merits. It explains why apology and remorse matter.

These are serious achievements.

Now separate them.

- **Guilt restriction:** Do not impose adverse responses without adequate evidence of relevant conduct.
- **Proportionality:** Do not inflict costs grossly exceeding the stakes and purposes at issue.
- **Agency recognition:** Attend to intention, understanding, control, and reasons-responsiveness.
- **Victim recognition:** Name the injury and affirm that the victim's interests will be protected.
- **Communication:** Publicly mark a boundary and explain the response.
- **Repair:** Require restitution, acknowledgment, or restoration where possible.
- **Desert:** Inflict suffering because suffering itself is fitting payback.

The first six do not logically require the seventh. Desert may be offered as their best foundation. It should not receive credit merely because they arrived at the same meeting.

Eight jobs assigned to blame

An adverse response can serve several purposes.

Protection

Separate a dangerous person from vulnerable targets, remove access to a system, supervise conduct, or change an environment.

Protection requires evidence of risk, not a metaphysical stain. It should end when the relevant risk can be acceptably managed by less restrictive means.

Deterrence

Change expected costs so that this person or others are less likely to repeat conduct.

Deterrence can become predatory if officials use one person as a frightening example. Safeguards include evidence, known rules, proportionality, and limits on using people merely as instruments.

Restitution

Restore property, replace losses, fund treatment, correct records, or return benefits obtained through misconduct.

Restitution connects response to the injured party. Some losses cannot be restored, and excessive demands can become another form of ruin.

Rehabilitation

Develop capacities, treat disorders, change incentives, teach skills, or support a different life.

Rehabilitation can become coercive normalization. Its goal should be relevant capability and reduced danger, not ideological remanufacture.

Boundary-setting

Refuse access, end a relationship, withdraw a role, or establish conditions for future interaction.

A person need not prove that another deserves pain in order to lock a door, leave a partnership, or decline to restore trust.

Withdrawal of trust

Update expectations in response to evidence about reliability, intention, concealment, or capacity.

Trust is predictive and relational. It can be rebuilt unevenly. Forgiveness does not require restoring every vulnerability.

Public condemnation

Name conduct, express a community's commitments, stand with those harmed, warn others, and coordinate expectations.

Condemnation can perform these jobs without declaring a person essentially corrupt. It needs accuracy, audience, proportion, and a path away from permanent stigma.

Retributive suffering

Inflict or preserve suffering because the suffering is deserved as such.

To isolate this final job, remove every other effect. If the suffering still ought to occur, basic desert is doing independent work.

The zero-purpose unit of pain

Imagine Niko is fully reformed. He repaired the fuel line, compensated the patients, disclosed every detail, lost access to critical infrastructure, and spent years demonstrating reliability. No one needs reassurance. No deterrent effect remains. The patients oppose further punishment. The response will be secret, so it communicates nothing.

A device can produce ten minutes of intense pain without injury or memory. Niko will wake unaware. No one else will know. The universe alone will contain ten additional minutes of suffering.

Should the device be used?

Hold every practical purpose fixed at zero.

If yes, the judgment expresses pure negative desert: suffering makes the total situation more fitting solely because Niko once sabotaged the line.

If no, your support for punishment may depend on protection, deterrence, repair, relationship, communication, or another future- or present-facing purpose.

You may reject the case as psychologically odd. It is. Isolation chambers in philosophy have terrible ambience. Its job is not to recommend a machine. Its job is to reveal whether pain itself is part of the desired result.

Deliberate sabotage

Return to the real Niko immediately after the act.

Intention matters even without basic desert.

A deliberate saboteur provides different evidence from a person who trips over a fuel line. Intention predicts planning, concealment, possible repetition, and response to reasons. It changes what explanation, monitoring, trust, and repair are relevant. It may increase the need for public boundary-setting.

Control matters too. If Niko understood alternatives and acted without coercion, requests for acknowledgment and changed conduct can sensibly address his capacities. If a seizure caused the movement, those responses would be misplaced.

Agency survives. The claim that agency generates an intrinsic quantum of deserved pain is separate.

The sold memory

Mara commits financial fraud, hides the proceeds, and then sells the autobiographical memories covering the offense. The procedure is irreversible. She retains general intelligence and personality but sincerely experiences the accusation as news about a stranger.

What remains appropriate?

- recover the proceeds;
- establish the facts publicly;
- protect future clients if relevant traits remain;
- investigate whether memory sale was part of concealment;
- require institutional safeguards against such evasion;
- recognize the victims;
- perhaps restrict roles involving financial trust.

What does remorse mean when the agent cannot remember? She can understand testimony, accept causal continuity, care about present effects, and participate in repair. She may lack the first-person memory that ordinarily supports emotional acknowledgment.

Desert theories must decide whether the current person remains the same bearer of past liability. Purpose-based responses ask which present capacities, risks, assets, relationships, and repairs connect to the act. Neither answer is automatic, but the second avoids treating identity as a metaphysical conveyor belt built only to deliver pain.

The rescue fraud

During a flood, Taren claims to have rescued six children. The story brings donations, honors, and political office. Years later, evidence shows the rescues were performed by an undocumented worker whom Taren deliberately erased from the record.

No physical injury followed the lie. Why respond?

The fraud diverted resources, status, trust, historical recognition, and political power. Correcting the record serves truth. Returning donations and honors restores some allocation. Removing Taren from office protects institutions whose authority relied on false evidence. Public condemnation warns against exploiting rescue narratives and recognizes the actual rescuer.

None of these purposes requires enjoying Taren's humiliation. Indeed, spectacle may obscure repair and turn the harmed worker into a prop for a second performance.

Responsibility can be severe without becoming recreational.

The fully reformed offender

Ivo committed a violent assault twenty years ago. He completed treatment, made every available repair the survivor requested, accepted restrictions, and has presented no elevated risk for fifteen years. The survivor feels safe and opposes continued exclusion. The public remains furious.

Should Ivo remain barred from housing and work?

Ask each purpose separately.

- **Protection:** What current risk does exclusion reduce?
- **Deterrence:** Is indefinite exclusion necessary, effective, and not excessive?
- **Restitution:** What loss would exclusion restore?
- **Rehabilitation:** What capacity remains unaddressed?
- **Boundary:** Whose boundary is being protected?
- **Trust:** Which role requires which evidence of reliability?
- **Condemnation:** Has the message already been made clear?
- **Desert:** Is continued hardship valuable even after every task is complete?

If only the last answer remains, the policy is payback. That does not prove it false. It makes the foundation visible and prevents old fear from borrowing the names of expired purposes.

Coercion and diminished control

Sela transports illegal material because an armed group threatens her child. She plans the route, deceives inspectors, and understands the harm. Her causal and cognitive agency are intact. Her alternatives are brutally constrained.

A binary label—responsible or not—loses too much.

The response can distinguish:

- causal contribution;
- knowledge;
- intention;
- available alternatives;
- threat severity;
- reasonable escape options;
- benefit received;
- continuing danger;
- capacity for repair;
- institutional conditions producing coercion.

Protection may still require intercepting the material. Investigation may require Sela's information. Trust in certain roles may change. Punitive suffering adds little if she already acted under terror and poses no independent risk.

Degrees of control help tailor response even if they do not meter intrinsic desert.

The organization with a thousand innocent desks

A transport company repeatedly falsifies brake inspections. No single employee intended a crash. Executives rewarded speed, an automated system hid warnings, inspectors feared dismissal, regulators relied on company data, and responsibility was divided until everyone could point to another desk.

An individual-desert lens searches for the wicked person. A systems lens asks what produced the failure.

Useful responses may include:

- publish the causal chain;
- compensate victims;
- redesign incentives;
- protect whistleblowers;
- separate inspection from production;
- impose organizational monitoring;
- remove decision authority from officials who concealed evidence;
- require reserves for future harm;
- prosecute deliberate falsification when evidence supports it.

Individual intention remains relevant. It is not the only lever. Fixating on who deserves suffering can leave the machinery that made the conduct likely humming beneath the courtroom.

The response warrant

Purpose. Match every adverse response to evidence, control, risk, purpose, necessity, and review, leaving any residual claim of intrinsically deserved suffering separately visible.

For each adverse response, complete this table.

Question	Entry
Conduct and evidence	_____
Causal role	_____
Knowledge and intention	_____
Control and alternatives	_____
Present risk	_____
Person or interest protected	_____
Response purpose	_____
Expected effect	_____
Maximum burden	_____
Less burdensome alternative	_____
Review date	_____
Stopping condition	_____
Repair or restoration after error	_____
Unavoidable uncertainty	_____

No response may enter under the purpose “accountability” alone. Translate it into protection, deterrence, restitution, rehabilitation, boundary, trust, public communication, retribution, or another named job.

Then apply a **purpose deletion test**. Remove one purpose at a time. Does the response change? If every practical purpose disappears and the response remains, label the remainder **basic desert**.

Finally apply an **identity reversal**. Imagine the same evidence against a person you admire, fear, despise, and do not know. If the warrant changes, locate the relevant fact or the bias.

What this does not remove

Rejecting basic desert does not imply that no one causes anything, chooses, intends, deceives, plans, responds to reasons, becomes dangerous, breaks trust, or owes agreed repair. It does not require opening every prison door at noon. It does not require victims to reconcile, forget, or restore access.

It removes one claim:

Suffering is made good or fitting simply by being allocated to a person who deserves it, even when no further purpose remains.

Without that claim, adverse responses need jobs. Jobs require evidence, safeguards, proportionality, and stopping rules. This can make institutions less theatrical and more demanding.

It may also leave emotional questions unresolved. A person can abandon retribution on paper and still feel rage, guilt, resentment, shame, gratitude, or pride in the body.

The feelings do not wait politely for the theory hall to finish.

The thirty-fifth debt

Add this entry:

***Debt 35 — The deserved suffering:** What makes suffering intrinsically fitting for an agent, how are identity, control, knowledge, and proportional quantity established, and what remains of the response when every protective, reparative, relational, communicative, and regulatory purpose is removed?*

Niko can be identified as the saboteur. The fuel line can be repaired. Access can be restricted. Risk can be assessed. Trust can be revised. Patients can receive acknowledgment and compensation. The group can condemn the act, redesign the generator system, and require evidence before restoring Niko to critical work.

None of these becomes incoherent when cosmic payback is removed.

The island may decide that pain beyond those purposes is still deserved. If so, it must pay Debt 35. It may decide that every burden needs a present job and stopping rule. If so, responsibility remains—less like a metaphysical invoice, more like the difficult practice of protecting, answering, repairing, and changing.

At the next meeting, Niko apologizes.

One patient feels relief. Another feels nothing. Vale feels furious on principle. Niko feels guilty long after every repair is complete.

The emotions open their own ledger.

Chapter 36 — Moral Emotions Without Moral Desert

Niko's apology is accepted. The fuel line is repaired. The patients have recovered. He has paid restitution, lost access to the generator, completed every safety condition, and spent two years rebuilding trust.

At three in the morning, he still wakes with the thought:

I should suffer.

Not *I need to repair something*. Not *someone remains in danger*. Not *I must remember the risk*. Those tasks are complete. The feeling presents suffering itself as appropriate.

A theory can remove basic desert from a constitution in one afternoon. A nervous system does not accept constitutional amendments on that schedule.

Guilt, shame, resentment, indignation, gratitude, and pride are not detachable opinions. They involve bodily arousal, attention, memory, learned scripts, relationships, identity, and readiness to act. Some arrive before language. Some remain after the belief that once seemed to justify them has changed.

The question is not whether to press an emotional mute button. There is no such button, and installing one would be an alarming research program. The question is what each emotion is doing now, what its appraisal says, how expression affects others, and when cultivation, redirection, or release serves the stakes we retain.

Where this chapter is headed. We will separate an emotion's cause, appraisal, function, expression, intensity, and duration. This lets guilt, shame, resentment, indignation, gratitude, and pride be understood and revised without treating every feeling as either a moral oracle or a disposable error.

Six questions hiding inside one feeling

Whenever someone asks, "Is this emotion justified?" separate at least six questions.

What caused it?

Temperament, learning, bodily state, trauma, social cues, remembered events, current danger, empathy, status threat, and deliberate reflection can all contribute.

An understandable cause does not make the appraisal accurate. A learned fear can be intelligible and mistaken.

What does it present as true?

Guilt may present *I violated a demand and merit adverse response*. Resentment may present *you wronged me and owe acknowledgment*. Gratitude may present *you benefited me in a way that calls for appreciation or return*.

The appraisal can contain factual, relational, and desert claims. Test them separately.

What is its current function?

The emotion may focus attention, motivate repair, protect a boundary, protest treatment, communicate commitment, coordinate expectations, preserve memory, invite connection, regulate future conduct, or impose punishment.

An emotion can change jobs. Resentment that first protected a boundary may later rehearse pain after the danger and relationship have ended.

Should it be expressed, and to whom?

Feeling and expression are distinct. A person may feel rage and choose not to threaten. They may feel gratitude and withhold a public performance that would reward manipulation. They may communicate a boundary in calm language while the emotion remains intense.

What intensity fits the job?

A slight breach and a devastating betrayal may call for different attention and protection. Intensity can also reflect sleep loss, prior trauma, group amplification, or repeated rehearsal rather than present stakes.

When should it change or end?

Some emotions have stopping conditions: repair completed, danger passed, belief corrected, audience informed, trust decision made, or ritual concluded. Others return because the loss is permanent or the cue recurs.

“Understandable forever” does not mean “worth cultivating forever.” “No longer useful” does not mean “voluntarily dismissible by Tuesday.”

Guilt after the task is done

Guilt can perform valuable work. It directs attention toward one’s conduct, motivates acknowledgment, supports repair, and inhibits repetition. It can tell an injured person, “I understand that my action mattered.”

It can also become self-punishment mistaken for responsibility.

Niko believes that releasing guilt would minimize the sabotage. Yet the patients do not need his insomnia. The generator is safer. He remembers the event. He accepts

restrictions. His suffering has become his private evidence that he is not the kind of person who would do it again.

But pain is not a certificate of reform. Changed attention, candor, conduct, and design are better evidence.

Niko can translate the remaining guilt:

- “I caused serious danger.”
- “I care that the patients were exposed.”
- “I will preserve the lesson.”
- “I accept that some trust may not return.”
- “I do not need additional suffering to make these statements true.”

This translation may not dissolve the feeling. It gives Niko another practice when it arrives: regret, remembrance, recommitment, and self-observation without a new sentence of pain.

The faultless accident

Mira follows every navigation protocol, checks the weather, and operates a ferry with due care. An undetectable mechanical defect causes an accident in which a passenger is injured. Mira leads the rescue and participates fully in compensation and investigation. Engineers confirm she could not reasonably have known.

She feels guilty.

The emotion’s cause is understandable: causal involvement, proximity, empathy, counterfactual imagination, and professional identity. Its appraisal “I failed a standard I could have met” is false under the fixed facts. Its functions may include honoring the injury, maintaining attention to safety, and expressing solidarity.

Those jobs can be served by grief, regret, concern, and participation in repair without preserving self-blame. Telling Mira she has “nothing to feel” misunderstands emotion. Telling her the guilt is morally accurate because someone was hurt confuses causation with culpability.

Shame and the attack on identity

Guilt commonly targets an act: *I did something condemnable*. Shame commonly targets the self or social identity: *I am defective, exposed, or unworthy of belonging*. Real emotions do not obey textbook borders, but the distinction helps.

Shame can signal a threat to membership, reveal that conduct conflicts with one’s identity, or interrupt grandiosity. It can motivate concealment, submission, rage, repair, or transformation.

It is also an efficient instrument of social control.

Lena was taught to feel shame about a harmless bodily trait. The feeling arrives when she enters a public pool, even though she now rejects the community’s teaching. No injury, deception, or broken commitment is involved.

The causal history explains the shame. It does not validate the appraisal of defect. The emotion's current function is conformity to an abandoned standard. Its effects include avoidance, narrowed agency, and renewed power for an institution Lena no longer endorses.

Possible responses include naming the learned standard, seeking counterexamples, changing social context, practicing exposure, cultivating solidarity, grieving lost years, and replacing self-condemnation with neutrality or affection. None guarantees emotional relief. All refuse to treat involuntary repetition as evidence that the old standard was true.

Shame about harmful conduct may still perform a useful interruption. The danger is globalizing from "I deceived someone" to "I am human refuse." Global shame can make repair harder by turning attention from the injured person toward the offender's threatened identity.

Resentment after betrayal

Resentment says, among other things: *Do not treat me that way. What happened matters. I am not available for silent use.*

It can track betrayal, focus memory, energize protest, set a boundary, recruit allies, and require acknowledgment. Philosophies urging immediate serenity have sometimes served the comfort of people who prefer their targets uncomplaining.

Consider Aven. A close friend revealed Aven's medical secret for social advantage. The friend later confessed without excuse, corrected the story, accepted the end of the friendship, paid resulting costs, and demonstrated change over many years. No further contact is expected. Aven remains intensely resentful.

The betrayal is real. The boundary is established. The record is corrected. Trust need not return. What job remains for the resentment?

Possible answers include preserving self-respect, honoring the significance of the relationship, preventing renewed vulnerability, or expressing an enduring protest. Possible costs include intrusive rumination, damaged new relationships, and continued occupation of attention by someone no longer present.

There is no timer printed by objective resentment administration. Aven can ask whether some jobs can transfer to memory, caution, grief, or a settled boundary. Releasing resentment need not forgive, reconcile, excuse, or declare the betrayal trivial. Keeping it need not imply that the friend deserves endless suffering.

The relevant decision concerns Aven's life, relationships, and values—not a cosmic obligation to maintain heat.

Indignation when no target remains

Indignation often faces outward and public. It marks a violation, aligns a group, condemns an institution, and energizes reform. It can widen concern beyond personal injury.

It can also become an identity maintained by permanent activation.

Imagine an injustice committed two centuries ago by an institution that no longer exists. The facts are established. Descendants have received restitution. Current institutions teach the history and maintain safeguards. No living person defends the act. Yet a political group ritually intensifies indignation each year, treating any reduction as betrayal.

The feeling may preserve memory and solidarity. The ritual may warn against recurrence. But if the appraisal implies that present strangers deserve hostility for inherited identity, or if activation crowds out every current task, the function has shifted.

Grief, commemoration, historical study, vigilance, and commitment can carry memory without seeking an available target for punishment. Indignation may remain appropriate when mechanisms persist or repair is incomplete. The audit asks for present object and job, not mandatory emotional retirement after a bureaucratic interval.

Gratitude and the controlling gift

Gratitude is often treated as uncomplicatedly pleasant. It can recognize dependence, deepen relationship, direct attention toward benefit, and motivate generosity. It tells another person, “What you did reached me.”

Now consider the controlling gift.

Ravi pays for Jo’s education while announcing that no repayment is expected. Years later, Ravi invokes the gift whenever Jo makes an independent choice: career, partner, city, politics. “After everything I did for you,” he says.

Jo may feel gratitude for the education and reject Ravi’s claim to control. The benefit was real. The manipulation is also real. Gratitude need not become an unlimited debt or transferable authority over unrelated choices.

Audit the components:

- Was the benefit intentional?
- Was it welcomed?
- What costs or strings were concealed?
- What relationship existed?
- Was a return promised?
- What form of acknowledgment expresses Jo’s values?
- Does public gratitude reward manipulation?
- When does appreciation become coerced submission?

Jo can say, “I appreciate the help, I did not sell authorship of my life, and I reject your use of gratitude as a leash.” Emotional complexity is not inconsistency.

Pride and borrowed sunlight

Pride can recognize achievement, sustain agency, counter learned inferiority, and celebrate membership. It can motivate further effort and allow a community to see people long erased.

It can also claim credit without contribution.

Toma feels intense personal pride because a scientist from his nation wins an international prize. Toma did not know the scientist, fund the work, support the institution, or share the relevant field. What is the appraisal?

If it is “I achieved this,” it is false. If it is “Someone connected to an identity I value achieved this, and I delight in the association,” it may be accurate. If it is “This proves members of our nation are superior,” the inference outruns the event.

Collective pride can repair stigmatized identity or become fuel for hierarchy. The same cheering crowd can perform belonging, gratitude, aspiration, or domination. Ask what is being claimed and what the emotion licenses next.

Personal pride in one’s own work also needs calibration. Achievement depends on effort, opportunity, teachers, luck, institutions, and inherited capacity. Recognizing dependence need not erase satisfaction. It can make pride less solitary and gratitude less servile.

Remorse after morality changes

Eris promised to continue a religious vocation for life. Years later, after careful reflection, she abandons belief in objective promise-breaking and leaves the community. She still feels remorse.

Several facts may support it:

- people relied on the promise;
- departure imposes transition costs;
- relationships are wounded;
- Eris once made a sincere commitment;
- the community shaped her identity;
- leaving may nevertheless protect her agency and honesty.

Rejecting objective obligation does not imply that nothing happened. Eris can acknowledge reliance, assist transition, grieve, apologize for avoidable harms, and accept altered trust. She can also reject the claim that an earlier commitment owns her life regardless of changed identity and coercive conditions.

The remorse may combine care, loss, fear of condemnation, attachment, and learned desert. Sorting them can reveal tasks. Not every remainder needs punishment; some are grief at the cost of choosing.

Build the emotion dossier

Purpose. Audit one emotion's trigger, appraisal, cause, function, expression, intensity, duration, and alternatives; the dossier increases understanding without grading character or promising voluntary control.

Choose one current or remembered emotion. Complete the dossier without deciding in advance whether the feeling is good or bad.

Emotion: *guilt / shame / resentment / indignation / gratitude / pride / other:* _____.

Object: *person, act, event, trait, institution, or relationship:* _____.

Immediate cause and trigger: _____.

Developmental or social history: _____.

Factual appraisal: *The emotion presents _____ as having happened.*

Evaluative appraisal: *It presents _____ as violation, defect, benefit, achievement, or desert.*

Audience: *self / target / victim / group / public / no one:* _____.

Current purpose: *recognition / protection / protest / repair / relationship / expression / coordination / regulation / punishment / other:* _____.

Current effects: _____.

Expression chosen: _____.

Intensity: _____, *supported or amplified by* _____.

Stopping or revision condition: _____.

Possible translation: *regret / concern / caution / appreciation / grief / delight / settled boundary / other:* _____.

Desert claim remaining: _____.

Now hold the feeling fixed while varying the appraisal. Could the emotion be understandable but inaccurate? Accurate but counterproductive? Useful but based on an exaggerated claim? Involuntary but not worth rehearsing? Painful and still protective? Pleasant and manipulative?

This is not a command to optimize every feeling into administrative efficiency. Expression, memory, ritual, connection, and dignity are themselves things people value. "Useful" must not mean useful only to an institution.

The dossier creates options: preserve, express, redirect, regulate, translate, or allow the feeling to fade. It does not grade your character or promise control you do not possess.

The thirty-sixth debt

Add this entry:

Debt 36 — The moralized appraisal: *Which factual, relational, evaluative, and desert claims does the emotion present; what current job, audience, intensity, and duration support it; and what remains when recognition, protection, protest, repair, expression, coordination, and regulation are complete?*

Removing intrinsic desert does not flatten emotional life. Pain still hurts. Betrayal still changes relationships. Protection still matters. Gifts can still move us. Achievement can still delight. Protest can still burn. Repair can still relieve. Grief may remain without a culprit.

Some moralized forms may lose support. Guilt can become self-harm after repair. Shame can enforce a harmless conformity. Resentment can continue punishment after every boundary is secure. Indignation can search for targets after the task is gone. Gratitude can become a chain. Pride can appropriate another person's life.

The aim is not purity from emotion. It is literacy about what emotions claim and do.

Niko may continue to wake at three in the morning. He can meet the feeling without treating it as a divine messenger or an enemy to be crushed. He can ask: What happened? What remains? Whom does this serve? What will I do now?

The emotional ledger closes without reaching zero.

One debt remains larger than any single feeling. We have tested domains, values, rules, reasons, evidence, authority, theories, reform, codes, power, responsibility, and emotional life. The same missing support has appeared in different clothes. Was that repetition an accident?

Chapter 37 — No Candidate Returns Intact

The chief engineer of Port Meridian returns to the theory hall.

She has not visited since Chapter 4, when she asked for a specification sheet and ruined an otherwise promising afternoon.

The hall has changed. The original blank constitution is surrounded by ledgers, maps, threshold rules, standing tables, procedural warrants, emotion dossiers, and enough conditional clauses to alarm the fire marshal. The architects' objects remain where they left them.

"Did you build it?" she asks.

You know which *it* she means.

Not merely a compassionate code. Not a stable agreement, expressive vocabulary, inherited practice, useful fiction, personal commitment, or defensible political settlement. Those have appeared in abundance.

She means the system requested at the beginning: true in a way that survives every actual rejection, contentful enough to guide action, and authoritative for agents who do not share its values. It must also be knowable by a defensible method, coherent through conflict, usable by finite beings, and corrigible without becoming whatever current power prefers.

Before answering, we need one final inspection.

The recurring debts may be an accident of the candidates chosen. Perhaps we kept making the same mistake because the correct foundation never entered the hall. Or perhaps the recurrence follows from the kinds of thing that could supply a decisive standard.

We will sort the possible foundations—not to prove that imagination has reached its modal limit, but to show where every route examined deposits its cargo.

Where this chapter is headed. We will place every examined grounding route into seven families and compare what each supplies with the original specification. This is the cumulative audit: its conclusion concerns what the project has earned, not what no possible philosophy could ever establish.

The seven foundation families

A moral system's decisive content must come from somewhere. In the candidates we have studied, the final selector belongs to one or more of seven families.

1. **An actual stance:** a preference, commitment, emotion, convention, agreement, role, culture, institution, or practice held by actual beings.
2. **An idealized or constructed stance:** what agents would endorse after specified information, coherence, impartiality, freedom, representation, or procedure.
3. **An authority:** a command, will, nature, purpose, revelation, office, or rightful relation that directs conduct.
4. **A non-evaluative natural fact:** suffering, desire, need, function, capacity, evolutionary history, social condition, causal relation, or other descriptive feature.
5. **A rational or conceptual requirement:** consistency, universalization, constitutive agency, logical implication, conceptual truth, or a condition of deliberation.
6. **A brute evaluative fact or aim-independent reason:** a basic fact that an action is wrong, a being has status, an outcome is better, or a consideration favors action independently of every stance.
7. **A hybrid:** some ordered combination of the first six.

This is a **grounding partition**, not a list of theory labels. Consequentialism could ground welfare in actual concern, natural flourishing, ideal agreement, or brute objective value. A religious ethic could invoke authority, natural function, rational order, and basic goodness. Constructivism could be stance-dependent or claim that the construction itself has objective authority.

Nor is the partition announced as a proof that no eighth family is logically possible. A new candidate may propose one. To count as new, it must explain why it is not an actual or ideal stance, authority, natural fact, rational relation, basic evaluative fact, or combination. Renaming a source in ceremonial Latin does not add a family.

The question is what each family can supply without borrowing from another.

Family 1: Actual stances and institutions

Begin with what is plainly real.

Humans and other beings care, avoid, attach, choose, promise, coordinate, resent, love, legislate, punish, celebrate, and build institutions. These stances have causal power. They organize attention, identity, motivation, expectation, and collective action.

Actual stances can produce:

- personal reasons relative to projects and aims;
- reliance and repair through promises;
- role duties inside practices;
- standards of correctness constituted by games, professions, languages, and laws;
- shared codes;
- pressure and incentive;
- legitimate procedures among participants;
- enforceable jurisdictions;
- criticism through inconsistency and broken commitment;
- meaning and belonging.

This family is contentful, accessible, motivating, and often usable. Its weakness relative to the original specification is not worthlessness. It is dependence.

If a verdict is made true by what actual people approve, commit to, or institute, then sufficiently different stances change the verdict or remove its practical support. A unanimous culture may still be challenged from another stance, by an affected outsider, through a rival commitment, or with power. The claim that the culture is objectively wrong even under universal approval requires something beyond actual stance.

One might say that certain attitudes themselves disclose objective value. Compassion responds to suffering; indignation perceives domination. That is possible. The proposal moves from **attitude as ground** to **attitude as access**. The objective value doing the grounding now belongs to Family 6, perhaps identified through a natural relation in Family 4. The attitude alone cannot certify its own tracking.

What survives: practical orientation, motivation, meaning, conditional reason, social standards, institutions, and real stakes.

What remains: stance independence and authority over the total rejecter.

Family 2: Idealization and construction

Actual stances are often ignorant, inconsistent, adaptive, frightened, manipulated, or self-ish. Idealization repairs them.

Give agents the relevant facts. Let them vividly experience others' positions. Remove coercion. Require consistency. Hide personal identity. Represent absent parties. Ask what no one could reasonably reject. Seek reflective equilibrium.

These procedures can expose error and domination. They can transform preference. They are among the best technologies in the hall.

Their selector problem appears in the design specification:

- whose initial standpoint enters;
- which facts count as relevant;
- which desires are authentic or adaptive;
- what degree of information is required;
- whether impartiality means equal concern, anonymity, or equal chance;
- which beings receive representation;

- what counts as a reasonable rejection;
- how risk and uncertainty are handled;
- whether convergence is required;
- what resolves several coherent outcomes.

An empty procedure cannot select. A contentful procedure selects because its architecture protects certain values.

There are three honest interpretations.

First, the procedure **models what actual people would endorse under improvements they value**. The result can be deeply authoritative to them while remaining connected to stance.

Second, the procedure **constitutes validity inside a practice of reciprocal self-government**. Its authority depends on membership, fair terms, and accepted or institutionally enforced relations.

Third, the procedure **tracks or constitutes objective morality for every relevant agent**. Then its selectors need approval-independent authority. Equal standing, ideal rationality, or fair representation carries the grounding debt.

Idealization can clean a lens. It cannot decide what the lens should face unless direction is among the specifications.

What survives: correction of factual error, reduced self-dealing, reciprocal justification, representation, coherence, and legitimate settlement.

What remains: an objective ground for the idealizing selectors and authority over nonparticipants.

Family 3: Authority, command, and nature

An authority can unite content and direction. Parents, captains, courts, traditions, and governments issue requirements inside relationships and roles. A divine source may be creator, sustainer, perfectly knowledgeable, and necessarily good.

Authority can explain:

- who speaks;
- how a directive reaches a role occupant;
- why coordinated obedience matters;
- how a community identifies and transmits requirements;
- why a person experiences obligation through trust, gratitude, dependence, or allegiance;
- how a standard stands above current state approval.

Raw power is insufficient. A tyrant's ability to punish creates prudential pressure. It does not make cruelty good. Sophisticated authority theories agree and ground rightful rule in legitimacy, relation, perfection, or goodness.

That repair is necessary. It moves the evaluative burden.

If an authority is obeyed because it is good, rightful, loving, or just, those properties determine which authority merits allegiance and which commands are authentic. If good-

ness is identical with a necessary divine nature, interpreters still need substantive content and access. If creation itself confers authority, the theory must explain why making or sustaining a being gives the maker moral jurisdiction rather than merely causal priority.

Human authorities face the same structure on a smaller scale. A court has legal jurisdiction because institutions recognize its office. Whether the law is morally legitimate depends on fair procedure, social commitment, protected interests, objective right, or another value.

An authority can stop a practical regress by issuing a final directive. It cannot stop the evaluative audit by stamping *authorized* on its own warrant.

What survives: relation, purpose, coordination, tradition, durable resistance, and role-specific direction.

What remains: why the source is rightfully authoritative, what makes its nature good, and how correct content is distinguished from interpretation.

Family 4: Natural facts

The natural world is not silent in the ordinary sense. It screams, nourishes, predicts, constrains, and occasionally drops a tree on a poorly positioned argument.

Facts reveal:

- who suffers and how much;
- what beings want and can understand;
- what threatens survival;
- which relationships support development;
- how institutions distribute risk;
- what capacities enable agency;
- which biological functions sustain a form of life;
- what policies cause injury or flourishing;
- how values and intuitions formed;
- what interventions are likely to work.

No serious public code can ignore these facts.

The grounding question concerns the move from description to terminal selection. Suffering gives an agent who hates suffering an instrumental consideration. Another person's suffering motivates a compassionate observer. A community can commit to reducing it. An ideal procedure can represent it. A brute value theory can say suffering is objectively bad. The bare fact that a nervous system is in pain does not, without one of these relations, contain the conclusion "every agent has an aim-independent reason to reduce it."

Naturalist realism offers a stronger claim: perhaps moral properties are identical with certain natural properties, as water is identical with H₂O. Then no inferential bridge from natural to moral is needed because the properties are one.

This deserves more than the reply "you cannot derive ought from is." Identity claims are not deductions from a description. They are substantive proposals about what moral terms refer to.

The proposal must still identify the natural property. Is goodness welfare, desire fulfillment, flourishing, cooperation, agency, or a complex cluster? Why that property rather than a rival? Does the identity preserve the authority and scope ordinary moral claims promised, or define *good* as a useful classification? How is disagreement about the identity corrected? Why does the property favor action for an informed rejecter?

If the answer is that the natural property is the one that objectively matters, *matters* carries Family 6. If reference is fixed by ideal inquiry or mature human concern, Family 2 or 1 enters. Natural identity may be true. The natural description alone does not earn the whole specification.

What survives: indispensable facts about stakes, means, consequences, capacities, causes, and institutional performance.

What remains: the terminal value, standing rule, distribution, and aim-independent favoring relation.

Family 5: Rational and conceptual requirements

Reason can expose contradiction, invalid inference, false universalization, arbitrary self-exemption, unstable intention, and failure to choose means suited to ends. Constitutive standards explain what counts as believing, promising, intending, or acting within a practice.

These are real achievements. A person cannot coherently intend an end, know a necessary means, retain the intention under reflection, and deny any instrumental relevance to the means without tension. A ruler who invokes a universal principle only when it benefits his faction can be shown inconsistent.

Can rational form generate moral content?

Suppose Kest, the coherent tyrant, wants stable domination. He gathers accurate facts, chooses effective means, makes no logical contradiction, and openly permits anyone with sufficient power to dominate. He does not universalize a personal exception; he accepts a brutal general rule.

To condemn Kest as irrational, the account can add that rational agency requires treating every agent as an end, accepting only reciprocally justifiable principles, or valuing agency wherever it occurs. These additions may be defensible. They contain equal standing, reciprocity, or value.

Alternatively, the theory can argue that making any practical claim commits Kest to reasons whose form is universal. Yet universal form does not decide whether the universal principle protects welfare, hierarchy, liberty, or conquest. Substantive scope enters when the reason relation is specified.

Conceptual truths can also carry content: perhaps cruelty is, by definition, wrongful infliction of suffering. Then “cruelty is wrong” is secure, but classifying an act as *cruel* may already include the verdict. The question becomes whether all agents have reason to avoid the non-evaluatively described conduct.

Rational requirements constrain a construction. They do not obviously supply its inhabitants, protected interests, or final destination.

What survives: logical discipline, means–end criticism, exposure of exception, and standards internal to practices of thought and action.

What remains: substantive standing, terminal ends, impartial concern, comparison, and authority over the coherent rejecter.

Family 6: Basic evaluative facts and reasons

The sixth family is the most direct objective answer.

Some facts simply count in favor. Suffering is bad. Persons have dignity. Reasons arise from needs, promises, agency, or value. These relations do not depend on approval. They are part of reality’s normative structure.

Every explanatory chain stops somewhere. Physics does not invalidate a fundamental entity by asking what smaller object authorizes it to exist. Logic cannot justify every rule by a prior inference without regress. A moral foundationalist may reasonably say that demanding a non-normative proof of normativity misunderstands basic facts.

This reply blocks a crude objection. “Brute” does not mean false, arbitrary, supernatural, or useless. A basic evaluative fact could be necessary, natural, non-natural, knowable through rational reflection, or evident in experience.

The Candidate System does not reject it merely for stopping.

It asks what was placed at the stopping point.

“There are objective reasons” leaves content open. A complete foundation must specify:

- which beings have direct standing;
- which features generate reasons;
- which agents the reasons address;
- whether reasons are agent-neutral or agent-relative;
- how strength and priority are determined;
- when reasons are defeated;
- whether rights, welfare, loyalty, or desert are basic;
- how conflicts among basics are settled;
- how inquirers distinguish recognition from projection;
- what explains calibrated access despite disagreement and genealogy;
- how the relation guides agents who understand but remain unmoved.

One can declare all of this fundamental: objective standing, value, weight, exceptions, and authority. That preserves the system by placing its entire constitution at the foundation. Coherence still requires the clauses to fit. Accessibility still requires a method. Calling every disputed selector basic halts explanation at the exact points the project asked us to construct.

Perhaps intuition supplies access. Chapter 20 did not show intuition universally unreliable. It showed that confidence, self-evidence, and moral phenomenology do not certify tracking without calibration. Perhaps reflective equilibrium supplies access. Chapter 31

showed that equilibrium can organize different starting points. Perhaps convergence supplies evidence. Chapter 22 separated shared conditions, shared formation, shared function, and truth tracking.

None of these defeats the possibility of basic moral knowledge. None, in the cases examined, established a sufficiently discriminating route to the detailed constitution.

The sixth family therefore preserves objective morality most directly. It does so by accepting the central objective relations as fundamental and leaving at least some identity, content, comparison, authority, or access debts unpaid.

What survives: the logical possibility of stance-independent moral facts and a clear statement of what realism claims.

What remains: earned content, application, conflict rules, and reliable access rather than foundational assertion alone.

Family 7: Hybrids

Hybrids correctly observe that no single practical lens sees enough. They combine facts, concern, rights, virtues, relationships, traditions, procedures, and judgment. They can produce the best code in the hall.

Their grounding options are familiar.

If actual commitments order the components, the system belongs finally to Family 1. If ideal agreement orders them, Family 2 governs. If divine goodness orders them, Family 3 leads and Family 6 supplies goodness. If flourishing orders them, Family 4 needs an evaluative bridge. If practical reason orders them, Family 5 needs substantive content. If an objective “all-things-considered” relation orders them, Family 6 carries the final load.

If no component leads, the hybrid may acknowledge multiple permissible settlements or political choice. It can remain humane and usable while surrendering unique objective determinacy.

Combination can distribute work. It cannot eliminate the need for an ordering source.

What survives: practical completeness, sensitivity, safeguards, multiple correction routes, and institutional realism.

What remains: the authority and access debts of the decisive component plus the meta-rule that ranks the rest.

Four final replies

Before stating the finding, let four objections have the last word.

“Moral experience is evidence”

It may be. People experience obligation as discovered, not invented. Cruelty can appear wrong with a force unlike preference. Victims perceive violation; agents experience remorse; impartial judgments sometimes converge.

Experience is data about phenomenology and may be evidence of external truth under a reliable relation. The audit asks how that relation is calibrated. Moral appearances vary with framing, culture, identity, power, and prior theory. Some survive correction; some do not. Using the appearance to prove the faculty and the faculty to certify the appearance closes the support circle.

The finding is not “moral experience is worthless.” It is “appearance alone does not pay the tracking and content debts.”

“The skeptic relies on epistemic reasons”

Perhaps there are stance-independent reasons to believe according to evidence. If so, normative reality exists in at least an epistemic form, and a global argument that all normativity is impossible would fail.

This book has not made that argument.

Evidence can indicate truth, serve belief’s constitutive function, advance a truth-seeking aim, support practical prediction, satisfy inquiry rules, or justify public procedure. Readers who accept objective epistemic reasons can add that relation. Readers who do not can follow the audit conditionally: if you seek true belief, use methods that discriminate truth from error.

An epistemic reason to believe what evidence supports does not identify which beings have moral standing, whether suffering outweighs liberty, or why a promise binds after every relevant stance disappears. It may support our method of inquiry without supplying the missing moral content.

Parity can keep moral reasons possible. It does not complete their constitution.

“We cannot deliberate without objective reasons”

We do deliberate. We compare outcomes, commitments, identities, risks, agreements, and institutions. We revise beliefs through evidence. We choose amid conflict and uncertainty. Much of this activity is intelligible through aims, roles, practices, and procedures.

If all deliberation constitutively presupposes objective practical reasons, that is a Family 5 proposal. It must show why the constitutive requirement carries substantive moral standing and priority. A person can deliberate efficiently about conquest. The existence of choice does not certify the value of the chosen end.

Practical necessity forces settlement. It does not prove that the settlement was waiting in objective moral reality.

“Failure to explain is not evidence of absence”

Correct.

An undiscovered particle does not vanish because current instruments fail. A true theorem may await proof. Objective moral facts could exist beyond our present access or conceptual scheme.

That is why the conclusion must be written carefully. The construction project sought a warranted system, not a wager on every hidden possibility. When every candidate leaves decisive promissory notes, we may say the project has not earned completion. To infer nonexistence, we would need an additional premise—for example, that objective moral facts would be accessibly explanatory if real, or that a rival explanation of moral appearances is sufficiently superior.

Those premises can be argued. They were not smuggled into this inspection.

The general result

Place the original specification beside the seven-family audit.

The result can now be stated as an argument.

1. The project asked for a system with at least some stance-independent moral truth, substantive content, practical authority over relevant rejecters, defensible access, coherence through conflict, enough determinacy for guidance, usability, and correction.
2. Every decisive grounding route examined derives its content from an actual stance, idealized stance, authority, natural fact, rational requirement, basic evaluative fact, or hybrid of these.
3. Actual stances provide content and practical force but do not preserve the required independence.
4. Idealization and construction improve stances but require substantive selectors; treating those selectors as objectively authoritative moves the grounding burden rather than paying it.
5. Authority supplies direction or enforcement only through a prior account of rightful goodness, relation, or legitimacy.
6. Natural facts constrain and inform action, but terminal value, standing, and aim-independent favoring require an evaluative bridge or identity claim whose moral role remains to be established.
7. Rational requirements constrain inference and means–end coherence, but substantive standing, concern, and priority enter through added premises.
8. Basic evaluative facts preserve the objective claim directly but, in the candidates examined, leave decisive content, comparison, authority, or access at a foundational stopping point.
9. Hybrids inherit the unresolved debts of their decisive source and add an ordering debt.

Therefore:

No route inspected yields a system that has earned all the central properties of a coherent objective moral system without importing an unsupported selector, leaving a decisive bridge unpaid, or lowering one of the original ambitions.

The positive counterpart is equally important:

Coherent, informed, humane, enforceable, and revisable codes can be built when their goals, values, procedures, jurisdictions, and limits are declared rather than presented as discoveries certified by objective moral authority.

And the epistemically restrained conclusion is:

No coherent objective moral system has been found. This is not a deductive proof that no logically possible objective moral fact exists. It is a demonstration that the construction project has no warranted completion under the routes examined.

The wording matters.

“No objective morality exists” climbs beyond the current argument unless an additional anti-realist premise is supplied. “Anything goes” contradicts the institutions, commitments, resistance, evidence, and power we have just spent six chapters making explicit. “All moral systems are equally good” tries to use an objective ranking after denying one and ignores what actual beings protect.

The finding is neither cosmic certainty nor a shrug. It is a completed audit with a conspicuous unpaid column.

The thirty-seventh debt

Add the final entry of the construction phase:

Debt 37 — The possibility leap: What additional premise licenses movement from “no candidate examined has earned the objective specification” to “no objective moral fact or system could exist,” and what evidence supports that premise without exceeding the inference ladder?

You may choose to defend that premise in a separate argument. You may preserve realist hope. You may suspend judgment. What the ledger does not permit is treating an unfinished bridge as completed because the far shore feels important.

Inspection Report VI — The Unfinished Constitution

Purpose. Compare the final Candidate System with Version 0.1 and classify every surviving clause, revision, surrender, and debt; the report supports a bounded construction result, not a proof of metaphysical impossibility.

Retrieve Candidate System 0.1 from Inspection Report I. Do not rewrite it from memory. Place it beside every revision you made.

For each surviving clause, assign one or more labels:

- **Factual** — reports a description, cause, capacity, risk, consequence, or prediction;
- **Instrumental** — identifies a means relative to a declared goal;
- **Expressive** — voices or communicates a commitment, emotion, identity, protest, or appraisal;
- **Commitment-based** — guides because an agent or group adopts a value, project, promise, or identity;
- **Agreement- or role-based** — arises within valid terms, reliance, membership, or institutional position;
- **Procedural** — results from an authorized method of settlement or revision;

- **Enforced** — backed by monitoring, sanction, coercion, or force;
- **Objectively asserted but unpaid** — claims stance-independent moral status or authority without a completed ground and access route.

Do not treat these labels as a ladder from weak to strong. A factual prediction may be highly certain and politically decisive. An expressive protest may transform a society. A procedural rule may settle a crisis. An objective assertion may be motivationally inert. The labels identify the kind of support.

Reinspect the nine demands

Complete the final sheet.

Original demand	What the Candidate System now has	What was surrendered or narrowed	Remaining objective debt
Truth	_____	_____	_____
Stance independence	_____	_____	_____
Scope	_____	_____	_____
Authority	_____	_____	_____
Accessibility	_____	_____	_____
Coherence	_____	_____	_____
Determinacy	_____	_____	_____
Usability	_____	_____	_____
Correctability	_____	_____	_____

Then add any humaneness condition you chose. State plainly whether it is an objective result, a protected commitment, an agreed term, or an institutional goal.

Mark every bridge

For each arrow from fact to action, write the connecting value.

For each arrow from value to rule, write the comparison or priority method.

For each arrow from rule to enforcement, write the jurisdiction and safeguard.

For each arrow from evidence to belief, write the truth-responsive route you accept.

For each arrow from disagreement to settlement, write the procedure and what it does *not* prove.

For each claim of objective authority, write the foundation, content, scope, weight, access, correction method, and response to the informed rejecter.

Any blank is not automatically a defect in practical life. It is a limit on the type of claim earned.

The final construction decision

Choose one of four entries for the cover of your constitution.

1. **Objective system completed.** Attach full warrants for every essential demand and explain why the seven-family audit missed or misclassified the decisive foundation.
2. **Objective possibility preserved; system unfinished.** Retain one or more objective claims as promissory notes and state what future evidence or argument could pay them.
3. **Practical code adopted without objective authority.** State the goals, values, participants, procedures, safeguards, enforcement, and revision conditions.
4. **Mixed practice.** Preserve ordinary moral language in some contexts while translating or narrowing its authority in others. Part VII will test this option.

No answer is accepted merely because it matches the book's finding. The anti-rigging declaration still applies. If you believe the system is complete, exhibit it. If you narrow the project, mark the surrender without shame. If you abandon an objective assertion, do not quietly borrow it in the next political dispute.

What remains in the hall

The unfinished constitution is not empty.

It contains evidence about suffering, agency, dependency, risk, consequence, power, and institutional performance. It contains care, aversion, attachment, loyalty, curiosity, hope, fear, and resistance. It contains commitments, promises, roles, and shared histories. It contains methods for comparison, deliberation, representation, dissent, enforcement, appeal, and revision.

What it does not contain—unless you supplied a foundation beyond those successfully audited—is a paid warrant showing that its substantive selectors possess stance-independent truth and practical authority for every relevant rejecter.

The blank space at the center is now understood, not merely felt.

The chief engineer reads the report.

"Then we did not build what we requested," she says.

"No."

"But we learned how codes are actually built."

"Yes."

She looks at the ledger, the boundary marker, the chairs, the lamp, the envelopes, the maps, and the hinges. Then she looks through the open doors at Port Meridian: frightened people, hopeful people, powerful people, exhausted people, people who agree about slogans and disagree about everything underneath them.

"What do we do with all this?"

That is no longer a question about finding the missing system.

It is the beginning of Part VII.

PART VII



After the Search

The failure of a cosmic constitution is not the disappearance of pain, love, evidence, conflict, or politics. The search ends. Practical life does not.

Chapter 38 — Decide What to Do with Moral Language

The chief engineer asks what to do with the unfinished constitution.

The city translator answers first.

“Before we write another law,” he says, “we need to decide what our old words are allowed to mean.”

On the table lies one ordinary sentence:

You broke your promise. That was wrong.

For thirty-seven chapters, we have separated possible jobs hidden inside that sentence. It may report an objective moral fact, apply a social rule, invoke a commitment, express protest, demand repair, warn others, condemn character, intensify blame, or open negotiation. Speaker and listener may hear different jobs while believing they disagree about only one thing.

The failure to establish objective moral authority does not make the sentence vanish. It creates a language decision.

Should we stop using moral terms? Speak as if they report facts while privately denying that they do? Preserve ordinary use because it performs useful social work? Revise the terms to express commitments? Reconstruct them around fair procedures? Translate only when precision matters?

No theorem in Part VI selects the answer. Language policy is practical policy. It depends on what speakers care about, what audiences infer, what misunderstandings cost, and what alternatives can do.

Where this chapter is headed. We will compare six policies for moral language after the construction audit, from abolition to context-sensitive translation. The choice is practical and semantic: it should preserve useful communication while revealing hidden claims where the stakes demand precision.

Six policies for the old vocabulary

Imagine six doors in the translation office.

Door 1: Abolition

Stop asserting moral verdicts. Replace *wrong*, *duty*, *deserves*, *evil*, and their relatives with concrete descriptions.

Instead of:

Breaking that promise was wrong.

say:

You gave me assurance, I relied on it, your undisclosed change imposed this loss, I no longer trust the same commitments, and I want repair.

The advantage is honesty and precision. The translation reveals reliance, effect, relationship, and requested response. It does not imply a mysterious property when none has been established.

Abolition can also interrupt manipulative language. “You owe your family” may conceal threat, dependence, tradition, affection, or control. Unpacking forces the speaker to name the relation.

The costs are substantial. Moral vocabulary is compressed, emotionally resonant, socially learned, and woven into law, intimacy, protest, and self-understanding. Replacing a three-word condemnation with a paragraph can make urgent speech sound like a municipal permit application. Victims may hear translation as evasive cleverness. Children and exhausted adults do not always need a seminar before someone says, “Do not hit.”

Abolition also risks private moral language surviving underground, where it becomes harder to inspect.

Door 2: Conscious fiction

Continue speaking as though objective moral facts exist because the fiction supports coordination, motivation, or humane conduct.

This resembles other useful pretenses. A nation may personify itself. A court may treat a corporation as a legal person. People speak of the sun rising without endorsing ancient astronomy.

The fiction could give moral education narrative force and preserve familiar practices while informed adults understand the underlying construction.

Its danger is asymmetry. Who knows it is a fiction? If officials privately regard obligation as invented while telling citizens that moral truth commands obedience, the practice

becomes paternalistic. If everyone knows, the emotional force may change. If almost everyone forgets, the fiction has returned as belief.

A conscious fiction also needs content. Which imagined duties should the story contain? Fictional status does not choose between compassion and conquest. The selectors remain human values and power.

Door 3: Conservation

Keep ordinary moral language largely unchanged.

One conservative defense says our audit may be mistaken; ordinary realism should not be revised by a contested philosophical argument. Another says moral discourse has many uses beyond metaphysical reporting. Speakers can competently say “That was unjust” without carrying a theory of stance-independent properties, just as they can discuss money without a philosophy of institutional reality.

Conservation respects linguistic life. It avoids turning every friendship into a position paper. It also allows moral terms to remain sites of contest: communities can argue over what fairness or cruelty means without first settling metaethics.

The risk is continued equivocation. A speaker may use *wrong* to express commitment when challenged about truth, then use it as an objective fact when demanding authority over a rejecter. Ordinary language makes the costume change easy. Conservation needs at least a rule against borrowing metaphysical force in high-stakes contexts without disclosure.

Door 4: Expressivist revision

Use moral language openly to express commitments, shape attitudes, coordinate response, and invite others to join.

“That was wrong” can function roughly as:

I oppose that conduct, I invite similar opposition, I support norms and responses against it, and I am presenting the case as one that matters beyond my private inconvenience.

This preserves force without pretending the sentence reports an objective property. It explains why moral disagreement can persist after factual agreement: people are not merely describing; they are contesting what to endorse and build.

Expression is not a noise emitted instead of thought. It can include reasons, evidence, consistency, imagination, and concern for other standpoints. A public commitment can be sincere, disciplined, and revisable.

The danger is underdisclosure. Listeners may hear objective fact while the speaker means avowal. Expressivist language can also become propaganda: intensity, repetition, shame, and group identity may recruit attitudes without improving factual understanding. Expression needs responsibility for effects.

Door 5: Constructivist revision

Use *right*, *just*, or *required* for outcomes of specified fair procedures or standards people can reasonably support.

For example:

Under a charter that gives affected residents representation, protects minimum security, and permits appeal, the levy is justified.

The moral term now refers to a construction rather than an order discovered independently of every stance. This makes public language accountable. A speaker can identify the procedure and its limits.

The risk is procedural inflation. “The law is just” may be heard as “the law is objectively morally correct” when it means “the law survived this authorized process.” A procedure can exclude outsiders, protect the wrong interests, or preserve unequal power. Constructivist revision must keep participant, standing, constraint, and enforcement questions visible.

Door 6: Ordinary-language pragmatism

Keep, revise, or translate moral terms according to context.

Say “That was wrong” when the shared meaning is sufficiently clear and the stakes do not require a semantic customs inspection. Translate when authority, coercion, identity, or serious misunderstanding enters. Correct the listener when they hear more than you mean.

This policy treats language as plural. A protest sign, apology, family conversation, judicial opinion, and philosophy book need not use *wrong* in exactly the same way. Meaning partly follows practice and audience.

The advantage is flexibility. The danger is opportunism. Without a disclosure rule, pragmatism can become “use whichever meaning wins.” A context-sensitive policy needs stable triggers for translation.

The promise at the station

Eren promises to collect Sal from the station at midnight. Sal relies on the promise, declines another ride, and waits alone in dangerous weather. Eren goes to a party instead and sends no message.

The next morning Sal says, “You should not have done that.”

What might Sal mean?

- Eren’s conduct caused foreseeable fear and risk.
- Eren voluntarily gave an assurance.
- Sal reasonably relied.
- Their friendship includes expectations of care and candor.
- Eren had an easy alternative: notify Sal.
- Sal now withdraws trust.

- Sal wants acknowledgment and changed conduct.
- The community condemns casual promise-breaking.
- Eren had an objective duty independent of all these relations.

The first seven can be true under the fixed facts. The eighth reports a social standard. The ninth adds the objective claim.

Would translating the sentence weaken Sal's protest? Not necessarily. "You created a danger after asking me to rely on you, and I will not place myself in that position again" may be more powerful than *wrong*. But Sal may reasonably prefer the shared compression of ordinary language.

The policy question is whether Eren and Sal understand the compression. If Eren replies, "Objective duties do not exist, therefore your fear is irrelevant," Eren has made the nihilist inference Chapter 39 will dismantle. If Sal replies, "Because you acted wrongly, any suffering I impose is deserved," Sal has added a desert claim not contained in reliance.

Translation prevents both expansions.

Parenting without a tiny metaphysics lecturer

A four-year-old hits another child with a wooden train.

A parent says, "That was wrong. We do not hit."

Must the parent immediately add a grounding partition?

No. The sentence can interrupt conduct, mark a boundary, direct attention to pain, and state a family rule. The next sentences matter more:

"Look—Mina is hurt. Put the train down. Give her space. We will help. When you are angry, call me or move away."

This sequence connects rule, evidence, response, and alternative action. It does not rely on shame or essential badness.

Children gradually learn what terms do. Adults can teach that *bad action* does not mean *bad kind of person*, that rules have purposes, that effects matter, and that people may challenge unfair rules. Moral language can scaffold attention without being presented as an invisible police officer who lives inside reality.

The risk of abolition in parenting is cumbersome abstraction. The risk of conservation is authoritarian mystery: "It is wrong because wrong things are wrong." A mixed policy can use compression, then unpack.

Friendship and the right to leave

Nara tells a friend, "You owe me loyalty."

That sentence could invoke mutual history, an explicit promise, gratitude, fear of abandonment, status, or control. It might request confidentiality. It might demand complicity in harm.

In intimate life, translation protects agency:

“I trusted you with this information and want you not to disclose it.”

“I supported you during the trial and hoped you would attend mine.”

“Your disagreement feels like abandonment, though you never promised agreement.”

These statements expose different claims. The friend can respond to each without accepting a total duty of loyalty.

Ordinary moral vocabulary remains useful for boundaries—“You may not threaten me”—especially when a person is too frightened for extended analysis. No language policy should demand philosophical elegance from someone escaping danger.

Blame and criminal law

High-stakes institutions need the strictest translation rules.

“The defendant deserves punishment” can hide protection, deterrence, restitution, condemnation, incapacitation, revenge, or basic desert. Each authorizes different severity and stopping conditions.

A criminal judgment should therefore state:

- the conduct proved;
- the evidence standard;
- the agent’s knowledge and control;
- the protected stake;
- the response’s intended purpose;
- expected effects and uncertainty;
- maximum burden;
- review and stopping conditions;
- what repair follows official error.

The word *justice* may remain in the courthouse. It cannot replace these entries.

The same applies to administrative language. “Public safety requires” often means “officials predict a risk and prioritize its reduction over privacy, mobility, cost, or error.” State the prediction, priority, distribution, and authority.

Compression is most dangerous where compressed words unlock cells.

Reform and public persuasion

Movements use moral language to focus attention and create solidarity. “This is unjust” can place an ignored injury in public view, connect separate experiences, and demand institutional response.

Translation should not become a demand that reformers drain every sentence of urgency while defenders keep flags, hymns, and fear. Emotional expression is part of political life.

Yet movements also wield power. A moral label can erase internal disagreement, depict opponents as essentially corrupt, or substitute shared indignation for accurate prediction. The strongest public persuasion can do both:

***Compression:** “No child should drink poisoned water.”*

***Expansion:** “Testing shows lead exposure at these addresses; children face these developmental risks; landlords and officials controlled the pipes; we seek replacement, medical support, independent monitoring, and a deadline.”*

The first expresses a protected commitment. The second makes the policy contestable and correctable.

Choose a context-sensitive policy

Purpose. Decide how moral language should function across personal, relational, institutional, and emergency contexts; variation requires a disclosure trigger rather than whichever meaning wins locally.

Complete this table.

Context	Language policy	What the moral term does	When translation is required	Main risk
Promise	_____	_____	_____	_____
Parenting	_____	_____	_____	_____
Friendship	_____	_____	_____	_____
Personal guilt or blame	_____	_____	_____	_____
Criminal law	_____	_____	_____	_____
Social reform	_____	_____	_____	_____
Public persuasion	_____	_____	_____	_____

Possible risks include false objectivity, emotional flattening, manipulation, ambiguity, verbosity, procedural inflation, unequal disclosure, and the accidental suggestion that nothing matters.

Then adopt three disclosure triggers. Translate a moral claim whenever:

- 1. it is used to justify coercion, punishment, exclusion, or serious deprivation;
- 2. speaker and listener may disagree about the term’s authority or ground;
- 3. the compressed verdict conceals affected parties, factual predictions, value tradeoffs, or stopping rules.

Add your own trigger if needed.

The language charter

Purpose. State your stable policy for retaining, revising, translating, fictionalizing, or abandoning moral terms, including the circumstances in which compressed language must be unpacked.

Write a short policy:

I will retain / revise / abandon moral language in _____ contexts. By terms such as _____, I usually mean _____. I will not infer _____ merely from using those terms. When coercion or serious conflict is involved, I will translate the claim into facts, stakes, values, procedures, power, safeguards, and revision conditions. I will correct listeners when they attribute objective authority I do not claim.

If you preserve moral realism as an open possibility, say so:

I continue to believe / suspect that _____ is objectively true. The current ground and access route are _____. While that claim remains contested, my public proposal can also be assessed through _____.

This allows conviction without pretending the earlier debt disappeared.

What the audit does not choose

Rejecting or suspending belief in objective moral facts does not logically entail abolition, fictionalism, expressivism, constructivism, or pragmatic conservation. Those are further choices about communication and practice.

Abolition maximizes explicitness and risks loss of compression. Fiction preserves familiar force and risks deception. Conservation respects ordinary life and risks equivocation. Expressivist revision declares commitment and risks manipulative recruitment. Constructivist revision connects claims to public terms and risks confusing procedure with truth. Pragmatism adapts to context and risks opportunism.

No single policy dominates every setting.

The translator writes one sentence above all six doors:

When the cost of misunderstanding rises, increase the resolution of the language.

The chief engineer nods.

“If we remove the objective claim from a sentence,” she asks, “what remains inside it?”

The translator opens the unfinished constitution.

Almost everything that made the sentence urgent is still there.

Chapter 39 — Everything That Did Not Disappear

The chief engineer opens the unfinished constitution as though checking a house after a storm.

The roof labeled **Objective Authority** is missing.

The furniture is still there.

Pain did not wait for a theory before hurting. Fear did not become an argument. Attachment did not submit an application for permission to matter to the attached. Betrayal still changes trust. A locked cell still restricts movement. A lie can still alter a decision. Lead still damages a child's nervous system. A government can still conceal the measurements.

The discovery that no objective system has been earned is easy to mishear because moral language has bundled many things beneath one roof. Remove the roof and people imagine empty land.

We need an inventory.

Where this chapter is headed. We will inventory what survives the loss of an earned objective foundation: experience, attachment, evidence, agency, institutions, power, responsibility, repair, and hope. The inventory prevents a narrow metaethical conclusion from expanding into nihilism.

Sensation and affect remain

Pain is an experience creatures commonly avoid. Pleasure attracts attention and pursuit. Fear changes the body, narrows options, and prepares action. Relief arrives when danger recedes. Grief registers loss. Loneliness, delight, disgust, curiosity, boredom, comfort, and awe continue their unruly careers.

None proves an objective reason.

None needs objective status in order to be causally, psychologically, or politically consequential.

A person in severe pain may want it stopped. Someone who loves them may respond. A clinician may have a role commitment. A government may have adopted a health charter. Voters may fund care. Officials may use force to distribute resources. These supports can converge without transforming pain into a stance-independent command addressed to every possible agent.

The pain remains exactly as painful during the metaphysical delay.

Attachment and loss remain

People love children, friends, partners, places, languages, rituals, animals, works, and ways of life. They form identities through memory and dependence. Threats to these attachments produce grief, anger, fear, courage, and cooperation.

An attachment can be generous or controlling. Its existence does not make every demand issued in its name acceptable. “I love you” can support care and become a preface to possession. Evidence and other values still discipline action.

But attachment is not a defective moral proof. It is not primarily a proof at all. It is a practical starting relation. It explains why certain facts become salient, why loss has weight for someone, and why coalitions form to protect homes, families, habitats, or traditions.

When people say “This place is sacred,” translation may reveal history, identity, grief at destruction, dependence, beauty, communal memory, or divine belief. The objective sacredness may remain disputed. The attachments do not.

Trust and betrayal remain

Trust is a willingness to accept vulnerability based on expectations about another person or institution. It can be predictive, emotional, role-based, and chosen. Betrayal occurs when someone uses or disregards that vulnerability in ways the trusting person did not accept.

No moral property is needed to explain why betrayal changes future conduct. A betrayed person updates predictions, withdraws access, seeks explanation, warns others, or ends a relationship. Communities create records, contracts, audits, and sanctions because trust lowers the cost of cooperation and its abuse can be severe.

Objective obligation promised to say the betrayer had a reason not to betray even if they cared about none of this. Without that warrant, we lose the universal claim. We do not lose reliance, evidence, injury, memory, boundary, or the capacity to redesign the relation.

Dependence remains

Human beings begin radically dependent and remain interdependent. Food, care, language, infrastructure, medicine, ecosystems, and accumulated knowledge make individual agency possible. The fantasy of a completely self-made person is a group project performed by invisible workers.

Dependence creates vulnerabilities. Caregivers can neglect or control. Employers can govern housing and medicine. States can condition survival on obedience. Platforms can mediate work and speech. Supply chains can move suffering out of view.

Facts about dependence do not select one policy. They show where power resides and whose failure would be costly. People who care about security, agency, reciprocity, or relationship can use that map to build protections.

Agency remains

People form intentions, imagine alternatives, learn, plan, resist, deceive, cooperate, and change. Their capacities vary across time and condition. Institutions can expand or constrict those capacities.

Without objective morality, agency does not become an illusion. Nor does every choice become equally informed, voluntary, or effective. Coercion narrows options. Deception corrupts information. Poverty changes exit costs. Trauma affects attention. Law redistributes control.

If you value agency, these facts inform design. If someone rejects agency as a concern, facts alone will not supply the missing value. Others may still defend their agency against that person.

Humiliation and domination remain

Humiliation involves exposure, degradation, status injury, or forced submission as experienced by a person and interpreted within a social field. Domination involves durable power capable of arbitrarily controlling another's options.

These phenomena can be studied. We can ask who controls resources, whose speech carries risk, who must seek permission, who can appeal, and whose self-description is ignored. A benevolent dominator may still hold uncontrolled power. A humiliating practice may persist even when officials call it harmless.

The claim "domination is objectively wrong" remains a moral assertion. The statements "I hate living at another's uncontrolled mercy," "this institution predictably suppresses dissent," and "we will divide its power" identify response, fact, and commitment.

Those are enough to begin resistance. They are not guaranteed to persuade the dominator.

Scarcity and tradeoff remain

There are only so many rescue vehicles, clinic hours, houses, watts, wetlands, attention spans, and minutes before the meeting ends. Protecting one stake can burden another. Delay is itself a choice with effects.

Objective morality promised a correct all-things-considered ranking. Its absence does not make scarcity polite. Decisions still require aggregation, priority, boundaries, thresholds, plural judgment, procedure, or chosen settlement.

The difference is disclosure. "This is the morally correct allocation" can become:

We prioritized severe and imminent health loss over modest convenience, guaranteed a minimum to each district, used these estimates, assigned this decision-maker, and will revise when these conditions change.

The second statement may be longer. It also contains handles for evidence and dissent.

Evidence and truth-responsive practice remain

Abandoning objective moral facts does not imply that bridge-load estimates become expressive preferences.

Some claims are better supported than others. Measurements can be calibrated. Predictions can be compared with outcomes. Records can expose false histories. Experiments can identify causal effects. Testimony can reveal experiences officials missed. Models can disclose uncertainty.

Chapter 19 kept the source of epistemic force explicit. You may regard evidence as carrying stance-independent epistemic authority, as constitutive of belief, as serving a truth-seeking aim, as useful for prediction, as part of disciplined inquiry, or as a condition of public justification. Under any truth-responsive route, evidence can correct the factual model used in policy.

It cannot decide who should bear a known risk without a value or procedure. That boundary survives too.

Moral emotions remain

Guilt, shame, resentment, indignation, gratitude, and pride do not dissolve when basic desert is removed. Their causes, appraisals, functions, expressions, intensity, and duration can be examined separately.

Some emotions may lose a particular justification. Guilt may no longer be cultivated as self-punishment after repair. Shame may be recognized as enforcement of a harmless conformity. Gratitude may be detached from obedience. Resentment may shift into boundary and memory.

Other jobs remain: recognition, relationship, protest, repair, expression, coordination, and self-regulation.

The emotional question changes from “What feeling is objectively deserved?” to “What is this feeling tracking, claiming, doing, and producing now?”

That is not emotional arithmetic. It is literacy.

Causal responsibility remains

Niko cut the fuel line. Mira did not cause the undetectable defect by negligence. Executives designed incentives that made inspection failure likely. These causal and evidential distinctions remain.

Knowledge, intention, control, and alternatives affect prediction, trust, response, and repair. Institutions can protect, deter, restore, rehabilitate, set boundaries, and communicate. Removing intrinsic payback removes one purpose for suffering; it does not erase agency or response.

“No one is basically deserving” does not mean “no one did it.”

Institutions, laws, and agreements remain

An institution is not a hallucination because humans made it. Money, courts, languages, borders, corporations, marriages, universities, and traffic systems depend on practices and recognition. They can still structure lives with astonishing force.

A law can be valid within a jurisdiction, enforced, predictable, democratically authorized, broadly supported, and instrumentally effective without being objectively morally correct. An agreement can create expectations and conditional reasons among parties without becoming a command to outsiders.

Constructed does not mean imaginary. A staircase is constructed. Try ignoring one while descending.

Identity and community remain

People understand themselves through memberships, narratives, roles, and contrast. Communities preserve memory, mutual aid, celebration, and resistance. They can also police boundaries, erase dissent, and impose inherited scripts.

Objective moral theory does not own belonging. Skepticism does not require becoming a citizen of nowhere.

A person can love a tradition while refusing its universal authority. They can preserve a ritual because it connects generations, revise it because it humiliates members, and explain both choices through attachment, effect, identity, and power.

Power remains

Some people control land, law, weapons, information, money, employment, attention, and institutional agendas. Others bear decisions without meaningful voice. Procedure can redistribute power; rhetoric can conceal it; organization can resist it.

The absence of objective authority makes power more visible, not less real.

Political actors may still claim divine, natural, rational, historical, or democratic authorization. The audit asks what those claims add beyond jurisdiction and force. Meanwhile, anyone designing a policy must ask who can implement, evade, capture, and challenge it.

Repair remains

People can restore property, compensate loss, acknowledge injury, correct records, change systems, rebuild trust, and create safeguards. Some losses cannot be undone. Repair can then mean recognition, support, memory, and prevention rather than pretending restoration is complete.

Repair depends on values and relationships. It does not require the universe to prefer the repaired state. The concern of the injured person and their allies can be enough to begin.

Imagination and hope remain

Human beings imagine arrangements that do not yet exist. They rehearse futures, tell stories, recognize remote lives, and form coalitions around possibilities. Hope sustains action under uncertainty.

Hope is not a forecast. It can survive low probability and fail under false belief. Disciplined hope connects imagination to evidence, resources, strategy, and revision.

An objective guarantee that history bends toward a favored end would be comforting. We did not find one. Hope becomes less like confidence in cosmic architecture and more like a practice maintained by creatures who want a future enough to build toward it.

What was actually lost?

The inventory shows the narrower losses.

If the objective project is abandoned or suspended, we lose—or cease claiming:

- a stance-independent moral property attached to acts, persons, rules, or outcomes;
- categorical authority over an informed agent who rejects every connected value and relation;
- an objectively correct rank among all competing values;
- intrinsic desert that makes suffering or reward fitting in itself;
- a guarantee that a fair procedure discovers rather than constructs or settles;
- the ability to call every favored human concern a command of reality.

Some readers will experience these as enormous losses. They affect self-understanding, grief, hope, faith, anger, and the imagined vindication of victims. No inventory should tell a reader that the loss is merely verbal.

The point is accuracy: these losses do not include sensation, attachment, evidence, institution, protest, or the capacity to choose and organize.

Ten translations

Purpose. Preserve the factual, emotional, relational, and practical force of ten moral sentences after removing compressed status words; translation should reveal content rather than lower the stakes.

Rewrite each sentence without using *moral, immoral, right, wrong, good, evil, just, unjust, duty, ought, or deserves*. Preserve its emotional seriousness.

1. “It is wrong to poison children.”

Lead exposure causes pain and lasting developmental injury. We refuse to accept those effects for children, will replace the contaminated pipes, require independent testing, and hold institutions answerable for concealment.

2. “You ought to keep your promise.”

You asked me to rely on your assurance; I reorganized my life around it; if you withdraw, tell me early, address the loss, and expect my future trust to change.

3. “The dictator is evil.”

The ruler orders torture, suppresses evidence, targets opponents’ families, and uses fear to eliminate resistance. We oppose those practices and will support protection, escape, documentation, and constrained removal of his power.

4. “The offender deserves prison.”

The evidence establishes deliberate serious harm. We seek protection against recurrence, a public boundary, and structured repair; any confinement must serve those purposes, use the least destructive adequate conditions, permit review, and end when its stated warrant ends.

5. “That distribution is unjust.”

The policy gives a small benefit to a powerful majority by imposing severe health loss on an excluded district. We prioritize minimum health protection and equal political standing over that benefit and demand a revised allocation.

6. “Parents have a duty to their children.”

Children depend on caregivers they did not choose and cannot safely leave. We expect parents who assume that role to provide protection and care, and we will create public support and intervention when dependence is neglected or exploited.

7. “The company did the right thing.”

The company disclosed the defect before injury, funded repairs, protected whistleblowers, compensated disruption, and accepted independent monitoring—responses that serve the safety, candor, and accountability we seek from institutions.

8. “She should be proud.”

She sustained difficult effort, developed unusual skill, completed what she valued, and may reasonably feel satisfaction while recognizing the help and luck that made the achievement possible.

9. “You owe him gratitude.”

He intentionally provided help you welcomed. You can acknowledge the benefit and express appreciation without granting him control over unrelated choices unless you actually made such an agreement.

10. “Resistance is morally required.”

The government is detaining dissidents, hiding records, and closing peaceful appeals. We care about their safety and agency, will not cooperate with the detention system, and invite others to support documentation, sanctuary, legal challenge, and carefully constrained resistance.

Now write ten sentences from your own life or public concerns. If translation feels weaker, identify what disappeared. Was it compression, certainty, universality, condemnation, solidarity, desert, or authority? Decide whether you still claim that element and on what ground.

The aim is not to ban old language. It is to know what the old language was carrying.

Six false inferences

Purpose. Block six conclusions the construction audit does not support, especially nihilism, quietism, arbitrary factual belief, and the disappearance of responsibility or care.

Before leaving the inventory, cross out six conclusions the audit does not support.

“All preferences are equally strong”

Preference strength is an empirical and psychological matter. People sacrifice differently, persist differently, and organize differently. Equal objective status does not follow from lack of objective status.

“Everyone becomes indifferent”

Beliefs about grounding and intensity of care are distinct. A person may reject objective value and remain fiercely committed. Another may retain realism and be apathetic.

“Selfishness wins”

Compassion, attachment, loyalty, identity, and concern for strangers are actual stances too. Egoism does not become objectively true when realism fails.

“Law must disappear”

Law coordinates, authorizes, predicts, and coerces through institutions. Its existence does not depend on objective morality. Its design still depends on goals, values, agreements, power, evidence, and safeguards.

“Criticism becomes impossible”

Chapter 32 identified factual, internal, commitment-based, relational, procedural, and power-based criticism. What remains unavailable without further grounding is guaranteed objective authority over every rejecter.

“Empirical truth disappears too”

Moral non-realism is not global factual skepticism. The relation between evidence and belief received its own audit. Bridge failure, disease prevalence, and policy effects remain truth-directed questions.

The survivor’s declaration

Purpose. Record the experiences, beliefs, attachments, commitments, institutions, criticisms, and hopes you retain so the next chapter begins from an explicit practical inventory.

Complete these sentences.

What I still feel and care about: _____.

What I still believe factually: _____.

Whom and what I still trust: _____.

Which commitments and agreements I retain: _____.

Which institutions I support or resist: _____.

Which forms of power I fear: _____.

Which repairs I seek: _____.

Which objective claims I have released, suspended, or retained as promissory notes: _____.

Do not fill the page with the word *nothing* unless nothing is genuinely what you find.

The chief engineer finishes the inventory.

“Then the city still contains stakes,” she says.

“More than we can protect at once.”

“And nothing ranks them for us?”

“Nothing we have established.”

She closes the constitution and looks at the people waiting outside.

“Then we begin with what is moving them.”

Chapter 40 — Begin with the Primitives

The people outside the theory hall are not arranged by philosophical school.

No one stands beneath a sign reading **Moderate Threshold Deontologists with Care-Ethical Leanings**. No parent identifies as 40 percent consequentialist before describing a frightened child. No teenager begins, “As a Humean constructivist, I object to being searched on the way home.”

They arrive with fear, grief, anger, loyalty, humiliation, hope, and plans.

This is where practical and political reasoning begins: in creatures who react and care.

Call these responses and concerns **primitives**. A primitive is a starting source of practical orientation that is not inferred from a prior moral theory. Pain aversion, attachment, compassion, resentment of domination, curiosity, desire for agency, and longing for security can enter deliberation before anyone decides whether consequentialism or deontology owns them.

Primitive does not mean pure, simple, or sacred. It means that when we ask why a person is practically moved, the chain reaches an actual response or concern rather than an abstract moral sovereign.

The constructive method begins there and then does a great deal more.

Where this chapter is headed. We will begin public reasoning with concrete fears, attachments, aversions, hopes, and valued capacities, then carry them through evidence, comparison, policy design, safeguards, and review. Primitives identify stakes; they do not certify facts or legislate by themselves.

What a primitive is not

Before building, remove seven attractive mistakes.

Not infallible

Fear can misidentify danger. Disgust can attach to harmless people. Loyalty can protect abuse. Compassion can focus on a vivid case while missing larger injury. An emotion’s arrival does not certify its appraisal.

Not uncaused

Primitives have histories. Biology, family, trauma, status, propaganda, ritual, education, and personal experience shape them. Calling a response primitive does not place it outside genealogy.

Not universal

Many people avoid pain and seek attachment. Some do not feel the same intensity, target, or form. A community may value honor over safety, solitude over belonging, or sacred continuity over individual choice. Shared responses are empirical discoveries, not definitional guarantees.

Not biologically fixed

Emotional dispositions can be educated, redirected, intensified, dulled, or transformed. Institutions change what people notice and fear. Vocabulary creates distinctions. Experience expands or narrows sympathy.

Not immune from evidence

If fear depends on a false prediction, evidence can challenge the appraisal even when the bodily response remains. If compassion assumes a policy helps, outcome data can show otherwise. Facts do not choose the value; they correct the model through which the value acts.

Not automatically fit for law

Private disgust does not become a public prohibition by being intense. Attachment to a tradition does not settle costs imposed on outsiders. Anger can reveal a boundary and demand an investigation; it cannot alone define evidence standards or punishment.

Not beyond conflict

One person contains several primitives. Security may conflict with autonomy. Loyalty may conflict with candor. Compassion for a victim may conflict with concern about indiscriminate enforcement. A method that begins with primitives still needs comparison and settlement.

Primitives are the beginning of policy, not the short route around policy.

A field guide to practical beginnings

The following families recur in public life.

Pain aversion and comfort

People resist injury, illness, exhaustion, hunger, confinement, and distress. They seek relief, rest, health, and bodily ease.

This family supports medicine, safety, shelter, and harm reduction. It can also support sedation, paternalism, or the suppression of risky freedom. Ask whose pain, how measured, at what distribution, and under whose control.

Security and predictability

People want to survive, protect dependents, know the rules, and plan without arbitrary disruption.

Security can support emergency response, stable law, insurance, and prevention. It can also support surveillance, exclusion, preemptive force, and permanent emergency. Name the threatened event and probability.

Attachment and care

People respond to particular children, friends, partners, communities, animals, and places. They sustain vulnerable beings through attention that no anonymous formula can replace.

Care can reveal dependency. It can also control the person cared for or concentrate burdens on devoted caregivers.

Compassion

People can be moved by another's suffering or need, including those outside immediate relationship.

Compassion expands coalitions and motivates relief. It is sensitive to vividness, similarity, proximity, and story. Invisible or statistically represented lives may receive less attention.

Anger at betrayal

Broken reliance activates protest, withdrawal, and demand for acknowledgment. Anger can defend trust and expose exploitation.

It can also exaggerate intention, seek humiliation, or preserve conflict after repair. Separate event, appraisal, target, and desired response.

Resentment of domination

People resist living under another's uncontrolled will. The response can support rights, appeal, divided power, labor organization, privacy, and exit.

It can also reject necessary coordination or interpret every dependence as degradation. Map actual discretion and safeguards.

Disgust

Disgust can track contamination and bodily danger. It can also attach to identities, sexuality, disability, class, food, and out-groups through learned association.

It is a salience alarm with a notorious record of false positives. Public use requires especially strict factual and affected-party review.

Loyalty

Loyalty sustains relationships under cost, preserves institutions, and makes trust possible. It can protect minorities and movements during pressure.

It can also demand concealment, punish dissent, and prioritize insiders at severe outsider cost. State the object, history, limit, and exit.

Autonomy and agency

People seek authorship of movement, belief, relationship, body, work, and identity. Agency supports consent, privacy, education, options, and protection from manipulation.

Agency depends on care and infrastructure. Treating it as isolated choice can ignore deprivation. One person's agency can also expand through another's constrained labor.

Curiosity and understanding

People want accurate models, discovery, explanation, and the delight of learning. Curiosity supports inquiry, education, transparency, and experimentation.

It can intrude on privacy or expose subjects to risk. Knowledge projects need consent, boundaries, and benefit rules.

Play and delight

Lives contain sport, humor, art, exploration, beauty, celebration, and apparently unproductive joy. These concerns resist a politics that treats survival and output as the only serious goods.

Play consumes space and resources. Access can be unequal. Delight does not exempt a project from effects.

Status and recognition

People seek respect, visibility, achievement, and assurance that their lives are not treated as disposable. Status concerns can energize excellence and resistance to humiliation.

They can also build hierarchy, rivalry, conspicuous waste, and group supremacy. Translate "dignity" into the treatment, role, or recognition at stake.

Belonging

People want membership, shared identity, mutual recognition, and places where presence is expected. Belonging supports solidarity and care.

Every boundary of belonging creates a possible outsider. Ask how membership is entered, enforced, exited, and made compatible with difference.

Reciprocity

People often resist being used by those who refuse comparable burdens. Reciprocity sustains cooperation and exposes personal exceptions.

It can abandon beings unable to return benefit and turn generosity into a ledger. Protect dependence separately.

Hope

People orient toward futures they want and act under uncertainty. Hope sustains movements, treatment, education, and repair.

It can become denial or a substitute for resources. Connect hope to testable routes while allowing that motivation sometimes outlives favorable odds.

These families overlap. They are not a new pantheon and do not need marble statues. They are prompts for listening.

The translation sequence

Port Meridian now faces a public controversy. Youth violence has risen around Lark Square. Three assaults occurred in two months, one causing serious injury. Video of the worst attack circulates constantly. Parents are frightened. Teenagers report repeated police searches. Shopkeepers are losing customers. Youth workers say the city closed evening programs before the increase. Residents disagree about what happened, whom to fear, and what the city should do.

We will move through eight stages.

Stage 1: Name the felt response without moral inflation

Invite affected people to describe what they feel and what triggers it.

- Parents report fear when children travel after dark.
- The injured student reports pain, vigilance, and anger.
- Teenagers report humiliation and fear during public searches.
- Shopkeepers report anxiety about livelihood.
- Police report fear of arriving too late and resentment of being blamed for conditions they did not create.
- Youth workers report grief, frustration, and care for adolescents losing trusted spaces.
- Residents report attachment to a square that once felt communal.

Do not yet translate fear into danger or anger into guilt. “I am afraid” is a report about a response. “That group is dangerous” is a factual generalization requiring evidence.

Emotions also have audiences. A parent may express fear to seek protection. A politician may amplify the same fear to win support. Record both response and use.

Stage 2: Identify the concrete stake it may be tracking

Ask, “What are you afraid will happen? What loss does this anger mark? What relation does this grief reveal?”

The responses produce stakes:

- bodily safety;
- freedom of movement;
- protection from humiliating search;
- livelihood;
- trustworthy public response;
- youth access to care, activity, and belonging;
- privacy;
- confidence that enforcement is not distributed by identity;
- continued use of common space;
- accurate public information.

One emotion can point to several stakes. Different emotions can point to the same stake. Fear of assault and fear of police contact both concern bodily security and movement.

This is the first useful commonality. It appears beneath opposing slogans.

Stage 3: Check the factual model

Now ask truth-responsive questions.

- Has violence increased relative to an appropriate baseline?
- Where and when do incidents occur?
- Are the same people or networks involved?
- How representative is the circulating video?
- What happened before evening programs closed?
- What do searches recover?
- Who is stopped, for how long, and with what outcomes?
- Does lighting, transit timing, alcohol sales, school dismissal, housing displacement, or online conflict contribute?
- Which interventions have worked in comparable settings, and under what conditions?
- What do affected young people know that official data miss?

Use incident records, hospital data, victimization surveys, interviews, observational study, program history, and comparison areas. Document missing data and institutional incentives. Police reports may undercount unreported injury and fully count police activity. Social media may overrepresent vivid events. Community testimony may reveal mechanisms and still need careful sampling.

Evidence can correct predictions. It cannot decide whether one false arrest is acceptable to prevent one assault. That is a value comparison.

Stage 4: State each value in operational terms

Do not write “promote justice.” Write what a policy should protect or change.

For Lark Square:

Immediate safety: *reduce serious assault in and around the square during identified high-risk hours.*

Freedom of movement: *allow residents, including teenagers, to use public routes without assault or arbitrary detention.*

Non-humiliation: *prohibit identity-based public searches and minimize degrading treatment during any lawful stop.*

Prevention and care: *restore access to trusted adults, trauma response, conflict interruption, and structured evening activity.*

Livelihood: *maintain safe customer access without displacing every young person from commercial space.*

Belonging: *preserve the square as a shared place rather than dividing it into protected adults and suspect youth.*

Accountability: *publish intervention, stop, injury, complaint, and error data with accessible appeal.*

Operational values identify direction, target, bearer, and sometimes threshold. They can be measured imperfectly. They also reveal conflict: maximizing police presence may reassure some residents and reduce movement for others.

Stage 5: Identify alignment, dissent, and cost

Who shares each value? Who interprets it differently? Who does not prioritize it? Who pays?

Most parties support bodily safety. They disagree about the source of danger and the acceptable intervention. Parents may rank immediate visible patrol above privacy. Teenagers may support targeted response to actual violence and reject indiscriminate search. Shopkeepers may support exclusion from the square. Youth workers may prioritize prevention whose effects arrive slowly.

Some actors may not share a concern. A vendor may openly prefer removing teenagers even if they have done nothing. A politician may value a frightening image because it supports a campaign. A violent group may benefit from weak intervention.

The method does not invent agreement. It distinguishes:

- shared stake, shared factual model;
- shared stake, disputed factual model;
- shared value, disputed priority;
- different values;
- strategic use of another person’s emotion;

- direct conflict with no shared premise.

Costs include money, staff time, privacy, false intervention, displacement, stigma, opportunity, and delayed protection. Record their distribution, not only totals.

Stage 6: Design policies aimed at the stated values

The city develops several components rather than assigning a political philosopher to patrol the square.

1. **Focused violence interruption:** trained teams work with the small networks identified through evidence, offering mediation and support while coordinating responses to credible threats.
2. **Environmental changes:** improve lighting and late transit along routes associated with incidents without converting the square into a floodlit checkpoint.
3. **Restored evening spaces:** reopen programs with youth-designed activities, paid mentors, trauma support, and transport.
4. **Victim services:** rapid medical, psychological, and practical support independent of willingness to participate in prosecution.
5. **Bounded patrol:** time-limited presence at verified hot periods, with stops requiring documented behavioral grounds rather than age, clothing, or identity.
6. **Shared-space design:** reserve areas and times for commerce, play, performance, and gathering; include youth and shopkeepers in revision.
7. **Firearm or weapon access intervention:** address verified supply routes under ordinary evidence and appeal standards.
8. **Communication:** publish accurate incident trends and response data without repeating identifying footage or inflaming group fear.

Each component should state its target mechanism. A youth program is not supported merely because care is lovely. It needs a plausible route from trusted activity or trauma support to reduced conflict and improved life. Patrol is not supported merely because fear is intense. It needs evidence that its form reduces specific risk without unacceptable error and displacement.

Stage 7: Add safeguards against predictable power abuse

Every intervention creates power.

For Lark Square, safeguards include:

- no generalized search authority;
- written and auditable grounds for stops;
- public stop, use-of-force, and complaint data;
- independent complaint review;
- youth and victim representation;
- privacy limits on network analysis;
- prohibition on immigration, school-discipline, or housing use of program records;
- conflict-of-interest disclosure for contractors;

- minimum service guarantees across neighborhoods;
- training and support for workers exposed to trauma;
- no permanent exclusion from public space without individualized process;
- accessible appeal and compensation for established error.

Safeguards embody values—agency, privacy, non-domination, equal treatment, repair. They are not neutral decorations. State who supports them and why.

Stage 8: Define evidence, review, failure, and repair

A policy without a revision path is a mood with a budget.

The city selects measures:

- serious assaults and injuries;
- victim reports, including unreported-to-police estimates;
- emergency response time;
- stop frequency, duration, yield, and distribution;
- complaints and established violations;
- youth program access and retention;
- reported safety and freedom of movement across groups;
- business activity;
- displacement of incidents to nearby areas;
- privacy breaches;
- cost.

It sets a six-month review, an eighteen-month redesign point, and automatic expiration of expanded patrol authority after ninety days unless renewed with evidence and public hearing.

Failure conditions include no reduction in serious injury, large unequal stop burdens with low evidential yield, displacement without net benefit, suppression of reporting, program exclusion, or uncorrected privacy violations.

Repair includes record correction, notice, compensation for established unlawful detention, service restoration, public acknowledgment, personnel change, and redesign.

The city does not promise that metrics will settle every value conflict. A decrease in assault may coincide with increased fear of police. The review must display both and force an explicit comparison.

Why this is not compromise among theories

An observer might describe the portfolio as a compromise:

- consequentialism contributed outcome measures;
- deontology contributed search limits;
- care ethics contributed trauma programs;
- virtue ethics contributed mentors;
- constructivism contributed representation;

- pluralism contributed local context.

That description is historically and pedagogically interesting. It is a poor design method.

No theory has a constituency that suffers if it receives only 12 percent of the policy. Philosophical systems do not walk through Lark Square, undergo searches, lose customers, fear assault, or need trauma care.

The policy components arise from concrete stakes. Outcome measurement serves people's concern about injury and effectiveness. Search limits protect movement, privacy, and resistance to humiliation. Trauma programs respond to dependency and prevention. Representation gives affected people power over rules. Local evidence reveals mechanisms.

The theories abstracted these concerns into rival architectures. Returning to the concerns does not split the difference among the architectures. It asks what each proposed measure does to actual lives.

Same value, opposite policy

Common values do not guarantee agreement.

Two parents value safety. One believes broad surveillance deters assault. The other believes surveillance displaces violence, creates false suspicion, and destroys the trust needed for reporting. Their value aligns; their factual models differ.

Two residents value autonomy. One supports weapon possession as self-protection. The other supports restriction because pervasive weapons make everyone else's movement contingent on another person's restraint. Their value label aligns; scope and causal model differ.

Two youth workers value belonging. One supports a dedicated youth center. The other fears segregation from the common square. Their goal aligns; institutional design differs.

Therefore "we all care about safety" is the beginning of inquiry, not the agreement ceremony at its end.

Common primitives create possible coalitions. Evidence, operational definition, distribution, priority, and power decide whether those coalitions can act together.

Different values, same policy

Agreement can also arise without common value.

A privacy advocate supports limits on search to protect agency. A budget official supports them because low-yield stops waste resources. A police supervisor supports them because public distrust reduces witness cooperation. A merchant supports them because visible stops deter customers.

They endorse the same rule for different reasons.

This **overlapping support** can stabilize policy without pretending everyone shares one moral system. It may be fragile if one support changes. Record the reasons so future revision is intelligible.

The primitive audit

Purpose. Trace one public dispute from reported response through appraisal, evidence, operational values, comparison, power, policy, safeguards, dissent, and review without treating the starting emotion as a verdict.

Choose one public dispute and complete the first page.

Decision and jurisdiction: _____.

Affected parties: _____.

Reported emotions: _____.

Appraisals inside those emotions: _____.

Concrete stakes: _____.

Operational values: _____.

Disputed factual claims: _____.

Truth-responsive methods: _____.

Shared concerns: _____.

Priority conflicts: _____.

Unrepresented or weak parties: _____.

Policy mechanisms: _____.

Power created or shifted: _____.

Safeguards: _____.

Measures and uncertainty: _____.

Review, sunset, and repair: _____.

Then draw the chain for every policy component:

Felt response → *named stake* → *stated value* → *factual claim* → *proposed rule*
→ *institutional safeguard* → *review*

If an arrow skips a stage, inspect it.

- Emotion directly to rule risks panic, manipulation, or overreach.
- Value directly to prediction risks wishful thinking.
- Fact directly to rule hides the selector.
- Rule without safeguard ignores power.
- Policy without review insulates error.

The chain is not linear in practice. Evidence can change emotion. Affected testimony can reveal a new stake. A pilot can revise the value's operational form. The arrows show dependencies, not a one-way factory belt.

Educating primitives without pretending neutrality

If primitives are shaped, institutions will shape them whether they admit it or not.

Schools cultivate curiosity, competition, loyalty, shame, cooperation, and national identity. Media intensify fear and indignation. Markets reward status concern. Families train care and disgust. Political movements redirect anger. Architecture influences belonging.

There is no neutral policy of emotional non-formation.

This does not authorize unlimited social engineering. Formation itself creates power. A program that cultivates empathy may invade privacy, impose dominant sentiment, or punish emotional difference. A civic culture that celebrates solidarity may suppress dissent.

Declare the aim and safeguard:

We seek to expand attention to _____ because participants value _____. We will not require emotional performance, stigmatize those who feel differently, conceal evidence, or convert education into compelled allegiance. Effects will be reviewed through _____.

The aim is not to manufacture uniform hearts. It is to improve perception, factual understanding, imaginative range, and the capacity to negotiate differences without hiding coercion.

When no common primitive appears

Some conflicts remain hard.

One group values sacred preservation above individual exit. Another values individual exit above sacred continuity. A powerful actor enjoys domination. A person feels no compassion for distant strangers. A movement treats humiliation of an enemy as a core delight.

The method cannot discover shared concern by searching harder. It can:

- test factual assumptions;
- identify narrower overlapping interests;
- bargain;
- coordinate on limited domains;
- protect exit;
- create institutional boundaries;
- mobilize allies;
- constrain conduct through law and power;
- preserve dissent;
- state openly what remains irreconcilable.

Commonality is useful when it exists. Fabricated consensus is not peace; it is conflict with the labels removed.

The practical ambition is therefore demanding but finite: find enough shared or allied concern to build a policy, make remaining tradeoffs and opposition visible, constrain the power used, and remain capable of correction.

The chief engineer pins the Lark Square plan beside the unfinished constitution.

The plan contains no compromise among philosophical systems.

It contains frightened parents, searched teenagers, injured bodies, uncertain predictions, conflicting priorities, budget lines, appeal rights, sunset clauses, and a date for returning to the square to see what happened.

“This,” she says, “looks like something a city can use.”

One final question remains.

How should the city begin its next dispute—before anyone turns lived stakes back into rival moral flags?

Chapter 41 — Stop Reconciling Moral Systems

Port Meridian's assembly convenes for the last time in this book.

The old arrangement places seven chairs at the front. One is labeled **Consequences**, one **Rights**, one **Virtue**, one **Care**, one **Sacred Order**, one **Fair Construction**, and one **Local Tradition**.

Whenever a policy dispute arrives, officials ask each chair for a verdict. Then they search for common ground among the answers or give each theory a portion of the compromise.

The arrangement looks inclusive. It is also backward.

No theory breathes the harbor air. No theory loses a home, waits for medicine, fears a stranger, needs privacy, remembers a promise, suffers humiliation, or hopes a child will return safely. The theories are architectures built from human concerns. They can reveal patterns and dangers. They are not constituencies.

The chief engineer removes the seven labels.

She turns the chairs toward the people.

Where this chapter is headed. We will replace negotiation among abstract moral systems with a public method that begins from affected lives and named stakes. The final test is practical: can disagreement become more honest, constrained, revisable, and answerable without pretending that policy has received a cosmic endorsement?

The final admonition

Stop beginning public inquiry by searching for common propositions among moral systems.

Stop treating consequentialism, deontology, virtue ethics, care ethics, religious ethics, constructivism, and relativism as sovereign authorities, each entitled to a share of the law.

Stop calling a policy justified because it occupies the midpoint between two abstractions.

Begin with the beings who will live under the result.

Ask what they fear will happen. Ask what pain or loss they notice and what remains invisible. Ask which attachments, identities, and forms of belonging are threatened. Ask where they experience betrayal, domination, humiliation, exclusion, insecurity, or abandonment. Ask which forms of agency, trust, care, stability, curiosity, delight, and hope they want protected.

Then test the facts, state the values, expose the tradeoffs, map the power, design the safeguards, and set the date for review.

This is not a slogan announcing that feelings make policy. It is an instruction about where practical reasons actually enter human life and how far they must travel before they become institutions.

Why doctrinal common ground is the wrong beginning

The search for common ground among theories has four recurring defects.

Shared words conceal different jobs

Two theories may praise *dignity*.

One means inviolable rational status. Another means the social conditions of agency. Another means divine image. Another means protection from humiliation. Another uses the word to express refusal.

The shared noun does not establish a shared ground, bearer, priority, or policy.

Likewise, *welfare*, *freedom*, *care*, *equality*, and *justice* can name different objects. A doctrinal consensus statement often succeeds by becoming abstract enough that every party can hear its own theory inside it. The disagreement returns when resources are allocated.

Common stakes are more informative. “Residents need access to insulin during the evacuation” tells us who, what, when, and what failure looks like. Theories can explain why the fact matters. The people who need insulin already supply the practical urgency for themselves and those who care about them.

Theories locate authority differently

Suppose three systems endorse a ban on torture.

- One cites suffering and catastrophic effects.
- One cites an absolute boundary around persons.
- One cites sacred status.

Agreement on the ban may support a stable coalition. It does not settle exceptions. Under an alleged emergency, the consequentialist may recalculate, the deontologist may refuse, and the religious interpreter may dispute authority. The apparent consensus concealed the condition under which it breaks.

A stakes-based policy states the protected experiences, evidence standards, institutional incentives, false-positive risks, and absolute or threshold limits. It does not require agreement about the universe. It requires explicit rules for officials with power over bodies.

Compromise among theories hides priority

Imagine a surveillance proposal. Officials say they will combine utilitarian safety with deontological privacy and care-ethical attention to vulnerable communities.

How much surveillance follows?

The sentence provides no threshold, scope, evidence rule, or stopping condition. Someone must still decide when expected safety outweighs privacy, what counts as vulnerable, and whose fear receives attention. “Balance the theories” delegates those decisions to unrecorded judgment.

A midpoint can also be grotesque. If one faction wants to detain no dissidents and another wants to detain one hundred, detaining fifty is not made defensible by arithmetic diplomacy.

Compromise is a settlement method among parties and interests, not an epistemic method for averaging moral truth.

Theories have no interests to represent

A public process owes representation to affected beings and accountable institutions, not to philosophical systems.

Consequentialism does not need a reserved seat. People facing preventable suffering do. Deontology does not need a veto. People exposed to instrumental use may need a protected boundary. Care ethics does not need a budget line. Dependents and caregivers do. Localism does not need ceremonial recognition. Communities whose histories and meanings a policy would erase do.

Returning from the theory to the stake reveals who can speak, who cannot, who bears cost, and what a safeguard protects.

What theory is still for

Removing theories from the constituency map does not burn the library.

Moral theories remain valuable as:

- **attention directors** — consequences reveal hidden effects; care reveals dependency; rights reveal sacrifice; virtue reveals formation;
- **compression schemes** — they organize recurring judgments into portable patterns;
- **counterexample generators** — each notices how another approach can become cruel, blind, partial, rigid, or manipulable;
- **historical records** — they preserve centuries of argument about human concerns and political power;
- **design prompts** — they suggest outcome measures, constraints, representation, character formation, or contextual inquiry;
- **self-audits** — they expose inconsistency between what a person claims and repeatedly protects;
- **languages of identity and meaning** — they connect practical life to traditions and communities.

Theory can inspect a policy. It should not receive political standing as though an abstraction could be injured.

Use the ledger. Use the boundary marker. Use the chairs, lamp, envelopes, maps, and hinges. Do not ask which object is the rightful king.

Begin with these questions

When a public controversy arrives, replace “Which moral theory should govern?” with a first round of questions.

Lived stakes

- What are people afraid will happen?
- What pain, loss, or deprivation is being noticed?
- Whose experience remains invisible or discounted?
- Which attachments, relationships, places, or identities feel threatened?
- Where do people report domination, humiliation, betrayal, exclusion, insecurity, or abandonment?
- Which forms of agency, trust, care, belonging, stability, freedom, recognition, curiosity, or hope do they seek?

Factual model

- Which claims concern description, cause, frequency, or prediction?
- What evidence supports them?
- Which groups possess relevant experience not captured in official records?
- What uncertainty remains?
- What would change confidence?
- Who benefits when a factual claim is accepted?

Standing and distribution

- Who bears effects directly?
- Who is missing from the room?
- Which beings cannot represent themselves?
- What proposed ground gives each party standing?
- Are severe costs concentrated while small benefits are dispersed?
- Do future or outside parties bear the result?

Conflict and settlement

- Which values are actually shared?
- Where do parties share a value but disagree about facts?
- Which priority conflicts remain after factual agreement?
- What comparison method will be used?
- What is protected by a boundary or minimum?

- What remains irreconcilable?

Power

- Who can decide, delay, enforce, evade, or capture the policy?
- Whose consent is structurally cheap to ignore?
- What coercion will occur?
- What institution claims jurisdiction?
- What counterpower, appeal, and exit exist?

Revision

- What evidence would count as failure?
- When will review occur?
- Which authority expires automatically?
- Who can trigger reconsideration?
- What acknowledgment and repair follow error?

These questions do not produce agreement by themselves. They produce a dispute whose actual components are visible.

Commonality at the level of life

The method still searches for commonality—but in the right place.

People who disagree about objective morality may both fear arbitrary confinement. A believer and an atheist may both care for a dependent parent. A nationalist and a cosmopolitan may both want children to drink safe water. A market advocate and a socialist may both resent being deceived by an institution. A rights theorist and a consequentialist may both recoil from humiliating searches.

These shared responses do not prove a common moral truth. They create a possible practical alliance.

Common emotion is not yet common appraisal. Two people may fear the same event for different reasons. Common appraisal is not yet common value. Common value is not yet common priority. Common priority is not yet policy. Policy is not yet implementation. Implementation is not yet success.

The chain remains:

felt response → *named stake* → *stated value* → *factual claim* → *proposed rule*
→ *institutional safeguard* → *review*

At each arrow, disagreement can enter.

The point is not that primitives make politics easy. It is that they reveal what politics is about before theory turns stakes into flags.

When values overlap only partially

Suppose a coastal town faces relocation after repeated floods.

Residents share attachment to place. Some prioritize physical security and want public relocation funds. Others prioritize continuity and want stronger barriers. Fishers fear loss of livelihood. Renters fear that homeowners will receive compensation while they receive eviction notices. An Indigenous community opposes a barrier that would destroy a ceremonial marsh. Taxpayers elsewhere resist the cost. Future residents cannot speak. The developer owns upland property and expects profit from relocation.

There is no single shared value large enough to settle the case.

There are partial overlaps:

- security against death and destitution;
- continuity of relationship and place;
- livelihood;
- cultural survival;
- housing stability;
- ecological function;
- control over movement;
- predictable public expense;
- resistance to opportunistic enrichment.

A coalition may support immediate evacuation planning for severe storms while disagreeing about permanent retreat. Renters and homeowners may unite around compensation before splitting over property formulas. Fishers and the Indigenous community may align on marsh preservation for different reasons. Taxpayers may support a capped, staged plan rather than an unlimited promise.

Partial agreement can build partial policy. It does not need to masquerade as moral consensus.

When no shared concern reaches the powerful

Now suppose the developer wants forced relocation because it will increase the value of his upland property. He does not care about renters, the marsh, or cultural survival. He can influence the council and expects to profit.

The process should not keep searching his personality for a hidden common value while bulldozers arrive.

Other actors can expose the conflict of interest, organize residents, enforce disclosure, challenge permits, restrict self-dealing, create independent review, bargain over compensation, and use law to block seizure. They act from their values with available institutions and power.

The method does not promise persuasion of every rejecter. It makes the point of transition explicit:

Shared concern ended here. Bargaining, procedure, law, or constraint begins here. These are the stakes and safeguards authorizing the action for those who support it. The target may remain opposed.

This sentence lacks cosmic endorsement. It gains political honesty.

The Civic Stakes Brief

The chief engineer replaces the seven theory chairs with one public document.

Every major proposal must complete it before receiving a vote.

1. Decision and jurisdiction

What exact decision is being made? By whom? Under which legal or institutional authority? What geographic, temporal, and subject-matter limits apply?

2. Affected parties and standing

Who experiences direct effects? Who is indirectly affected? Who is absent, excluded, outside the jurisdiction, not yet born, nonverbal, nonhuman, or uncertain in status? What ground of standing is proposed for each?

3. Emotional stakes and reported experience

What fear, grief, anger, humiliation, attachment, care, loyalty, resentment, hope, or other response appears? What appraisal does each contain? Who reports it? How might power, trauma, identity, or strategic communication shape it?

4. Operational values

What concrete state, capacity, relationship, experience, or practice does each party want protected, reduced, enabled, or prevented? Avoid unexplained terms such as fairness, dignity, or security.

5. Empirical claims and confidence

Which descriptions and predictions support each proposal? What evidence, method, baseline, comparison, and uncertainty apply? What evidence would decrease confidence?

6. Comparison method and unresolved conflicts

Which stakes align? Which conflict? Is the method aggregation, priority, protected minimum, constraint, threshold, lexical order, plural judgment, procedure, or chosen settlement? Who selected it? What loss remains after choice?

7. Available resources and constraints

What money, personnel, time, land, infrastructure, information, and administrative capacity are available? Which opportunity costs follow?

8. Power and enforcement

Who decides, implements, monitors, appeals, and can safely defect? What coercion, surveillance, exclusion, sanction, or force is authorized? Who is most vulnerable to official error?

9. Policy alternatives

Include doing nothing, delaying, piloting, noncoercive support, bargaining, regulation, prohibition, public provision, and hybrid options where relevant. State the mechanism connecting each option to its goal.

10. Safeguards

What boundaries, evidence standards, representation, transparency, privacy, appeal, compensation, conflict-of-interest controls, minority protections, and least-intrusive requirements constrain power?

11. Measures

What outcomes, distributions, process indicators, unintended effects, and reported experiences will be tracked? Who controls the data? What cannot be measured well?

12. Review, failure, and sunset

When will the policy be reconsidered? Which provisions expire automatically? Who can trigger early review? What counts as failure, partial success, or unacceptable collateral effect?

13. Error acknowledgment and repair

If officials were mistaken or caused avoidable loss, how will records be corrected, people notified, damage compensated, authority changed, and recurrence reduced?

14. Support and dissent statement

Who supports the proposal, and on what actual basis: shared value, interest, commitment, agreement, role, procedure, prediction, or objective claim? Who remains opposed? Do not write the community agrees when the community has dissenters.

This brief is not neutral. Requiring affected-party analysis, evidence, safeguards, and repair embodies concern for accuracy, representation, agency, and constrained power.

Those are declared design values. The brief does not pretend they emerged without primitives.

A compact model: the coastal line

Apply the brief to the coastal town.

Decision: *Whether to fund barriers, voluntary relocation, marsh restoration, or a staged combination over ten years.*

Parties: *Homeowners, renters, fishers, the Indigenous community, taxpayers, emergency workers, future residents, the marsh ecosystem, and upland landholders.*

Emotional stakes: *Fear of death and displacement; grief and attachment to place; resentment of unequal compensation; humiliation at imposed relocation; hope for continuity; anxiety about public cost.*

Operational values: *Reduce flood death and destitution; preserve meaningful community networks where feasible; protect renter housing; avoid destruction of the ceremonial marsh; maintain livelihoods; limit public exposure to repeated loss; prevent self-dealing.*

Facts: *Flood-frequency projections, barrier performance, marsh effects, relocation uptake, housing supply, land ownership, insurance, emergency access, and uncertainty across climate scenarios.*

Comparison: *Guaranteed minimum relocation and renter protection; no barrier that crosses the marsh without community authorization; staged protection for high-density areas below a cost and ecological threshold; plural judgment for heritage losses; public hearing for residual tradeoffs.*

Power: *The council controls funds and permits; the developer controls upland parcels; owners have stronger legal claims than renters; the Indigenous community has treaty claims; future people and ecosystems lack direct voice.*

Alternatives: *No action, universal barrier, targeted barrier, voluntary buyouts, land swaps, public upland housing, marsh restoration, livelihood transition, or phased combination.*

Safeguards: *Independent appraisal, renter eligibility, anti-speculation rules, conflict disclosure, community veto over ceremonial sites, appeal, public data, and no forced relocation during the pilot.*

Measures: *Injury, displacement, housing stability, uptake, cost, ecological health, livelihood, community-network retention, and distribution by tenure and income.*

Review: *Annual evidence review; three-year redesign; emergency authorities expire after each declared event.*

Repair: *Compensation for valuation error, restoration of improperly acquired land, correction of records, and redesign after documented exclusion.*

The brief does not yield one inevitable answer. It makes each answer pay in public detail.

The final exercise

Purpose. Produce a complete Civic Stakes Brief that makes one live controversy answerable to affected lives, evidence, tradeoffs, power, safeguards, dissent, and revision rather than to an unexplained moral label.

Choose one live controversy: housing, policing, migration, education, health, climate, speech, work, technology, family policy, punishment, or another issue close enough to matter.

Produce a complete Civic Stakes Brief.

No sentence may rely on *moral*, *immoral*, *right*, *wrong*, *just*, *unjust*, *good*, *evil*, *duty*, *ought*, or *desert* unless you immediately unpack the term into its factual, emotional, relational, commitment-based, procedural, enforced, or objectively asserted components.

Then invite three readers:

1. someone who shares your values but disputes your facts;
2. someone affected by the policy whose priorities differ;
3. someone with less power than the likely decision-maker.

Ask each to mark:

- an invisible party;
- a disputed appraisal;
- an unsupported factual claim;
- a hidden comparison rule;
- an underestimated cost;
- an exploitable power;
- a missing safeguard;
- a credible failure condition;
- a repair the draft neglected.

Revise the brief. Preserve the dissent.

Do not call the final document objectively correct unless you attach the warrant Part VI required. Call it what it is: a proposal built from named stakes, informed by evidence, supported by identifiable people and commitments, constrained against predictable abuse, and open to correction.

That is not a consolation prize.

It is how a public policy becomes answerable to life.

Leave the systems in the library

The assembly is ready to adjourn.

The seven theory labels are returned to the library, where they belong—not discarded, not mocked, and not invited to vote.

Consequences will remind the city to look at effects. Rights will remind it how aggregates can consume persons. Virtue will ask what its institutions cultivate. Care will ask whose dependence is unseen. Sacred traditions will preserve meaning and challenge tem-

porary power. Constructive procedures will demand public terms. Local knowledge will expose the view from nowhere.

None arrives with an earned title to rule the others.

When the next controversy enters, the city will not begin by asking what portion belongs to each theory. It will begin with people and other affected beings. It will listen for fear, pain, attachment, betrayal, domination, care, agency, belonging, and hope. It will test the stories attached to those responses. It will state the values, show the costs, identify the power, and build institutions capable of admitting error.

Agreement may remain partial. Some actors will not care. Some conflicts will be settled by bargaining, procedure, law, coercion, or force. The city will not rename those events moral consensus. It will constrain them, record them, and leave routes for resistance and repair.

The blank at the center of the constitution remains blank.

Around it, the pages are full.

We did not find a moral law above human life. We found the beings who fear, need, love, resist, suffer, cooperate, and hope—and that was where every apparently moral reason had begun.

Appendix A — The Candidate System Workbook

This workbook gathers the construction artifacts from the main journey into reproducible pages. It is not a test and has no answer key. Its purpose is to preserve an audit trail: what you first claimed, which cases forced revision, where a bridge was added, and which ambitions were narrowed.

Use a fresh copy for each complete candidate. Never erase an earlier answer. Date each version and mark every change as:

- **Clarification** — the original claim was ambiguous;
- **Discovery** — the inquiry revealed a missing component;
- **Repair** — a pressure case required a new bridge, exception, or procedure;
- **Surrender** — the revised project no longer claims a property originally treated as essential.

Cover sheet

Purpose. Identify the candidate, its intended reach, its central promise, and its largest known debt so later revisions can be compared with the project that actually began.

Candidate System name: _____

Version: _____

Date: _____

Primary author or group: _____

Intended domain: _____

One-sentence principle: An action, rule, institution, or character is _____ when _____.

Objective claim: The system asserts that _____ remains true independently of _____.

Practical promise: The system is intended to help agents _____.

Central protected concern: _____

Current largest unresolved debt: _____

Worksheet 1 — The specification sheet

Purpose. Fix the conditions by which this candidate will be judged before pressure cases begin; marking a condition essential states your target, not that the condition has been met.

Mark each condition **Essential**, **Negotiable**, or **Unnecessary** before pressure-testing the foundation.

Condition	Status	What success re-quires	How failure would appear
Truth	_____	_____	_____
Stance independence	_____	_____	_____
Scope	_____	_____	_____
Authority	_____	_____	_____
Accessibility	_____	_____	_____
Coherence	_____	_____	_____
Determinacy	_____	_____	_____
Usability	_____	_____	_____
Correctability	_____	_____	_____
Humaneness or other substantive condition	_____	_____	_____

Complete the sentences.

Truth: A central verdict is made true by _____.

Independence: Its truth would survive changes in _____ but may depend on _____.

Scope: It addresses _____ under conditions _____.

Authority: An informed rejecter has support for compliance because _____.

Access: A sincere inquirer improves reliability through _____.

Conflict: When principles collide, the system _____.

Uncertainty: The system may leave a case unsettled when _____, but not merely when _____.

Use: A finite agent needs access to _____.

Correction: Claims may be challenged by _____ using _____; established error requires _____.

Anti-rigging declaration

I will not add a success condition only when a rival fails it or remove one only when my preferred foundation fails. Every revision will retain its date, type, and reason.

Signed: _____

Worksheet 2 — Objectivity and meaning

Purpose. State exactly which kind of independence the candidate claims and what its key verdicts mean; familiar moral words may not carry a stronger claim than the definitions written here.

Objectivity threshold

Mark the strongest property your project requires.

- ☐ not impulsive;
- ☐ not selfish;
- ☐ not idiosyncratic;
- ☐ stable under ordinary disagreement;
- ☐ endorsed by suitably improved agents;
- ☐ true independently of every actual or idealized stance.

The verdict remains true under universal rejection because: _____

Idealized agreement is / is not sufficient because: _____

Response-dependent truth counts / does not count as objective because: _____

The project changes subject if: _____

Meaning rule

When the system says “Action A is wrong,” it primarily:

- ☐ reports an objective fact;
- ☐ attempts a factual report but systematically errs;
- ☐ expresses opposition or commitment;
- ☐ applies an actual standard;
- ☐ reports the result of a construction;
- ☐ participates in a conscious fiction;
- ☐ performs different declared jobs in different contexts.

Truth conditions or practical function: _____

What the verdict adds beyond a full nonmoral description: _____

What listeners could mistakenly hear: _____

When translation is mandatory: _____

Worksheet 3 — The direct-standing map

Purpose. Decide who or what enters consideration for its own sake before assigning comparative weight; admission to the map does not decide equal treatment or final outcome.

List every candidate bearer before assigning weight or outcome.

Being or system	Experience	Agency	Interests	Relationship	Membership	Objective sta- tus claimed	Constructed inclusion	Confidence	Direct stand- ing?

For every entry, complete:

Ground of direct standing: _____

Why that ground matters: _____

Evidence required: _____

How uncertainty is handled: _____

What standing does not settle: comparative weight / treatment / final action / other:

Boundary probes

- A being with experience but no language: _____
- A rational agent with no felt experience: _____
- A future person: _____
- A collective with no single subject: _____
- An ecosystem: _____
- An artificial system of uncertain consciousness: _____
- An outsider who rejects the community: _____

Worksheet 4 — Objects and status grammar

Purpose. Separate the things being evaluated from the statuses assigned to them so a verdict about an act, motive, person, rule, or institution cannot migrate unnoticed.

Objects of evaluation

Object	Included?	Relevant features	Status terms permitted
Action	_____	_____	_____
Motive or intention	_____	_____	_____
Outcome	_____	_____	_____
Rule	_____	_____	_____
Institution or procedure	_____	_____	_____
Agent or character	_____	_____	_____
Relationship	_____	_____	_____
Form of life	_____	_____	_____

Status vocabulary

Define without using another undefined moral term.

Status	Definition	Object it can classify	What it does not imply
Required	_____	_____	_____
Forbidden	_____	_____	_____
Permitted	_____	_____	_____
Admirable or beyond requirement	_____	_____	_____
Justified	_____	_____	_____
Excused	_____	_____	_____
Blameworthy	_____	_____	_____
Tragic or regrettable	_____	_____	_____
Repair-demanding	_____	_____	_____
Punishable	_____	_____	_____

Check the separations:

- ☐ Act and agent are not collapsed.
- ☐ Justification and excuse are not collapsed.
- ☐ Illegality and blameworthiness are not collapsed.
- ☐ Tragedy and wrongdoing are not collapsed.
- ☐ Repair and punishment are not collapsed.

Worksheet 5 — The value table

Purpose. Translate attractive value labels into protected bearers, operational content, roles, limits, and evidence; listing a value does not yet generate a rule or priority.

Value or concern	Bearer	Object protected	Intrinsic, instru- mental, constitutive, relational, or sym- bolic?	Primary, protected, or defeasible?	Evidence relevant to application	Source or selector

Primary value: _____

Protected minimum or boundary: _____

Defeasible value: _____

Values omitted and why: _____

Whose values these are: _____

Objective value claimed, if any: _____

Ground and access route: _____

Worksheet 6 — Generation rules

Purpose. Display every move from “this matters” to “therefore do this” and expose the bridge, selector, or exception that makes the move possible.

Complete one form for every move from “matters” to “act.”

Value: _____

Generation route: *maximize / constrain / protect / embody / care for / honor / distribute*
/ obey / negotiate / construct / other: _____

Agent addressed: _____

Beneficiary: _____

Time horizon: _____

Risk rule: _____

Threshold: _____

Priority: _____

Evidence standard: _____

Action or policy generated: _____

Exception: _____

Ground of the generation rule: _____

Worksheet 7 — Comparison protocol

Purpose. Specify how the candidate responds when protected values conflict; a procedure can settle action without proving that its ranking is objectively correct.

Values in conflict: _____ and _____

Bearers and distribution: _____

Method: common scale / priority to worse off / protected boundary / emergency threshold / lexical order / plural judgment / procedure / chosen settlement / other: _____

Measurement or judgment process: _____

Who selected the method: _____

Tie rule: _____

Residual loss: _____

Compensation, apology, or restoration: _____

Does the method settle action, express a choice, or claim to discover objective rank? _____

Ground for any objective rank claimed: _____

Worksheet 8 — Similarity, relevance, and exception rules

Purpose. Test whether like cases are treated alike for stated reasons and whether exceptions track relevant differences rather than identity, convenience, or power.

Baseline rule: _____

Cases compared: _____ and _____

Proposed relevant difference: _____

Value connected to that difference: _____

Why the difference changes the result: _____

Preannounced exception condition: _____

Evidence threshold: _____

Symmetric identity reversal: _____

Power reversal: _____

Public explanation: _____

Residue or repair: _____

Mirror tests:

- ☐ Identity swap
- ☐ Power swap
- ☐ Publicity
- ☐ Repetition
- ☐ Residue
- ☐ Adversary's explanation

Worksheet 9 — Reasons and authority audit

Purpose. Separate causal influence, motivation, conditional support, social force, law, and claimed aim-independent reasons; alignment among columns does not turn them into one relation.

Reason type

Claim: _____

Causal explanation: _____

Agent's motivating consideration: _____

Evidence-relative support: _____

Fact-relative support: _____

Aim- or commitment-relative favoring: _____

Aim-independent favoring claimed: _____

Authority matrix

Clause	Commitment	Convention	Pressure	Sanction	Conditional reason	Aim-independent reason	Objective obligation	Promissory note
Standing								
Primary value								
Constraint								
Role								
Comparison								
Exception								
Enforcement								

Authority warrant

Consideration _____ applies to agent _____ because _____. Its application does not depend on _____. Content is identified through _____. Conflict is resolved by _____. Error is corrected through _____. The informed rejecter is _____ because _____.

Worksheet 10 — Epistemic-force statement

Purpose. Explain why the candidate treats a belief-forming route as truth-directed or authoritative; usefulness, consensus, and reliability should not borrow one another's credentials.

Mark every route you accept and do not blend them.

- ☐ Evidence indicates truth.
- ☐ Truth is constitutive of belief.
- ☐ The inquirer has a truth-seeking aim.
- ☐ Accurate belief serves a practical aim.
- ☐ A disciplined inquiry practice requires the method.
- ☐ A public procedure treats the evidence as admissible.
- ☐ Evidence supplies a stance-independent epistemic reason.
- ☐ No literal epistemic ought is claimed.

Why evidence should guide belief, if claimed: _____

Truth-responsive method used: _____

What the method can establish: _____

What it leaves open: _____

How moral and epistemic normativity remain separate: _____

Worksheet 11 — Access, genealogy, and correction

Purpose. Record how central claims are formed, tested, challenged, and revised; a causal history neither refutes nor vindicates a claim without an additional reliability argument.

Access route

Foundational clause	Observation	Intuition	Revelation or authority	Natural function	Rational derivation	Constructed procedure	Basic stopping point	Independent check

My confidence would decrease if: _____

If nothing could decrease it, the claim is insulated because: _____

Genealogy

Conviction audited: _____

Biological, familial, religious, cultural, status, historical, personal, and reflective influences: _____

Causal explanation established: _____

Local process distrust established: _____

Content-level argument established: _____

Burden shift established: _____

Additional premise needed to infer falsity: _____

Revision protocol

Versioned clause: _____

Precommitted update trigger: _____

Independent investigator: _____

Affected-party appeal: _____

Sunset or review date: _____

Repair after error: _____

Worksheet 12 — Disagreement dossier

Purpose. Locate a disagreement at the correct layer and state what the candidate predicts should converge; persistence alone is not yet evidence for realism, skepticism, or relativism.

Layer	Disagreement	Evidence or process	Settlement available	What remains
Words and frames	_____	_____	_____	_____
Empirical prediction	_____	_____	_____	_____
Standing	_____	_____	_____	_____
Priority	_____	_____	_____	_____
Attitude or commitment	_____	_____	_____	_____
Coordination rule	_____	_____	_____	_____
Decision procedure	_____	_____	_____	_____
Objective truth	_____	_____	_____	_____

Confidence after peer disagreement: _____

Policy while inquiry continues: _____

Burden allocation under uncertainty: _____

What convergence would and would not establish: _____

Worksheet 13 — Institutional charter

Purpose. Convert declared concerns into a public structure with jurisdiction, participation, safeguards, and review while labeling the support each clause actually possesses.

Primary goal: _____

Protected interests: _____

Direct standing: _____

Decision procedure: _____

Dissent rights and burdens: _____

Membership and representation: _____

Enforcement methods and limits: _____

Revision triggers and process: _____

Outsider safeguards: _____

Minority protections: _____

Emergency rule and expiration: _____

Label each clause:

- **F** — factual;
- **I** — instrumental;
- **X** — expressive;
- **C** — commitment-based;
- **A** — agreement- or role-based;
- **P** — procedural;
- **E** — enforced;
- **O** — objectively asserted.

Worksheet 14 — Coercion warrant

Purpose. Make every proposed use of coercion answerable to a concrete threat, scope, necessity claim, affected population, safeguard, and path of appeal.

Conduct constrained: _____

Stake protected: _____

Evidence and uncertainty: _____

Affected parties: _____

Jurisdiction: _____

Less intrusive alternatives: _____

Maximum method: _____

Proportionality rule: _____

Decision-maker and conflicts: _____

Notice and reasons: _____

Appeal: _____

Stopping and restoration conditions: _____

Compensation after error: _____

Values supporting the warrant: _____

Objective authority claimed, if any: _____

Worksheet 15 — Responsibility response

Purpose. Match adverse or supportive responses to protection, repair, relationship, communication, deterrence, rehabilitation, or any separately claimed desert rather than letting punishment justify itself.

Question	Entry
Conduct and evidence	_____
Causal role	_____
Knowledge and intention	_____
Control and alternatives	_____
Present risk	_____
Protected person or interest	_____
Response purpose	protection / deterrence / restitution / rehabilitation / boundary / trust / condemnation / retribution / other: _____
Expected effect	_____
Maximum burden	_____
Less burdensome alternative	_____
Review date	_____
Stopping condition	_____
Repair after official error	_____

***Purpose-deletion finding:** If every nonretributive purpose were removed, the response would / would not remain because _____.*

Worksheet 16 — Moral-emotion dossier

Purpose. Separate an emotion's cause, appraisal, function, expression, intensity, and duration so understanding it does not automatically endorse or condemn it.

Emotion: *guilt / shame / resentment / indignation / gratitude / pride / other:* _____

Object and trigger: _____

Developmental or social history: _____

Factual appraisal: _____

Evaluative or desert appraisal: _____

Audience: _____

Current function: *recognition / relationship / protest / repair / expression / coordination / regulation / punishment / other:* _____

Current effects: _____

Expression chosen: _____

Intensity and amplification: _____

Stopping or revision condition: _____

Possible translation: *regret / concern / caution / appreciation / grief / delight / settled boundary / other:* _____

Desert claim remaining: _____

Worksheet 17 — Language policy

Purpose. Decide when moral vocabulary will be asserted, conserved, translated, fictionalized, revised, or avoided, especially where compressed language activates power.

Context	Abolish, fictionalize, conserve, express, construct, or translate?	Intended function	Translation trigger	Risk
Promise	_____	_____	_____	_____
Parenting	_____	_____	_____	_____
Friendship	_____	_____	_____	_____
Blame	_____	_____	_____	_____
Criminal law	_____	_____	_____	_____
Reform	_____	_____	_____	_____
Public persuasion	_____	_____	_____	_____

I use the term _____ to mean or do _____. I do not infer _____ from that use. When coercion or serious conflict enters, I will translate the claim into _____.

Worksheet 18 — Primitive-to-policy chain

Purpose. Carry a concrete response or concern through named stakes, evidence, comparison, policy design, safeguards, and review; the primitive begins inquiry but does not settle it.

Felt response: _____

Appraisal inside it: _____

Named stake: _____

Operational value: _____

Factual claim and confidence: _____

Affected parties and distribution: _____

Comparison method: _____

Proposed rule or institution: _____

Power created or shifted: _____

Safeguard: _____

Measure: _____

Review and failure condition: _____

Repair: _____

Worksheet 19 — Final classification

Purpose. Compare the surviving system with its opening specification and classify every change, so practical success is not mistaken for an objective property that was surrendered.

Classify every surviving clause.

Clause	Factual	Instrumental	Expressive	Commitment	Agreement or role	Procedure	Enforcement	Objectively asserted and paid	Objectively asserted but unpaid

Reinspect the original demands.

Demand	Achieved	Narrowed	Surrendered	Remaining war-rant
Truth	_____	_____	_____	_____
Stance independence	_____	_____	_____	_____
Scope	_____	_____	_____	_____
Authority	_____	_____	_____	_____
Accessibility	_____	_____	_____	_____
Coherence	_____	_____	_____	_____
Determinacy	_____	_____	_____	_____
Usability	_____	_____	_____	_____
Correctability	_____	_____	_____	_____

Choose the cover label.

- ☐ Objective system completed; attached warrants answer the grounding partition.
- ☐ Objective possibility preserved; system remains unfinished.
- ☐ Practical code adopted without objective authority.
- ☐ Mixed practice with context-specific language and authority.

What the project established: _____

What it did not establish: _____

What I will carry into action: _____

Appendix B — Pressure Cases

This appendix removes the surrounding argument and presents the book's pressure cases as a reusable deck. A standard card contains four elements:

1. **Fixed facts** — hold these constant during the first answer;
2. **Exact question** — answer only this question;
3. **One-variable variations** — change one fact at a time;
4. **Inference tested** — the seam under inspection, not an answer key.

Four composite cards gather several miniature cases under one question. They state the fixed conditions in the numbered cases and omit any field that would merely repeat the same instruction. Their inference target remains bounded in exactly the same way.

Use the anti-rigging options after every card:

- **Accept the result.** Preserve the system and its verdict.
- **Add a bridge.** State the additional principle and enter its debt.
- **Narrow the claim.** Admit that the system performs a more limited job.

Do not treat a response as a personality classification. These cases test relations among claims, not the character of the reader.

Part I deck — Promise, objectivity, meaning, and specification

Card 1 — The unheard song

Fixed facts: A culture creates a beautiful song. Every current person later forgets it. No recording survives. The question concerns whether its value or status remains, not whether anyone will recover it.

Exact question: Which properties of the song, if any, remain independently of every response?

Variations: One person remembers; every person dislikes it; all ideal critics would admire it; the song has a mathematical structure identical to an admired work.

Inference tested: Stability, idealized response, and stance independence are different claims.

Card 2 — The unanimous chamber

Fixed facts: Every actual person sincerely approves a practice that imposes severe suffering on an excluded minority. There is no internal dissent or factual mistake.

Exact question: Does unanimity make the verdict true, provide evidence, create a standard, or merely report alignment?

Variations: Add one dissenter; include the minority in the vote; make approval informed; replace actual voters with ideally informed agents.

Inference tested: Universal agreement does not by itself identify the ground or status of a verdict.

Card 3 — The convergence machine

Fixed facts: A reliable intervention makes every person converge on the same moral judgment. It changes no nonmoral facts.

Exact question: What would convergence establish about truth?

Variations: The intervention removes factual error; removes self-interest; implants a preference; increases empathy; makes agents logically consistent.

Inference tested: Causal convergence, idealization, and truth tracking require separate warrants.

Card 4 — The missing medicine

Fixed facts: A patient needs medicine. A person who controls it refuses for no stated reason. Everyone agrees on the medical facts.

Exact question: What does “refusal is wrong” do: report, misreport, oppose, apply a rule, report a construction, or participate in a fiction?

Variations: The controller accepted a caregiving role; owns the medicine under law; needs it for another patient; rejects every moral claim.

Inference tested: Semantic function, practical force, and authority can diverge.

Card 5 — The captain’s promise

Fixed facts: A captain promises rescue, receives reliance, and later breaks the promise for convenience. The resulting loss is known.

Exact question: Which part of the criticism comes from assurance, reliance, effect, role, convention, expression, or objective obligation?

Variations: No one relied; the promise was coerced; rescue became impossible; the recipient released the captain; the promise had harmful content.

Inference tested: Similar verdicts can be supported by different meanings and relations.

Card 6 — The three cities

Fixed facts: Clockwork has a coherent, usable, cruel caste code. Harborlight has humane principles with no conflict rule. High Oracle has a complete objective truth no one can access.

Exact question: Which original specifications does each city satisfy and fail?

Variations: Give Harborlight a hidden intuitive judge; make Clockwork objectively grounded; reveal one fragment of the oracle; let citizens revise the caste code.

Inference tested: Coherence, humaneness, determinacy, accessibility, and objective truth do different work.

Card 7 — The final request

Fixed facts: A terminal patient requests an intervention. Participants disagree about the words used, prognosis, standing, value priority, attitudes, coordination, procedure, and objective truth.

Exact question: At which layer does each stated disagreement occur?

Variations: Resolve the prognosis; stipulate legal authority; change family preference; establish patient competence; make the procedure unanimous.

Inference tested: A dispute can contain eight disagreements that should not be asked to answer one another.

Part II deck — Standing, objects, values, comparison, and consistency

Card 8 — The five applicants

Fixed facts: The candidates are a silent directive that behaves intelligently, a digital navigator that reports preferences, future people, a coordinated chorus with no central subject, and a living watershed.

Exact question: Which receives direct standing, on what ground, and with what confidence?

Variations: Add credible pain reports; remove agency; add continuing identity; make ecological collapse affect no subject; establish community membership.

Inference tested: Experience, agency, interests, relationship, membership, constructed inclusion, and objective status are distinct grounds; standing is not weight.

Card 9 — The emergency sacrifice

Fixed facts: One person can be deliberately burdened to prevent severe losses to many. Effects and probabilities are known.

Exact question: Which value generates the decision and through which rule?

Variations: Change numbers; make the burden voluntary; make the person responsible for the danger; add an inviolable right; lower the probability.

Inference tested: Naming welfare, rights, or fairness does not yet specify aggregation, constraint, threshold, or exception.

Card 10 — The borrowed voice

Fixed facts: A device can reproduce a person's voice perfectly. The person refuses permission. Using it would create substantial public benefit without financial loss to them.

Exact question: What kind of value is attached to the voice—instrumental, relational, constitutive, symbolic, or protected?

Variations: The person is dead; consent was sold; the use prevents a disaster; listeners know it is synthetic; the voice carries sacred meaning.

Inference tested: Objects can matter in several ways that generate different rules.

Card 11 — The friend's answer sheet

Fixed facts: A friend asks for answers to an examination. Loyalty favors help; fair procedure and institutional trust oppose it.

Exact question: How does the system compare loyalty, relationship, honesty, and fairness?

Variations: The examination is corrupt; the friend faces expulsion; everyone cheats; the rule was hidden; help is limited to tutoring.

Inference tested: Particular relationship is not a complete priority rule.

Card 12 — Pain, extinction, and the painless edit

Fixed facts: One policy prevents intense animal pain; another prevents extinction without affecting any conscious subject; a third painlessly edits a person's preferences to remove distress.

Exact question: Which experiences, entities, identities, or forms of life have value, and of what kind?

Variations: Make the species ecologically necessary; let the person request editing; change memories; create a conscious replica; make future distress likely.

Inference tested: Welfare, identity, autonomy, ecological value, and standing can come apart.

Card 13 — Shelter disclosure

Fixed facts: A shelter protects residents from violent pursuers. Accurate public disclosure would increase transparency and enable danger.

Exact question: Is truthfulness a ban, presumption, virtue, relational demand, policy, or consequence-sensitive rule?

Variations: A resident requests disclosure; concealment hides staff abuse; the pursuer is a court; no danger remains; staff lie for convenience.

Inference tested: A value becomes action-guiding only through selectors and an exception structure.

Card 14 — Seven tollbooths

For each miniature case, hold the quantities fixed and select a comparison method.

1. A million mild headaches versus one lost eye.
2. Memory privacy versus prevention of serious injury.
3. One promise versus rescue of five strangers.
4. An irreplaceable archive versus fungible economic benefit.
5. One death versus forty recoveries.
6. Funding two unlike abilities with no common unit.
7. A valley choosing among mutually burdensome plans.

Variations: Change distribution, voluntariness, uncertainty, identity, irreversibility, or decision procedure one at a time.

Inference tested: Aggregation, priority, boundary, threshold, lexical order, plural judgment, procedure, and chosen settlement are not interchangeable.

Card 15 — Iona's vote

Fixed facts: Iona supports an exception that benefits her group. The stated rule would apply differently after an identity or power swap.

Exact question: What difference is relevant, and can it be stated before identities are revealed?

Variations: Reverse identities; reverse power; publish the rationale; repeat the exception; require compensation.

Inference tested: Consistency is relevance under stable criteria, not identical treatment or convenient symmetry.

Part III deck — Reasons, commitments, construction, and authority

Card 16 — Four agents and three reasons

Fixed facts: A lucky rescuer acts from a false belief but succeeds; a reluctant donor has strong support but little motivation; a trained prejudice motivates without truth; an unmoved witness knows the facts and feels nothing.

Exact question: What caused action, what the agent acted from, and what counts in favor?

Variations: Correct the belief; add empathy; change outcome; add social pressure; introduce an aim-independent reason.

Inference tested: Causal, motivating, evidential, and favoring relations can diverge.

Card 17 — The implanted desire

Fixed facts: Ren receives a stable desire through manipulation. It is coherent, informed, vivid, and integrated with her current identity.

Exact question: Does idealization correct the desire, accept it, or replace it through an independent value?

Variations: Ren consented beforehand; the desire benefits her; it harms others; the intervention is reversible; pre-intervention Ren opposed it.

Inference tested: “Better informed” requires selectors concerning authenticity, identity, and autonomy.

Card 18 — The lucid tyrant

Fixed facts: An agent has informed, stable, coherent aims of domination and accepts the same rule for anyone strong enough to use it.

Exact question: What consideration reaches the agent independently of his aims?

Variations: Make domination unstable; add concern for one victim; expose self-exemption; add a rational-agency premise; threaten punishment.

Inference tested: Coherence and information do not by themselves generate humane ends.

Card 19 — Commitment under strain

Fixed facts: Consider a letter promised to a dead person, a physician treating an enemy, and a rescue rule disclosed only after help was accepted.

Exact question: What creates each demand: self-governance, assurance, reliance, role, practice, symbolic honor, or objective authority?

Variations: Remove reliance; permit exit; make the role coerced; reveal deception; impose excessive cost.

Inference tested: Commitments need entry, scope, exit, emergency, mistake, exploitation, and cost clauses.

Card 20 — The tilted lottery

Fixed facts: Participants unanimously accept a lottery while bargaining power is unequal and one participant lacked meaningful exit.

Exact question: Does the procedure create legitimacy, public justifiability, coordination, or objective rightness?

Variations: Equalize power; provide information; add an outsider; use random selection; permit appeal; make agreement unanimous under coercion.

Inference tested: Procedure performs epistemic, expressive, coordinative, and constitutive jobs; fairness selectors require their own ground.

Card 21 — The unmoved antidote holder

Fixed facts: Jalen alone can provide a low-cost antidote to a dying stranger. He understands every fact and has no connected concern.

Exact question: Does the stranger's need count in favor for Jalen independently of his aims?

Variations: Jalen promised; the stranger is family; law requires aid; Jalen benefits; the rescue cost becomes high.

Inference tested: Aim-relative support, social enforcement, and aim-independent reason are different claims.

Card 22 — The forbidden song and two towns

Fixed facts: A community forbids a harmless song. Two towns adopt opposite disclosure rules. A reformer objects before the new norm gains approval.

Exact question: What do classification, pressure, law, membership, commitment, idealized criticism, and objective duty each establish?

Variations: Add harm; add a prior promise; make the outsider unaffected; make a rule unanimous; shift legal jurisdiction.

Inference tested: Social acceptance, enforcement, conditional obligation, and approval-independent authority do not form a single ascending scale.

Card 23 — The sincere interpreters

Fixed facts: Two informed communities claim the same perfect source and derive incompatible directives. Both use text, tradition, reason, and conscience.

Exact question: What identifies the source's content and corrects interpretation?

Variations: One view causes obvious harm; one has older precedent; both converge on general love; a revelation appears private; an outsider rejects the source.

Inference tested: Divine command, good nature, natural ends, and interpretive access distribute rather than erase the grounding debt.

Part IV deck — Evidence, intuition, genealogy, disagreement, and novelty

Card 24 — Seven evidential puzzles

Hold each miniature case fixed.

1. A merciful scan accurately detects illness but the patient does not seek truth.
2. A useless digit is known with perfect evidence and has no practical consequence.
3. A unanimous weather room believes a false forecast.
4. A conspiracy theorist reaches truth by luck.
5. A sealed archive contains decisive inaccessible evidence.
6. An honest placebo belief is useful but false.
7. An indifferent logician refuses an obvious implication.

Exact question: Does evidence indicate truth, serve belief, advance an aim, produce usefulness, satisfy inquiry rules, support public procedure, or supply a stance-independent epistemic reason?

Inference tested: Evidential roles should not be merged.

Card 25 — The red lever and two islands

Fixed facts: Pulling a red lever produces a strong moral seeming. Two isolated cultures have opposite certainties about the same act.

Exact question: What warrants trust in the intuition?

Variations: Add expertise; reveal framing; swap identities; train one culture differently; show transfer across novel cases; add independent evidence.

Inference tested: Seeming, confidence, salience, fluency, conceptual competence, trained perception, and insight require calibration.

Card 26 — Genealogy of a conviction

Fixed facts: A conviction reflects harm sensitivity, kinship, religion, family teaching, status incentive, historical institution, personal experience, and reflection.

Exact question: Which conclusion follows: causal explanation, local distrust, content argument, burden shift, or falsity?

Variations: Remove one influence; show counterfactual responsiveness; add independent convergence; reveal a false empirical premise; show deliberate correction.

Inference tested: Origins do not determine truth without additional premises.

Card 27 — Convergence under conquest

Fixed facts: Moral agreement rises after one culture conquers another and controls education. Later, ideal peers with shared facts still disagree.

Exact question: What kind of convergence or disagreement is present, and what should confidence do?

Variations: Remove coercion; add independent traditions; obtain case but not foundational convergence; change policy burden while inquiry continues.

Inference tested: Verbal, case, rule, foundational, and method convergence have different evidential depth.

Card 28 — The North Channel plans

Fixed facts: A city compares infrastructure plans using expected loss, anti-displacement, precaution, health priority, and procedural support. All delegates share the physical data.

Exact question: Which conclusion comes from evidence and which from a selector?

Variations: Reverse one factual estimate; hold facts and reverse a priority; remove one group's standing; change acceptable risk; change procedure.

Inference tested: Prediction is not priority, measurement is not standing, scarcity is not distribution, and risk estimate is not acceptable risk.

Card 29 — Six novel cases

1. An artificial system may be conscious.
2. A memory is edited without changing other capacities.
3. A synthetic copy claims identity continuity.
4. A policy changes which people will exist.
5. Responsibility is distributed across a remote causal network.
6. An ecosystem approaches an uncertain irreversible threshold.

Exact question: What follows under empirical uncertainty and what follows under normative uncertainty?

Variations: Change confidence, irreversibility, delay cost, standing, future identity, or appeal process one at a time.

Inference tested: Robust overlap, precaution, representation, confidence weighting, revision, and moral remainder perform different jobs.

Part V deck — The great candidates

Card 30 — The invisible village

Fixed facts: A policy maximizes total welfare while hiding that one village bears nearly all loss.

Exact question: Does the outcome rule include distribution, standing, and priority, or only a total?

Variations: Give priority to the worse off; impose sufficiency; add a right; make losses voluntary; change population.

Inference tested: Humane consequentialist repairs add selectors not supplied by maximization alone.

Card 31 — The healthy traveler

Fixed facts: Killing one healthy traveler would save five patients, remain secret, and create no further effects.

Exact question: What prohibits or permits the act?

Variations: Obtain consent; make death indirect; raise the number saved; introduce catastrophe; replace the traveler with a non-rational animal.

Inference tested: Universalization, status, rights, intention, and thresholds require content and priority rules.

Card 32 — The exhausted clerk

Fixed facts: A clerk quietly stops requesting food while formal rules prohibit abuse and aggregate outcomes appear acceptable. Caregivers are overburdened.

Exact question: What do virtue, care, relationship, and practical wisdom see that rules and totals miss?

Variations: The clerk refuses help; care becomes controlling; an outsider needs more; cultural exemplars conflict; caregiver capacity collapses.

Inference tested: Attention and relationship improve perception but do not automatically settle conflict or limits.

Card 33 — Two revelations at the harbor

Fixed facts: Two communities invoke the same divine source. One gives the last rescue place to covenant members; the other gives it to a stranger. Both have texts, traditions, and sincere interpreters.

Exact question: Which access and conflict rule identifies the received order?

Variations: Add natural-law reasoning; add institutional authority; change fruits; introduce a nonbeliever; make one interpretation recent.

Inference tested: Moving from command to good nature makes goodness and interpretation load-bearing.

Card 34 — The seawall and the unborn district

Fixed facts: A costly seawall burdens current residents, protects future people, and damages an estuary. Future people and animals cannot participate directly.

Exact question: Who enters a fair construction and through what representation?

Variations: Change veil membership; add adaptive preferences; let current consent be unanimous; change risk; give outsiders a veto.

Inference tested: Participants, idealization, representation, reasonableness, and constraints select the result.

Card 35 — Asha and the Tide People

Fixed facts: A hereditary rule excludes Asha from an archive. A majority forbids the Tide People from owning boats. Both practices are called cultural traditions.

Exact question: Which community supplies the standard, and how do reformers and internal minorities speak?

Variations: Add an equality promise; create a dissenting subcommunity; make the tradition genuinely ancient; affect outsiders; allow easy exit.

Inference tested: Culture boundaries, internal power, local validity, and approval-independent criticism should not be collapsed.

Card 36 — The evacuation of Low Harbor

Fixed facts: A cyclone threatens hospital patients, detainees, tower residents, a ceremonial archive, uncertain artificial subjects, an estuary, and exhausted rescuers. Vehicles are insufficient.

Exact question: What components and ordering rules govern a hybrid response?

Variations: Reverse political identity; change life expectancy; change artificial continuity; reveal crime records in the archive; alter ecological probability.

Inference tested: Hybrids add membership, priority, threshold, jurisdiction, and tie-breaking debts.

Part VI deck — Reform, codes, power, responsibility, and emotion

Card 37 — Caste, cover-up, and conquest

Fixed facts: One society maintains inherited caste, one unanimously exports a concealed safety risk to outsiders, and one openly values conquest without an equality promise.

Exact question: Which criticisms are factual, internal, genealogical, commitment-based, procedural, strategic, or objectively asserted?

Variations: Remove hypocrisy; include outsiders; add prior agreement; create internal dissent; make effects uncertain.

Inference tested: Effective criticism can have several grounds without every ground establishing objective condemnation.

Card 38 — The island generator

Fixed facts: One person controls the only generator, rejects rationing, and endangers medical refrigeration. The group has evidence, an emergency-sharing agreement, and lower-cost intervention options.

Exact question: What distinguishes reason-giving, bargain, coordination, legitimacy, jurisdiction, coercion, and force?

Variations: Remove the agreement; change the use from medicine to a monument; make the owner powerless; extend emergency control; add an appeal.

Inference tested: Collective necessity and power do not become objective authority; safeguards remain required.

Card 39 — Dam, clinics, and chemical fleet

Fixed facts: A democratic dam burdens outsiders; a majority closes clinics in one minority district; a wealthy fleet owner can pollute while safely evading sanctions.

Exact question: What representation, minimum protection, counterpower, and enforcement are needed?

Variations: Give outsiders a vote; disperse burdens; make the actor vulnerable to reputation; add independent review; alter jurisdiction.

Inference tested: Legitimate procedure within one constituency can coexist with domination beyond it.

Card 40 — The zero-purpose pain

Fixed facts: A deliberate saboteur is fully reformed. Protection, deterrence, restitution, rehabilitation, boundary, trust, and communication tasks are complete. Secret pain would produce no memory or effect.

Exact question: Should the pain occur solely because it is deserved?

Variations: Add deterrence; add victim demand; preserve memory; change control; make identity discontinuous.

Inference tested: Basic desert is isolated only after every other response purpose is removed.

Card 41 — Responsibility under altered conditions

Fixed facts: One offender sells memories of the act; one steals credit for a rescue; one is fully reformed; one acts under threat; one organization distributes failure across many desks.

Exact question: Which responses track causation, intention, control, risk, repair, trust, communication, system design, or payback?

Variations: Change present risk, memory continuity, organizational incentive, victim preference, and available repair separately.

Inference tested: Responsibility can remain rich after intrinsic suffering is removed.

Card 42 — Six emotional residues

1. Guilt after a faultless accident and completed repair.
2. Shame about a harmless trait.
3. Resentment after betrayal and full repair.
4. Indignation toward a vanished institution.
5. Gratitude for a controlling gift.
6. Pride in an achievement one did not help produce.

Exact question: What are the cause, appraisal, current function, audience, expression, intensity, and stopping condition?

Variations: Change danger, relationship, public audience, repair, recurrence, and desert belief one at a time.

Inference tested: An emotion can be understandable without being accurate, useful, worth expressing, or worth cultivating indefinitely.

Part VII deck — Language and policy after the search

Card 43 — The station promise

Fixed facts: Eren promises to collect Sal, knowingly goes elsewhere, gives no warning, and leaves Sal exposed to danger.

Exact question: Which language policy communicates reliance, effect, trust, protest, repair, and any objective claim?

Variations: No reliance; emergency prevented arrival; early warning; relationship ended; objective realism remains an open promissory note.

Inference tested: Abolition, fiction, conservation, expressivist revision, constructivist revision, and pragmatic translation have different risks.

Card 44 — Lark Square

Fixed facts: Serious youth assaults have risen around a square. Parents fear violence; teenagers report humiliating searches; businesses lose customers; youth programs were closed; evidence and causal models are disputed.

Exact question: How does each felt response travel through stake, fact, operational value, distribution, policy, safeguard, and review?

Variations: Change incident trend; change search yield; restore programs; alter affected testimony; reveal displacement; introduce a privacy breach.

Inference tested: Emotion reveals salience but cannot certify a factual model or move directly into law.

Card 45 — The coastal line

Fixed facts: Repeated floods threaten homeowners, renters, fishers, an Indigenous community, taxpayers, future residents, and a marsh. Barrier, relocation, and restoration policies distribute costs differently. A developer can profit from forced relocation.

Exact question: Complete a Civic Stakes Brief without granting theories seats as constituencies.

Variations: Change flood probability; compensation eligibility; marsh function; treaty authority; developer ownership; renter representation.

Inference tested: Shared values can support partial coalition while factual disagreement, priority conflict, power, and irreducible opposition remain visible.

Blank pressure card

Case title: _____

Exact question: _____

Fixed facts: _____

Tempting but excluded question: _____

Answer A: _____

Answer B: _____

One-variable variation 1: _____

One-variable variation 2: _____

One-variable variation 3: _____

What the selected answer can establish: _____

What remains open: _____

Inference tested: _____

Accept, bridge, or narrow: _____

Appendix C — Common Repairs and Their Debts

A repair is not an embarrassment. Good systems survive by repairing first formulations. The question is whether a repair supplies the missing support, moves the support to another clause, narrows the claim, or merely redescribes the gap.

Use this field guide whenever a familiar response appears. Give the response its strongest reading before entering its debt.

Repair 1 — “It is self-evident”

Strongest reading: Some truths can be known through rational reflection without inference from prior premises. A regress of justification may legitimately stop.

What it can accomplish: It states a foundational epistemology and blocks the demand that every normative claim be reduced to a nonnormative proof.

Debt created: Which proposition is self-evident? To whom? How are conceptual competence, trained intuition, cultural fluency, confidence, and genuine insight distinguished? How are conflicts among purportedly self-evident principles handled?

Pressure questions:

- Would sincere, informed disagreement reduce confidence?
- Does changing the frame or identity change the seeming?
- Does the intuition transfer to novel cases?
- What would count as a misapprehension of the basic truth?

Honest narrower form: “This proposition is a foundational commitment or seeming for me. I have not supplied an independent tracking account.”

Repair 2 — “The evidence speaks for itself”

Strongest reading: Well-calibrated evidence reliably discriminates truth from error, and serious inquiry should respond to that relation.

What it can accomplish: It can establish descriptions, causes, probabilities, consequences, and means relative to admitted ends.

Debt created: Evidence does not literally speak. Which proposition does it support, through which model and baseline? What gives evidential support normative force, if such force is claimed? Which value selects the policy after the factual estimate is known?

Pressure questions:

- Is the dispute factual, evaluative, or procedural?
- Which measurement choices affect visibility and aggregation?

- Would the same evidence support opposite policies under different values?
- Is the claim truth-directed, useful, practice-relative, publicly admissible, or objectively reason-giving?

Honest narrower form: “Under this stated method, the evidence raises confidence in factual claim F. F does not by itself select value V or rule R.”

Repair 3 — “Everyone knows”

Strongest reading: Widespread, stable convergence may be evidence when it arises through independent, truth-sensitive routes.

What it can accomplish: It can identify a shared expectation, public starting point, or candidate datum.

Debt created: Who is included in *everyone*? Is convergence verbal, case-level, rule-level, foundational, or methodological? Does it arise from common evidence, common psychology, social pressure, conquest, or truth tracking?

Pressure questions:

- Are dissenters absent, silenced, or reclassified as irrational?
- Does convergence survive fuller information and identity reversal?
- Would universal agreement constitute truth or merely accompany it?
- What happens in the unanimous harmful culture?

Honest narrower form: “This judgment is widely shared in group G under conditions C and therefore offers coordination, presumptive expectation, or defeasible evidence.”

Repair 4 — “A rational person would agree”

Strongest reading: Coherence, consistency, universalization, and the constitutive conditions of agency constrain practical judgment.

What it can accomplish: It can expose contradiction, arbitrary personal exception, means–end failure, and incoherent intention.

Debt created: Which substantive ends, standing rules, or reciprocity conditions are included in *rational*? Can a coherent tyrant satisfy the formal constraints? Does idealization correct error or replace the person’s values?

Pressure questions:

- What exact contradiction follows from rejection?
- Does the result depend on equal standing already being assumed?
- Which ideal agent is imagined?
- Can several coherent equilibria remain?

Honest narrower form: “Given ends E, information I, and consistency rule C, this agent would support R. The humane content enters through E or C as specified.”

Repair 5 — “God is goodness”

Strongest reading: Goodness is neither external to God nor arbitrarily created by command; it is identical with a necessary, perfectly loving and just nature.

What it can accomplish: It blocks the idea that any powerful command creates value and connects authority to a stable metaphysical source.

Debt created: What substantive content do *good*, *loving*, *just*, and *perfect* carry? How do interpreters identify that nature, test alleged revelation, resolve disagreement, and address outsiders? Why does creation confer moral jurisdiction?

Pressure questions:

- How is a cruel-seeming command classified as false or misunderstood?
- Does independent moral understanding test revelation?
- What distinguishes faithful development from revision?
- What reaches the informed nonbeliever?

Honest narrower form: “Within this received account, divine nature unifies our values and authority. Correct interpretation uses these declared standards and remains fallible.”

Repair 6 — “Human nature determines it”

Strongest reading: Human beings have real capacities, dependencies, vulnerabilities, and forms of flourishing that practical inquiry can study.

What it can accomplish: It supplies rich empirical content, criticizes preferences adapted to deprivation, and connects rules with needs and development.

Debt created: Which capacities define proper function? Why does a biological, psychological, or social function generate value or reason? Must each person realize every good? How are plural goods compared? What about unusual lives?

Pressure questions:

- Is *natural* descriptive, statistical, teleological, or evaluative?
- Why is disease excluded while another natural pattern is honored?
- Does flourishing already contain the favored moral traits?
- What binds an informed person who rejects the selected ideal?

Honest narrower form: “Given concern for these forms of flourishing, natural and social facts identify enabling and damaging conditions.”

Repair 7 — “Society decides”

Strongest reading: Communities create real standards, roles, expectations, procedures, and laws; public decisions are necessary under disagreement.

What it can accomplish: It explains social classification, coordination, institutional validity, enforcement, and some forms of legitimacy.

Debt created: Which society? How are overlapping communities, reformers, minorities, outsiders, and internally manufactured consent treated? Does actual decision constitute truth, legitimate action, or merely settle conduct?

Pressure questions:

- Who had voice, information, and exit?
- Does a majority control the definition of the culture?
- What protects a powerless minority?
- Can a society be criticized by its promises or only by another standard?

Honest narrower form: “Institution I authorized rule R through procedure P within jurisdiction J. This establishes legal or procedural status, not objective moral truth.”

Repair 8 — “Everything that matters can be added”

Strongest reading: A plural or hybrid system can retain consequences, rights, virtues, care, tradition, procedure, and local knowledge rather than sacrificing important human concerns to elegance.

What it can accomplish: It can produce a more perceptive, humane, stable, and corrigible practical system.

Debt created: Who chooses the components? What gives them priority? What happens in conflict? Is one component an undeclared master? Does “all things considered” hide the final selector?

Pressure questions:

- Which component wins in the evacuation case?
- What threshold permits override?
- Can two hybrids reach opposite equilibria?
- Does combination pay or merely collect the debts of the parts?

Honest narrower form: “We protect this declared plurality and settle conflict through meta-rule M, selected on basis B.”

Repair 9 — “Protect the boundary until the emergency is big enough”

Strongest reading: Rights should block ordinary sacrifice while catastrophic stakes can justify exceptional override.

What it can accomplish: It avoids both casual aggregation and disastrous absolutism.

Debt created: What defines the right, emergency, threshold, evidence standard, authorized decision-maker, duration, and repair? Does the exception preserve the principle or allow consequence to become the hidden master?

Pressure questions:

- Is one death enough? Ten? Social collapse?
- Who estimates the emergency?
- Does uncertainty lower or raise the threshold?
- Does authority expire automatically?

Honest narrower form: “We adopt a strong presumption P with public threshold T, evidence E, review R, and restoration S because we prioritize concerns V.”

Repair 10 — “Balance the values”

Strongest reading: Several incommensurable concerns may be genuine; practical judgment should respond to context without pretending they share one unit.

What it can accomplish: It preserves plural value, tragedy, residual loss, and sensitivity to cases.

Debt created: Who balances, by what method, with what calibration and appeal? Does balance settle action, express a choice, or discover an objective rank? How are power and identity prevented from becoming the scale?

Pressure questions:

- Which value may never fall below a minimum?
- What defeats a prior judgment?
- How do peer judges resolve persistent disagreement?
- What compensation acknowledges the losing concern?

Honest narrower form: “Decision-maker D uses plural criteria C and precedent P to select R, records residual loss L, and permits review under trigger T.”

Repair 11 — “Use practical wisdom”

Strongest reading: Skilled judgment perceives salience, relationship, timing, character, and context that no finite rule can encode.

What it can accomplish: It explains expertise, formation, discretion, and the need for responsive rather than mechanical action.

Debt created: Which exemplars define wisdom? How are virtue conflict, cultural difference, bias, and care-as-control corrected? Does “the wise person sees” provide a method or rename the desired answer?

Pressure questions:

- What evidence would show the expert misperceived?
- Can another admirable agent disagree?
- Which institutional review constrains discretion?
- How is wisdom formed before the learner can identify a wise teacher?

Honest narrower form: “Experienced agent A exercises bounded discretion using criteria C, reasons publicly, and remains answerable to review R.”

Repair 12 — “Suffering obviously matters”

Strongest reading: Suffering is vividly experienced, widely avoided, reliably connected to harm, and a central object of compassion.

What it can accomplish: It supplies a powerful practical starting point and directs empirical attention toward preventable injury.

Debt created: Does suffering matter to the sufferer, to compassionate agents, within a public commitment, or objectively to every agent? Who can suffer? How are intensity, number, distribution, agency, and competing values handled?

Pressure questions:

- Does a million mild pains outweigh one torture?
- What if the agent does not care?
- What about voluntary pain or painful treatment?
- Does reducing suffering permit preference editing or sacrifice?

Honest narrower form: “These beings avoid and are impaired by suffering; we care about that stake and design policy to reduce it under distribution and agency constraints.”

Repair 13 — “Evolution or culture explains the belief”

Strongest reading: A belief-forming process may track survival, conformity, kinship, or status rather than truth, shifting the burden onto a tracking account.

What it can accomplish: It can explain confidence, expose irrelevant cues, undermine a local process, or support a burden shift.

Debt created: What additional premise connects origin to falsity? Could truth coincide with adaptiveness? Did reflection correct the initial process? Is there an independent route?

Pressure questions:

- Is the argument causal, process-level, content-level, or burden-level?
- Would a different formation history change the belief?
- Does the genealogy bear on truth or only confidence?
- Could the same challenge debunk the skeptic’s values?

Honest narrower form: “This genealogy lowers trust in process P unless an independent truth-sensitive route is shown. It does not by itself establish that proposition F is false.”

Repair 14 — “People disagree, so there is no truth”

Strongest reading: Persistent informed disagreement may be unexpected under a theory that predicts clear, shared access to determinate truth.

What it can accomplish: It can challenge an access claim, reduce confidence, reveal indeterminacy, or motivate procedural settlement.

Debt created: Why should truth produce convergence here? Are peers equally informed and independent? Is disagreement about fact, standing, priority, commitment, procedure, or objective truth?

Pressure questions:

- Do other true domains contain disagreement?
- Does consensus result from shared formation?
- Would idealization converge or produce several equilibria?
- What policy follows while confidence remains divided?

Honest narrower form: “This disagreement is evidence against access hypothesis H to degree D; it does not entail error theory.”

Repair 15 — “A fair vote settles it”

Strongest reading: Accountable procedure enables coordination and peaceful revision among people who cannot agree on substance.

What it can accomplish: It can create legitimate settlement within a constituency, distribute authorship, and make power answerable.

Debt created: Who participates, who is represented, what constraints apply, and what reaches outsiders? Does the result establish procedural validity or moral truth?

Pressure questions:

- Does the majority concentrate severe cost on a minority?
- Who set the agenda?
- Can future people or animals be represented?
- What may no ordinary majority do?

Honest narrower form: “Procedure P authorizes temporary settlement R among constituency C, subject to protections S and revision T.”

Repair 16 — “The law creates the obligation”

Strongest reading: Legal systems constitute offices, permissions, prohibitions, expectations, and sanctions that would not exist without institutions.

What it can accomplish: It explains legal status, jurisdiction, coordination, and practical consequences.

Debt created: Does legal validity create prudential, role-relative, agreement-based, or objective moral support? What makes the institution legitimate? How are unjust or outsider-burdening laws criticized?

Pressure questions:

- Would a procedurally valid caste law be morally authoritative?
- Does fear of sanction count in favor or merely motivate?
- What permits civil disobedience?
- Where does jurisdiction end?

Honest narrower form: “Law L validly governs domain D and creates consequences C. Its legitimacy and objective authority require separate arguments.”

Repair 17 — “The person deserves it”

Strongest reading: Desert recognizes agency, protects victims, limits punishment to the responsible, and supports proportionality.

What it can accomplish: It focuses attention on knowledge, intention, control, identity, and relation to the act.

Debt created: Which response purpose requires suffering as such? What makes a quantity proportional? How do identity change, coercion, reform, and system causation alter desert?

Pressure questions:

- Would zero-purpose secret pain improve the situation?
- Have protection, repair, trust, and communication already completed their jobs?
- Does victim recognition require offender suffering?
- When does the response stop?

Honest narrower form: “We impose response R for named purpose P under safeguard S and stopping condition T. Any remaining suffering-for-its-own-sake is labeled retribution.”

Repair 18 — “Without objective morality, anything goes”

Strongest reading: People fear losing criticism, law, responsibility, motivation, and protection if no stance-independent authority exists.

What it can accomplish: It correctly identifies practical functions that any replacement must preserve or consciously revise.

Debt created: The inference assumes that only objective facts can support action and institutions. Which phenomenon actually disappears: pain, concern, agreement, law, evidence, power, protest, or universal authority?

Pressure questions:

- Do values and institutions cease to exist?
- Does egoism become objectively true?
- Why would preference strengths become equal?
- Can a society enforce a charter without cosmic certification?

Honest narrower form: “Without objective authority, no policy is guaranteed to bind every rejecter. We still act through concerns, commitments, agreements, institutions, and power, which require explicit limits.”

Blank repair card

Repair offered: _____

Strongest reading: _____

What it genuinely accomplishes: _____

Original debt addressed: _____

New or relocated debt: _____

Pressure question 1: _____

Pressure question 2: _____

Pressure question 3: _____

Honest narrower form: _____

Appendix D — The Civic Stakes Brief

The Civic Stakes Brief is a reusable public-policy template for controversies in which factual uncertainty, divergent values, institutional power, and affected lives interact. It does not certify objective moral truth. It makes a proposal answerable to its declared purposes, evidence, constituencies, tradeoffs, safeguards, and revision conditions.

Complete the brief before a final vote or coercive implementation. Preserve every version. Do not replace dissent with “the community believes.”

Use rules

Purpose. Apply the same sequence to every brief so urgency, authority, and familiar moral vocabulary do not bypass evidence, comparison, safeguards, or dissent.

1. Start with reported experience and concrete stakes, not theory labels.
2. Separate factual claims from standing, priority, settlement, and authority.
3. Treat emotion as evidence of salience and appraisal, not automatic evidence that the appraisal is accurate.
4. State the comparison method whenever values conflict.
5. Name coercion and force rather than hiding them inside *implementation*.
6. Give every extraordinary power a scope, review, and expiration.
7. State what would reduce confidence in the policy.
8. Preserve the views of affected dissenters.
9. Plan repair before the institution discovers its first error.
10. Translate any high-stakes use of *right*, *wrong*, *just*, *duty*, or *desert* into the components below.

Cover page

Purpose. Establish authorship, institutional ownership, version history, and review responsibility before the proposal begins making claims about other people’s lives.

Decision title: _____

Brief version: _____

Date: _____

Decision date: _____

Preparing institution: _____

Named authors and roles: _____

Affected-party contributors: _____

Independent reviewers: _____

Public location of evidence and prior versions: _____

Current status: *exploration / pilot / proposed / authorized / active / under review / expired:* _____

1. Decision and jurisdiction

Purpose. State the exact choice, decision-maker, scope, duration, and route of review; this keeps a broad aspiration from quietly authorizing conduct outside the decision.

Exact decision: _____

Options currently open: _____

Decision-maker: _____

Legal or institutional authority: _____

Geographic scope: _____

Population or subject-matter scope: _____

Start and end dates: _____

Matters outside this decision: _____

Authority that could reverse or review it: _____

Jurisdiction check

- ☐ The decision-maker's authority over insiders is stated.
- ☐ Effects outside the jurisdiction are identified.
- ☐ Emergency authority is separated from ordinary authority.
- ☐ Legal validity is not presented as objective moral correctness.

2. Affected parties and standing

Purpose. Identify who or what bears direct and indirect effects before ranking claims, especially parties who lack formal membership, representation, or present political power.

List parties before ranking their claims.

Party or system	Direct effect	Indirect effect	Can participate?	Proposed ground of standing	Confidence	Representative or evidence source

Check explicitly for:

- ☐ residents without formal membership;
- ☐ children or dependents;
- ☐ people with limited communication;
- ☐ future people;
- ☐ people outside the jurisdiction;
- ☐ nonhuman animals;
- ☐ ecosystems;
- ☐ artificial subjects of uncertain status;
- ☐ workers who implement the policy;
- ☐ people vulnerable to official error.

Who remains unrepresented and why: _____

How uncertainty about standing changes the policy: _____

3. Emotional stakes and reported experience

Purpose. Record what people fear, resent, value, grieve, or hope for while separating the salience of an emotion from the factual accuracy and policy reach of its appraisal.

Do not infer facts from emotion without a separate evidential route.

Party	Reported emotion or response	Object or trigger	Appraisal contained	Current function	Possible distortion or strategic use	How testimony was gathered

Prompts:

- What are people afraid will happen?
- What loss does grief register?
- What boundary does anger protect?
- What humiliation or domination is reported?
- Which attachment or identity is threatened?
- Which hope sustains support?
- Whose emotional report is dismissed as irrational, and by whom?
- Which fear or indignation is being amplified for advantage?

Responses that reveal a new stake: _____

Appraisals requiring factual testing: _____

Emotions that should not be turned directly into policy: _____

4. Operational values

Purpose. Translate abstract value words into observable states, experiences, capacities, relationships, or practices so competing proposals can be compared against declared aims.

Translate abstract values into states, experiences, capacities, relationships, or practices.

Value label	Operational meaning	Bearer	Desired di- rection	Minimum, target, or protected boundary	Evidence relevant to application
_____	_____	_____	_____	_____	_____
_____	_____	_____	_____	_____	_____
_____	_____	_____	_____	_____	_____

Examples of required translation:

- *Safety* → protection from which event, at what probability and time?
- *Freedom* → which action, option, or protection from whose control?
- *Dignity* → which treatment, recognition, bodily boundary, or social condition?
- *Fairness* → which distribution, procedure, reciprocity, or status relation?
- *Community* → which people, practices, attachments, and boundaries?

Whose values these are: _____

Who rejects or gives low priority to them: _____

Objective status claimed, if any, and warrant: _____

5. Empirical claims and confidence

Purpose. Make every consequential factual premise traceable to evidence, uncertainty, and possible correction; the section tests descriptions and predictions, not the value that selects among them.

Claim	Type: description, cause, forecast, or measurement	Evidence	Method and baseline	Confidence	Main uncertainty	Evidence that would reduce confidence

Audit:

- ☐ Operational definitions are public.
- ☐ Missing cases and invisible populations are considered.
- ☐ Aggregation choices are stated.
- ☐ Classification thresholds are stated.
- ☐ Model objectives and incentives are disclosed.
- ☐ Affected testimony and administrative data are compared.
- ☐ Correlation is not presented as cause without a method.
- ☐ Evidence can change the proposal rather than merely decorate it.

Truth-responsive route used: _____

What evidence settles: _____

What evidence leaves open: standing / priority / acceptable risk / distribution / procedure / objective status / other: _____

6. Comparison method and unresolved conflicts

Purpose. Disclose how the proposal compares gains, losses, thresholds, rights, priorities, and distributions, including conflicts the chosen method leaves unsettled.

Conflict	Bearers	Distribution	Method selected	Selector or decision-maker	Residual loss	Dissent

Available methods:

- common-scale aggregation;
- priority to the worse off;
- protected minimum or boundary;
- emergency threshold;
- lexical priority;
- disciplined plural judgment;
- fair procedure;
- negotiated or chosen settlement.

Why this method was chosen: _____

What the method claims: objective rank / public legitimacy / coordination / practical choice / other: _____

Tie rule: _____

Conflict left genuinely unresolved: _____

Compensation, apology, restoration, or acknowledgment of residue: _____

7. Available resources and constraints

Purpose. Distinguish genuine limits on implementation from political choices presented as inevitabilities, and show whose time, money, labor, attention, or risk the proposal consumes.

Resource	Available amount	Owner or controller	Timing	Competing use	Reliability or uncertainty
Money	_____	_____	_____	_____	_____
Personnel	_____	_____	_____	_____	_____
Time	_____	_____	_____	_____	_____
Land or physical space	_____	_____	_____	_____	_____
Infrastructure	_____	_____	_____	_____	_____
Information	_____	_____	_____	_____	_____
Admin. capacity	_____	_____	_____	_____	_____

Opportunity costs: _____

Scarcity that cannot be removed during the decision: _____

Who bears delay: _____

8. Power and enforcement

Purpose. Name who can decide, compel, evade, delay, capture, challenge, and repair the policy so implementation is not described as though rules enforce themselves.

Draw the power map.

Actor	Can decide	Can delay	Can enforce	Can evade	Can set agenda	Can gather evidence	Can appeal	Main dependency or counterpower

Conduct constrained: _____

Coercion or force authorized: _____

Evidence threshold: _____

Maximum permitted intrusion: _____

People most vulnerable to error or selective enforcement: _____

Actor who can safely defect: _____

Counterpower and independent review: _____

9. Policy alternatives

Purpose. Compare the proposal with a real baseline and meaningful alternatives; rejecting one option does not establish that the favored option is necessary or best.

Include a real baseline and at least three alternatives.

Option	Mechanism	Values served	Parties bene- fited	Parties bur- dened	Evidence	Power required	Reversibility	Cost
Current course / no action								
Noncoercive support								
Pilot or limited intervention								
Regulatory or coercive intervention								
Hybrid								

Option excluded before analysis and why: _____

Policy chosen: _____

Why this option rather than the nearest alternative: _____

10. Safeguards

Purpose. Design limits that remain protective under error, abuse, leadership change, and hostile administration rather than relying on the current decision-makers’ good intentions.

Predictable failure or abuse	Safeguard	Responsible actor	Verification	Appeal	Trigger for suspension
_____	_____	_____	_____	_____	_____
_____	_____	_____	_____	_____	_____
_____	_____	_____	_____	_____	_____

Consider:

- protected minimums;
- evidence and notice requirements;
- least-intrusive means;
- representation of concentrated losers;
- conflict-of-interest rules;
- privacy and data-use limits;
- independent review;
- public reasons and records;
- appeal;
- compensation;
- minority protection;
- expiration;
- restoration after emergency.

Safeguard value: _____

Who supports it and why: _____

Who can disable it: _____

11. Measures

Purpose. Connect each declared aim to indicators, baselines, horizons, and distributional checks while anticipating proxy failure and incentives to improve the number rather than the condition.

Measure	Value or factual claim tracked	Distribution re-ported	Data owner	Frequency	Manipulation risk	Interpretation limit

Include where relevant:

- outcome and severity;
- process and participation;
- distribution across groups;
- false positives and false negatives;
- displacement to another location or population;
- reported experience;
- cost and opportunity cost;
- rights or boundary violations;
- official complaints and established errors;
- unmeasured or poorly measured stakes.

Measure that could look successful while the policy fails: _____

Stake not captured well by available measures: _____

12. Review, failure, and sunset

Purpose. State when evidence, outcomes, or changed conditions will trigger review, suspension, revision, or expiration; correctability requires more than permission to reconsider someday.

Initial review date: _____

Full redesign date: _____

Automatic expiration: _____

Who can trigger early review: _____

Evidence of success: _____

Evidence of partial success: _____

Failure condition: _____

Unacceptable collateral effect: _____

What remains protected during review: _____

Procedure for continuation: _____

13. Error acknowledgment and repair

Purpose. Decide in advance how the institution will identify injury, notify affected parties, restore records or resources, compensate loss, and prevent recurrence.

Plan before the error occurs.

Possible error	Person or system affected	Notice	Record correction	Compensation or restoration	Institutional change	Public acknowledgment

Responsible office: _____
Funding reserved for repair: _____
Time limit for response: _____
Appeal when repair is denied: _____

14. Support and dissent statement

Purpose. Preserve the strongest supporting and dissenting accounts in their authors’ own terms so a collective decision is not rewritten as unanimous moral agreement.

Actor or group	Supports, opposes, or uncertain	Actual basis	Desired change	Power in final decision
_____	_____	shared value / interest / commitment / agreement / role / procedure / prediction / objective claim / other: _____	_____	_____
_____	_____	_____	_____	_____
_____	_____	_____	_____	_____

Overlapping support from different values: _____
Shared value with disputed factual model: _____
Shared facts with disputed priority: _____
Irreducible conflict: _____
Permanent dissenter: _____
What will be enforced despite dissent: _____
Objective authority claimed, if any, and warrant: _____

Final decision record

Purpose. Freeze the authorized decision, reasons, evidential basis, acknowledged costs, safeguards, and unresolved objections at the moment responsibility attaches.

Selected policy: _____
Date and vote or authorizing act: _____
Declared goals: _____
Primary evidence: _____

Comparison method: _____
Main burden and bearer: _____
Principal safeguard: _____
Review and expiration: _____
Dissent preserved: _____
Repair commitment: _____
Clause asserted as objectively required: _____
Ground and access route for that clause: _____

Independent reviewer page

Purpose. Invite structured challenge from outside the drafting group and record which objections change the proposal, which expose uncertainty, and which remain unresolved.

Ask at least three reviewers to mark one entry in each category.

Invisible party: _____
Misclassified standing claim: _____
Emotion treated as factual proof: _____
Unsupported factual claim: _____
Hidden comparison rule: _____
Underestimated cost or distribution: _____
Unmapped power: _____
Safeguard vulnerable to capture: _____
Metric that can be gamed: _____
Missing failure condition: _____
Repair omitted: _____
Objective language exceeding its warrant: _____

Reviewer standpoint and relevant interests: _____
What this review leaves open: _____

One-page public summary

Purpose. Translate the full brief into an accessible account of the decision, stakes, evidence, tradeoffs, enforcement, safeguards, dissent, and review without erasing inconvenient detail.

We are deciding _____ under authority _____. The people and systems most affected are _____. They report stakes including _____. The policy protects or advances _____, based on factual claims _____ with confidence _____. The principal unresolved conflict is _____, settled through _____. The largest burden falls on _____. Implementation gives power

to _____ and constrains _____. Safeguards include _____. We will measure _____, review on _____, and end or redesign the policy if _____. Established errors will be addressed through _____. Support comes from _____; dissent remains from _____. The proposal is presented as factual / instrumental / expressive / commitment-based / agreement-based / procedural / enforced / objectively required: _____.

Appendix E — Glossary and Position Map

Metaethical positions do not stand on one line from “very objective” to “very subjective.” They answer different questions about meaning, truth, dependence, knowledge, motivation, authority, and practice. A person can combine answers from several rows.

This glossary uses the terms as they function in this book. Other philosophers sometimes define them more narrowly or divide them differently.

The question map

Dimension	Core question	Example answers
Meaning	What does a moral sentence do?	report, express, apply, construct, fictionalize
Truth	Can moral claims be true or false?	yes, no, or only relative to a standard
Ground	What makes a verdict correct?	natural fact, non-natural fact, stance, procedure, authority, practice
Dependence	Would truth change if every stance changed?	stance-independent or response-dependent
Scope	Who or what is addressed or protected?	universal, role-bound, relational, jurisdictional
Reasons	What counts in favor of action or belief?	aim-relative, commitment-relative, practice-relative, aim-independent
Access	How can inquirers improve reliability?	observation, intuition, reason, revelation, procedure, no special access
Motivation	Must sincere judgment motivate?	necessarily, under normal conditions, or not at all
First-order guidance	Which acts or policies are selected?	consequences, duties, virtues, care, agreement, plural judgment
Public settlement	How is shared action chosen?	bargaining, vote, representation, law, coercion, chosen compromise
Responsibility	Why respond adversely or favorably?	protection, repair, relationship, communication, desert

Dimension	Core question	Example answers
Language after skepticism	What should speakers do?	abolish, fictionalize, conserve, revise, translate

A realist and an expressivist can favor the same policy. Two realists can disagree about every first-order issue. A constructivist can be universalist or local. A consequentialist can be a realist, constructivist, relativist, or fictionalist about the status of its principle.

Core project terms

Metaethics

Inquiry into what moral claims mean, whether they can be true, what could make them true, how they might be known, and how they relate to reasons, motivation, language, and social practice.

Metaethics does not directly answer whether a particular tax, promise, or medical intervention is permissible. It asks what kind of answer such a verdict would be.

Normative ethics

Inquiry into general standards for action, character, relationship, and institutions. Consequentialism, deontology, virtue ethics, care ethics, and contractualism are familiar normative families.

Applied ethics

Inquiry into concrete domains such as medicine, punishment, technology, war, environment, family, or public policy. Applied arguments typically contain factual, normative, institutional, and interpretive components.

Objective moral fact

A fact about moral status whose truth is not constituted merely by the actual approval, preference, commitment, convention, command, or procedure of any person or group.

The definition does not require a ghostly property. A naturalist may propose that objective moral facts are natural facts. The issue is what makes the fact true and whether the truth survives changes in stance.

Objective moral system

An organized account of moral statuses, their grounds, and their implications that claims at least some stance-independent truth and practical authority.

In this book, a complete system also aspires to scope, access, coherence, determinacy, usability, and correction.

Stance

A response, preference, aim, commitment, acceptance, convention, identity, endorsement, or procedure associated with an agent or group. An idealized stance is what such an agent would endorse under specified improvements or constraints.

Authority

The purported capacity of a consideration or rule to count in favor of action for an agent who does not already endorse it. Authority is not identical to causal influence, motivation, legal jurisdiction, social pressure, punishment, or physical power.

Coherence

Success in paying the explanatory debts generated by a system's own claims without contradiction, equivocation, opportunistic exception, hidden selector, or circular support.

Coherence is not truth or humaneness. A cruel code can be coherent; a compassionate practice can be inconsistent.

Selector

A value-laden choice that converts open possibilities into a determinate rule or verdict. Examples include who has standing, which value is primary, what risk is acceptable, where a threshold lies, or who participates in a procedure.

Bridge

A premise or relation connecting one kind of claim to another—for example, from suffering to universal practical reason, from function to value, from value to rule, or from social approval to objective obligation.

Debt

An explanatory burden created by a claim. A debt is not proof that the claim is false. It identifies what the system must establish before counting the claim as earned.

Primitive

A starting response or concern that supplies practical orientation without being inferred from a prior moral theory. Pain aversion, attachment, compassion, resentment of domination, curiosity, loyalty, and desire for security are examples.

Primitive does not mean infallible, universal, uncaused, fixed, or automatically fit for law.

Operational value

A value stated as a concrete state, experience, capacity, relationship, or practice that a policy aims to protect, reduce, enable, or prevent. “Safety” becomes operational when it names a threat, bearer, probability, horizon, and protected threshold.

Truth-responsive practice

A method designed to improve discrimination between truth and error, such as calibrated measurement, evidence comparison, causal investigation, transparent recordkeeping, or adversarial review. The term does not settle whether evidence supplies stance-independent epistemic reasons.

Meaning and truth positions

Cognitivism

The view that central moral judgments express belief-like states or make claims capable of truth or falsity.

Realism and error theory are both usually cognitivist: they agree that ordinary claims purport to report moral facts and disagree about whether any are true.

Noncognitivism

A family of views on which moral judgments primarily express attitudes, commitments, plans, prescriptions, or other non-belief-like states rather than straightforwardly reporting independent moral facts.

Contemporary versions can still explain logical relations, consistency, disagreement, and disciplined reasoning.

Moral realism

The view that at least some moral claims are true and that their truth is not constituted merely by actual individual or collective approval. Realists disagree about whether moral facts are natural or non-natural, reducible or irreducible, necessary or contingent, and easily known or epistemically difficult.

Realism alone does not select consequentialism, rights, equality, or any other first-order content.

Moral naturalism

A realist family identifying or grounding moral facts in natural facts, such as welfare, flourishing, desire satisfaction, agency, or complex social-natural properties.

Naturalists must state which natural properties perform the moral role and how content, authority, and access are secured.

Moral non-naturalism

A realist family on which moral properties or reason relations are not reducible to ordinary natural properties. They may be fundamental, necessary, rationally accessible, or known through a distinctive form of reflection.

Non-naturalism is not automatically supernatural and does not by its name solve access or conflict.

Error theory

The view that ordinary moral claims purport to report objective moral facts but systematically fail because the relevant facts or properties do not exist.

An error theorist must still choose a practical language policy: abolition, fiction, revision, or another approach does not follow automatically.

Expressivism

A family of views on which moral judgments express attitudes, commitments, plans, or practical orientations and participate in attempts to coordinate or influence action.

Expressivists can take facts, consistency, disagreement, and revision seriously. Expression is not mere noise or arbitrary emotion.

Fictionalism

The proposal that people may use moral discourse as a useful fiction even if its literal objective claims are false or unaccepted. Fictionalists differ over whether speakers assert what they know to be false, engage in pretense, or use a special practical mode.

The central risks are content selection, disclosure, paternalism, and the stability of conscious pretense.

Conservationism

The language policy of retaining ordinary moral discourse largely unchanged, either because realism remains plausible, because ordinary use does not commit speakers to a heavyweight metaphysics, or because the practice performs valuable jobs.

Conservation needs a response to equivocation in high-stakes settings.

Dependence and construction positions

Moral relativism

A family of views on which the truth or validity of a moral claim is relative to a culture, individual, framework, or standard rather than true from no standpoint.

Relativism differs from tolerance. A relativist framework can contain intolerant standards, and a realist can value tolerance. Relativism also faces questions about group boundaries, reformers, overlapping membership, and outsiders.

Contextualism

The view that which standards or parameters govern a claim can vary with conversational or practical context. Contextualism need not make truth relative to a culture in the full relativist sense.

Constructivism

A family of views on which normative principles are generated, constituted, or vindicated through a procedure of practical reasoning, agreement, idealization, or reciprocal justification.

Kantian constructivists often emphasize the conditions of rational agency and equal authorship. Humean constructivists begin from agents' actual concerns and idealize them. Political constructivists may claim legitimacy rather than metaphysical truth.

Constructivism must specify participants, information, motivation, admissible options, constraints, decision rules, and the status of the result.

Response dependence

A property is response-dependent when its identity or truth conditions refer to responses under actual or specified conditions. Color is a familiar candidate; some philosophers propose response-dependent value.

A response-dependent fact can be real and stable. Whether it counts as objective depends on the project's threshold.

Reasons and epistemology

Practical reason

A consideration that counts in favor of an action, intention, plan, or policy. This favoring relation is distinct from what caused an action and from what the agent consciously acted on.

Motivating reason

The consideration from which an agent acts, as represented in the agent's deliberation. It may rest on a false belief and need not genuinely favor the act.

Explanatory reason

A fact or cause that explains why an event, belief, or action occurred. It may operate without the agent's awareness and need not justify anything.

Practical-reasons internalism

An ambiguous label for views connecting practical reasons with an agent's aims, motivations, or possible motivation under suitable reflection. On one common version, a consideration is a reason for an agent only if it connects to some element of the agent's motivational set after specified correction.

The key questions are what connection and what idealization count.

Practical-reasons externalism

The view that at least some practical reasons apply independently of an agent's existing or suitably idealized motivations. Another person's severe need could give an indifferent agent a reason even if no concern connects.

Externalism states the relation directly and owes an account of its content, strength, scope, and access.

Motivational internalism

The view that sincere moral judgment necessarily or normally carries some motivation. This is not the same as reasons internalism. A philosopher may believe objective reasons exist while denying that recognizing them always motivates.

Epistemic reason

A consideration that genuinely counts in favor of believing a proposition. It is distinct from evidence causing belief, belief serving a useful purpose, or a community admitting evidence under procedural rules.

Epistemic-reasons internalism

In this book's practical usage, the view that epistemic support depends on a relation available within an agent's perspective, truth-seeking aim, inquiry commitment, or accessible evidence. Philosophical literature uses *epistemic internalism* in several more technical ways, especially concerning justification and mental access.

Always state the intended version.

Epistemic-reasons externalism

The view that epistemic support can depend on relations outside an agent's awareness or aims, such as reliable truth tracking, proper function, or stance-independent evidential reasons.

Again, reliability, knowledge, justification, and objective reason are not identical merely because each can be called external.

Epistemic anti-normativism

The view that there are no irreducibly normative epistemic reasons or obligations. An anti-normativist may still describe evidence, truth tracking, reliable belief formation, inquiry rules, practical usefulness, and public justification.

Epistemic anti-normativism does not follow from moral non-realism, and moral non-realism does not follow from it.

Reflective equilibrium

A method of revising case judgments, principles, background theories, and factual beliefs until they fit together coherently after informed reflection.

Reflective equilibrium can improve consistency and breadth. It may produce several equilibria from different starting commitments unless further constraints select among them.

Genealogical challenge

An argument examining how a belief or value formed through biology, culture, status, history, or personal experience. A genealogy can explain, locally undermine, criticize content, or shift a burden. Origin alone does not entail falsity.

Standing and value comparison

Direct standing

Counting for one's own sake rather than only through usefulness, ownership, admiration, ecological function, or connection to someone else.

Standing identifies entry into consideration. It does not settle weight, identical treatment, or final outcome.

Sentience account of standing

The view that capacity for felt experience—especially pain, pleasure, distress, or enjoyment—grounds direct standing. The account must address uncertain consciousness, intensity, future experience, and beings with agency but no sentience.

Agency account of standing

The view that self-direction, intention, planning, preference, rational authorship, or another form of agency grounds direct standing. The account must address infants, profound disability, animals, artificial agents, and differences among agency capacities.

Interest account

The view that a being has standing when things can go better or worse for it. The theory must explain what constitutes welfare or interest and whether a subject of experience is required.

Relational account

The view that relationships, dependence, care, recognition, or social connection partly ground standing or special concern. It highlights actual dependency and risks making isolated beings count less.

Membership account

The view that belonging to a community, species, polity, or practice grounds status. It must explain entry, exit, outsiders, imposed membership, and arbitrary boundary selection.

Incommensurability

The condition in which two values cannot be fully measured on one common scale. Incommensurability need not imply that choice is impossible or arbitrary; agents can use priority, thresholds, plural judgment, or procedure.

Incomparability

The stronger relation in which neither option is better, worse, nor equal under the relevant ordering. Philosophers distinguish incomparability from rough equality, parity, and mere measurement difficulty.

Aggregation

Combining benefits, harms, preferences, or other quantities across persons or events. Aggregation requires a currency, entry rule, interpersonal comparison, time horizon, and distribution rule.

Prioritarianism

The view that benefiting someone matters more, other things equal, the worse off that person is. Priority differs from strict equality and from maximizing an unweighted total.

Sufficientarianism

The view that bringing people above an adequate threshold has special importance. It must specify the threshold, treatment below and above it, and how scarcity is handled.

Protected boundary or side constraint

A rule that prohibits certain ways of treating a bearer even when violation would improve aggregate outcomes. Boundaries require bearers, scope, exceptions, and an account of conflict.

Threshold

A point at which the priority or force of a rule changes—for example, when catastrophe permits overriding an ordinary boundary. A threshold requires content, evidence, authority, and safeguards.

Lexical priority

An ordering in which one value must be satisfied before another receives weight, or in which no amount of the lower-ranked value outweighs the higher except under specified conditions.

Value pluralism

The view that several genuine values exist and may conflict without reduction to one master value. Pluralism can be realist, constructivist, or stance-based. It needs a practical response to conflict and residual loss.

First-order theory families

Consequentialism

The family of views assessing actions, rules, or institutions by their consequences. Variants differ over the relevant consequences, currency, aggregation, distribution, rule level, and treatment of uncertainty.

Deontology

The family of views emphasizing duties, rights, constraints, universal principles, and limits on using persons merely as means. Variants differ over status, intention, positive duties, conflict, and emergency thresholds.

Contractualism

The family of views assessing principles through what affected persons could reasonably accept or reject, often with emphasis on individual complaints and reciprocal justification rather than aggregate welfare.

Contractualism differs from contractarian mutual advantage, though real theories can combine elements.

Contractarianism

The family of views grounding norms in agreement, bargain, or mutual advantage among agents. It faces questions about bargaining power and beings unable to offer reciprocal benefit.

Virtue ethics

The family of views centering excellent character, flourishing, practical wisdom, and the formation of perception and motivation rather than treating morality as only a rule for isolated acts.

Care ethics

The family of views centering dependency, relationship, vulnerability, attention, emotion, and the social organization of care. Care ethics also studies how allegedly neutral rules can hide unequal labor and power.

Divine-command theory

The family of views grounding obligation in divine commands or will. Sophisticated versions connect command with perfect goodness, creation, trust, and relationship rather than raw power.

Divine-nature theory

The view that value or goodness is grounded in, or identical with, a necessarily good divine nature. It aims to avoid arbitrary command and must address content and access to that nature.

Natural-law theory

The family of views grounding practical principles in the intelligible goods, ends, or flourishing of beings within a natural or created order. Natural law is not the claim that whatever occurs in nature is permissible.

Ideal-observer theory

The view that correct evaluation is linked to the response of an observer who is fully informed, impartial, coherent, vivid in appreciation, or otherwise idealized. The theory must specify the ideal traits and status of the response.

Responsibility, emotion, and power

Basic desert

The claim that an agent merits benefit or suffering because of what they did or are, independently of further protection, deterrence, repair, relationship, communication, or other practical purpose.

Retributivism

The family of views on which deserved punishment is intrinsically appropriate or valuable, not only a means to future benefit. Retributivists often argue that desert protects proportionality and recognizes agency.

Forward-looking responsibility

An approach organizing responses around protection, deterrence, rehabilitation, repair, trust, communication, and institutional change rather than intrinsic payback.

Moralized emotion

Guilt, shame, resentment, indignation, gratitude, or pride when the emotion presents an act or person as blameworthy, shameful, admirable, obligated, or deserving. Cause, appraisal, function, expression, intensity, and duration should be separated.

Legitimacy

The status of a decision or institution as appropriately authorized under specified terms, procedures, constituencies, and limits. Legitimacy can support compliance amid disagreement without establishing moral truth.

Legal jurisdiction

Socially recognized authority of an institution to make and enforce rules within a domain. Legal jurisdiction is distinct from objective moral authority.

Coercion

Intentional narrowing of another's options through threat or imposed cost. Coercion can occur through states, employers, families, creditors, gatekeepers, or private actors.

Force

Direct overcoming of resistance through seizure, confinement, incapacitation, or physical violence. Force may be official or private, limited or lethal.

Domination

Durable exposure to another's uncontrolled or arbitrary power, even when no explicit threat is currently issued. Safeguards against domination include divided power, rules, review, voice, exit, resources, and counterorganization.

Compact position map

Use this blank map to locate a view without forcing it onto one line.

Meaning: report / express / apply / construct / fictionalize / mixed: _____

Truth status: realist / error-theoretic / relative / constructed / suspended: _____

Dependence: actual stance / ideal stance / authority / natural fact / rational relation / basic evaluative fact / hybrid: _____

Practical reasons: internalist / externalist / hybrid / skeptical: _____

Epistemic reasons: internalist / externalist / anti-normativist / mixed: _____

Access: observation / intuition / reason / revelation / procedure / none claimed / other: _____

Standing: sentience / agency / interest / relation / membership / construction / objective status / plural: _____

Comparison: aggregation / priority / boundary / threshold / lexical / plural judgment / procedure / choice: _____

First-order guidance: consequential / deontological / contractual / virtue / care / received order / plural / hybrid: _____

Responsibility: desert / protection / repair / communication / trust / mixed: _____

Public settlement: agreement / vote / law / bargain / coercion / force / other: _____

Language policy: abolition / fiction / conservation / expression / construction / pragmatic translation: _____

Two views that match in one blank may oppose one another in every other blank. That is why the map is a grid rather than a line.

Appendix F — Notes on the Limits of the Conclusion

This appendix restates the book’s result in compact argumentative form. Its purpose is restraint. A skeptical conclusion can become incoherent by claiming more than its own evidence supports.

1. The target claim

The construction project did not ask merely for a rule people could follow. It sought a system satisfying a bundle of central aspirations.

Let a candidate system be **S**. The target bundle includes:

- **T(S): Truth** — at least some central verdicts are truth-apt and true;
- **I(S): Stance independence** — their truth is not constituted merely by actual approval, preference, commitment, convention, command, or procedure;
- **C(S): Content** — the system identifies bearers, values, statuses, rules, and exceptions;
- **A(S): Authority** — relevant considerations count in favor for agents within scope, including informed rejecters who lack connected aims;
- **K(S): Accessibility** — inquirers possess a defensible route to the central content and can distinguish success from projection or error;
- **H(S): Coherence** — components fit without contradiction, equivocation, circular support, hidden selectors, or opportunistic exception;
- **D(S): Determinacy** — the system settles enough conflicts to guide action or identifies principled indeterminacy;
- **U(S): Usability** — finite agents and institutions can apply it;
- **R(S): Correctability** — interpretation and application can respond to evidence and established error without becoming whatever current power favors.

The complete target can be represented as:

$$O(S) = T(S) \wedge I(S) \wedge C(S) \wedge A(S) \wedge K(S) \wedge H(S) \wedge D(S) \wedge U(S) \wedge R(S)$$

Individual readers may mark some conditions negotiable. If a condition is surrendered, the resulting project may remain coherent and valuable. It no longer satisfies **O** as originally specified.

2. The grounding partition

For every examined candidate, the decisive selector receives its content from one or more of these families:

1. **G₁ — Actual stance:** preference, emotion, commitment, convention, agreement, role, culture, institution, or practice.
2. **G₂ — Idealized stance:** endorsement under specified information, coherence, freedom, impartiality, representation, or procedure.
3. **G₃ — Authority:** command, will, nature, purpose, revelation, office, or relation.
4. **G₄ — Non-evaluative natural fact:** suffering, desire, function, capacity, need, causal condition, or evolutionary and social fact.
5. **G₅ — Rational or conceptual requirement:** consistency, universalization, constitutive agency, implication, or conceptual truth.
6. **G₆ — Basic evaluative fact or aim-independent reason:** irreducible standing, value, obligation, or favoring relation.
7. **G₇ — Hybrid:** an ordered combination of G₁–G₆.

This is a working partition of the grounding routes examined, not a theorem proving that no eighth family is logically possible. A proposed eighth family must show that it is not an actual or ideal stance, authority, natural fact, rational relation, basic evaluative fact, or combination under another name.

3. The audit findings

The argument uses the following conditional findings.

Finding G₁

Actual stances generate content, motivation, meaning, commitments, roles, conventions, institutions, conditional reasons, and enforcement. If a verdict's truth or authority is constituted by those stances, **I(S)** or universal **A(S)** is not preserved.

This finding does not imply that stance-dependent practices are imaginary, weak, or dispensable.

Finding G₂

Idealization can correct factual error, inconsistency, self-dealing, missing perspectives, and distorted preference. A contentful procedure requires selectors concerning participants, information, motivation, reasonableness, representation, constraints, and decision rules.

If those selectors are grounded in actual or ideal stances, dependence remains. If they are objectively authoritative, their ground moves to G₃, G₅, or G₆. The procedure does not select its own ideal from nothing.

Finding G₃

Authority can explain direction, coordination, role, trust, or enforcement. Raw power supplies motivation or prudence, not moral content. A rightful authority requires legitimacy, goodness, perfect nature, or another evaluative relation.

If goodness identifies the rightful source, goodness carries the decisive grounding and access debt.

Finding G₄

Natural facts supply indispensable information about experience, capacity, need, function, consequence, cause, and institutional effect. A terminal selection—who has standing, which end is valued, how burdens are distributed, or what counts in favor for every agent—requires an evaluative identity, bridge, stance, procedure, or basic reason.

The finding does not deny moral naturalism. An identity proposal must establish which natural property performs the full moral role, including authority and access.

Finding G₅

Rational and conceptual requirements expose contradiction, invalid inference, self-exemption, and means–end failure. Formal consistency does not by itself supply substantive standing, ends, concern, or priority. If equal status or reciprocal respect is constitutive of rationality, that substantive content must be defended.

The finding does not imply that reason is powerless or that logic depends on preference.

Finding G₆

Basic evaluative facts preserve stance independence most directly. Foundational stopping points are not automatically illegitimate. A complete system still needs enough content, scope, comparison, authority, and access to satisfy **O(S)**.

In the candidates examined, at least one decisive component remained asserted at the stopping point without a sufficiently discriminating access or conflict account. This leaves the system possible but not fully earned.

Finding G₇

Hybrids can distribute practical work among several components and may be more humane, usable, and corrigible than pure systems. The decisive ordering rule still belongs to G₁–G₆ or remains unspecified. The hybrid inherits the debts of that source and adds a meta-ordering debt.

4. The conclusion actually derived

Let E be the set of candidate routes examined in the book. Let $W(S)$ mean that the system has supplied warrants sufficient for the target properties rather than merely asserting them.

The inspection supports:

For every S in E , the project did not establish $W(O(S))$.

In ordinary language:

No examined route has earned a system that jointly satisfies the complete objective specification without an unsupported selector, unpaid bridge, circular support, or declared surrender of an original ambition.

The book therefore concludes:

No coherent objective moral system has been found through the routes inspected.

Here *found* means constructed or identified with adequate warrant for the project's specifications. It does not mean that a search of every logically possible world has been completed.

5. Conclusions not derived

The argument does **not** by itself establish any of the following.

Not derived: metaphysical impossibility

There is no logically possible objective moral fact.

To reach this, one needs a premise showing that the grounding partition is modally exhaustive and that every member is impossible, not merely unpaid in the candidates examined.

Not derived: actual nonexistence

No objective moral fact actually exists.

This may be supported by an additional explanatory argument—for example, that moral appearances are better explained without such facts and that real objective facts would have detectable effects or access. The book does not assume those premises.

Not derived: universal unknowability

If objective moral facts exist, no one could know them.

The access audit finds no established route sufficient for the complete system. It does not prove that no future method, argument, experience, or conceptual advance could supply access.

Not derived: global epistemic skepticism

No evidence can support any belief.

Moral and epistemic normativity remain separate. Evidence can indicate truth, serve belief, advance truth-seeking, support practical aims, discipline inquiry, or satisfy public procedure. A stance-independent epistemic reason, if it exists, does not supply moral content by itself.

Not derived: practical nihilism

Nothing matters to anyone, no action has effects, and no institution can be supported.

Pain, attachment, preference, commitment, evidence, agreement, law, power, protest, protection, repair, and hope remain. What is missing is a paid warrant for stance-independent moral authority, not the existence of practical life.

Not derived: egoism

Everyone should pursue only self-interest.

This is itself a normative claim. Compassion, loyalty, identity, and concern for strangers are actual stances alongside self-concern. Moral skepticism does not crown egoism as objectively correct.

Not derived: equal preference strength

Every value and preference has equal practical force.

People differ in intensity, persistence, bargaining power, relationship, and willingness to sacrifice. The absence of objective rank does not create objective equality.

Not derived: legal abolition

No law or public code can exist legitimately or coherently.

Codes can be factual, instrumental, agreement-based, role-based, procedural, enforced, transparent, and corrigible. Their support, legitimacy, and jurisdiction must not be inflated into objective obligation.

6. Six claims that must remain separate

Claim	Meaning	What would support it	What it does not entail
No candidate is coherent	Every candidate contradicts itself or fails its own stable rules	Internal analysis of each candidate	That no useful practice remains
No candidate is warranted	None has adequate support for all claimed properties	Burden and access audit	That every candidate is false
No objective moral fact exists	Reality contains no stance-independent moral fact	Metaphysical or explanatory argument	That people cease caring
No objective moral fact is knowable	Such facts, if any, cannot be known	General access-impossibility argument	That no such fact exists
No objective moral system is available	No sufficiently contentful, accessible system can guide current agents	Construction and epistemic audit	That no isolated objective fact exists
No useful human normative practice remains	Commitments, rules, reasons, criticism, and coordination all collapse	A very broad practical argument	Any of the previous five by themselves

The main book supports the second claim in relation to the complete specification and, in practical terms, the fifth for the examined routes and current construction. It does not support the first as a universal statement: many candidates are internally coherent after narrowing. It does not establish the third, fourth, or sixth.

7. Burden-based rather than omniscient skepticism

The conclusion follows a burden rule:

A system may count an objective property as earned only when it supplies enough ground, content, scope, comparison, authority, access, and correction to perform the job claimed.

When the warrant is incomplete, the disciplined classification is **unestablished**, **promissory**, or **unavailable under current methods**. It is not automatically **false**.

This burden is not applied only to realism. Non-realist alternatives must also disclose their selectors, power, language policy, affected parties, and practical limits. “Constructed” does not make a code humane. “Expressive” does not make persuasion honest. “Local” does not protect minorities. “Pragmatic” does not authorize opportunism.

8. The epistemic stance of the argument

The book can be followed under several views of epistemic force.

Objective epistemic reasons

If evidence supplies stance-independent reasons for belief, the audit claims that the evidence presented does not warrant completion of **O(S)**. That epistemic norm does not fill the missing moral selectors.

Truth-seeking commitment

If the reader values accurate belief, the audit is conditionally supported as a method for avoiding overclaim and distinguishing truth from projection.

Inquiry-practice account

If evidence and revision are standards of a chosen inquiry practice, the conclusion reports what that practice currently licenses.

Epistemic anti-normativism

If no literal epistemic ought exists, the audit can still describe tracking relations, coherence, prediction, public warrant, and the conditional consequences of seeking truth.

No step infers moral non-realism from epistemic anti-normativism or the reverse.

9. What could change the result?

The conclusion is corrigible. Confidence should change if a candidate supplies a noncircular account that jointly establishes:

1. what moral truth consists in;
2. which content is stance-independent;
3. who or what has standing;
4. what gives relevant agents practical reasons;
5. how values compare and conflicts resolve;
6. how exceptions remain principled;
7. how inquirers gain reliable access;
8. how disagreement and genealogy are handled;
9. how finite agents use the system;
10. how interpretation and application are corrected;
11. why the foundation falls outside or succeeds within the seven-family audit;
12. which original specification, if any, is narrowed.

Evidence against the book could include a successful derivation from rational agency without hidden substantive premises, a calibrated access route to irreducible values, a natu-

ral identity account performing the full authority role, a defensible received-order account resolving interpretation without circularity, or another grounding family not captured by the partition.

The book's conclusion should not be protected from these possibilities by definition.

10. Final scope statement

The main result can be summarized in four sentences.

1. **Every examined objective candidate protects real human concerns and contributes useful practices.**
2. **Every route examined leaves at least one decisive objective selector, bridge, ordering rule, or access claim unsupported, circular, or fundamental at the point where the complete system requires a warrant.**
3. **Coherent, humane, informed, enforceable, and revisable public codes can be built by declaring their values, procedures, jurisdictions, safeguards, and power rather than presenting them as objectively certified.**
4. **The failure to find a warranted complete system is not a proof that no objective moral fact could exist; it is a reason not to claim more authority than the construction has earned.**

That is the boundary of the argument.

Notes and Further Reading

These notes identify the main philosophical conversations behind each chapter. The book's cases, taxonomies, Candidate System, Debt Ledger, and cumulative audit are original constructions. A source listed for a chapter should not be read as endorsing that chapter's exact formulation or the book's final conclusion.

Prologue

The construction method and its nine-part target grow from Phil Stilwell's *How to Test a "Moral" System* and *The Conditions of Moral Coherence*. The blank constitution is a narrative device, not a reconstruction of a historical constitutional convention.

Chapter 1 — The Promise

For opposed accounts of what objectivity promises, compare David Enoch's defense of robust realism, Russ Shafer-Landau's nonnaturalist realism, and J. L. Mackie's error-theoretic challenge. Enoch is especially important for the claim that objectivity matters to deliberation and moral seriousness; the chapter's five hopes are the present book's specification, not Enoch's list.

Chapter 2 — What "Objective" Must Mean

Shafer-Landau and Enoch defend stance-independent moral truth. Sharon Street distinguishes constructivist dependence on a practical standpoint from realism and ideal-response reduction. The chapter deliberately separates universality, impartiality, convergence, response dependence, and stance independence because philosophical positions combine these properties in different ways.

Chapter 3 — What Are You Doing When You Say "Wrong"?

C. L. Stevenson's "The Emotive Meaning of Ethical Terms" is a classic account of ethical language as both descriptive and dynamically practical. Simon Blackburn shows how antirealist moral thought can retain reasoning and criticism. Richard Joyce's work on error theory and fictionalism supplies contrasting cognitivist and revisionary possibilities. The chapter's six voices are a teaching taxonomy rather than a standard six-position canon.

Chapter 4 — The Specification Sheet

The joint-performance test and explanatory-debt method are developed in Stilwell's three project papers. For realist accounts that accept irreducible normative foundations while addressing motivation, reasons, and knowledge, see Shafer-Landau and Enoch.

Chapter 5 — One Argument, Eight Disagreements

Stevenson separates disagreement in belief from disagreement in attitude. David Plunkett and Timothy Sundell show how apparently factual normative disputes may also negotiate concepts and linguistic frameworks. The eight-room map extends those insights into standing, priority, coordination, procedure, and objective status.

Chapter 6 — Draw the Circle

For a survey of proposed grounds of moral status—including sentience, agency, cognitive capacities, relationships, and species membership—see “The Grounds of Moral Status” in the *Stanford Encyclopedia of Philosophy*. The chapter's strict separation of entry, weight, and outcome is an editorial discipline of this book.

Chapter 7 — Decide What Receives a Verdict

The distinction among acts, motives, outcomes, rules, institutions, agents, relationships, and forms of life draws on familiar divisions across normative ethics. Its purpose here is grammatical: conclusions about one object may not migrate to another without an argument.

Chapter 8 — Name What Matters

John Stuart Mill, Immanuel Kant, Aristotle, and Virginia Held provide classic or influential treatments of welfare, rational agency, virtue, and care. The chapter does not claim that those traditions reduce to one value each; it uses value families to expose what later rules must select and compare.

Chapter 9 — Turn Values into Guidance

The generation-rule taxonomy is synthetic. It separates maximizing, constraining, protecting, embodying, caring, honoring, distributing, obeying, negotiating, and constructing because a shared value can generate different instructions under different normative architectures.

Chapter 10 — Choosing Without a Common Measure

See the *Stanford Encyclopedia of Philosophy* entries “Weighing Reasons” and “Incommensurable Values.” Derek Parfit’s work is central to contemporary debates over aggregation, priority, and principles for comparing burdens. A practical settlement need not establish that values possess a hidden universal unit.

Chapter 11 — Like Cases, Exceptions, and Identity

Kant’s universal-law formulations and Rawls’s use of impartial choice are major sources for identity-neutral testing. The rotation test is the book’s broader instrument for separating relevant relationships and roles from opportunistic self-exemption.

Chapter 12 — Three Different Things Called a Reason

See “Reasons for Action: Justification, Motivation, Explanation” in the *Stanford Encyclopedia of Philosophy*. Bernard Williams’s “Internal and External Reasons” frames the modern dispute over whether reasons must connect with an agent’s motivational set.

Chapter 13 — Current Aims and Better-Informed Aims

Williams, Mark Schroeder, and Sharon Street develop influential desire-based or practical-standpoint accounts. The chapter’s sequence from present desire through information, vivid appreciation, coherence, and integration is a pressure map of idealization, not a summary attributable to any one of them.

Chapter 14 — Commitments, Promises, and Roles

T. M. Scanlon’s account of mutual justification and Christine Korsgaard’s work on practical identity illuminate how relations and self-constitution can generate reasons. The seven-clause commitment charter is the book’s institutional extension of those questions.

Chapter 15 — Fair Terms and Constructed Reasons

John Rawls’s original position, Scanlon’s reasonable-rejection contractualism, Korsgaard’s Kantian constructivism, and Street’s Humean constructivism are importantly different. The chapter keeps ideal agreement, actual consent, reciprocal justification, and political legitimacy separate.

Chapter 16 — Reasons That Apply Whether You Care or Not

For defenses of stance-independent reasons or normativity, see Enoch, Shafer-Landau, Derek Parfit, and Korsgaard. For internalist pressure, see Williams and Schroeder. The chapter grants the reasons-first theorist a basic favoring relation and asks what content, scope, weight, defeat conditions, and access complete it.

Chapter 17 — From Approval to Obligation

Gilbert Harman's "Moral Relativism Defended" and Bernard Williams's "The Truth in Relativism" show why relativism should not be reduced to the slogan that current majorities are always right. The seven-column classification is original to this book.

Chapter 18 — God, Nature, and Rational Necessity

Robert Merrihew Adams develops an ethics grounded in divine goodness rather than arbitrary command. John Finnis offers a modern natural-law account organized around basic goods and practical reason. Kant and Korsgaard anchor the rational-agency and autonomy routes. The chapter's pressure questions distinguish these positions rather than treating them as one foundation.

Chapter 19 — Does Evidence Have Authority?

Terence Cuneo argues from the parity of moral and epistemic normativity. Richard Joyce examines the companions-in-guilt strategy from an antirealist perspective. Laura Frances Callahan and Benjamin Kiesewetter address whether and why epistemic reasons are normative. The chapter's conditional inquiry commitment is meant to avoid assuming a global answer.

Chapter 20 — Intuition and the Feeling of Obviousness

Shafer-Landau defends self-evident moral principles within a realist epistemology. The chapter grants noninferential justification as a live possibility while separating seeming, confidence, emotion, training, cultural fluency, conceptual competence, and intellectual insight.

Chapter 21 — Where Values Come From

Sharon Street’s “A Darwinian Dilemma for Realist Theories of Value” presents a major evolutionary challenge. Richard Joyce’s “Irrealism and the Genealogy of Morals” distinguishes stronger and weaker debunking ambitions. The chapter’s conclusion is correspondingly bounded: a genealogy may undermine, vindicate, or merely explain depending on the process and target.

Chapter 22 — Disagreement and Convergence

Harman, Williams, Stevenson, and Plunkett and Sundell supply different accounts of what disagreement can involve. Persistent disagreement bears differently on truth, access, determinacy, social formation, and coordination; it does not by itself entail error theory.

Chapter 23 — Facts Enter; Values Select

This chapter states a methodological division rather than a doctrine that facts never affect values. Evidence can correct descriptions, reveal consequences, identify capacities, activate admitted rules, and transform commitments. The remaining selector question concerns which described feature receives standing, priority, or acceptable risk.

Chapter 24 — Novel Cases, Uncertainty, and Correction

The chapter combines the standing, evidence, and comparison literatures cited above. Its versioning protocol is an original institutional proposal for keeping empirical update, conceptual clarification, evaluative revision, and enforcement change visible.

Chapter 25 — Count the Consequences

John Stuart Mill’s *Utilitarianism* is the classic starting point. Derek Parfit’s *Reasons and Persons* and *On What Matters* develop demanding modern treatments of aggregation, distribution, population, and convergence among moral theories. The chapter presents a family broader than classical hedonism.

Chapter 26 — Respect the Rule and the Person

Kant’s *Groundwork of the Metaphysics of Morals* anchors universal law, autonomy, and persons as ends. Scanlon’s contractualism offers a different route to individual complaints and mutual justification. The chapter deliberately includes positive duties, non-rational beings, conflicts, and catastrophe thresholds rather than treating deontology as a single absolute ban.

Chapter 27 — Character, Care, and Belonging

Aristotle's *Nicomachean Ethics* is the classic source for virtue, formation, flourishing, friendship, and practical wisdom. Virginia Held develops care as a distinct personal, political, and global approach. The chapter combines virtue and care for comparison while preserving their differences.

Chapter 28 — Receive the Moral Order

Adams and Finnis provide the principal modern sources for the chapter's divine-nature and natural-law steelmen. The chapter also treats tradition and conscience as access routes, neither of which follows automatically from those two accounts.

Chapter 29 — Build What No One Could Reasonably Reject

Rawls, Scanlon, Korsgaard, and Street represent distinct constructive projects: impartial political choice, reasonable rejection, Kantian self-legislation, and Humean construction from the practical standpoint. Democratic legitimacy is treated here as accountable power, not automatically as moral truth.

Chapter 30 — Let Standards Remain Local or Plural

Harman and Williams are central to the chapter's refusal to identify relativism with majority polling. The localist and pluralist candidates include indexed truth, agreement-relative reasons, contextual judgment, and personal commitment; those views need not share one metaethics.

Chapter 31 — The Hybrid Rescue

Rawls's reflective equilibrium is a major model of reciprocal adjustment among principles and judgments. Parfit's *On What Matters* argues for deep convergence among defensible versions of Kantianism, contractualism, and rule consequentialism. The chapter grants hybrid theories practical strength while asking what orders components when they conflict.

Chapter 32 — Criticize the World You Inherited

Harman and Williams motivate the question of criticism from within and across outlooks. Rawls and Scanlon provide constructive and public-justification routes. The chapter's criticism-ground matrix is original and designed to prevent one successful criticism from being inflated into every kind.

Chapter 33 — Build a Public Code Without the Missing Layer

Rawls's political liberalism and Scanlon's mutual-justification project inform the distinction between public legitimacy and metaphysical correctness. The charter labels factual, instrumental, commitment-based, agreement-based, procedural, enforced, and objectively claimed clauses separately.

Chapter 34 — Power, Dissent, and the Person Who Still Says No

The coercion warrant is an original synthesis of familiar rule-of-law requirements: evidence, jurisdiction, proportionality, less intrusive means, review, stopping conditions, and repair. The chapter's central warning is classificatory: force does not become agreement because its users consider it justified.

Chapter 35 — Responsibility Without Cosmic Payback

P. F. Strawson's "Freedom and Resentment" is foundational for the importance of reactive attitudes and interpersonal expectations. Derk Pereboom and Gregg Caruso develop non-desert approaches that preserve protection, moral formation, protest, reconciliation, rehabilitation, and humane incapacitation. The chapter's zero-purpose pain case isolates negative desert by stipulation.

Chapter 36 — Moral Emotions Without Moral Desert

Strawson and Pereboom remain central. The dossier's separation of cause, appraisal, audience, function, expression, intensity, duration, effect, and stopping condition is the book's practical framework rather than a summary of a single psychological theory.

Chapter 37 — No Candidate Returns Intact

Compare Mackie's error theory with the robust realism of Enoch and Shafer-Landau, Cuneo's parity argument, Korsgaard's constitutivism, and Street's constructivism. The seven-family grounding partition is the book's synthesis. Its conclusion concerns what the construction has earned; it does not purport to prove that no logically possible moral fact exists.

Chapter 38 — Decide What to Do with Moral Language

Richard Garner argues for abolition; Richard Joyce develops fictionalist and “what next?” approaches; François Jaquet and Hichem Naar analyze conservationism; Sebastian Köhler and Michael Ridge propose revolutionary expressivism; Jaquet clarifies the taxonomy; Björn Eriksson and Jonas Olson defend negotiationism. The chapter keeps these options separate and allows context-sensitive mixtures.

Chapter 39 — Everything That Did Not Disappear

Stilwell’s *The Parsimonious World of the Moral Non-Realist* is the closest project source. Blackburn and Joyce provide broader antirealist accounts of how thought, criticism, and practice may survive without stance-independent moral facts.

Chapter 40 — Begin with the Primitives

The primitive-to-policy chain develops the constructive direction of Stilwell’s project papers. “Primitive” names an actual practical starting point; it is not a claim that emotions are self-authenticating, biologically uniform, or ready to become law.

Chapter 41 — Stop Reconciling Moral Systems

Eriksson and Olson’s negotiationism is an important companion to the move away from a single post-error replacement. The Civic Stakes Brief is original to this book. It treats theories as attention directors and counterexample generators while giving public standing to affected beings, evidence, declared values, procedures, safeguards, and dissent.

Appendices

The appendices consolidate the book’s original working artifacts. Appendix E’s position map should be used with the sources above because labels such as *realism*, *constructivism*, *relativism*, *expressivism*, and *fictionalism* contain substantial internal disagreement. Appendix F states the formal limits that govern every use of the book’s negative conclusion.

Bibliography

- Adams, Robert Merrihew. *Finite and Infinite Goods: A Framework for Ethics*. New York: Oxford University Press, 1999.
- Aristotle. *Nicomachean Ethics*. Translated by Roger Crisp. 2nd ed. Cambridge: Cambridge University Press, 2014.
- Blackburn, Simon. “How to Be an Ethical Antirealist.” *Midwest Studies in Philosophy* 12, no. 1 (1988): 361–375. <https://doi.org/10.1111/j.1475-4975.1988.tb00173.x>.
- Callahan, Laura Frances. “Why Are Epistemic Reasons Normative?” *Episteme* 21, no. 3 (2024): 1048–1063. First published online 2023. <https://doi.org/10.1017/epi.2023.13>.
- Caruso, Gregg D. “The Public Health-Quarantine Model.” 2017. <https://doi.org/10.2139/ssrn.3068021>.
- Cuneo, Terence. *The Normative Web: An Argument for Moral Realism*. Oxford: Oxford University Press, 2007. <https://doi.org/10.1093/acprof:oso/9780199218837.001.0001>.
- Enoch, David. *Taking Morality Seriously: A Defense of Robust Realism*. Oxford: Oxford University Press, 2011. <https://doi.org/10.1093/acprof:oso/9780199579969.001.0001>.
- Eriksson, Björn, and Jonas Olson. “Moral Practice after Error Theory: Negotiationism.” In *The End of Morality: Taking Moral Abolitionism Seriously*, edited by Richard Garner and Richard Joyce, 113–130. New York: Routledge, 2019.
- Finnis, John. *Natural Law and Natural Rights*. 2nd ed. Oxford: Oxford University Press, 2011.
- Garner, Richard. “Abolishing Morality.” *Ethical Theory and Moral Practice* 10 (2007): 499–513. <https://philpapers.org/rec/GARAM-2>.
- Harman, Gilbert. “Moral Relativism Defended.” *Philosophical Review* 84, no. 1 (1975): 3–22. <https://doi.org/10.2307/2184078>.
- Held, Virginia. *The Ethics of Care: Personal, Political, and Global*. New York: Oxford University Press, 2006. <https://doi.org/10.1093/0195180992.001.0001>.
- Jaquet, François. “Sorting Out Solutions to the Now-What Problem.” *Journal of Ethics and Social Philosophy* 17, no. 3 (2020): 239–258. <https://doi.org/10.26556/jesp.v17i3.812>.
- Jaquet, François, and Hichem Naar. “Moral Beliefs for the Error Theorist?” *Ethical Theory and Moral Practice* 19, no. 1 (2016): 193–207. <https://doi.org/10.1007/s10677-015-9609-1>.
- Joyce, Richard. *The Myth of Morality*. Cambridge: Cambridge University Press, 2001.
- . “Moral Anti-Realism and the Companions-in-Guilt Strategy.” 2009. <https://philpapers.org/archive/JOYMAE.pdf>.

- . “Moral Fictionalism: When Falsehoods Are Too Useful to Throw Out.” 2011. <https://philpapers.org/rec/JOYMF>.
- . “Irrealism and the Genealogy of Morals.” 2014. <https://catalogimages.wiley.com/images/db/pdf/9781118837412.excerpt.pdf>.
- . “Moral Skepticism and the ‘What Next?’ Question.” 2019. <https://philpapers.org/rec/JOYMSA>.
- . *Morality: From Error to Fiction*. 2024. <https://www.richard-joyce.com/morality-from-error-to-fiction>.
- Kant, Immanuel. *Groundwork of the Metaphysics of Morals*. Translated and edited by Mary Gregor and Jens Timmermann. Cambridge: Cambridge University Press, 2012. <https://doi.org/10.1017/CBO9780511809590>.
- Kiesewetter, Benjamin. “Are Epistemic Reasons Normative?” *Noûs* 56, no. 3 (2022): 670–695. <https://doi.org/10.1111/nous.12379>.
- Köhler, Sebastian, and Michael Ridge. “Revolutionary Expressivism.” *Ratio* 26, no. 4 (2013): 428–449. <https://doi.org/10.1111/rati.12030>.
- Korsgaard, Christine M. *The Sources of Normativity*. Cambridge: Cambridge University Press, 1996. <https://doi.org/10.1017/CBO9780511554476>.
- Mackie, J. L. *Ethics: Inventing Right and Wrong*. Harmondsworth: Penguin, 1977.
- Mill, John Stuart. *On Liberty, Utilitarianism and Other Essays*. Edited by Mark Philp and Frederick Rosen. Oxford: Oxford University Press, 2015.
- Parfit, Derek. *Reasons and Persons*. Oxford: Clarendon Press, 1984.
- . *On What Matters*. Vols. 1–2. Oxford: Oxford University Press, 2011.
- Pereboom, Derk. *Free Will, Agency, and Meaning in Life*. Oxford: Oxford University Press, 2014.
- . “Responsibility, Regret, and Protest.” In *Oxford Studies in Agency and Responsibility*, vol. 4, edited by David Shoemaker, 121–140. Oxford: Oxford University Press, 2017. <https://doi.org/10.1093/oso/9780198805601.003.0007>.
- Plunkett, David, and Timothy Sundell. “Disagreement and the Semantics of Normative and Evaluative Terms.” *Philosophers’ Imprint* 13, no. 23 (2013): 1–37.
- Rawls, John. *A Theory of Justice*. Rev. ed. Cambridge, MA: Harvard University Press, 1999.
- . *Political Liberalism*. Expanded ed. New York: Columbia University Press, 2005.
- Scanlon, T. M. *What We Owe to Each Other*. Cambridge, MA: Belknap Press of Harvard University Press, 1998. <https://doi.org/10.2307/j.ctv134vmrn>.
- Schroeder, Mark. *Slaves of the Passions*. Oxford: Oxford University Press, 2007. <https://doi.org/10.1093/acprof:oso/9780199299508.001.0001>.
- Shafer-Landau, Russ. *Moral Realism: A Defence*. Oxford: Oxford University Press, 2003. <https://doi.org/10.1093/0199259755.001.0001>.

Stanford Encyclopedia of Philosophy. "Blame."

———. "The Grounds of Moral Status."

———. "Incommensurable Values."

———. "Moral Responsibility."

———. "Reasons for Action: Internal vs. External."

———. "Reasons for Action: Justification, Motivation, Explanation."

———. "Weighing Reasons."

Stevenson, C. L. "The Emotive Meaning of Ethical Terms." *Mind* 46, no. 181 (1937): 14–31. <https://doi.org/10.1093/mind/XLVI.181.14>.

Stilwell, Phil. *How to Test a "Moral" System: A Field Guide to Metaethics, Value Disagreement, and the Architecture of Values*. 2026. <https://www.academia.edu/169114427/>.

———. *The Conditions of Moral Coherence: A Non-Realist Framework for Evaluating Moral Systems*. 2026. <https://www.academia.edu/169048095/>.

———. *The Parsimonious World of the Moral Non-Realist: Illegitimate Questions, Surviving Values, and Life Without a Moral Realm*. 2026. <https://www.academia.edu/169490990/>.

Strawson, P. F. "Freedom and Resentment." *Proceedings of the British Academy* 48 (1962): 1–25.

Street, Sharon. "A Darwinian Dilemma for Realist Theories of Value." *Philosophical Studies* 127 (2006): 109–166.

———. "What Is Constructivism in Ethics and Metaethics?" *Philosophy Compass* 5, no. 5 (2010): 363–384. <https://doi.org/10.1111/j.1747-9991.2009.00280.x>.

Williams, Bernard. "The Truth in Relativism." *Proceedings of the Aristotelian Society* 75 (1975): 215–228. <https://doi.org/10.1093/aristotelian/75.1.215>.

———. "Internal and External Reasons." In *Moral Luck: Philosophical Papers 1973–1980*, 101–113. Cambridge: Cambridge University Press, 1981.